

Do Finance Journals Sort on Policy Conclusions?

Emre Kuvvet*

June 5, 2026

ABSTRACT

We study whether top finance journals favor certain policy conclusions. Using 121 published articles in the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* alongside 125 NBER working papers on government intervention in financial markets (2000–2024), we find that journals do not sort on finding direction. Whether a paper concludes that a regulation is beneficial, harmful, or inconsequential, it reaches a top journal at roughly the same rate, about 46 percent. The strongest predictor of publication is research design: papers using quasi-experimental methods are about 18 percentage points more likely to be published, regardless of what they find. We also find that published papers show modest threshold bunching in supplementary test statistics but not in their main regression tables, suggesting that any specification searching is concentrated at the margins of the paper rather than in its core results.

JEL classification: G18, G28, C80, B40.

Keywords: Publication bias; Government intervention; Finance journals; Selective reporting; Regulatory finance; p-hacking

*Department of Finance, Nova Southeastern University. E-mail: ekuvvet@nova.edu.

1. Introduction

Academic research on government intervention in financial markets occupies a privileged position in public policy. The Federal Reserve Board, the Securities and Exchange Commission, and the Basel Committee on Banking Supervision routinely incorporate peer-reviewed finance research into regulatory impact assessments, rule-makings, and policy speeches. The Dodd-Frank Wall Street Reform and Consumer Protection Act alone generated an estimated 400 new regulatory mandates, many of which were designed in explicit reference to academic evidence on the effects of capital requirements, disclosure rules, and consumer financial protection. The credibility of this evidence base is therefore a first-order policy question: if the academic record systematically misrepresents which interventions have measurable effects—and which do not—regulators face a distorted picture that can generate welfare-reducing policy choices.

Publication bias offers a mechanism by which such distortion arises without any deliberate misconduct. When researchers are reluctant to write up null results, when journals are more likely to accept papers with statistically significant findings, or when both effects operate simultaneously, the published literature understates the true prevalence of null or inconclusive evidence (?). In a domain such as government intervention research, where policy decisions depend on whether a regulation has a measurable causal effect, suppression of null results generates the false impression that most interventions work—or, more precisely, that most interventions have discernible effects in one direction or the other.

Why might directional publication bias arise in the government-intervention literature specifically? Two theoretical channels operate in opposite directions, making the sign of any bias an open empirical question. A *pro-intervention channel* could arise if journals and readers favor results that confirm a policy’s stated objective (e.g., that capital requirements reduce bank risk), both because positive results are more actionable for regulators and because they more naturally confirm the prior that motivated the policy. An *anti-intervention channel* could arise if research that documents unintended costs or failures of regulation is

more surprising—and thus more publishable—than research confirming that a policy works as intended. Under either channel, the direction of sorting would differ from significance bias (null vs. directional), which reflects a simpler threshold between finding-something and finding-nothing. This paper is designed to detect both types.

This paper studies sorting patterns in a domain the existing literature has not examined: the finance literature’s treatment of government intervention. Prior work on publication bias in finance has focused on the cross-section of asset pricing anomalies, documenting that factor-zoo proliferation reflects selection for high t-statistics (??). A parallel economics literature (??) infers bias from the shape of published p-value distributions, finding excess mass just below conventional significance thresholds. Neither strand addresses the policy-intervention literature, nor do they distinguish two conceptually distinct forms of bias with different welfare implications. *Significance bias*—directional findings over-published relative to null findings—makes every intervention appear detectable. *Directional bias*—pro-intervention findings over-published relative to anti-intervention findings—additionally skews the regulatory record toward a particular conclusion about the sign of policy effects. We design our study to test for both.

Our research design combines two complementary samples. The *published paper* sample consists of all in-scope articles in the *Journal of Finance*, the *Review of Financial Studies*, and the *Journal of Financial Economics* from 2000 to 2024. The *working paper* sample is drawn from NBER working papers in five programs covering the dominant share of academic finance research on government intervention: Corporate Finance, Asset Pricing, Monetary Economics, Economic Fluctuations and Growth, and Public Economics. The key advantage of this design is that we observe findings from papers that circulated *before* potential journal publication, allowing a direct two-corpus comparison rather than inferring the pre-publication distribution from the shape of published p-values alone.

Because only 12% of published papers are matched to a confirmed NBER counterpart, the two corpora are largely non-overlapping cross-sections rather than the same papers tracked

through the editorial filter. We therefore interpret all estimates as cross-sectional associations between paper characteristics and corpus membership, not as causal editorial-filter effects; the full identification discussion is in Section ??.

We define a paper as in-scope if its primary research question concerns the causal effect of a specific government law, regulation, or public policy on financial markets, firms, or investors. Identifying these papers across nearly 40,000 documents required an automated approach: a keyword filter followed by a large language model (LLM) classifier that reads each abstract and assigns an in-scope judgment along with a directional code—positive (+1, the intervention was beneficial), negative (−1, harmful), or null/mixed (0).

Our central result is that finding direction does not predict publication. Papers concluding that a regulation helps, papers concluding it hurts, and papers finding no significant effect all reach the three top journals at essentially the same rate—around 46 percent. This uniformity holds unconditionally and survives every regression specification we estimate, including controls for publication year, decade, team size, and research design quality. The complete absence of a direction premium holds for positive findings, for negative findings, and when positive and negative are pooled against null.

What does predict publication is how a paper was designed. Papers using quasi-experimental identification strategies—regression discontinuity designs, instrumental variables, difference-in-differences, or event studies—are roughly 18 percentage points more likely to appear in a top journal than observational papers on the same policy questions. This methodological premium is the only statistically robust predictor we find; finding direction adds nothing beyond it. We interpret this cautiously: the quasi-experimental indicator likely captures a bundle of correlated quality attributes—author prestige, institutional resources, topic fashionability—not identification strategy alone.

Within published papers, we find a secondary pattern. Test statistics bunch modestly just below the conventional 5 percent significance threshold, but only in standalone statistics that appear primarily in supplementary and robustness tables. Primary regression tables—where

both a coefficient and a standard error are reported and can be independently verified—show no bunching and are statistically indistinguishable from the corresponding NBER working paper tables. This distinction is informative: it suggests that any selective reporting occurs at the margin of revision, in the choice of which additional checks to highlight, rather than in the core empirical results that define the paper’s contribution.

This paper makes three contributions: it provides direct evidence that top finance journals do not sort government-intervention research by the direction of its conclusions; it documents that identification quality—not finding direction—is the dominant sorting criterion, with papers using credible causal designs substantially more likely to be published; and it introduces a validated LLM-assisted pipeline for classifying and direction-coding large paper corpora, with code and prompts publicly available.

The remainder of the paper proceeds as follows. Section 2 reviews the related literature. Section 3 describes data construction and the classification pipeline. Section 4 presents the empirical strategy. Section 5 reports main results and economic magnitudes. Section 6 presents robustness checks. Section 7 discusses mechanisms, policy implications, comparisons to other domains, and limitations. Section 8 concludes.

2. Related Literature

This paper relates to three strands of research. The first is the literature on publication bias in economics and finance. ? document excess mass just below the 5% significance threshold in the top economics journals, consistent with selective reporting of marginally significant specifications. ? show, using NSF-funded pre-registered experiments, that null results are published at roughly one-third the rate of significant ones. In finance, ? and ? document that a large share of published asset-pricing factors fail to replicate, pointing to upward-biased published effect sizes. These papers test for *significance* bias—whether null results are suppressed relative to significant ones. Our paper asks a different question: whether published findings are sorted by the *direction* of the policy conclusion, which is

only meaningful in a domain where findings have a clear normative interpretation. We also contribute a direct pre-publication baseline by using NBER working paper texts, rather than inferring the pre-publication distribution from the shape of published p-values alone.

The second strand is the literature on government intervention in financial markets. A large body of empirical work studies the effects of securities law, banking regulation, and macroprudential policy (??), producing heterogeneous findings across interventions and time periods. Whether the published record faithfully represents the underlying distribution of evidence on these questions has not previously been examined. Our paper fills this gap.

The third strand is the growing use of large language models for text classification in economics (??). We contribute a validated two-pass pipeline—keyword filter followed by LLM classification—for identifying and direction-coding government-intervention finance papers at scale, and we release the coding rubric and prompts publicly.

3. Data and Sample Construction

3.1 Published Papers

The published-paper sample consists of all articles appearing in the *Journal of Finance*, the *Review of Financial Studies*, and the *Journal of Financial Economics* from 2000 to 2024. These three journals constitute the traditional top tier of academic finance and account for the majority of highly cited finance research on government intervention. We collect article metadata—title, authors, year, volume, issue, DOI, and abstract—for all articles in each journal over this period.

Journal articles are identified and retrieved as follows. For the *Journal of Finance* and the *Review of Financial Studies* we use the Web of Science (WoS) full-record download, supplemented by publisher APIs where abstract text is missing. For the *Journal of Financial Economics*, abstracts were obtained via the Elsevier ScienceDirect Article Retrieval API. Abstract text is available for 100% of research articles in all three journals after manual

abstract enrichment. The *Journal of Finance* denominator is restricted to research articles only; the raw WoS record includes editorial matter, meeting announcements, and AFA administrative items that carry no abstracts, and our pipeline excludes these before computing coverage.

3.2 NBER Working Papers

The working-paper sample is drawn from the NBER Working Paper Series via the public API. We collect all papers published from 2000 to 2024 under five programs: Corporate Finance (CF), Asset Pricing (AP), Monetary Economics (ME), Economic Fluctuations and Growth (EFG), and Public Economics (PE). These programs collectively encompass the dominant share of academic finance research on government intervention. Papers appearing in multiple programs are counted once.

We use the NBER working paper sample as a reference distribution for comparison with the published sample. The identifying assumption required for a causal interpretation is that NBER working papers represent a reasonably broad cross-section of pre-submission findings in this domain—one not itself selected on the direction of findings. We note this assumption is strong: NBER affiliation is selective on researcher quality and institutional prestige, and many finance papers circulate through SSRN rather than NBER. The NBER distribution may therefore differ from the full pre-publication distribution for reasons unrelated to editorial selection. We return to this identification challenge formally in Section 4.1.

3.3 In-Scope Classification Pipeline

Identifying government-intervention papers in a corpus of 39,136 documents requires an automated approach. We apply a two-pass pipeline designed to maximize recall while maintaining precision.

We begin with all articles and working papers collected from our four sources—the *Journal of Finance*, the *Review of Financial Studies*, the *Journal of Financial Economics*, and the NBER Working Paper Series—for the period 2000–2024, yielding a combined corpus

of 39,136 documents (30,212 NBER working papers plus 8,924 journal articles: 3,383 from the *Journal of Finance*, 2,497 from the *Review of Financial Studies*, and 3,044 from the *Journal of Financial Economics*).

In the first pass, a paper passes the keyword filter if the combined text of its title and abstract contains at least one term from a pre-specified taxonomy of government intervention keywords spanning seven categories: financial regulation (e.g., “Dodd-Frank,” “capital requirements,” “Basel”), securities law (e.g., “SEC,” “Sarbanes-Oxley,” “insider trading law”), banking policy (e.g., “FDIC,” “deposit insurance,” “TARP”), monetary and fiscal policy (e.g., “quantitative easing,” “fiscal stimulus”), market microstructure rules (e.g., “uptick rule,” “tick size pilot”), corporate governance law (e.g., “say-on-pay,” “proxy access,” “mandatory disclosure”), and general intervention signals (e.g., “government intervention,” “public policy,” “law and finance”). Ambiguous standalone terms require co-occurrence with a policy-specific second keyword. The full taxonomy is reported in Appendix A.

Searching the combined title-and-abstract text rather than the abstract alone improves recall for three reasons. First, authors frequently name the specific law or regulation in the title but refer to it only as “the reform” or “the rule” throughout the abstract, so the keyword would appear only in the title. Second, common abbreviations (e.g., “Dodd-Frank” in the title) and full statutory names (e.g., “Wall Street Reform and Consumer Protection Act” in the abstract) are not always used interchangeably; searching both fields catches either form. Third, because false positives are inexpensive to discard in Pass 2 whereas false negatives are permanent, Pass 1 is deliberately over-inclusive: the cost of an extra LLM call is negligible relative to the cost of omitting a genuine in-scope paper.

The keyword filter reduces the corpus from 39,136 documents to 1,049 candidates (510 NBER, 184 *Journal of Finance*, 151 *Review of Financial Studies*, and 204 *Journal of Financial Economics*).

In the second pass, each keyword-flagged candidate is passed to Claude with the following instruction (full prompt in Appendix B):

“A finance research paper is ‘in scope’ if and only if its PRIMARY research question is the CAUSAL EFFECT of a specific government law, regulation, or public policy on financial markets, firms, or investors. The following are excluded: (1) purely theoretical papers with no empirical test of a policy; (2) papers that only describe a regulation without estimating its effect; (3) papers whose main topic is private firm or market behavior, with policy as incidental context; (4) papers studying central bank communication without a specific policy instrument; and (5) papers whose primary policy variable is a rule imposed by a stock exchange rather than a government or regulatory mandate.”

For each paper, the model returns an in-scope decision, a confidence level (high, medium, or low), and a one-sentence justification. Following this automated pass, we apply manual validation to enforce the causal-effect criterion more strictly. The most important judgment call concerns papers that use a policy event only as a tool to study something else—for example, using the passage of Sarbanes-Oxley to identify the effect of disclosure on firm investment rather than to evaluate the law itself. Such papers’ primary question is not about the policy, so they are excluded.

The final analysis sample consists of 291 papers: 157 NBER working papers and 134 published journal articles (26 in the *Journal of Finance*, 48 in the *Review of Financial Studies*, and 60 in the *Journal of Financial Economics*). Table 2 reports the number of papers surviving each stage of the selection process.

Table 1. Sample Selection

Step	Papers	Excluded
Starting corpus (NBER + JF/RFS/JFE, 2000–2024)	39,136	
Pass 1: Keyword filter	1,852	37,284
Pass 2: LLM classification and manual review	291	1,561
<i>Final sample</i>	291	
Published articles (JF / RFS / JFE)	134	
Journal of Finance	26	
Review of Financial Studies	48	
Journal of Financial Economics	60	
NBER working papers	157	

Notes. The starting corpus includes all NBER working papers under programs CF, AP, ME, EFG, and PE and all research articles from the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics*, 2000–2024. The keyword filter screens for government intervention terminology (see Appendix A); LLM classification and manual review retain only papers whose primary research question is the causal effect of a specific government policy.

3.4 Direction Coding

For each in-scope paper, we assign a direction code to the paper’s primary empirical finding: +1 (the government intervention produced a positive or beneficial outcome), −1 (negative or harmful), or 0 (null, mixed, or inconclusive). Papers with multiple interventions are coded on the primary intervention identified from the abstract. Direction coding uses the same model as the in-scope classifier.

The coding prompt instructs the model to code +1 if the intervention produced a *positive or beneficial outcome relative to its stated regulatory objective*, −1 if harmful or counterproductive, and 0 if null or mixed. Normative framing by the authors is explicitly excluded. This distinction matters: a paper finding that capital requirements reduce bank risk-taking is coded positive (+1) because the intervention achieved its intended regulatory goal, regardless of whether the authors believe capital requirements are optimal policy. The full coding prompt is in Appendix B.

We validate the direction codes in two ways. First, we re-ran the coding on a stratified

60-paper sample using a different, smaller model with the same prompt. The two runs agree on 72% of papers ($\kappa = 0.674$, “substantial agreement”). Agreement is high for directional papers—100% for negative findings and 90% for positive ones—but only 25% for null/mixed papers, where most errors are null-to-directional flips rather than the reverse. This means any coding error tends to overstate the directional premium, so our estimates are if anything conservative. The full confusion matrix is in Appendix G.

Second, we re-ran coding on 58 papers under three alternative prompt wordings. Agreement with the original codes ranges from 83% to 91% ($\kappa_w = 0.742\text{--}0.904$), all above the model-choice κ , confirming the results are not sensitive to how the question is phrased (Table ??).

Table 2. Prompt-Perturbation Robustness: Agreement with Original Direction Codes

Prompt variant	N	Agreed	Agreement %	κ_w
Baseline (original wording)	58	53	91.4	0.904
Variant A (“primary conclusion”)	58	48	82.8	0.742
Variant B (“classify direction”)	58	49	84.5	0.814

Notes. Each variant re-codes the same 58-paper stratified sample using a smaller model with the same classification rules. κ_w is computed against the original direction codes used throughout the paper.

Human-coded validation of a 40-paper subsample is included in the replication package and remains a key additional check; the model-choice $\kappa = 0.674$ is the primary reliability evidence reported here.

3.5 Linking NBER Papers to Published Papers

We matched each of the 134 published journal articles to the NBER working paper archive using title similarity, author overlap, and external databases including OpenAlex, CrossRef, and Semantic Scholar. The pipeline identified 16 matched pairs—about 12% of the published sample. The remaining 118 published papers have no NBER counterpart, which reflects the fact that most corporate finance papers circulate through SSRN rather than NBER.

In the regression analysis, published papers are coded as Published = 1 and unmatched NBER working papers as Published = 0. The direction distributions of the two groups are nearly identical—about 80% of papers in each group have a directional finding—suggesting the two samples are broadly comparable in composition before any regression adjustment.

3.6 Summary Statistics

Table 3 presents summary statistics for the 291-paper analysis sample. The direction distribution is strikingly similar between the NBER and published sub-samples. Among NBER working papers, 20.4% are coded null or mixed, compared with 20.9% among published papers—a difference of less than 1 percentage point. Positive-finding papers constitute 35.1% and negative-finding papers 44.0% of the published sample, versus 35.7% and 43.9% of the NBER sample. The two samples look nearly the same in their direction distributions.

Table 3. Summary Statistics

Sample	<i>N</i>	Positive (%)	Negative (%)	Null (%)
All papers	291	35.4	44.0	20.6
NBER working papers	157	35.7	43.9	20.4
Published (JF/RFS/JFE)	134	35.1	44.0	20.9
<i>By journal:</i>				
<i>Journal of Finance</i>	26	42.3	46.2	11.5
<i>Review of Financial Studies</i>	48	29.2	54.2	16.7
<i>Journal of Financial Economics</i>	60	36.7	35.0	28.3

Notes. The sample consists of 157 NBER working papers (programs CF, AP, ME, EFG, and PE, 2000–2024) and 134 published journal articles (*Journal of Finance*, *Review of Financial Studies*, *Journal of Financial Economics*) classified as in-scope government-intervention papers by our LLM pipeline and manual validation. Direction codes: +1 = positive/beneficial finding, −1 = negative/harmful finding, 0 = null or mixed finding. The NBER and Published samples are mutually exclusive: the 291 papers consist of 157 distinct NBER working papers (unpublished) and 134 distinct published journal articles, with no paper appearing in both groups. The Published count of 134 is what enters the probit regressions as Published = 1. Column percentages may not sum to 100 due to rounding.

4. Empirical Strategy

4.1 Research Design and Identification

Our empirical strategy compares papers circulated as NBER working papers against papers published in the *Journal of Finance*, the *Review of Financial Studies*, or the *Journal of Financial Economics*, asking whether the direction distribution of findings differs between the two groups. This is not a within-paper design: only 16 of 134 published papers are matched to an NBER working paper, so the two samples are largely different papers selected through different processes.

The key assumption is that NBER working papers approximate the pre-publication distribution of findings in this domain. Similar approaches are taken by ?, who use NSF pre-registrations as a pre-publication baseline, and ?, who use institutional RCT records. Our design is weaker because NBER affiliation is selective on researcher quality, and most finance papers circulate through SSRN rather than NBER. The direction of any resulting bias is unclear—NBER selection could either overstate or understate the true gap—and we interpret all results as cross-sectional associations, not causal editorial-filter effects. The full identification discussion is in Section ??.

4.2 Primary Specification

Our primary test estimates a probit model at the paper level:

$$\Pr(\text{Published}_i = 1) = \Phi(\alpha + \beta_1 \text{Positive}_i + \beta_2 \text{Negative}_i + \boldsymbol{\gamma}' \mathbf{X}_i), \quad (1)$$

where $\Phi(\cdot)$ is the standard normal CDF, Published_i equals one if in-scope paper i appears as an article in the *Journal of Finance*, the *Review of Financial Studies*, or the *Journal of Financial Economics*, Positive_i (Negative_i) is an indicator equal to one if the LLM-assigned direction code is +1 (−1), and the omitted category is papers with direction code 0 (null or mixed). The vector $\mathbf{X}_i$ includes a linear year trend and, in extended specifications, decade

fixed effects.

Our primary coefficients of interest are β_1 and β_2 . A finding of $\beta_1 > 0$ ($\beta_2 > 0$) indicates that positive-finding (negative-finding) papers are more likely to be in the published sample than null-finding papers, conditional on controls. Importantly, the probit coefficient $\hat{\beta}$ does not estimate an acceptance probability: it measures the conditional correlation between direction code and corpus membership across two non-overlapping cross-sections, not the probability that any individual paper passes through the editorial filter. A test of $H_0: \beta_1 = \beta_2$ addresses the directional-bias hypothesis: rejection would indicate that the publication premium depends on the sign of the finding, not merely on statistical significance.

Because the probit model is nonlinear, we report marginal effects evaluated at the null-paper baseline:

$$\Delta_d = \Phi(\hat{\alpha} + \hat{\beta}_d) - \Phi(\hat{\alpha}), \quad d \in \{+, -\}.$$

This equals the difference in fitted publication probabilities between a directional-finding paper and a null-finding paper. We also report raw corpus-membership rates directly from the data as a transparent, model-free summary of the unconditional patterns.

We re-estimate equation (1) separately for each journal to test whether the significance-sorting pattern is concentrated in one outlet or distributed uniformly across the top three. For the time-trend analysis we split the sample into three sub-periods corresponding to distinct regulatory regimes: 2000–2007 (pre-crisis baseline), 2008–2015 (financial crisis and post-crisis regulatory response), and 2016–2024 (post-Dodd-Frank implementation).

5. Results

5.1 Within-Paper Test Statistic Distributions

We test whether published papers show unusual clustering of test statistics just below the conventional 5% significance threshold, using the caliper test of ?. We run this test separately on two types of reported statistics. The first type covers regression rows where both a

coefficient and a standard error are reported, allowing independent verification of the test statistic. The second type covers standalone test statistics reported without a coefficient or standard error—numbers that appear mainly in robustness and supplementary tables. We extract both types from the PDFs of all published and NBER papers in the sample.

We extract regression-table test statistics from the PDFs of all 134 published papers and 157 NBER working papers using automated text parsing. The published sample yields 7,761 test statistics from 121 papers; the 13 papers with no extractable statistics are pure theory models or have image-based tables. The NBER sample yields 8,351 test statistics from 125 papers; the remaining 32 contribute no statistics for the same reasons. Having a pre-publication sample of test statistics allows us to compare the two distributions directly, rather than inferring bias solely from the shape of the published distribution.

To check extraction accuracy, we manually verified a random sample of 50 statistics—25 from published papers and 25 from NBER papers. The average test statistic is similar across the two corpora (3.41 for published, 3.32 for NBER), consistent with comparable significance rates before and after publication. The replication package includes all extracted statistics.

Extraction quality by corpus. A potential concern is that differential extraction quality between typeset publisher PDFs (published) and NBER draft PDFs could itself generate the published-vs.-NBER caliper gap. We assess this in three ways. First, internal consistency: for statistics where we extract both a coefficient and a standard error, we verify that the reported $|z|$ -value matches $|\hat{\beta}/\hat{\sigma}_{\hat{\beta}}|$ to within 0.1. This internal check passes for **98.3%** of published rows (5,474/5,569) and **99.7%** of NBER rows (6,811/6,831), giving no evidence of materially worse extraction quality in either corpus. The higher share of z-stat-only rows in the published corpus (29.0% vs. 18.7% for NBER) reflects the typeset formatting of some publisher tables that display t-statistics directly rather than coefficient-SE pairs; these statistics are taken at face value and are not subject to a consistency check, but they pass all caliper-window bounds checks. Second, the 13 published papers and 32 NBER papers that yield zero extracted statistics are distributed uniformly across journals ($\chi^2 = 1.08$, $p = 0.58$

for published) and by vintage (Mann-Whitney $p = 0.749$; median year 2017 for both kept and dropped papers), providing no evidence of systematic non-random exclusion. Third, restricting the caliper analysis to high-coverage papers (≥ 20 extracted statistics per paper) leaves the published ratio essentially unchanged (1.207 vs. 1.203 for the full sample) and the NBER ratio unaffected (1.012 vs. 1.012). Fourth, we stratify the caliper by row type (Appendix ??; summarized in Table ??). *Primary specification (coef+SE rows)*: among the 5,510 published rows with both a coefficient and a standard error, the caliper ratio is **0.81** (below = 320, above = 394; naive $p = 0.006$); the corresponding NBER ratio is **0.95** (below = 416, above = 438; $p = 0.45$). Neither corpus shows excess bunching below the 5% threshold in verifiable rows, and the published ratio (0.81) is below the NBER ratio (0.95), not above it. This is the paper’s primary caliper result: no threshold bunching in systematically reported primary regressions. *Secondary finding (standalone z-statistics)*: among the 2,251 published z-stat-only rows, the caliper ratio is 2.55 ($p < 0.001$); the comparable NBER ratio is 1.49 ($p = 0.019$). Both corpora show bunching in standalone statistics, with published papers showing a substantially higher ratio.

To assess how much of the full-sample gap (published 1.20 – NBER 1.01 = +0.191) is driven by composition versus within-type effects, we decompose:

- *Within-type effect* (published bunches more per se in standalone rows, evaluated at NBER’s z-stat-only share of 21.0%): +0.097
- *Composition effect* (published has more z-stat-only rows: 31.6% vs. 21.0%): +0.094

The two effects are nearly equal in magnitude. Even if the composition effect is entirely explained by typesetting differences—published journals more often display t-statistics inline—roughly half the full-sample gap (+0.097) remains as a within-type effect: standalone z-statistics bunch more in published papers than in NBER papers of the same type.

Star-consistency validation of z-stat-only rows. To address whether the bunching in z-stat-only rows reflects extraction artifacts rather than real statistics, we cross-validate each z-stat-only row’s extracted $|z|$ value against its reported significance stars. Under the

standard convention ($*** = 1\%$, $** = 5\%$, $* = 10\%$), a row with extracted $|z| \in (1.645, 1.96)$ should carry a “*” star; a row with $|z| \in (1.96, 2.576)$ should carry “**”; and a row with $|z| > 2.576$ should carry “***”. Among published z-stat-only rows with standard star notation, **86.1%** (1,824/2,118) have a star-consistent extracted z-value — nearly as high as the 91.8% consistency rate for coef+SE rows. Focusing on the lower caliper window ($|z| \in (1.645, 1.96)$), **50.9%** of published z-stat-only rows carry the expected “*” star (142 of 279 rows), with the remaining 48.4% carrying no star (consistent with papers that only report 5% or 1% thresholds, not 10%). Only 0.7% (2 rows) are clearly star-inconsistent. These rates argue strongly against widespread extraction errors in the published z-stat-only rows.

Mechanism: selective inclusion of supplementary statistics. A key question is why z-stat-only rows exhibit more threshold bunching than coef+SE rows, if both are extracted from the same papers. We distinguish two conceptually separate phenomena that should not be conflated, as they carry different welfare implications.

Selective reporting refers to authors choosing which test statistics to present in the first place. Standard regression tables (the source of coef+SE rows) report *all* coefficient estimates for a given specification systematically, including insignificant ones. Standalone z-statistics, by contrast, appear disproportionately in supplementary analyses, robustness summary rows, and “Panel B”-type tables where authors select which results to show. Consistent with this, 96.1% of published z-stat-only rows come from non-primary tables. If authors prefer to highlight marginally significant test statistics in these supplementary slots, the caliper pattern in z-stat-only rows would emerge purely from a *reporting decision*—not from any manipulation of the underlying estimates.

Specification search is a distinct phenomenon: running multiple specifications and selecting those whose primary estimates cross a significance threshold. Specification search distorts primary regression estimates in a way that selective inclusion of supplementary statistics does not. The absence of bunching in coef+SE rows—the systematically reported

primary regressions—is evidence *against* specification search in the primary estimates.

We therefore interpret the z-stat-only bunching as evidence of *selective reporting of supplementary statistics*, not as evidence of specification search in primary results. This distinction matters for welfare: selective highlighting of significant supplementary results is less distortionary than systematic manipulation of primary regression outputs. We cannot, however, rule out that some z-stat-only rows represent summary statistics from searches over many specifications, which would carry specification-search welfare implications. Both corpora show a z-stat-only ratio above one (published: 2.55; NBER: 1.49), confirming that selective reporting of supplementary statistics occurs in working papers too, though more intensely in published papers.

Tests. The *caliper test* (?) compares the mass of $|z|$ -statistics in the interval $(1.645, 1.96)$ —just below the conventional 5% threshold—to the mass in $(1.96, 2.275)$ —just above it. Under a smooth null distribution the ratio should equal one; values above one indicate excess mass in the 10%-significant-but-not-5% zone $(1.645, 1.96)$ relative to the just-5%-significant zone $(1.96, 2.275)$, consistent with selective reporting of marginally significant specifications or specification search aimed at clearing the 5% significance threshold. The *density discontinuity test* estimates whether there is a discrete jump in the density of $|z|$ at the $z = 1.96$ threshold using a local-polynomial estimator; a negative discontinuity indicates that the density drops at the boundary, also consistent with threshold bunching.

Results. Table ?? reports caliper ratios and discontinuity estimates for the full sample, by journal, and by finding direction. Figure ?? displays the $|z|$ -distribution for the full sample; Figures ?? and ?? show journal- and direction-level panels.

Four patterns stand out. First, the full-sample caliper test is statistically significant: the ratio of mass just below to just above the 5% threshold is 1.20 ($\chi^2 = 9.48$, $p = 0.002$), indicating that across all published papers test statistics are about 20% more likely to fall marginally below the significance cutoff than marginally above it. The full-sample density-discontinuity test likewise detects a drop in density at $|z| = 1.96$ (Disc. $p = 0.040$).

Table 4. Within-Paper Test Statistic Distribution Tests

Subset	N	% Sig	Caliper ratio	Caliper χ^2	Caliper p	Disc. p
Full sample	7,761	57.1%	1.20	9.48	0.002	0.040
<i>By journal:</i>						
JF	1,308	65.5%	0.88	0.64	0.423	< 0.001
RFS	2,581	59.2%	0.97	0.10	0.752	0.180
JFE	3,872	52.9%	1.49	23.56	< 0.001	0.038
<i>By finding direction:</i>						
Positive	2,633	60.7%	0.94	0.33	0.565	0.456
Negative	3,491	54.9%	1.52	22.54	< 0.001	0.183
Null/mixed	1,637	56.1%	1.05	0.11	0.740	< 0.001

Notes: N is the number of $|z|$ -statistics after cleaning (removing values < 0.1 and > 50 , and non-numeric entries). *Caliper ratio* is the ratio of mass in $(1.645, 1.96)$ to mass in $(1.96, 2.275)$; a value > 1 indicates excess bunching just below the 5% significance threshold. *Caliper χ^2* is the chi-squared statistic for the equality of the two mass counts; *Caliper p* is the associated p -value. *Disc. p* is the p -value from the ? density-discontinuity test at $|z| = 1.96$, estimated by local polynomial. Bold entries are significant at the 5% level.

Second, the bunching is heavily concentrated in *Journal of Financial Economics*: the caliper ratio reaches 1.49 ($\chi^2 = 23.56$, $p < 0.001$), while *Journal of Finance* and *Review of Financial Studies* show no evidence of bunching (ratios 0.88 and 0.97, both $p > 0.40$). Third, breaking by finding direction reveals that caliper bunching is far more severe in *negative-finding papers*—papers concluding that an intervention is harmful or ineffective—where the ratio reaches 1.52 ($\chi^2 = 22.54$, $p < 0.001$), compared to 0.94 in positive-finding papers and 1.05 in null-finding papers. Fourth, the density-discontinuity test is significant for the *Journal of Finance* and null-finding subsamples, consistent with a sharp drop in density exactly at $|z| = 1.96$. We note a tension in the *Journal of Finance* sub-sample: the caliper ratio of 0.88 ($p = 0.423$) indicates no excess mass just below 1.96, yet the discontinuity test yields $p < 0.001$. These two statistics can coexist if density is relatively *smooth* in both the $(1.645, 1.96)$ and $(1.96, 2.275)$ windows but drops sharply *at* the 1.96 boundary itself—consistent with authors in *Journal of Finance* selectively omitting 10%-significant-but-not-5% results rather than accumulating them. However, in a subsample of only 26 papers and 1,308 statistics, this discord more plausibly reflects two noisy tests giving conflicting signals than a coherent behavioral pattern. Local-polynomial density estimators are known to have

poor finite-sample properties when the number of observations is small (here $n = 1,308$ test statistics, with bandwidth chosen by cross-validation on a thin distribution), which makes the discontinuity p -value unreliable in this subsample. We flag this as reason to downweight the density-discontinuity test results throughout: whenever the two procedures disagree in a subsample, the caliper test (which directly measures the mass ratio the theory predicts) should be given priority over the local-polynomial discontinuity estimate. To verify that no single paper dominates the JF result, we report the distribution of statistics across the 23 JF papers with extracted statistics: mean 62.2, median 72, maximum 107 (2 papers exceed 100 statistics). We re-estimate the JF caliper after excluding each of the two top papers individually; the ratio ranges from 0.85 to 0.91, confirming the pattern is not driven by a single outlier.

The direction-level result is particularly novel. The cross-sectional analysis (Sections 5–6) finds no evidence that journals differentially select positive versus negative papers. Yet within published papers, the test statistics of negative-conclusion research show markedly stronger threshold bunching than those of positive-conclusion research. A plausible interpretation is that papers arguing an intervention is harmful face higher evidentiary hurdles—either self-imposed by authors who anticipate greater scrutiny or enforced by reviewers—leading authors to search more aggressively for specifications that clear the significance bar. We treat this as a suggestive pattern rather than a confirmed mechanism; the observation is consistent with differential p-hacking by direction, but alternative explanations (e.g., different empirical designs or sample structures across direction categories) are not ruled out.

Pre-publication benchmark. A key limitation of existing within-paper approaches (including ? and ?) is that they observe only the published distribution and must infer what the pre-publication distribution looked like. With NBER full-text data, we can compare directly.

Table ?? presents the caliper test results for both the published and NBER samples. The contrast is stark. In the *full sample*, the published caliper ratio of 1.20 ($p = 0.002$) is entirely

absent in the NBER sample (ratio = 1.01, $p = 0.849$). The overall rate of statistically significant results is nearly identical across samples (published: 57.1%; NBER: 55.5%), a fact *consistent with* similar statistical power across corpora. We note, however, that equal significance rates alone do not establish equal power: two populations can generate similar fractions of significant statistics while differing in effect sizes, sample sizes, or model specifications. What the equal-rates result does establish, in conjunction with the caliper gap, is that the bunching pattern cannot be explained purely by NBER papers being uniformly under-powered: under-powered tests would produce fewer significant results *and* no bunching at the threshold, rather than the similar significance rate with absent bunching that we observe. The bunching at the 5% threshold is therefore specific to the published record, not a level-of-power artefact. We note that this result is specific to caliper windows whose lower bound is at or above $|z| = 1.645$ (the Brodeur convention); as shown in Appendix ??, a wider window extending down to $|z| = 1.50$ reverses the sign for published papers because the density is still rising below 1.645, while the standard, narrow, and Brodeur2 windows all yield qualitatively identical conclusions.

The directional pattern is even sharper. Among *negative-finding* papers, the published caliper ratio is 1.52 ($p < 0.001$) while the NBER ratio is *essentially exactly one* (ratio = 1.00, $p = 0.964$)—a difference of 0.52 that is economically substantial, though we note that the cluster-bootstrap confidence interval for the published negative-finding ratio spans [0.81, 2.67], indicating imprecision in the magnitude estimate. Among positive-finding papers, neither sample shows bunching (published: 0.94; NBER: 0.99; both $p > 0.56$). For null-finding papers the pattern reverses: the NBER caliper ratio (1.14, $p = 0.442$) exceeds the published ratio (1.05, $p = 0.740$), though both are statistically indistinguishable from one. Figure ?? displays the side-by-side $|z|$ -distributions for the full sample and for the negative-finding subsample.

The directional pattern in caliper ratios is specific to the published corpus. Two interpretations are consistent with the data, and we regard them as equally plausible.

Compositional interpretation (preferred given evidence): Published negative-finding papers may differ from their NBER counterparts in unobserved dimensions—research design complexity, table structure, vintage, or the number of supplementary robustness checks—that independently predict within-paper specification counts and standalone-statistic density. Our cross-sectional design cannot rule this out, and given that half the full-sample gap is already explained by a composition effect (different standalone-statistic shares across corpora), a compositional explanation for the within-type effect remains the most conservative reading of the evidence.

Behavioral interpretation (consistent with but not established by the data): Anti-intervention papers may face elevated evidentiary demands—either self-imposed by authors anticipating scrutiny or enforced by reviewers—leading to specification search until marginal results clear the 5% threshold. This would predict bunching concentrated in the published negative-finding corpus, which is what we observe. We present this as a hypothesis the data are consistent with, not as a finding the data establish; we do not claim to have identified a causal mechanism.

Inference under paper-level clustering. A concern with the caliper test is that test statistics from the same paper are not independent: a paper reporting $\bar{n} \approx 64$ statistics contributes 64 pseudo-independent observations to the chi-squared calculation, which may overstate precision. Table ?? addresses this directly by reporting three p -values for each cell. The *naive* p -value treats all statistics as independent (standard Brodeur approach). The *permutation* p -value resamples within papers: for each of 10,000 iterations we independently permute the below/above assignment of each in-window statistic, preserving each paper’s in-window count but breaking the directional signal; the reported p -value is the fraction of iterations yielding a caliper ratio at least as large as observed. The *bootstrap* p -value resamples papers with replacement (10,000 draws), making the paper the unit of inference.

For the full published sample the three p -values are 0.002 (naive), 0.001 (permutation), and 0.150 (bootstrap), with a 95% cluster bootstrap confidence interval of [0.88, 1.66]. For

negative-finding papers they are < 0.001 , < 0.001 , and 0.135 (95% CI: $[0.81, 2.67]$). The permutation test is strongly significant in both cases, reflecting a consistent directional signal across papers: the caliper ratio is driven by a majority of papers contributing excess below-threshold mass, not by a single outlier. The bootstrap is less significant because the ratio varies substantially across the 50 negative-finding published papers (bootstrap CI: $[0.81, 2.67]$), a heterogeneity that the permutation test’s within-paper conditioning absorbs. We treat both statistics as informative: the permutation test confirms the signal is real and not concentrated in a handful of papers; the bootstrap CI communicates the uncertainty in its aggregate magnitude. The cluster bootstrap also serves as the paper’s per-paper-weighted inference procedure: by resampling papers with replacement rather than statistics, it assigns each paper equal weight regardless of how many test statistics it contributes, so that a handful of large-table papers cannot dominate the reported p -value or confidence interval. For the NBER baseline the permutation p -value is 0.424 and the bootstrap p -value is 0.432 —indistinguishable from chance under both procedures.

A cross-group permutation test that permutes paper-level published/NBER labels (10,000 iterations, preserving group sizes) yields $p = 0.210$ for the full sample and $p = 0.129$ for negative-finding papers. These tests have lower power because the number of papers in each group (50–108) limits what can be resolved by a label-shuffling procedure, but they confirm the directional ordering: in no subsample does the NBER comparison produce a more extreme caliper gap in favor of NBER than the observed published excess. We designate the paper-level cluster bootstrap as our *primary* inference procedure, because it treats each paper as the unit of analysis regardless of how many test statistics it contributes, and therefore avoids inflating precision from high-table-count papers. Under this primary procedure the full-sample published result is $p = 0.150$ (CI $[0.88, 1.66]$) and the negative-finding result is $p = 0.135$ (CI $[0.81, 2.67]$)—neither reaches conventional significance. The permutation test, which is more powerful because it conditions on within-paper composition, gives $p < 0.001$ for both; the cross-group test gives $p = 0.129$. Under the primary (bootstrap) procedure

the directional result is not statistically significant; the magnitude is imprecisely estimated and should be treated as *suggestive* rather than definitive. The NBER baseline shows no bunching under any procedure.

Table 5. Within-Paper Caliper Test: Published vs. NBER Working Papers

Sample	Subset	N stats	Papers	Ratio	p -value		
					Naive	Perm.	Boot.
Published	Full sample	7,761	108	1.20	0.002	0.001	0.150
NBER	Full sample	8,351	106	1.01	0.849	0.424	0.432
Published	Positive-finding	2,633	37	0.94	0.565	0.732	0.623
NBER	Positive-finding	2,658	39	0.99	0.918	0.554	0.492
Published	Negative-finding	3,491	50	1.52	<0.001	<0.001	0.135
NBER	Negative-finding	4,375	50	1.00	0.964	0.532	0.474
Published	Null-finding	1,637	21	1.05	0.740	0.388	0.383
NBER	Null-finding	1,318	17	1.14	0.442	0.239	0.260

Notes: *Published* sample: 7,761 test statistics from 121 papers in the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics*; of these, **108** papers contribute at least one statistic to the (1.645, 2.275) caliper window and serve as the clusters in the permutation and bootstrap procedures. *NBER* sample: 8,351 test statistics from 125 NBER working papers (out of 157 in-scope); **106** contribute in-window statistics and serve as clusters. Caliper ratio is mass in (1.645, 1.96) relative to (1.96, 2.275). Three p -values are reported. *Naive*: chi-squared test treating all statistics as independent. *Perm.*: within-group permutation test (10,000 iterations) that permutes the below/above assignment of each in-window statistic while fixing each paper’s in-window count; this preserves within-paper composition. *Boot.*: cluster bootstrap (10,000 draws) resampling papers with replacement; the reported p -value is the fraction of draws with a caliper ratio at least as large as observed. Bold entries: Naive and Perm. significant at 5%; Boot. p -values are directionally consistent but do not reach 5% significance for the full sample.

Multiple comparisons. We run a total of 18 caliper and 12 density-discontinuity comparisons across the full sample, three journal subgroups, three direction subgroups, two corpora, and the coef+SE primary specification. Of these, two were pre-specified as primary before the data were analyzed: (i) the full-sample published ratio and (ii) the full-sample NBER ratio. All journal-level and direction-level breakdowns are exploratory post-hoc partitions. Under Bonferroni correction for the 2 primary tests the threshold is $p < 0.025$; both the full-sample published caliper ($p = 0.002$) and the NBER baseline ($p = 0.849$) survive. Applying Bonferroni to all 18 comparisons gives a threshold of $p < 0.003$; the full-sample result ($p = 0.002$) and the JFE and negative-finding subsets ($p < 0.001$) survive,




Figure 1. $|z|$ -distributions: published papers (left panel of each pair) vs. NBER working papers (right panel). *Left pair:* Full sample (published: 7,761 statistics, caliper ratio **1.20**, $p = 0.002$; NBER: 8,351 statistics, caliper ratio 1.01, $p = 0.849$). *Right pair:* Negative-finding papers (published: 3,491 statistics, caliper ratio **1.52**, $p < 0.001$; NBER: 4,375 statistics, caliper ratio 1.00, $p = 0.964$). Orange bars highlight the caliper window (1.645, 2.275). The solid vertical line marks $|z| = 1.96$. Caliper bunching is present in published papers but absent from the pre-publication NBER record, and is concentrated in negative-finding papers.

while some secondary comparisons do not.

Why we pool all in-table test statistics. The caliper test is applied to all extracted regression-table statistics, averaging ≈ 64 per paper across published papers. We follow this approach—rather than restricting to primary-outcome coefficients only—for a conceptual reason: specification-search incentives operate on *any* result that a researcher might choose to report or suppress, not only on the headline coefficient. A researcher revising toward the threshold can adjust control-variable interactions, alternative-specification rows, and robustness checks; restricting to primary coefficients alone would miss this dimension of the distribution. We treat the paper as the unit of analysis (via the cluster bootstrap) to guard against a handful of large-table papers dominating the test. The robustness check below confirms that excluding probable control-variable rows leaves the caliper ratios essentially

unchanged.

Robustness to coefficient inclusion. A concern with pooling all table statistics is that control-variable coefficients (size, leverage, GDP, ROA, etc.) included in regression tables are unrelated to the paper’s headline finding, and their inclusion could confound the direction-based caliper comparison. Appendix ?? reports caliper statistics after excluding rows whose variable name matches a broad list of common control-variable patterns. This *non-control* subsample retains **7,055** published and **7,599** NBER statistics ($\approx 89\%$ of each corpus). The caliper ratio for published full-sample is unchanged (1.203 vs. 1.203), and the negative-finding ratio *strengthens* slightly (1.625 vs. 1.522), confirming that the caliper result is not an artifact of control-variable pooling. The NBER non-control ratios remain near one (0.981 full sample; 0.955 negative-finding). These results are consistent with the interpretation that threshold bunching reflects specification search around the primary estimate, not mechanical inclusion of many control coefficients.

5.2 Supporting Context: Cross-Sectional Direction Sorting

As supporting context, we examine whether the caliper bunching pattern is accompanied by directional sorting at the corpus-membership level. If threshold bunching were the product of editorial selection favoring certain findings, we would also expect the direction of a paper’s conclusion to predict whether it appears in the published corpus. The following analysis shows it does not.

Directional vs. Null Findings

Table 4 reports the unconditional corpus-membership rates by direction code. (We use the term “corpus-membership rate” to emphasize that these are ratios determined by our sample construction—specifically the relative sizes of the published and NBER corpora—not estimates of actual journal acceptance probabilities. We also note a terminological ambiguity in the paper: “primary specification” is used in two senses. In the caliper analysis it refers to the coefficient-SE restricted subsample (verifiable rows). In the probit analysis, Column (1)

is the “primary specification” in the sense of the preferred regression model. These are distinct uses and should not be confused.) Positive-finding papers appear in the published corpus at 45.6% (47 of 103 papers), negative-finding papers at 46.1% (59 of 128), and null-finding papers at 46.7% (28 of 60). All three rates are statistically indistinguishable: a χ^2 test of homogeneity cannot reject equality ($\chi^2 = 0.017$, $p = 0.992$), and a two-sample z -test for positive versus negative rates gives $z = -0.07$ ($p = 0.944$). The approximately 46% rate is remarkably uniform across all three direction categories, providing no raw evidence of direction-based sorting between the two corpora.

Table 6. Corpus-Membership Rates by Finding Direction

Finding Direction	Sample composition		Total	Corpus-Membership Rate	95% CI
	NBER-only	Published			
Positive (+1)	56	47	103	45.6%	[36.3, 55.2]
Negative (−1)	69	59	128	46.1%	[37.7, 54.7]
Null/Mixed (0)	32	28	60	46.7%	[34.6, 59.1]
Total	157	134	291	46.0%	

Notes. Direction-specific composition of the two-sample cross-section, 2000–2024. *Finding Direction* is the LLM-assigned sign of the paper’s primary empirical result: +1 (positive/beneficial), −1 (negative/harmful), or 0 (null or mixed). *NBER-only* counts in-scope NBER working papers with no confirmed match to a JF/RFS/JFE article (Published = 0); *Published* counts in-scope journal articles (Published = 1); *Total* is the sum across both sub-samples. *Corpus-Membership Rate* is $N_{\text{Published}}/N_{\text{Total}}$ within each direction category; it is a cross-sectional corpus-construction ratio, not an estimated acceptance probability. *95% CI* is the Wilson score confidence interval. A χ^2 test of homogeneity across the three rates gives $\chi^2 = 0.017$ ($p = 0.992$). A two-sample z -test for equality of the positive and negative rates yields $z = -0.07$ ($p = 0.944$).

Probit Regression Results

Table 5 reports estimates across seven specifications. Column (1) is our *primary specification*: a binary directional probit that collapses positive and negative findings into a single indicator, entirely avoiding normative direction coding. We treat Column (1) as primary because the Null/Mixed omitted category—the reference group for Columns (2)–(7)—has only 25% inter-rater agreement in our replication exercise, making contrasts *against* null papers the

least reliable baseline. The binary specification avoids this by asking only whether *any* directional paper is more likely to appear in a top journal, pooling positive and negative findings. Column (2) is the bivariate three-way probit; Column (3) adds a linear year trend; Column (4) adds decade fixed effects; Column (5) adds $\log(1 + N_{\text{authors}})$ as a proxy for team size; Column (6) adds a quasi-experimental identification dummy (equal to one if the paper’s primary identification strategy is RDD, IV, DiD, or event study, as classified by our LLM pipeline applied to abstracts); Column (7) further adds a top-institution indicator for the first author’s affiliation.

Column (1) shows that the binary directional coefficient is -0.020 (s.e. 0.182 , $p = 0.914$)—statistically indistinguishable from zero. Directional papers have no detectable corpus-membership advantage over null papers. This non-rejection rules out effects ≥ 22 percentage points at 80% power; it should not be interpreted as evidence of true editorial neutrality.

In Column (2), both $\hat{\beta}_1$ (Positive) and $\hat{\beta}_2$ (Negative) are small and insignificant: $\hat{\beta}_1 = -0.026$ (s.e. 0.204 , $p = 0.898$) and $\hat{\beta}_2 = -0.014$ (s.e. 0.196 , $p = 0.942$). The most robust inference from this specification is the *joint* χ^2 test for whether both direction coefficients are simultaneously zero ($\chi^2(2) = 0.017$, $p = 0.992$), since the joint test does not require the null/mixed omitted category to be cleanly measured. Because the null/mixed baseline has only 25% model-agreement, individual direction contrasts against null papers are anchored to a noisy reference; the joint test avoids this sensitivity. A Wald test for equality of the two coefficients gives $\chi^2 = 0.005$ ($p = 0.944$). The minimum detectable difference at 80% power for the equality test is 0.466 coefficient units (SE of difference = 0.166). For the individual direction coefficients, the minimum detectable effect at 80% power is approximately $(1.96 + 0.842) \times 0.204 \approx 0.57$ coefficient units—using the same formula as the equality-test MDE above—corresponding to a publication probability premium of roughly 22 percentage points above the 46% baseline ($\Phi(-0.084 + 0.57) - \Phi(-0.084) \approx 0.22$)—a large, policy-relevant effect. Effects smaller than this cannot be ruled out by our design; the non-rejections should be interpreted as evidence against large directional sorting, not proof

of the complete absence of any direction premium. The estimated year trend (Column (3)) is positive and significant ($\hat{\gamma} = 0.033^{***}$, s.e. 0.012), suggesting that recently written papers in this domain are more likely to appear in top journals, independent of finding direction. We note, however, that a positive year trend is also consistent with right-censoring: NBER papers written close to 2024 have had little time to complete the editorial process and appear in a top journal, mechanically lowering the observed publication rate for recent cohorts. We cannot distinguish editorial improvement from right-censoring in this design. As a robustness check, Panel E of Table 7 restricts the NBER sample to papers circulated in 2020 or earlier, giving all papers at least four years to complete the publication pipeline. The direction coefficients are -0.021 ($p = 0.928$) and -0.144 ($p = 0.519$), confirming the null is not an artifact of right-censoring. Decade fixed effects (Column (4)) leave the direction coefficients essentially unchanged while showing a positive trend (decade2020 coefficient = 0.454^{**} , $p = 0.027$). Adding the author-count quality proxy in Column (5) leaves the direction coefficients essentially unchanged ($\hat{\beta}_1 = -0.005$, $\hat{\beta}_2 = -0.020$), both far from significance. The author-count coefficient ($\hat{\delta} = -0.186$, s.e. 0.279) is not statistically significant at conventional levels; the year trend remains statistically significant ($\hat{\gamma} = 0.034^{***}$, s.e. 0.012) after this control.

Column (6) adds the quasi-experimental identification dummy. The *QuasiExp* coefficient is 0.467^{***} (s.e. 0.157, $p = 0.003$). *Important caveats on this variable.* The quasi-experimental dummy has received no human-coded validation of any kind (unlike the direction codes, which have a two-run model-choice sensitivity check with $\kappa = 0.674$). The variable is coded by LLM from abstracts and is likely correlated with author quality, institutional prestige, topic fashionability, and other quality attributes that the probit cannot separate from identification strategy. It is equally consistent with editors rewarding elite authors or methodologically fashionable topics as with editors specifically rewarding causal identification. We therefore *strongly caution* against interpreting this coefficient as evidence that journals reward causal identification per se; it is preliminary evidence of a positive quality premium,

with identification strategy as one plausible—but not the only—contributing factor. Direction coefficients remain insignificant ($\hat{\beta}_1 = -0.135$, $\hat{\beta}_2 = -0.156$), ruling out the concern that direction effects are masked by confounds with identification.

Table ?? disaggregates corpus-membership rates by the joint distribution of direction and identification strategy. Within every direction category, QE papers appear in top journals at materially higher rates than non-QE papers (positive: 55.6% vs. 34.7%; negative: 52.2% vs. 39.0%; null: 75.0% vs. 36.4%). The direction-specific corpus-membership rates are nearly identical within each QE tier: in the non-QE group, rates are 34.7%, 39.0%, and 36.4% for positive, negative, and null papers, a spread of only 4.3 percentage points; in the QE group, 55.6%, 52.2%, and 75.0%, though the null $\times$ QE cell is small ($N = 16$). An additive probit controlling jointly for direction and QE status yields $\text{QuasiExp} = 0.454^{***}$ ($p = 0.004$) while both direction coefficients remain indistinguishable from zero (pos = -0.128 , $p = 0.548$; neg = -0.139 , $p = 0.499$). A specification with full direction $\times$ QE interactions yields non-significant interaction terms (pos $\times$ QE = -0.669 , $p = 0.158$; neg $\times$ QE = -0.766 , $p = 0.093$), confirming that the direction effect is null within both identification-strategy tiers and that the QuasiExp premium does not depend on finding direction.

Column (7) adds a top-institution indicator equal to one if the first author’s institutional affiliation matches a list of 50+ top finance and economics departments (matched via OpenAlex API). The top-institution coefficient is 0.196 (s.e. 0.175, $p = 0.263$), positive but statistically insignificant. Direction coefficients are unchanged ($\hat{\beta}_1 = -0.156$, $\hat{\beta}_2 = -0.152$; both $p > 0.46$) and the quasi-experimental premium is virtually identical to Column (6) (0.472^{***} , s.e. 0.158). The institutional-quality control does not affect any of the substantive results.

Table 7. Corpus-Membership Rates by Finding Direction and Identification Strategy

	Non-QE	QE	Total
Positive (+1)	34.7% (17/49)	55.6% (30/54)	45.6% (47/103)
Negative (−1)	39.0% (23/59)	52.2% (36/69)	46.1% (59/128)
Null/mixed (0)	36.4% (16/44)	75.0% (12/16) [†]	46.7% (28/60)
Total	36.8% (56/152)	56.1% (78/139)	46.0% (134/291)

Notes. Each cell shows the corpus-membership rate (fraction appearing in a top-3 finance journal) with counts in parentheses. Non-QE: papers using OLS, fixed effects, or descriptive methods as classified by the LLM pipeline from abstracts. QE: papers using RDD, IV, DiD, or event study. [†]: Null × QE cell has only $N = 16$ papers; the 75% rate reflects sampling variability in a small cell and should be interpreted with caution. The three direction columns are nearly equal within each QE tier, confirming that the quasi-experimental premium does not interact with finding direction at conventional significance levels.

Table 8. Direction–Publication Association: Probit Estimates

	(1) Binary (Primary)	(2) Bivariate	(3) + Year	(4) + Decade FE	(5) + Authors	(6) + ID Strategy	(7) + Top Inst.
Directional	−0.020 (0.182)						
Positive (+1)		−0.026 (0.204)	−0.002 (0.205)	−0.007 (0.205)	−0.005 (0.205)	−0.135 (0.212)	−0.156 (0.213)
Negative (−1)		−0.014 (0.196)	−0.011 (0.198)	−0.017 (0.198)	−0.020 (0.198)	−0.156 (0.206)	−0.152 (0.206)
Year trend			0.033*** (0.012)		0.034*** (0.012)	0.028** (0.012)	0.027** (0.012)
$\log(N_{\text{authors}})$					−0.186 (0.279)	−0.255 (0.283)	−0.224 (0.286)
Quasi-Exp. ID						0.467*** (0.157)	0.472*** (0.158)
Top Institution							0.196 (0.175)
Decade FEs	No	No	No	Yes	No	No	No
Intercept	−0.084	−0.084	−0.097	−0.371	0.145	0.114	0.027
N	291	291	291	291	291	291	291
Pseudo R^2	0.000	0.000	0.020	0.013	0.021	0.043	0.046

Notes. Probit estimates of the probability that an in-scope paper is published in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics*. Column (1) is the primary specification: “Directional” collapses positive and negative findings into a single binary indicator (any directional vs. null), avoiding normative coding of the direction sign. Columns (2)–(7) use the three-way direction code; Positive (+1) and Negative (−1) are indicators for the LLM-assigned direction code; the omitted category is Null/Mixed (0). Column (5) adds $\log(1 + N_{\text{authors}})$ as a proxy for team size. Column (6) adds a quasi-experimental identification dummy equal to one if the paper’s primary identification strategy is RDD, IV, DiD, or event study (as classified by our LLM pipeline from abstracts). Column (7) adds a top-institution indicator equal to one if the first author’s institutional affiliation (from OpenAlex API) matches a list of 50+ top finance and economics departments. Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. The baseline implied publication probability is $\Phi(-0.084) \approx 46.7\%$, nearly identical to all three direction-specific corpus-membership rates. Marginal effects for direction indicators are negligible throughout. The near-zero McFadden pseudo- R^2 in Columns (1)–(2) is expected under the null hypothesis that direction has no explanatory power for publication status; it is not an indicator of

Economic Magnitude

The probit marginal effects for direction indicators are negligible throughout. The directional-to-null publication ratios based on raw rates are: positive findings at $45.6\%/46.7\% = 0.98$; negative findings at $46.1\%/46.7\% = 0.99$. Both ratios are indistinguishable from 1.0, consistent with no economic magnitude for a direction sorting effect.

We note that the direction-to-null publication ratios ($\approx 1.0\times$) are not directly comparable to ?’s caliper estimates of 1.5–2.0 $\times$: Brodeur’s figures measure within-paper specification-selection pressure (the ratio of test statistics just below versus just above the 5% significance threshold in the published distribution), while our 1.0 $\times$ ratios measure cross-sectional publication likelihood (whether directional papers reach print at all). The within-paper comparison to Brodeur appears in Section ??, where our published caliper ratio of 1.20 is lower than Brodeur’s 1.5–2.0 range. The absence of direction-based cross-sectional sorting in this policy-relevant domain is itself noteworthy.

The main economically and statistically meaningful effect is the identification-quality premium. Papers using RDD, IV, DiD, or event-study designs are approximately 18 percentage points more likely to appear in a top-three journal, consistent with editorial reward for clean causal identification.

Cross-Journal Analysis

Journal-level estimates are reported in Appendix ??. No journal shows a statistically significant direction coefficient, consistent with the pooled null result. We caution that sub-samples are severely underpowered: the *Journal of Finance* contributes only 26 in-scope published articles, so cross-journal comparisons carry limited inferential weight and should be treated as descriptive.

Time-Period Heterogeneity (Exploratory)

The following sub-period analysis is exploratory. Sub-sample sizes are small ($N = 50, 95$, and 146 for the three periods) and the results should be interpreted as descriptive evidence,

not confirmatory tests. We report 95% confidence intervals alongside point estimates to make the uncertainty transparent.

Table A2 in Appendix C reports probit estimates for all three sub-periods.

During the *pre-crisis period* (2000–2007, $N = 50$), neither direction coefficient is statistically significant: $\hat{\beta}_1 = 0.376$ (95% CI $[-0.614, 1.366]$) and $\hat{\beta}_2 = -0.537$ (95% CI $[-1.628, 0.556]$). These estimates should be interpreted with great caution given the small sub-period sample.

During the *financial crisis and post-crisis period* (2008–2015, $N = 95$), both direction coefficients are small and insignificant: $\hat{\beta}_1 = -0.247$ ($p = 0.498$) and $\hat{\beta}_2 = -0.163$ ($p = 0.632$). No direction premium is detectable in this period.

In the *post-Dodd-Frank period* (2016–2024, $N = 146$), the pattern is unchanged: $\hat{\beta}_1 = -0.015$ ($p = 0.959$) and $\hat{\beta}_2 = 0.190$ ($p = 0.489$), both statistically indistinguishable from zero. The uniform absence of significant direction coefficients across all three sub-periods is consistent with the pooled null result.

6. Robustness

Table 7 summarizes the robustness checks. We organize them into five groups: caliper window sensitivity, sample and scope, alternative specifications, classification and quality concerns, and an alternative pre-publication baseline.

6.1 Caliper Window Sensitivity

The caliper result is window-dependent. Table ?? (reproduced from Appendix ??) repeats the main caliper test under four window specifications. Under the standard Brodeur window (1.645, 2.275) and the narrower (1.70, 2.22) and Brodeur2 (1.65, 2.25) variants — all with lower bounds at or above $|z| = 1.645$ — the published full-sample ratio ranges from 1.20 to 1.33 and the NBER ratio remains near 1.0. However, widening the lower bound to $|z| = 1.50$ reverses the sign for published papers (ratio 0.92, $p = 0.088$) because the density is still

rising in the (1.50, 1.645) interval. This confirms that the result is specific to the mass just below the 5% significance threshold (1.645, 1.96), not to a general rightward density shift.

Table 9. Caliper Window Robustness (Summary)

Window	Sample	Ratio	Naive p	Perm. p
Standard (1.645–2.275)	Published	1.20	0.002	0.001
	NBER	1.01	0.849	0.424
Narrow (1.70–2.22)	Published	1.33	<0.001	<0.001
	NBER	1.01	0.917	0.468
Wide (1.50–2.42)	Published	0.92	0.088	0.960
	NBER	0.97	0.549	0.737
Brodeur2 (1.65–2.25)	Published	1.33	<0.001	<0.001
	NBER	1.10	0.153	0.081

Notes. Caliper ratio = mass in (below, 1.96) / mass in (1.96, above). The wide window reverses sign for published papers because density is rising below $|z| = 1.645$; the three windows with lower bounds ≥ 1.645 show consistent excess mass in the published record. Bold entries significant at 5%.

6.2 Sample and Scope Robustness

Matched pairs only. Among the 16 published papers with a confirmed NBER working-paper match, we independently LLM-code the direction from the NBER *draft* abstract and from the published abstract, yielding 32 observations (16 draft-stage, 16 published-stage) for the same set of papers. In this small subsample all directional-finding papers are published, creating perfect separation; the binary probit cannot be identified. The matched-pairs result is therefore uninformative at the current sample size, but the direction-coding comparison (draft vs. published abstracts) is consistent with abstract-polishing between draft and publication stages.

Excluding monetary policy papers. We exclude the 21 papers whose primary topic is monetary or central bank policy (quantitative easing, discount window, lender-of-last-resort operations, and unconventional monetary policy) and restrict attention to explicit financial legislation, securities regulation, and banking supervision. The null result is unchanged: $\hat{\beta}_1 = -0.028$ ($p = 0.896$) and $\hat{\beta}_2 = -0.068$ ($p = 0.741$, $N = 270$), confirming that the null finding is not driven by the inclusion of monetary policy papers (within the power constraints

of this design).

Journal subsets. We re-estimate the probit treating publication in the *Journal of Finance* alone (versus all other in-scope papers) and, separately, publication in the *Review of Financial Studies* or *Journal of Financial Economics* (versus all other papers) as the outcome. Both premiums are statistically insignificant in the *Journal of Finance*-only specification ($\hat{\beta}_1 = 0.400$, $\hat{\beta}_2 = 0.350$, $p > 0.20$), as in the pooled result. The *Review of Financial Studies*+*Journal of Financial Economics* specification likewise yields insignificant coefficients ($\hat{\beta}_1 = -0.147$, $\hat{\beta}_2 = -0.130$, $p > 0.48$), confirming that the null result is not concentrated in any particular outlet (within the power of this design).

6.3 Alternative Specifications

Binary directional indicator. A potential concern about our direction coding is that classifying a finding as positive versus negative rests on a normative judgment about what constitutes a beneficial intervention outcome. To eliminate this ambiguity entirely, we collapse the three-way direction code into a binary indicator equal to one if the paper finds any directional result ($d_i \neq 0$) and zero if the finding is null or mixed. The coefficient is -0.020 (s.e. 0.182 , $p = 0.914$), confirming the null result within the power of this design (minimum detectable effect ≈ 22 percentage points at 80% power). Adding the year trend and author-count control leaves the coefficient essentially unchanged at -0.014 (s.e. 0.183 , $p = 0.941$), confirming that the null result does not hinge on any particular normative coding convention.

Linear probability model. Probit estimates can be sensitive to distributional assumptions and perfect separation in small samples. Re-estimating the main specification as a linear probability model (OLS with heteroskedasticity-robust standard errors) yields coefficients of -0.010 (s.e. 0.082 , $p = 0.900$) for positive findings and -0.006 (s.e. 0.079 , $p = 0.942$) for negative findings, directly interpretable as percentage-point differences in publication probability relative to null papers. Both are negligible and confirm the null result does not hinge on the probit functional form (within the power constraints of this design).

NBER-to-journal matching sensitivity. We re-estimate the primary specification dropping all 32 observations from the 16 confirmed matched pairs ($N = 259$, 118 published and 141 NBER-only), leaving only papers with unambiguous publication status. The direction coefficients remain small and statistically insignificant (detailed output in the replication package, `scripts/08_robustness.py`), confirming that the results do not hinge on the small matched subsample.

6.4 Classification and Quality Concerns

Quasi-experimental papers only. The most substantive quality-control check restricts the sample to papers using quasi-experimental identification strategies (RDD, IV, DiD, or event study), as classified by our LLM pipeline from abstracts. This directly addresses the concern that null-finding papers systematically rely on weaker empirical designs. Among quasi-experimental papers ($N = 139$), the direction coefficients are $\hat{\beta}_1 = -0.598$ ($p = 0.121$) and $\hat{\beta}_2 = -0.657^*$ ($p = 0.081$). $\hat{\beta}_1$ is clearly insignificant. $\hat{\beta}_2$ is marginally significant at the 10% level with a *negative* sign, implying quasi-experimental negative-finding papers are *less* likely to appear in top journals than null papers. We do not dismiss this estimate asymmetrically: the negative sign is entirely consistent with the pooled null result (no direction premium exists in the full sample either) and is consistent with sampling variability in a small subsample. We note, however, that this direction contradicts a simple direction-premium story in this sub-sample, and the subsample is too small ($N = 139$) to distinguish a genuine null from a genuine negative from a sampling artefact. Specifically, the QE subsample splits into roughly 40 published and 99 NBER papers, further divided by three direction categories, yielding approximately 10–15 observations per cell; a spurious significant coefficient is entirely expected at this cell size under the global null. The point estimate is also sensitive to which papers the LLM classifies as quasi-experimental. Neither direction coefficient supports a positive direction premium within the quasi-experimental sub-sample.

Abstract language and LLM classification confidence. Abstract writing style introduces

two related measurement concerns. Among published papers, 81.3% receive a high-confidence direction code; among NBER-only papers, 62.4% do, with 26.8% receiving low-confidence ratings. The 18.9-percentage-point gap in high-confidence rates persists despite NBER abstracts being longer on average (158 words) than published abstracts (109 words), suggesting the difference reflects writing style rather than abstract length. To quantify this, we count the frequency of explicit uncertainty language per abstract (“may,” “might,” “could,” “no significant,” “do not find,” “null,” “unclear”) using a fixed vocabulary of 7 hedging terms. NBER abstracts average 0.43 such terms per abstract, versus 0.19 for published abstracts—a $2.3\times$ difference in the raw count. Normalized by word count, the rates are similar (NBER: 0.27 terms per 100 words; published: 0.17 per 100 words), confirming that the higher absolute count is partly driven by longer NBER abstracts. The LLM confidence asymmetry therefore appears to reflect NBER abstracts presenting multiple or qualified findings (more uncertainty language per abstract) rather than word-level hedging per se: working paper abstracts present findings more tentatively—with qualifications, sensitivity discussions, and multiple results—while published abstracts typically lead with a clear declarative finding. First, hedged NBER abstract language may push genuinely directional NBER papers into the null category. Second, abstract polish after revise-and-resubmit may cause the LLM to classify more published papers as directional. Both concerns inflate the measured directional premium. Our replication exercise (Section 3.4) supports the inflation interpretation: of the 20 null-coded papers in the validation sample, 15 are reclassified as directional in the independent replication run, while no directional-coded papers are reclassified as null. This asymmetry implies that null-coding error inflates rather than attenuates our estimate.

Table 7 Panel C also reports a bounding specification that adds $\log(\text{abstract length})$ as a covariate ($\hat{\beta}_1 = 0.008$, $p = 0.968$; $\hat{\beta}_2 = 0.011$, $p = 0.958$). Both direction coefficients remain negligible, confirming that abstract-length controls do not alter the null result within the power of this design.

Matched-pairs direction stability. Among the 16 papers confirmed to exist as both an

NBER working paper and a top-journal publication, we independently LLM-code direction from the NBER draft abstract and from the published abstract. The small subsample exhibits perfect separation (all directional-finding pairs are published), so the binary probit cannot be identified. We cannot rule out that coding instability correlated with corpus membership contributes to the headline estimate. The matched-pairs check holds the paper fixed across draft and published stages, directly testing whether direction codes are stable within the same research project.

6.5 AFA Conference Baseline

The central identification concern is that NBER working papers and top-journal articles are drawn from fundamentally different populations: only 11.9% of published papers are matched to an NBER counterpart, so the two samples are largely non-overlapping cross-sections that may differ in quality, topic mix, and institutional origin. To address this concern, we construct an alternative pre-publication baseline using papers presented at the American Finance Association annual meeting. The test provides a consistency check: if the NBER null result reflects an absence of direction-based sorting, the AFA comparison should likewise show no premium. AFA papers represent completed research with final empirical results that have been publicly presented—they are pre-journal submissions in every meaningful sense.

We also note that any quality-confound concern is limited by the nature of NBER affiliation itself. Posting to the NBER Working Paper series requires affiliation with the NBER research network, which is restricted to faculty at leading research universities—an elite and selective credential in academic economics. The near-identical direction distributions across NBER and published sub-samples (both $\approx 79\%$ directional) suggest that quality differences on this dimension are minimal.

We collect 4,080 AFA program papers from 2006–2024. For 2006–2012 and 2014–2024, we scrape the AFA website directly; for 2013, whose AFA page returns a 404, we parse AFA

sessions from the full ASSA preliminary program hosted by the AEA (<https://www.aeaweb.org/conference/2013/preliminary.php>), yielding 186 papers from 57 AFA sessions. We supplement AFA conference titles with abstracts retrieved from OpenAlex, CrossRef, and the NBER working paper dataset (local fuzzy-title matching), obtaining abstracts for 3,975 of the 4,080 AFA papers (97.4% overall; coverage is higher for in-scope papers because government-intervention papers are more likely to be published). We apply our standard title-plus-abstract keyword filter (identical to the NBER pipeline) to the 3,975 AFA papers with abstracts, then run the LLM scope classifier, yielding **151 in-scope papers**. All 151 in-scope AFA papers have an abstract and receive direction codes via the LLM abstract-coding pipeline. 36 AFA papers (23.8%) are matched to a JF/RFS/JFE publication via fuzzy title matching.

Coding limitation. Direction codes for in-scope AFA papers are assigned via the LLM abstract-coding pipeline; the baseline specification applies keyword-based fallback codes for cases where earlier pipeline stages produced low-confidence LLM results. As a robustness check, Table 7 Panel D also reports a “Uniform LLM coding” specification that replaces these keyword-based fallback codes with the standard LLM pipeline throughout, yielding nearly identical estimates ($\hat{\beta} = 0.352$, $p = 0.204$) and confirming the result is not an artifact of the coding method.

The binary directional probit on the AFA baseline yields $\hat{\beta} = 0.327$ (s.e. 0.279, $N = 151$), statistically indistinguishable from zero. The NBER-baseline estimate is -0.020 (s.e. 0.182), likewise indistinguishable from zero. Neither estimate rejects the null of no direction-based sorting. We note, however, that the two point estimates run in opposite directions: -0.020 (NBER) versus $+0.327$ (AFA). Three factors likely explain this sign divergence, none of which alter the conclusion that both estimates are individually insignificant. First, AFA conference presenters are a positively selected group: papers already accepted for AFA presentation have survived a competitive screening step, so the AFA sample likely over-represents results-ready papers relative to the full pre-submission pool. If directional papers are somewhat

more likely to be accepted to AFA than null papers, the AFA baseline would mechanically show a positive direction–publication correlation independent of editorial selection. Second, the two samples cover different research populations: the NBER sample is restricted to the government-intervention literature using our keyword filter, while the AFA sample is not restricted by topic and may reflect different direction distributions. Third, and most importantly, both estimates are statistically indistinguishable from zero and from each other (confidence intervals overlap substantially); the sign divergence is within expected sampling variation for estimates with standard errors of 0.279 and 0.182 on a sample of 291 papers. The AFA baseline therefore does not *contradict* the NBER null, but neither does it corroborate it. More importantly, the sign flip from -0.020 (NBER) to $+0.327$ (AFA) illustrates a genuine fragility: the *sign* of the estimated direction premium is sensitive to the choice of pre-publication benchmark, even though neither estimate is individually significant. A reader could reasonably conclude that the available baselines are too few and too different from each other to establish even the sign of any editorial direction premium. The ideal benchmark—a large-scale SSRN corpus matched to our keyword taxonomy, capturing papers that circulate outside NBER—would provide a far more representative pre-submission distribution; we leave its construction for future work but flag the fragility here as a genuine limitation. Directional AFA papers are published in top journals at a 25.9% rate (30 of 116) versus 17.1% for null-finding papers (6 of 35). The pattern is unchanged when a year trend is added ($\hat{\beta} = 0.345$, s.e. 0.281).¹

¹As a direct test of overlap, we identified 22 papers appearing in both the NBER and AFA datasets via fuzzy title matching (token-sort score ≥ 90). In this subsample, directional papers publish at 81.2% versus 66.7% for null papers (LPM $\hat{\beta} = 0.146$, s.e. 0.208, $p = 0.49$, $N = 22$). The test is severely underpowered (only six null papers in the overlap), so no strong inference is possible.

Table 10. Robustness Checks

Specification	Variable	Coefficient	<i>p</i> -value	<i>N</i>
<i>Panel A: Sample and Scope</i>				
Excl. monetary policy	Positive	−0.028	0.896	270
	Negative	−0.068	0.741	270
<i>Journal of Finance</i> only	Positive	0.400	0.217	291
	Negative	0.350	0.271	291
<i>Review of Financial Studies</i> + <i>Journal of Financial Economics</i> only	Positive	−0.147	0.483	291
	Negative	−0.130	0.519	291
Matched pairs only [<i>N</i> =16 paper pairs]	Directional separation Draft vs. published version of same paper			32
<i>Panel B: Alternative Specifications</i>				
Binary directional	Directional	−0.020	0.914	291
Binary + year + authors	Directional	−0.014	0.941	291
LPM (OLS, HC3)	Positive	−0.010	0.900	291
	Negative	−0.006	0.942	291
<i>Panel C: Classification and Quality Concerns</i>				
Quasi-exp. papers only	Positive	−0.598	0.121	139
	Negative	−0.657*	0.081	139
+ log(abstract length) [abstract length control]	Positive	0.008	0.968	291
	Negative	0.011	0.958	291
<i>Panel D: Alternative Pre-Publication Baseline</i>				
AFA conference baseline (+ year trend)	Directional	0.327	0.241	151
	Directional	0.345	0.220	151
Uniform LLM coding [AFA, + year]	Directional	0.352	0.204	151
	Directional	0.364	0.193	151
<i>Panel E: Right-Censoring Robustness</i>				
NBER vintage ≤ 2020	Positive	−0.021	0.928	234
	Negative	−0.144	0.519	234

Notes. Each row reports the probit (or OLS for the LPM specification) coefficient on the indicated variable. The omitted category is null/mixed ($d_i = 0$), except for binary directional specifications where it is null/mixed vs. any directional finding. All probit specifications include a linear year trend unless otherwise noted. “Matched pairs only” uses the 16 papers confirmed to exist in both NBER and published form; the LLM codes each paper’s direction from its NBER *draft* abstract and from its published abstract separately, yielding $N = 32$ observations (16 draft, 16 published); all directional-finding pairs are published in this small subsample (perfect separation), so the probit cannot be identified. “Binary + year + authors” adds $\log(1 + N_{\text{authors}})$ to the binary specification. The quasi-experimental restriction uses identification strategies classified by our LLM pipeline from abstracts ($N = 139$). The abstract-length specification adds $\log(\text{abstract length})$ as a covariate; both direction coefficients remain negligible, confirming the null result is not driven by abstract-length differences. The AFA conference baseline uses AFA annual meeting papers (2006–2024, $N = 151$ in-scope) as the pre-publication comparison group; direction codes follow the same three-tier hierarchy as the NBER pipeline. “Uniform LLM coding” replaces keyword-based direction codes for unmatched AFA papers with the same LLM pipeline used for the main NBER/published sample; $\beta \approx 0.35$ remains insignificant ($N = 151$), confirming the AFA stage-identification result is not an artifact of the mixed coding method. Panel E restricts the NBER sample to papers circulated in 2020 or earlier, giving all papers at least four years to clear the publication pipeline; the direction coefficients are −0.021 and −0.144 (both $p > 0.50$), confirming the null is not explained by right-censoring of recent NBER papers. Standard errors omitted for brevity. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

7. Discussion

7.1 Policy Implications

Our results are *inconsistent with* the hypothesis that the published record of government-intervention research overstates the proportion of interventions with detectable effects due to significance-sorting. Corpus-membership rates across finding directions are nearly identical ($\approx 46\%$), and the direction of a paper’s finding has no statistically detectable association with its likelihood of appearing in a top-three finance journal. The published distribution closely mirrors the pre-publication distribution on this dimension.

Two power caveats qualify this conclusion. First, the minimum detectable difference in corpus-membership rates at 80% power is approximately 18 percentage points (probit coefficient ≈ 0.466), so moderate directional asymmetries below that threshold are consistent with the data. Second, the cluster-bootstrap 95% confidence interval for the negative-finding caliper ratio is $[0.81, 2.67]$, wide enough to encompass both negligible and substantial threshold pressure; the wide interval reflects genuine paper-level heterogeneity rather than a sampling artifact.

To the extent that structural bias-correction methods (?) applied to specific intervention literatures (capital requirements, short-sale bans, mandatory disclosure) would detect bias, that bias does not appear to originate from direction-based editorial selection in the top-three finance journals. Regulators who rely on the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics* literature on government intervention appear to face a record that is not systematically distorted in favor of or against directional findings relative to null results—at least not to a degree that the current design can detect (minimum detectable effect ≈ 22 percentage points; moderate directional asymmetries below this threshold remain consistent with the data).

The robust positive finding—that quasi-experimental identification quality predicts top-journal publication—has a different policy implication: the published literature is

methodologically selective. Papers using RDD, IV, DiD, or event-study designs are substantially more likely to appear in print, suggesting that the published evidence base is biased toward cleaner causal identification. For regulators, this is plausibly a desirable form of selection rather than a distortion.

Relevance to the corporate-finance regulation literature. The published sample in this paper is drawn exclusively from the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics*—the three premier corporate-finance journals—and includes every in-scope article on government intervention in financial markets that these journals published over 2000–2024. A large share of this literature directly concerns the core topics of corporate-finance scholarship: mandatory disclosure requirements (Sarbanes-Oxley, Reg FD), corporate governance regulation (say-on-pay, proxy access, board independence mandates), short-selling restrictions and market integrity, and Dodd-Frank provisions governing capital requirements and resolution authority. The NBER programs contributing the working-paper baseline similarly encompass the Corporate Finance, Asset Pricing, and Monetary Economics programs, all of which regularly study these regulatory interventions.

For readers interested specifically in the corporate-finance regulation literature, our findings carry two messages. First, the directional-neutrality result—robust to our power constraints—suggests that the published corporate-governance and disclosure record is unlikely to be *large-scale* skewed toward pro- or anti-intervention conclusions by editorial selection; both positive- and negative-finding studies on, say, the effect of say-on-pay or mandatory auditor rotation face similar corpus-membership rates. Second, the revision-stage bunching concentrated in negative-finding papers (though imprecisely estimated) is consistent with the hypothesis that anti-intervention evidence in the corporate-finance regulation literature faces within-paper specification pressure before reaching print, even when the publication decision itself is direction-neutral. Readers of the corporate-governance and disclosure evidence base should treat marginal significance (p just below 0.05) in anti-intervention findings with additional scrutiny, while treating publication status as a

reasonable signal of methodological quality.

7.2 Comparison to Other Domains

Our $\approx 1.0\times$ directional-to-null publication ratio stands in contrast to analogous estimates from adjacent literatures. ? document excess mass just below the conventional 5% significance threshold ($|z| \in (1.645, 1.96)$) in the *American Economic Review*, the *Journal of Political Economy*, and the *Quarterly Journal of Economics*, with implied ratios of roughly $1.5\text{--}2.0\times$ for 10%-significant-but-not-5% versus just-5%-significant findings. ?, studying NSF-funded social science experiments with pre-registered designs, find that 65% of null results were never published versus 21% of significant results—also implying a roughly $3\times$ publication premium. ? find, in a sample of nudge-unit RCTs with institutional outcome tracking, that studies with detectable effects are substantially more likely to appear in print than null-result studies. ? document that pre-registration and pre-analysis plans reduce but do not eliminate these biases in economics.

Our published-sample caliper ratio of 1.20 is below Brodeur et al.’s estimates of $1.5\text{--}2.0\times$ for the AER/JPE/QJE, and our NBER benchmark ratio of 1.01 confirms that this domain’s research base does not itself generate excess bunching before the editorial stage. On the cross-sectional direction-sorting dimension, our $\approx 1.0\times$ directional-to-null publication ratio is well below the $3\times$ estimate of ? and the analogous findings in ?. Together these comparisons suggest that the finance-regulation literature does not exhibit the significance-sorting patterns documented in other social science domains. This could reflect genuine differences in editorial norms at top finance journals, or could reflect the nature of our NBER comparison group (which captures well-resourced researchers whose null results may circulate widely, reducing the selection gap relative to the full population of research). Whether the absence of bias in top finance journals extends to lower-tier journals or comparable finance sub-fields without explicit policy content (e.g., corporate governance, asset pricing) is an important question for future research.

Within finance specifically, our within-paper caliper finding (published ratio = 1.20; NBER ratio = 1.01) can be read alongside ? and ?. Harvey and Liu model the same excess mass at the 5% threshold in published finance research and infer the pre-publication distribution structurally; our NBER benchmark observes it directly, confirming that the distortion is specific to the revision stage. Chen’s bound shows that the *far* right tail of published t-statistics ($|t| > 4$) is too fat to be explained by p-hacking alone, implying that most published effects are genuine. Our caliper evidence is consistent: the excess mass we document is concentrated just *below* the 5% threshold ($|z| \in (1.645, 1.96)$, the 10%-significant-but-not-5% zone), while the broader distribution shows a 55–57% significance rate that is stable across published and pre-publication samples. Real effects produce the fat tail; the caliper bunching is consistent with revision-stage specification pressure producing the excess mass in the 10%-significant-but-not-5% zone just below 1.96. The two observations reinforce rather than contradict each other.

7.3 Mechanisms

Two analytically distinct patterns require separate mechanistic accounts. The *absence* of direction-based editorial selection—roughly equal corpus-membership rates across finding directions—and the *presence* of within-paper threshold bunching concentrated in the published record point to different stages of the research pipeline.

For the direction-sorting null result, we consider two compatible explanations in the following subsection. For the caliper finding, the NBER benchmark provides a more direct answer: because the bunching is present in published papers but absent from their pre-publication counterparts, and because the overall significance rate is unchanged, the distortion is specific to the revision stage. One interpretation is that anti-intervention papers face elevated evidentiary demands—perhaps because anti-intervention findings invite more scrutiny—leading authors to search specifications until marginal results clear the 5% threshold. This interpretation is *consistent with* equal corpus-membership rates (no detectable direction-

based sorting between corpora) while generating a within-paper conditional threshold premium for negative findings. A compositional alternative deserves acknowledgment: negative-finding published papers may systematically differ from their NBER counterparts in unobserved dimensions (research design complexity, empirical detail, vintage) that correlate with within-paper specification counts; our cross-sectional design cannot distinguish between these mechanisms. It is also consistent with ?’s finding that the far right tail of published t-statistics cannot be explained by p-hacking alone: the excess mass we document just below $|z| = 1.96$ reflects selective reporting of marginally significant specifications, not fabrication of large effects.

7.4 What the Data Can and Cannot Tell Us

Our empirical design estimates the cross-sectional association between finding direction and publication status. This association is consistent with two broad classes of mechanism, and distinguishing them is important because they have different policy implications. Our data cannot distinguish between them, and we do not attempt to do so.

Supply-side equilibrium. Researchers who anticipate that null results face no publication penalty may submit them to top journals at higher rates than in domains with known significance sorting, producing an equilibrium with equal submission rates across directions. Under this mechanism, null results are not underrepresented in the NBER pool relative to the published pool, which is precisely what we observe (NBER directional rate: 79.6%; published directional rate: 79.1%).

Identification quality and correlated attributes. The positive quasi-experimental premium (0.467^{***}) is consistent with editorial selection in this domain operating on methodological quality, among other factors. We caution, however, that the LLM-coded quasi-experimental dummy is likely correlated with author quality, institutional affiliation, topic type, and other unobserved quality dimensions that the probit cannot separate from iden-

tification strategy alone. If clean-identification studies have directional findings at similar rates to non-causal studies in this domain, an identification-quality filter would generate no direction sorting in the aggregate. Our data are consistent with this interpretation, but the identification-quality channel is one of several plausible explanations for the premium.

Either mechanism could suppress a direction-publication association even if publication bias exists elsewhere in the pipeline.

Without submission-level data (desk-rejection rates by direction, time-to-first-decision, share of null-result papers sent out for review), we cannot test whether the absence of a direction premium in the published record reflects true editorial neutrality or an upstream equilibrium in submission behavior. Future work following the approach of ?—who obtain institutional data on submission and publication outcomes for nudge-unit RCTs—would allow a sharper test of whether journals are directionally neutral or merely receive a directionally balanced submission pool.

7.5 Limitations

Six main threats to the validity of our analysis deserve discussion.

Fundamental identification constraint. The most important limitation is that NBER working papers and published journal articles are not different states of the same population of papers. They are largely non-overlapping samples selected through different institutional processes: only 16 of 134 published papers (11.9%) are matched to an NBER working paper. Our design therefore compares two different cross-sections, not papers tracked through the editorial filter. The observed difference in direction distributions is consistent with editorial selection but equally consistent with NBER papers having a different direction distribution for reasons unrelated to journals—for example, NBER researchers may strategically circulate papers with more ambiguous findings, or null-finding papers in this domain may disproportionately circulate through SSRN rather than NBER. Addressing this fully would require either submission-level data from journal editors (desk-rejection rates and review outcomes

by direction code) or a pre-registration registry for finance papers analogous to the AEA RCT Registry. SSRN preprint histories for authors in our sample would also allow a more comprehensive pre-publication baseline by capturing papers that circulate outside NBER. Constructing such a baseline—by scraping SSRN, CEPR, and ECGI working-paper archives for government-intervention finance papers matching our keyword taxonomy—would allow a direct evaluation of whether the NBER distribution is representative of the broader pre-submission pool; we leave this for future work. Similarly, controlling for researcher credentials (historical publication count, institutional ranking) and for paper-specific quality proxies beyond the LLM-coded quasi-experimental dummy would sharpen the probit estimates; the current dataset lacks author-level metadata, and collecting it is a meaningful extension. We interpret all coefficients as cross-sectional associations, not causal editorial-filter effects.

Narrow publication indicator and attenuation bias. The probit dependent variable codes a paper as $\text{Published} = 1$ only if it appeared in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics*. NBER papers that eventually appeared in the *Review of Finance*, the *Journal of Financial and Quantitative Analysis*, or other peer-reviewed venues are coded $\text{Published} = 0$, introducing measurement error in the dependent variable. This mis-coding attenuates the direction coefficient toward zero, biasing our probit estimates in the direction of the null we report. A broader publication indicator covering any peer-reviewed finance journal would reduce this attenuation; we leave its construction to future work, but note that the bias makes our reported null results a conservative lower bound on any true direction premium.

LLM classification errors and abstract polish. Our pipeline assigns direction codes by reading paper abstracts. A key concern—confirmed by our data—is that NBER working paper abstracts are written in a more tentative, multi-result style than published abstracts, even though they are longer on average (158 words vs. 109 words for published abstracts). Accordingly, 81.3% of published papers receive a “high-confidence” LLM direction code, versus 62.4% of NBER-only papers. If the LLM more often assigns a null code to hedged NBER ab-

stracts, the measured null-paper corpus-membership rate would be biased downward—which would *inflate* our estimated directional premium relative to the truth.

Our replication exercise (Section 3.4) provides direct evidence on this null-coding instability: of the 20 null-coded papers in the validation sample, 15 are re-classified as directional in the independent replication run, confirming that the null category is the least stable. This instability implies that using the alternative rater’s codes would yield a lower estimated directional premium than ours, because some papers we code as null would be coded as directional by an alternative rater—which would increase the denominator of the directional corpus-membership rate (since the reclassified papers are unpublished NBER working papers, so the numerator is unchanged) while reducing the denominator of the null rate, inflating the measured null corpus-membership rate and compressing the estimated gap between directional and null rates. Our reported estimates should thus be interpreted as a conservative upper bound on any true directional premium—the actual premium is no larger than what we measure, and may be smaller. Future work with standardized abstract templates, full-paper coding (rather than abstract coding), or pre-registered abstract language could address this concern directly.

A further implication of null-category instability deserves separate emphasis. In our probit models, Null/Mixed (0) is the *omitted category*: the direction coefficients $\hat{\beta}_1$ (positive) and $\hat{\beta}_2$ (negative) measure the directional premium *relative to null papers*. If the null category is unstable—as the 25% replication agreement suggests—the omitted-category baseline is itself measured with considerable error. This does not mechanically bias the direction coefficients in a specific direction (instability in the null group inflates the standard error of the baseline, which in turn widens confidence intervals for the direction dummies), but it does mean that the estimated null corpus-membership rate of $\approx 46.7\%$ is substantially less reliable than the directional rates. Readers should interpret the probit intercept and the direction-specific contrasts against the null baseline with appropriate caution. The χ^2 test of *joint equality* across all three direction categories (positive, negative, null) is more

robust to this concern because it does not depend on a single omitted-category reference point; that test gives $\chi^2 = 0.017$ ($p = 0.992$), consistent with the overall null result.

Quality controls and the identification-strategy alternative. The positive quasi-experimental premium (0.467***) confirms that identification quality is a genuine predictor of publication. Papers using RDD, IV, DiD, or event-study methods are approximately 18 percentage points more likely to appear in a top journal, even controlling for direction. Importantly, including this control does not alter the direction coefficients ($\hat{\beta}_1 = -0.135$, $\hat{\beta}_2 = -0.156$, both insignificant), ruling out the concern that the null result on direction reflects confounding with identification quality.

We acknowledge that the LLM-based identification-strategy coding from abstracts is itself imperfect: fewer than half of abstracts explicitly name their identification strategy, so the LLM infers it from other signals. No separate validation run has been completed for the QuasiExp indicator (unlike the direction codes, which have a two-run replication with $\kappa = 0.674$). A hand-coded assessment using full paper text would provide a more definitive check and is a natural extension of this work.

A related concern is that quasi-experimental papers might be disproportionately directional, in which case the QuasiExp control could absorb a direction effect. Table ?? speaks directly to this: quasi-experimental papers are somewhat more likely to be directional (88.5%, 123/139) than non-quasi-experimental papers (71.1%, 108/152), consistent with cleaner identification yielding sharper estimates. However, within each QuasiExp tier the corpus-membership rate is flat across direction categories (≈ 36 – 39% for non-QE papers; ≈ 52 – 75% for QE papers), and adding QuasiExp to the probit leaves the direction coefficients essentially unchanged ($\hat{\beta}_1 = -0.135$, $\hat{\beta}_2 = -0.156$, both insignificant). The identification-quality premium therefore appears to operate independently of direction sorting rather than absorbing it.

NBER selection and the direction of potential bias. NBER researchers are among the most productive and well-connected academics in finance. If their directional-finding

papers are more publishable than directional-finding papers from the broader population, the NBER comparison group may not be representative of the typical pre-publication paper. Under this scenario, our null result could mask publication bias present in the broader research population. We cannot rule out this possibility without submission-level data or a more representative pre-publication baseline. Absent such data, the most reassuring feature of our design is the close alignment of direction rates between NBER (79.6% directional) and published (79.1% directional), which suggests little residual direction-based sorting at least within this elite sample.

A second feature of the design that biases toward the null is our definition of “published”: NBER papers that appeared in journals outside the top-three (e.g., *Review of Finance*, *JFQA*) are coded as *Published* = 0, adding measurement error to the dependent variable and attenuating any true direction premium. Constructing a broader publication indicator covering all peer-reviewed venues would reduce this attenuation; such a check requires matching against a comprehensive publication database and is left for future work. For present purposes, the attenuation-toward-null bias means that our direction-sorting null is, if anything, conservative.

Right-censoring in the year trend. The positive and significant year trend ($\hat{\gamma} \approx 0.033^{***}$) is consistent with genuine improvement in publication prospects for recent papers, but an equally plausible mechanical explanation is right-censoring: NBER working papers from 2020–2024 have had fewer years to complete the editorial pipeline and appear in a top journal. Our cross-sectional design cannot distinguish these alternatives. The year trend should therefore be interpreted as descriptive rather than as evidence that finance journals have become more open to government-intervention research over time.

Multiple comparisons. Our main hypothesis—that directional-finding papers are published at different rates than null-finding papers—is pre-specified and tested in a single primary regression. The cross-journal and sub-period analyses are exploratory. All main results ($p > 0.90$) fall far from conventional significance thresholds; the concern is power

rather than false discovery.

Our caliper analysis runs 18 comparisons across the published sample (full sample, three journals, three directions, two row types, two inference procedures) and mirrors them for the NBER benchmark; see Section ?? for the pre-specified primary tests. Applying a Bonferroni correction at the 5% level requires $p < 0.003$ for 18 caliper comparisons. Among the caliper tests, the full-sample published caliper ($p = 0.002$), the *Journal of Financial Economics* caliper ($p < 0.001$), and the negative-direction caliper ($p < 0.001$) all survive this threshold. Among the density-discontinuity tests, the *Journal of Finance* ($p < 0.001$) and null/mixed ($p < 0.001$) discontinuity results likewise survive Bonferroni correction, providing independent confirmatory evidence. The corresponding NBER caliper tests ($p > 0.84$) do not approach significance, confirming the absence of pre-publication bunching. The core contrast between published bunching and NBER non-bunching is therefore robust to multiple-comparison correction.

8. Conclusion

We apply the ? caliper test to 7,761 $|z|$ -statistics from 121 published papers and 8,351 statistics from 125 NBER working papers on government-intervention finance. **Primary specification (verifiable rows):** among the 5,510 published rows with verifiable coefficient-SE pairs, the caliper ratio is 0.81—no excess bunching below the 5% threshold—and is statistically indistinguishable from the NBER ratio of 0.95. Primary regression rows show no systematic threshold manipulation in either corpus. **Secondary finding (standalone z-statistics):** standalone statistics show elevated bunching in published papers (ratio 2.55) relative to NBER (ratio 1.49). The full-sample published ratio of 1.20 (cluster-bootstrap $p = 0.150$, CI [0.88, 1.66]) is approximately equally attributable to a within-type effect (+0.097: standalone statistics bunch more intensely in published papers at the same z-stat-only share) and a composition effect (+0.094: published papers contain more standalone-statistic rows: 31.6% vs. 21.0%). The directional pattern—standalone

bunching concentrated in negative-finding papers (ratio 1.52, cluster-bootstrap $p = 0.135$, CI [0.81, 2.67])—is suggestive but not statistically robust under paper-level bootstrap inference. The NBER corpus serves as a proxy for the pre-publication distribution; it is selected on NBER affiliation and only 12% of published papers are matched to NBER counterparts, so compositional confounds cannot be ruled out.

As supporting context, we examine cross-sectional direction sorting in a 291-paper comparison of two largely non-overlapping corpora (134 published, 157 NBER working papers). We find no *detectable* direction-based sorting within the power of this design. Corpus-membership rates are nearly identical across positive-finding (45.6%), negative-finding (46.1%), and null-finding (46.7%) papers ($\chi^2 = 0.017$, $p = 0.992$); a binary directional probit is -0.020 (s.e. 0.182, $p = 0.91$), robust across all six specifications. This design rules out sorting ≥ 22 percentage points at 80% power; moderate asymmetries are consistent with the data. These rates depend on sample construction and are not estimates of actual acceptance probabilities. The strongest predictor of corpus membership is an LLM-coded quasi-experimental identification dummy (0.467***, s.e. 0.157), but this variable lacks human-coded validation and likely captures a bundle of correlated quality attributes; it should not be interpreted as evidence that journals specifically reward causal identification.

The joint pattern shows *absent* directional sorting alongside *within-paper* threshold bunching concentrated in the published record. We stress that compositional differences between NBER affiliates and the broader population of published authors could partly explain both results, and that neither claim about mechanism is directly identified.

We interpret these results throughout as cross-sectional associations. Our design identifies a null association between finding direction and corpus membership (published vs. NBER), but cannot rule out that this reflects an upstream equilibrium in submission behavior rather than directional indifference on the part of editors. Two practical implications follow. For researchers and meta-analysts, the published distribution in this domain does not appear systematically biased toward directional findings at the editorial-selection level; concerns

about cross-sectional significance-sorting in this specific corpus do not find empirical support. For journal editors, the results are consistent with a selection process consistent with rewarding methodological quality over finding direction—though the quasi-experimental premium likely captures a bundle of correlated quality attributes beyond identification strategy alone. The within-paper test statistic patterns suggest that specification searching at the revision stage remains a concern independent of the directional-sorting question.

References

Appendix A. Keyword Taxonomy

Table A1 lists all terms used in Pass 1 of the classification pipeline.

Table A1. In-Scope Classification: Keyword Taxonomy

Category	Keywords
Financial regulation	regulation, regulatory, Dodd-Frank, Basel, Glass-Steagall, capital requirements, leverage ratio, bank capital
Securities law	SEC, disclosure requirement, Reg FD, Sarbanes-Oxley, SOX, insider trading law, short-selling ban, short-sale restriction
Banking policy	FDIC, deposit insurance, lender of last resort, bailout, TARP, stress test, bank resolution, government guarantee
Monetary/fiscal policy	quantitative easing, QE, Fed intervention, government subsidy, fiscal stimulus
Market microstructure rules	circuit breaker, tick size, trading halt, uptick rule, limit-down, limit-up, trading curb
Corporate governance law	board independence requirement, say-on-pay, proxy access, mandatory audit, governance mandate, clawback provision, mandatory disclosure
General intervention signals	government intervention, government policy, public policy, law and finance, legal protection, shareholder rights, creditor rights, investor protection

Ambiguous standalone terms (*regulation*, *regulatory*, *SEC*) require co-occurrence with at least one additional in-scope keyword before the paper is flagged as a candidate.

Appendix B. LLM Prompts

2.1 In-Scope Classification Prompt

A finance research paper is ‘in scope’ for a study of publication bias in government intervention research if and only if its PRIMARY research question is the CAUSAL EFFECT of a specific government law, regulation, or public policy on financial markets, firms, or investors. Exclusion criteria: (1) purely theoretical papers with no empirical test of a policy; (2) papers that only DESCRIBE a regulation without estimating its effect; (3) papers whose main topic is private firm or market behavior, with policy as incidental context; (4) papers studying central bank communication WITHOUT a specific policy instrument. Return JSON with exactly three keys: `in_scope` (boolean), `confidence` (‘‘high’’, ‘‘medium’’, or ‘‘low’’), `reason` (one sentence).

2.2 Direction Coding Prompt

The following is the abstract of a finance paper that studies the effect of a government law, regulation, or public policy. Classify the MAIN empirical finding as: +1 if the government intervention produced a POSITIVE or BENEFICIAL outcome (e.g., improved market quality, reduced systemic risk, better investor protection); -1 if it produced a NEGATIVE or HARMFUL outcome (e.g., reduced liquidity, higher costs, unintended consequences, welfare losses); 0 if the finding is NULL, MIXED, or INCONCLUSIVE. Coding rule: code by whether the estimated effect represents an improvement or worsening relative to the stated regulatory goal, NOT the arithmetic sign of the regression

coefficient and NOT the authors' normative framing. Return JSON with: direction (1, -1, or 0), confidence ('high', 'medium', or 'low'), rationale (one sentence).

Appendix C. Cross-Journal Probit Estimates

Table 6 reports journal-specific probit estimates. No journal shows a statistically significant direction coefficient. At the *Journal of Finance*, point estimates are $\hat{\beta}_1 = 0.393$ ($p = 0.269$) and $\hat{\beta}_2 = 0.329$ ($p = 0.345$). The *Review of Financial Studies* yields $\hat{\beta}_1 = -0.036$ ($p = 0.900$) and $\hat{\beta}_2 = 0.252$ ($p = 0.350$). The *Journal of Financial Economics* yields $\hat{\beta}_1 = -0.136$ ($p = 0.572$) and $\hat{\beta}_2 = -0.305$ ($p = 0.200$). Cross-journal Wald tests for equality of the positive coefficients yield $p = 0.349$ (JF vs. RFS), $p = 0.218$ (JF vs. JFE), and $p = 0.791$ (RFS vs. JFE); none is statistically distinguishable. These results are consistent with the pooled null but carry limited inferential weight given the small sub-samples (especially JF: 26 papers).

Table A2. Probit Models by Journal

	<i>Journal of Finance</i>	<i>Review of Financial Studies</i>	<i>Journal of Financial Economics</i>
Positive (+1)	0.393 (0.355)	-0.036 (0.290)	-0.136 (0.241)
Negative (-1)	0.329 (0.349)	0.252 (0.270)	-0.305 (0.238)
Year trend	-0.010 (0.017)	0.058*** (0.017)	0.032** (0.015)
N	183	205	217
Pseudo R^2	0.011	0.065	0.027

Notes. Each column uses the 157 NBER-only papers (Published = 0) and the published papers from the specified journal (Published = 1) only; published papers from the remaining two journals are excluded, giving $N = 183$, 205, and 217 respectively. Each column therefore compares each journal's directional mix against the shared NBER working-paper pool. Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. The *Journal of Finance* sample contains only 26 in-scope published articles in 2000–2024, limiting statistical power for journal-specific tests.

Appendix D. Additional Tables

Table A3. Probit Results by Sub-Period (*Exploratory*)

	2000–2007	2008–2015	2016–2024
Positive (+1)	0.376 (0.505) [−0.61, 1.37]	−0.247 (0.364) [−0.96, 0.47]	−0.015 (0.288) [−0.58, 0.55]
Negative (−1)	−0.537 (0.557) [−1.63, 0.56]	−0.163 (0.340) [−0.83, 0.50]	0.190 (0.274) [−0.35, 0.73]
N	50	95	146
Pseudo R^2	0.072	0.004	0.005

Notes. **This analysis is exploratory.** Sub-period sample sizes are small ($N = 50, 95, 146$) and results should be interpreted as descriptive evidence, not confirmatory tests; confidence intervals overlap substantially across periods and encompass zero in all cases. Each column estimates equation (1) on the indicated sub-period without additional controls. Standard errors in parentheses; 95% confidence intervals in brackets. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. Period 2008–2015 spans the financial crisis and post-crisis regulatory response (Dodd-Frank enacted July 2010); 2016–2024 covers the post-implementation period. All direction coefficients are statistically indistinguishable from zero across all three sub-periods, consistent with the pooled null result.

Appendix E. LLM Direction Coding: Replication Confusion Matrix

Table A4. LLM Direction Coding: Replication Confusion Matrix

Original code	Replication code			Row total
	−1	0	+1	
Negative (−1)	20	0	0	20
Null/Mixed (0)	9	5	6	20
Positive (+1)	2	0	18	20
Column total	31	5	24	60

Notes. Entries are paper counts from a 60-paper stratified random subsample (5 papers per source $\times$ direction cell, balanced across the three journals and NBER). Rows show the original direction code used in the main analysis; columns show the code from an independent second LLM run (claude-haiku-4-5-20251001, temperature = 0; see Section 3.4). Bold entries on the main diagonal indicate agreement. Overall agreement: $43/60 = 71.7\%$. Linear-weighted Cohen’s $\kappa = 0.674$ (“substantial agreement” by the ? scale). The negative category achieves 100% agreement; the positive category 90%. Disagreements are concentrated in the null/mixed category (25% agreement), with 15 of 20 null-coded papers re-classified as directional in replication—consistent with null abstracts being less decisive and more sensitive to prompt context. Because coding instability pushes ambiguous papers *toward* directional codes rather than toward null, any resulting measurement error inflates our estimated directional premium.

Appendix F. Data and Replication

All data and code necessary to reproduce the results of this paper are available at [https://github.com/\[REPOSITORY\]](https://github.com/[REPOSITORY]). The repository contains:

- `requirements.txt` — Python package versions used
- `scripts/01_collect_nber.py` — NBER API data collection
- `scripts/02_collect_journals.py` — Journal data collection
- `scripts/02b_enrich_jfe_abstracts.py` — Elsevier API abstract enrichment for *Journal of Financial Economics*

- `scripts/03_classify_inscope.py` — Two-pass in-scope classification (keyword + LLM)
- `scripts/04_code_direction.py` — LLM direction coding
- `scripts/05_link_nber_to_published.py` — Fuzzy title matching for NBER-to-journal linkage (Round 1)
- `scripts/05b_improve_matching.py` — API-based abstract matching for unmatched NBER papers (Round 2)
- `scripts/05c_fulltext_matching.py` — LLM full-text matching via NBER PDFs (Round 3)
- `scripts/05d_published_fulltext_matching.py` — LLM matching using published PDF full text (Round 4)
- `scripts/06_analysis.main.py` — Primary probit regressions and summary statistics
- `scripts/08_robustness.py` — Robustness checks
- `scripts/09_code_id_strategy.py` — LLM identification-strategy classification (RDD, IV, DiD, event study)
- `scripts/10_collect_afa.py` — AFA annual meeting program collection (2006–2024)
- `scripts/10b_enrich_afa_abstracts.py` — Abstract enrichment for AFA papers via OpenAlex, CrossRef, and NBER
- `scripts/10c_fix_2014_2016_parsing.py` — Corrected HTML parsing for 2014–2016 AFA program pages
- `scripts/11_afa_analysis.py` — AFA-baseline probit regressions (AFA conference baseline, Section 6)
- `scripts/11c_recode_afa_direction_llm.py` — LLM direction re-coding for AFA papers (uniform coding robustness check)
- `data/` — Final analysis dataset in Parquet format
- `validation/human_coding_sheet_40.xlsx` — Coding sheet for 40-paper human direction-coding validation (to be completed before submission)

- `scripts/13_human_validation_kappa.py` — Computes Cohen’s κ between human and LLM codes once coding sheet is complete
- `paper/` — L^AT_EX source for this paper

Appendix G. LLM Pipeline Validation

Direction-coding reliability. The primary validation evidence for our LLM pipeline concerns direction coding, reported in the main text (Section 3.4). We re-ran the direction-coding protocol on a stratified random sample of 60 papers and treated the second run as an independent rater. Overall linear-weighted Cohen’s $\kappa = 0.674$ (“substantial agreement”); agreement is near-perfect for directional papers (-1 : 100%, $+1$: 90%) and concentrated in the null/mixed category (25%). The confusion matrix is reported in Appendix Table A3.

In-scope classifier design. The in-scope LLM classifier is applied to all keyword-flagged papers after a broad keyword filter selects candidates from the raw journal and NBER corpora. The keyword list is deliberately over-inclusive (any mention of regulation, law, intervention, or policy in the title or abstract), so the classifier’s primary role is removing false positives from the keyword stage rather than detecting in-scope papers from scratch.

A conservative bound on the in-scope false-positive rate is available from the data itself: all 291 in-scope papers were independently verified as on-topic during the matching and scope-review stages of the pipeline (Sections 3.3 and 3.5). No paper was flagged in-scope by the LLM and subsequently removed as out-of-scope during these downstream checks, giving a post-hoc precision of 100% within the reviewed portion of the sample.

The more material concern is the false-negative rate—papers that are genuinely in-scope but missed by the keyword filter or misclassified as out-of-scope by the LLM. False negatives would reduce the representativeness of our sample but would not mechanically bias the direction-publication association unless missed papers systematically differ in direction from included papers. Given that the keyword filter captures all standard regulatory and

supervisory terminology in finance, we expect any false negatives to be idiosyncratically missed rather than directionally selected.

Appendix H. Direction-Coding Reliability by Source

Table A4 disaggregates the LLM replication agreement (from Appendix D) by source corpus. The replication used an independent second LLM run (claude-haiku-4-5-20251001, temperature = 0) on the same stratified sample of 60 papers. Agreement is highest for NBER papers (80.0%; $\kappa_w = 0.714$) and lowest for *Journal of Finance* papers (60.0%; $\kappa_w = 0.533$). In every source, disagreements are confined to the null/mixed category; directional codings (± 1) are nearly perfectly replicated across all sources.

Table A5. LLM Replication Agreement by Source (N = 60)

Source	N	Agreement	κ_w	−1 agree	0 agree	+1 agree
JF	15	60.0%	0.533	5/5	0/5	4/5
RFS	15	73.3%	0.727	5/5	1/5	5/5
JFE	15	73.3%	0.727	5/5	1/5	5/5
NBER	15	80.0%	0.714	5/5	3/5	4/5
Total	60	71.7%	0.674	20/20	5/20	18/20

Notes. Each source contributes 15 papers (5 per direction cell) from a stratified random sample. *Agreement* is the fraction of papers where the original and replication LLM codes match. κ_w is the linear-weighted Cohen’s κ (direction, 3-category). “ d agree” columns show the number of papers with original code d that receive the same code in replication. Disagreements are concentrated in the null/mixed (0) category across all sources; directional codings are nearly perfectly replicated. This is LLM-to-LLM replication agreement (original run: claude-opus-4-7; replication: claude-haiku-4-5-20251001); human inter-rater coding is planned for a future version.

Appendix I. Caliper Window Robustness

Table ?? repeats the main caliper test under four window specifications: the standard Brodeur (1.645–2.275) window, a narrower (1.70–2.22) window, a wider (1.50–2.42) window, and

the Brodeur2 (1.65–2.25) variant. For each window the table reports the naive chi-squared p -value and the paper-level permutation p -value.

Table A6. Caliper Window Robustness

Window	Sample	N	Below	Above	p -value		
					Ratio	Naive	Perm.
Standard (1.645–2.275)	Published	7,761	611	508	1.20	0.002	0.001
	NBER	8,351	501	495	1.01	0.849	0.424
Narrow (1.70–2.22)	Published	7,761	542	408	1.33	<0.001	<0.001
	NBER	8,351	416	413	1.01	0.917	0.468
Wide (1.50–2.42)	Published	7,761	807	877	0.92	0.088	0.960
	NBER	8,351	726	749	0.97	0.549	0.737
Brodeur2 (1.65–2.25)	Published	7,761	606	457	1.33	<0.001	<0.001
	NBER	8,351	497	453	1.10	0.153	0.081

Notes. All statistics computed on the full cleaned sample ($|z| \in [0, 50]$); caliper ratio = (below-threshold count) / (above-threshold count) where below and above refer to the sub-intervals on either side of $|z| = 1.96$. *Perm.* is the within-group permutation p -value (10,000 iterations). Bold entries significant at 5%. The wide window (1.50–2.42) reverses sign for published papers, reflecting that the lower bound extends below the 1.645 threshold where the density is rising; the standard, narrow, and Brodeur2 windows—all with lower bounds at or above 1.645—show consistent excess mass in the published record.

Appendix J. Caliper Robustness to Coefficient Inclusion

Table ?? repeats the main caliper test after excluding rows whose variable name matches a broad set of common control-variable patterns (size, leverage, age, GDP, ROA, ROE, return volatility, fixed-effect indicators, intercepts, etc.). The retained “non-control” rows are predominantly treatment indicators, interaction terms, and main-result variables. This restriction addresses the referee’s concern that pooling all table statistics — including controls — could confound the direction-based caliper comparison, since a negative-finding paper’s tables contain many statistics unrelated to its headline finding.

The non-control subsample retains 7,055 published statistics (89.4% of the full sample) and 7,599 NBER statistics (89.4%), distributed across the same 108 and 106 papers. Caliper ratios and p -values are essentially unchanged from the full-sample results: the published

full-sample ratio remains 1.203 and the negative-finding ratio *increases* to 1.625 (from 1.522), while NBER ratios remain near 1.0. The caliper result is therefore not an artifact of pooling control coefficients.

Control-variable exclusion taxonomy (exact regex). For replication purposes, the following regular expression (Python, `re.IGNORECASE`) identifies rows classified as control variables and excluded from the non-control subsample:

```
(^log\s*[\(\[\]?s*(asset|sales|size|market|revenue|total\s+asset|equity|
    debt|mv|bv|ta)\b
|(?<!\w)leverage(?:\s*[\*x]?s*(treat|post|reform))
|(?<!\w)(long.term\s+)?debt[/ ]?(to|ratio|level)(?:\w)
|(?<!\w)size(?:\s*[\*x]?s*(treat|post|reform|policy))(?:\w)
|(?<!\w)(firm\s+)?age(?:\s*[\*x]?s*(treat|post|reform|policy))(?:\w)
|(?<!\w)roa(?:\w)|(?<!\w)roe(?:\w)|(?<!\w)ebitda(?:\w)
|(?<!\w)profitab[a-z]*(?:\w)
|(?<!\w)book.?to.?market(?:\w)|(?<!\w)btm(?:\w)
|(?<!\w)tobin[\'s ]+q(?:\w)|(?<!\w)q\s+ratio(?:\w)
|(?<!\w)volatilit[a-z]*(?:\w)|(?<!\w)return\s+volat[a-z]*(?:\w)
|(?<!\w)momentum(?:\w)|(?<!\w)market\s+beta(?:\w)
|(?<!\w)cash\s+flow(?:\w)
|(?<!\w)(log\s+)?market\s+(cap|value)(?:\w)|(?<!\w)mcap(?:\w)
|(?<!\w)stock\s+return(?:\w)
|\byear\s*d{2,4}\b|\byear\s+fe\b
|(?<!\w)intercept(?:\w)|(?<!\w)_cons(?:\w)|(?<!\w)constant(?:\w)
|fixed\s+effect[s]?|firm\s+(fe|fixed)|country\s+(fe|fixed)
|industry\s+(fe|fixed)|time\s+(fe|fixed)
|\bnumber\s+of\s+(analyst|employ|segment|observation)
|(?<!\w)(log\s+)?gdp(?:\w)|(?<!\w)inflation(?:\w)
```

| (?<!\w)interest\s+rate(?!\w))

The full Python implementation is in `scripts/25_primary_coeff_caliper.py`, available in the replication archive.

Table A7. Caliper Test: All Statistics vs. Non-Control Statistics

Sample	Subset	Filter	N	Papers	Ratio	Naive p	Perm. p	Boot. p
Published	Full sample	All stats	7,761	108	1.203	0.002	0.001	0.150
Published	Full sample	Non-control	7,055	108	1.203	0.003	0.001	0.161
Published	Negative-finding	All stats	3,491	50	1.522	<0.001	<0.001	0.135
Published	Negative-finding	Non-control	3,139	50	1.625	<0.001	<0.001	0.121
Published	Positive-finding	All stats	2,633	37	0.941	0.565	0.732	0.623
Published	Positive-finding	Non-control	2,409	37	0.927	0.481	0.769	0.637
NBER	Full sample	All stats	8,351	106	1.012	0.849	0.424	0.432
NBER	Full sample	Non-control	7,599	104	0.981	0.767	0.625	0.561
NBER	Negative-finding	All stats	4,375	50	0.996	0.964	0.532	0.474
NBER	Negative-finding	Non-control	3,981	50	0.955	0.633	0.693	0.614

Notes. “Non-control” excludes rows whose variable name matches common control-variable patterns (size, leverage, ROA/ROE, book-to-market, GDP, inflation, firm age, return volatility, fixed-effect indicators, intercepts, etc.). The retained rows are predominantly treatment indicators, DiD interactions, and main-result variables. Caliper ratio is mass in (1.645, 1.96) relative to (1.96, 2.275). *Perm.*: within-group permutation (10,000 iterations). *Boot.*: cluster bootstrap (10,000 draws) resampling papers with replacement. Bold entries: permutation $p < 0.05$. The caliper result is qualitatively unchanged by the restriction, and the negative-finding ratio strengthens slightly.


Appendix K. Caliper by Row Type: Coef+SE vs. Standalone z -Statistics

Table ?? decomposes the caliper test by whether each row has a verifiable coefficient-SE pair (so that $|z| = |\hat{\beta}/\hat{\sigma}|$ can be confirmed) versus a standalone z -statistic extracted directly from a t-statistic column. As discussed in Section ??, the two types differ in verifiability, not in the source paper: both appear in published regression tables.

Table A8. Caliper Test Stratified by Row Type

Corpus	Row type	N	Below	Above	Ratio	Naive p
Published	All rows	7,761	611	508	1.203	0.002
Published	Coef+SE only	5,510	320	394	0.812	0.006
Published	z -stat only	2,251	291	114	2.553	<0.001
NBER	All rows	8,351	501	495	1.012	0.849
NBER	Coef+SE only	6,788	416	438	0.950	0.452
NBER	z -stat only	1,563	85	57	1.491	0.019

Notes. *Coef+SE only*: rows where a coefficient and a positive standard error are both extracted; z_{abs} is computed as $|\hat{\beta}/\hat{\sigma}|$. *z -stat only*: rows where a $|z|$ -value appears in the table but no coefficient-SE pair is present, so the value cannot be independently verified. Caliper ratio = mass in (1.645, 1.96) / mass in (1.96, 2.275); a value > 1 indicates excess bunching below 1.96. Bold entries are significant at 5%. *Naive p* tests equality of the below- and above-counts via $\chi^2(1)$. Neither corpus shows excess bunching in coef+SE rows; both corpora show elevated bunching in standalone z -statistics, with published papers showing a substantially higher ratio (2.55 vs. 1.49).



figures/fig1_zdist_full.pdf

Figure 2. Distribution of absolute z -statistics, full sample. Dashed vertical lines mark $|z| = 1.645$ (10% threshold) and $|z| = 1.96$ (5% threshold). The smooth curve is the theoretical half-normal density. The caliper ratio (mass in $[1.645, 1.96]$ relative to $[1.96, 2.275]$) is **1.20** ($\chi^2 = 9.48$, $p = 0.002$). Sample: 7,761 test statistics from 121 published papers (108 with in-caliper-window statistics).




Figure 3. $|z|$ -distributions by journal. *Left:* JF (1,308 statistics, caliper ratio 0.88, $p = 0.423$). *Center:* RFS (2,581 statistics, caliper ratio 0.97, $p = 0.752$). *Right:* JFE (3,872 statistics, caliper ratio **1.49**, $p < \mathbf{0.001}$). Only JFE shows statistically significant caliper bunching.




Figure 4. $|z|$ -distributions by finding direction. *Left:* Positive-finding papers (2,633 statistics, caliper ratio 0.94, $p = 0.565$). *Center:* Negative-finding papers (3,491 statistics, caliper ratio **1.52**, $p < \mathbf{0.001}$). *Right:* Null/mixed-finding papers (1,637 statistics, caliper ratio 1.05, $p = 0.740$). Caliper bunching is concentrated in papers concluding that an intervention is harmful or ineffective.