

# Do Finance Journals Pick Sides? Evidence from Government Intervention Research

Emre Kuvvet\*

June 18, 2026

## ABSTRACT

Do top finance journals favor research that finds government regulations work, that finds them harmful, or that finds no clear effect? Using 134 articles in three leading finance journals (the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics*) alongside 157 NBER working papers on government-intervention research from 2000 to 2024, we find no evidence that journals favor any one type of conclusion. Papers concluding that a regulation helped, hurt, or had no measurable effect reach the published group at roughly the same rate, about 46 percent. Two other features do predict publication: research-design rigor and the prestige of the author’s home institution. Papers using credible causal methods—natural experiments, instrumental variables, difference-in-differences, regression discontinuity, or event studies with a causal interpretation—are about 18 to 20 percentage points more likely to appear among the published, and papers from top-ranked institutions are about 24 to 26 points more likely. Internal validation suggests the design-rigor finding’s magnitude is sensitive to coding choices, so we treat its direction (positive) as more reliable than its precise size. The two effects are nearly uncorrelated ( $r = 0.015$ ) and survive added controls, but our design cannot rule out that they partly reflect unobserved differences in paper quality rather than editorial selection. We also find subtle statistical clustering near the conventional significance threshold in published papers—but only in supplementary tables, not in their main regression results—which is consistent with selective reporting concentrated in robustness checks rather than headline findings.

*JEL classification:* G18, G28, C80, B40.

*Keywords:* Publication bias; Government intervention; Selective reporting; p-hacking

---

\*Department of Finance, Nova Southeastern University. E-mail: ekuvvet@nova.edu.

## 1. Introduction

Academic research on government intervention in financial markets occupies a privileged position in public policy. The Federal Reserve Board, the Securities and Exchange Commission, and the Basel Committee on Banking Supervision routinely incorporate peer-reviewed finance research into regulatory impact assessments, rule-makings, and policy speeches. The Dodd-Frank Wall Street Reform and Consumer Protection Act alone generated an estimated 400 new regulatory mandates, many of which were designed in explicit reference to academic evidence on the effects of capital requirements, disclosure rules, and consumer financial protection. The credibility of this evidence base is therefore a first-order policy question: if the academic record systematically misrepresents which interventions have measurable effects—and which do not—regulators face a distorted picture that can generate welfare-reducing policy choices.

Publication bias offers a mechanism by which such distortion arises without any deliberate misconduct. When researchers are reluctant to write up null results, when journals are more likely to accept papers with statistically significant findings, or when both effects operate simultaneously, the published literature understates the true prevalence of null or inconclusive evidence (Franco, Malhotra and Simonovits, 2014). In a domain such as government intervention research, where policy decisions depend on whether a regulation has a measurable causal effect, suppression of null results generates the false impression that most interventions work—or, more precisely, that most interventions have discernible effects in one direction or the other.

Directional publication bias could arise through two opposing channels. A *pro-intervention channel* arises if journals favor results that confirm a policy’s stated objective, both because positive results are more actionable for regulators and because they confirm the prior that motivated the policy. An *anti-intervention channel* arises if unintended-cost research is more surprising—and thus more publishable—than confirmatory work. Both differ from *significance bias* (non-null vs. null/inconclusive), which reflects a simpler

threshold between finding-something and finding-nothing.

Existing finance work on selection bias focuses on the cross-section of asset-pricing anomalies, documenting that factor-zoo proliferation reflects selection for high  $t$ -statistics (Harvey, Liu and Zhu, 2016; Hou, Xue and Zhang, 2020; Jensen, Kelly and Pedersen, 2023); a parallel economics literature (Brodeur, Lé, Sangnier and Zylberberg, 2016; Andrews and Kasy, 2019) infers bias from published  $p$ -value distributions. Neither strand addresses the policy-intervention literature or separates significance bias from directional bias. In our setup, these two biases map onto distinct tests: the across-corpus probit with separate Positive and Negative indicators (equation (1)) speaks to both—the  $\beta_1 = \beta_2$  test for directional bias, the joint  $\beta_1 = \beta_2 = 0$  test for significance bias—and the within-paper caliper test (Section 5.1) speaks to selective reporting near the significance threshold.

Our research design combines two complementary samples. The *published paper* sample consists of all in-scope articles in the *Journal of Finance*, the *Review of Financial Studies*, and the *Journal of Financial Economics* from 2000 to 2024. The *working paper* sample is drawn from all NBER working papers published from 2000 to 2024 (30,212 papers across every program), with in-scope papers identified through the same keyword filter and LLM classification applied to the journal sample. No program restriction is imposed; topic relevance alone determines inclusion. The key advantage of this design is that we observe findings from papers that circulated *before* potential journal publication, allowing a direct two-corpus comparison rather than inferring the pre-publication distribution from the shape of published  $p$ -values alone.

Because only 12% of published papers are matched to a confirmed NBER counterpart, the two corpora are largely non-overlapping cross-sections rather than the same papers tracked through the editorial filter. We therefore interpret all estimates as cross-sectional associations between paper characteristics and corpus membership, not as causal editorial-filter effects; the full identification discussion is in Section 7.4.

We define a paper as in-scope if its primary research question concerns the causal effect

of a specific government law, regulation, or public policy on financial markets, firms, or investors. Identifying these papers across nearly 40,000 documents required an automated approach: a keyword filter followed by a large language model (LLM) classifier that reads each abstract and assigns an in-scope judgment along with a directional code—positive (+1, the intervention was beneficial), negative (−1, harmful), or null/mixed (0).

Our central result is that finding direction does not predict published-versus-NBER corpus membership at the scale our design can detect. Within each direction category, the published share of our 291-paper sample is essentially the same—around 46 percent (published/(published + NBER), not a journal acceptance rate). The uniformity survives every regression specification, including controls for publication year, decade, team size, and research design quality.

What does predict published-side membership is research design. Papers using quasi-experimental identification strategies (RDD, IV, DiD, natural experiments, or event studies with a causal interpretation; defined in Section 5.2) are 18–20 percentage points more likely to sit on the published side. This is the only statistically robust predictor we find beyond finding direction. We interpret it cautiously: the QE indicator likely captures a bundle of correlated quality attributes (resources, framing, data access), not identification strategy alone.

Within published papers, we find a secondary pattern. Test statistics bunch modestly just below the conventional 5 percent significance threshold, but only in standalone statistics that appear primarily in supplementary and robustness tables. Primary regression tables—where both a coefficient and a standard error are reported and can be independently verified—show no bunching and are statistically indistinguishable from the corresponding NBER working paper tables. This distinction is informative: it suggests that any selective reporting occurs at the margin of revision, in the choice of which additional checks to highlight, rather than in the core empirical results that define the paper’s contribution.

The paper makes three contributions. First, it provides direct evidence that top finance

journals do not sort government-intervention research by the direction of its conclusions at scales our design can detect. Second, it documents that identification quality and institutional prestige—not finding direction—are the two largest cross-sectional predictors of corpus membership: conditional on observed controls, QE papers are about 18–20 pp more likely to appear in the published corpus and top-institution papers about 24–26 pp more likely, with  $r = 0.015$  between the two indicators and both premia surviving sequential controls. Third, it introduces a validated LLM-assisted pipeline for classifying and direction-coding large paper corpora, with code and prompts publicly available.

Section 2 reviews the literature; Section 3 describes the data and classification pipeline; Section 4 presents the empirical strategy; Sections 5 and 6 report main results and robustness checks; Section 7 discusses mechanisms and limitations; Section 8 concludes.

## 2. Related Literature

We use four terms with distinct meanings throughout, and each maps to a specific empirical test. *Significance bias* refers to the selective suppression of statistically null results, giving the published record an inflated share of non-null findings regardless of their substantive direction; in our analysis it corresponds to a joint test that the Positive and Negative coefficients in equation (1) are both zero (equivalently, the binary “Directional vs. null” coefficient in Column (1) of Table 8). *Directional bias* refers to editorial selection based on the substantive sign of a paper’s conclusion—whether a policy is found beneficial or harmful—conditional on the finding being non-null; it corresponds to the equality test  $\beta_1 = \beta_2$  in our primary Column (2) specification. *Selective reporting* refers to within-paper researcher behavior: the choice of which specifications, robustness checks, or standalone test statistics to highlight in the final manuscript based on proximity to a significance threshold; in our analysis it corresponds to the caliper test in Section 5.1. *Editorial selection* is our umbrella term for any mechanism by which the submission-to-publication pipeline shapes the composition of the published literature. The probit tests speak to significance and directional bias; the caliper

test speaks to within-paper selective reporting; the significance-bias literature provides the methodological benchmark for both.

This paper relates to three strands of research. The first is the literature on significance bias in economics and finance. Brodeur, Lé, Sangnier and Zylberberg (2016) document excess mass just above the  $|z| = 1.96$  significance cutoff (equivalently, on the  $p < 0.05$  side of the threshold) in the top economics journals—the classical p-hacking pattern. Our caliper analysis uses the same windows but the inverse ratio convention; we describe both directions explicitly in Section 5.1. Franco, Malhotra and Simonovits (2014) show, using NSF-funded pre-registered experiments, that null results are published at roughly one-third the rate of significant ones. In finance, Harvey, Liu and Zhu (2016) show that given the scale of data mining in the factor literature, conventional significance thresholds are far too low and most claimed factors are likely false discoveries under multiple testing; Hou, Xue and Zhang (2020) confirm this empirically, finding that between 65% and 82% of published anomalies (depending on the test and sample) fail to clear conventional significance thresholds under uniform replication standards, pointing to upward-biased published effect sizes. These papers test for *significance* bias—whether null results are suppressed relative to significant ones. Our paper asks a different question: whether published findings are sorted by the *direction* of the policy conclusion, which is only meaningful in a domain where findings have a clear normative interpretation. We also contribute a direct pre-publication baseline by using NBER working paper texts, rather than inferring the pre-publication distribution from the shape of published p-values alone.

The second strand is the literature on government intervention in financial markets. A large body of empirical work studies the effects of securities law, banking regulation, and macroprudential policy (La Porta, Lopez-de Silanes, Shleifer and Vishny, 1998; Berger and Bouwman, 2017; Agarwal, Lucca, Seru and Trebbi, 2014), producing heterogeneous findings across interventions and time periods. Whether the published record faithfully represents the underlying distribution of evidence on these questions has not previously been examined.

Our paper fills this gap.

This literature is particularly susceptible to directional publication bias for three reasons. *Normative salience*: unlike most empirical finance research, the direction of findings here carries clear policy implications, and editors, referees, and authors arrive with directional priors that can generate asymmetric selection (Andrews and Kasy, 2019). *Ideological sorting*: the academic finance community has historically leaned skeptical of government intervention; under that prior, Andrews and Kasy (2019) predicts that positive (+1, pro-intervention) findings face higher scrutiny than negative (−1, anti-intervention) findings, generating selection toward over-publication of negative results. *Policy-venue citation pressure*: papers in the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* are routinely cited in regulatory proceedings, so directional findings are more actionable than null results. These mechanisms jointly predict directional sorting while leaving the sign of any bias theoretically ambiguous; our paper provides the first direct test using a pre-publication benchmark.

The third strand is the growing use of large language models for text classification in economics (Ash and Hansen, 2023; Korinek, 2023). We contribute a validated two-pass pipeline—keyword filter followed by LLM classification—for identifying and direction-coding government-intervention finance papers at scale, and we release the coding rubric and prompts publicly.

### 3. Data and Sample Construction

#### 3.1 Published Papers

The published-paper sample consists of all articles appearing in the *Journal of Finance*, the *Review of Financial Studies*, and the *Journal of Financial Economics* from 2000 to 2024. These three journals constitute the traditional top tier of academic finance and account for the majority of highly cited finance research on government intervention. We collect article

metadata—title, authors, year, volume, issue, DOI, and abstract—for all articles in each journal over this period.

Journal articles are identified and retrieved as follows. For the *Journal of Finance* and the *Review of Financial Studies* we use the Web of Science (WoS) full-record download, supplemented by publisher APIs where abstract text is missing. For the *Journal of Financial Economics*, abstracts were obtained via the Elsevier ScienceDirect Article Retrieval API. Abstract text is available for 100% of research articles in all three journals after manual abstract enrichment. The *Journal of Finance* denominator is restricted to research articles only; the raw WoS record includes editorial matter, meeting announcements, and AFA administrative items that carry no abstracts, and our pipeline excludes these before computing coverage.

### 3.2 NBER Working Papers

The working-paper sample is drawn from the NBER Working Paper Series via the public API. We collect *all* working papers published from 2000 to 2024 across every NBER program (30,212 papers in total) and identify in-scope government-intervention papers through the keyword filter and LLM classification described in Section 3.3. No program restriction is applied at the collection stage; the topic filter alone determines inclusion. Government intervention finance research is concentrated in five programs— Corporate Finance (CF), Asset Pricing (AP), Monetary Economics (ME), Economic Fluctuations and Growth (EFG), and Public Economics (PE)—but papers assigned to neighboring programs (e.g., Economic History or International Finance and Macroeconomics) are automatically captured if they pass the scope filter.

We use the NBER working paper sample as a reference distribution for comparison with the published sample. A causal editorial-selection interpretation would require two distinct conditions: (i) NBER selection into working-paper circulation is itself direction-neutral, so that the NBER distribution is not pre-filtered on finding direction; and (ii) conditional

on direction, the NBER sample approximates the relevant top-journal submission pool on other factors jointly related to direction and publication (e.g., research-design quality, topic, author network, data access, vintage). Both conditions are strong. NBER affiliation is selective on researcher quality and institutional prestige, and many finance papers circulate through SSRN rather than NBER; the NBER distribution may therefore differ from the full pre-publication distribution for reasons unrelated to editorial selection. We also note that “pre-publication” here means circulated as an NBER working paper at some point—NBER posting dates establish working-paper circulation but do not establish that this circulation preceded journal submission, since papers may be posted to NBER before, alongside, or after a submission decision. We therefore read all results in this paper as cross-sectional associations between corpus membership and paper characteristics, not as causal estimates of editorial filtering; we return to this identification challenge formally in Section 4.1.

### 3.3 In-Scope Classification Pipeline

Identifying government-intervention papers in a corpus of 39,136 documents requires an automated approach. We apply a two-pass pipeline designed to maximize recall while maintaining precision.

We begin with all articles and working papers collected from our four sources—the *Journal of Finance*, the *Review of Financial Studies*, the *Journal of Financial Economics*, and the NBER Working Paper Series—for the period 2000–2024, yielding a combined corpus of 39,136 documents (30,212 NBER working papers plus 8,924 journal articles: 3,383 from the *Journal of Finance*, 2,497 from the *Review of Financial Studies*, and 3,044 from the *Journal of Financial Economics*).

In the first pass, a paper passes the keyword filter if the combined text of its title and abstract contains at least one term from a pre-specified taxonomy of government intervention keywords spanning seven categories: financial regulation (e.g., “Dodd-Frank,” “capital requirements,” “Basel”), securities law (e.g., “SEC,” “Sarbanes-Oxley,” “insider trading

law”), banking policy (e.g., “FDIC,” “deposit insurance,” “TARP”), monetary and fiscal policy (e.g., “quantitative easing,” “fiscal stimulus”), market microstructure rules (e.g., “uptick rule,” “tick size pilot”), corporate governance law (e.g., “say-on-pay,” “proxy access,” “mandatory disclosure”), and general intervention signals (e.g., “government intervention,” “public policy,” “law and finance”). Ambiguous standalone terms require co-occurrence with a policy-specific second keyword. The full taxonomy is reported in Appendix A.

Searching the combined title-and-abstract text rather than the abstract alone improves recall for three reasons. First, authors frequently name the specific law or regulation in the title but refer to it only as “the reform” or “the rule” throughout the abstract, so the keyword would appear only in the title. Second, common abbreviations (e.g., “Dodd-Frank” in the title) and full statutory names (e.g., “Wall Street Reform and Consumer Protection Act” in the abstract) are not always used interchangeably; searching both fields catches either form. Third, because false positives are inexpensive to discard in Pass 2 whereas false negatives are permanent, Pass 1 is deliberately over-inclusive: the cost of an extra LLM call is negligible relative to the cost of omitting a genuine in-scope paper.

The keyword filter reduces the corpus from 39,136 documents to 1,849 candidates (1,313 NBER, 182 *Journal of Finance*, 153 *Review of Financial Studies*, and 201 *Journal of Financial Economics*).

In the second pass, each keyword-flagged candidate is passed to Claude with the following instruction (full prompt in Appendix B):

*“A finance research paper is ‘in scope’ if and only if its PRIMARY research question is the CAUSAL EFFECT of a specific government law, regulation, or public policy on financial markets, firms, or investors. The following are excluded: (1) purely theoretical papers with no empirical test of a policy; (2) papers that only describe a regulation without estimating its effect; (3) papers whose main topic is private firm or market behavior, with policy as incidental context; (4) papers studying central bank communication without a specific policy*

*instrument; and (5) papers whose primary policy variable is a rule imposed by a stock exchange rather than a government or regulatory mandate.”*

For each paper, the model returns an in-scope decision, a confidence level (high, medium, or low), and a one-sentence justification. At this stage the model classifies 343 of the 1,849 keyword-flagged candidates as in scope (NBER: 187, *Journal of Finance*: 44, *Review of Financial Studies*: 50, *Journal of Financial Economics*: 62) and the remaining 1,506 as out of scope.

Following the automated pass, all 343 LLM in-scope classifications are reviewed manually by the author against the paper’s abstract, with the PDF introduction consulted whenever the abstract is ambiguous about whether the policy is the primary research question. The review proceeds in two directions.

(i) *Re-examining LLM in-scope decisions.* The most important judgment call concerns papers that use a policy event only as an identification device to study a downstream outcome that the regulation did not target. We operationalize the boundary at the level of the *outcome variable*:

- A paper is *in scope* if its primary outcome is something the regulation explicitly aimed to affect (its stated regulatory objective)—e.g., for Sarbanes-Oxley, audit quality, internal-control disclosure, financial-reporting accuracy, restatement frequency, or fraud detection. The paper’s contribution is then a direct test of “did SOX achieve its disclosure objective?”
- A paper is *out of scope* if its primary outcome is downstream of the regulated mechanism but is *not* a stated regulatory objective—e.g., for Sarbanes-Oxley, firm investment, M&A activity, cost of capital, or innovation. The paper’s contribution is then “what is the effect of disclosure (or governance, or liability) on outcome  $Y$ , using SOX as an instrument?” These are mechanism papers whose policy variable is incidental.

The criterion turns on whether the regulator stated the outcome as a target, not on whether

the outcome plausibly responds to the regulation. Borderline cases—where authors describe both a policy-effect and a mechanism interpretation—are resolved by the position of the outcome in the abstract: if the abstract’s headline result names a stated regulatory objective, the paper is in scope; if the abstract frames the regulation as the source of variation for a non-objective outcome, the paper is out of scope. A second class of overrides drops papers that test a policy whose empirical content turns out to be a private market response (e.g., voluntary cross-listing decisions after a directive). This re-examination produces 52 net exclusions relative to the LLM (NBER: 30, *Journal of Finance*: 18, *Review of Financial Studies*: 2, *Journal of Financial Economics*: 2).

(ii) *Reconsidering LLM out-of-scope decisions in the NBER corpus.* Because false-negative LLM exclusions are difficult to detect without a second pass, the author re-codes 803 NBER papers that the LLM flagged as low-confidence out-of-scope. Of these, 126 are returned to the in-scope pool; 106 survive the same Pass (i) criteria and enter the analysis sample. The published-side LLM decisions are not subjected to a symmetric refilter because the journal denominators (8,924 articles) are small enough that the keyword filter retains all plausible candidates at the first stage.

Borderline cases are resolved by the following rule: if the policy variable is the right-hand-side treatment in the paper’s headline regression, the paper is in scope; if the policy is used as an instrument or a covariate to study a non-policy outcome, it is out of scope. The full per-paper override log—LLM decision, manual decision, and reason when they differ—is available in the replication package.

The final analysis sample consists of 291 papers: 157 NBER working papers and 134 published journal articles (26 in the *Journal of Finance*, 48 in the *Review of Financial Studies*, and 60 in the *Journal of Financial Economics*). Table 1 reports the number of papers surviving each stage of the selection process.

**[INSERT TABLE 1 HERE]**

### 3.4 Direction Coding

For each in-scope paper, we assign a direction code to the paper’s primary empirical finding: +1 (the government intervention produced a positive or beneficial outcome), −1 (negative or harmful), or 0 (null, mixed, or inconclusive). Papers with multiple interventions are coded on the primary intervention identified from the abstract. Both the in-scope classifier and the direction coder use Claude Haiku (claude-haiku-4-5-20251001; Anthropic).

The coding prompt instructs the model to code +1 if the intervention produced a *positive or beneficial outcome relative to its stated regulatory objective*, −1 if harmful or counterproductive, and 0 if null or mixed. Normative framing by the authors is explicitly excluded. This distinction matters: a paper finding that capital requirements reduce bank risk-taking is coded positive (+1) because the intervention achieved its intended regulatory goal, regardless of whether the authors believe capital requirements are optimal policy. The full coding prompt is in Appendix B.

We validate the direction codes in two ways. First, we re-ran the coding on a stratified 60-paper sample (5 papers per source  $\times$  direction cell) using the same model (claude-haiku-4-5-20251001) in an independent second call. The two runs agree on 71.7% of papers ( $\kappa_w = 0.674$ , “substantial agreement”). Agreement is high for directional papers—100% for negative findings and 90% for positive ones—but only 25% for null/mixed papers; among the 20 originally null-coded papers, 15 were re-coded as directional and only 5 stayed null. These flips are not evenly distributed across corpora: 13 of the 15 originally null-coded *published* articles flipped to directional, against 2 of 5 NBER papers. Because the sign of the resulting bias in our direction premium depends on which corpus the miscoded papers come from—NBER directional-as-null inflates the premium, while published directional-as-null deflates it—the replication exercise establishes that null is the least stable category but does *not* establish a clean sign for the residual measurement error. The full confusion matrix is in Appendix C.

Second, we re-ran coding on 58 papers under three alternative prompt wordings. Agree-

ment with the original codes ranges from 83% to 91% ( $\kappa_w = 0.742\text{--}0.904$ ), all above the model-choice  $\kappa_w$ , confirming the results are not sensitive to how the question is phrased (Table 2).

[INSERT TABLE 2 HERE]

Human-coded validation of a 40-paper subsample is included in the replication package and remains a key additional check; the model-choice  $\kappa_w = 0.674$  is the primary reliability evidence reported here.

### 3.5 Linking NBER Papers to Published Papers

We matched each of the 134 published journal articles to the NBER working paper archive using title similarity, author overlap, and external databases including OpenAlex, CrossRef, and Semantic Scholar. The pipeline identified 16 matched pairs—about 12% of the published sample. The remaining 118 published papers have no NBER counterpart, reflecting that most corporate finance papers circulate through SSRN rather than NBER.

**Accounting rule for the 16 matched pairs.** We adopt a single auditable rule for the 16 matched-pair NBER drafts and apply it consistently across every table and regression in the paper:

- (i) The 134 *Journal of Finance*/*Review of Financial Studies*/*Journal of Financial Economics* articles enter the analysis sample with Published = 1 (the journal-published version is the one counted).
- (ii) The NBER comparison group of 157 papers contains *only* NBER working papers whose published version did not subsequently appear in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics* during 2000–2024. The 16 NBER drafts of papers that did reach a top-three journal are excluded from the 157 by construction.

(iii) No paper-author-content unit is therefore double-counted. The 291-paper sample partitions cleanly into  $134 + 157$  disjoint units, with the 16 matched-pair NBER drafts appearing in neither cell—they are represented only by their journal-published incarnation in the 134-paper Published group.

This rule has three immediate implications. First, the dependent variable in the probit ( $\text{Published} \in \{0, 1\}$ ) is well-defined and mutually exclusive across the 291 observations (matching the Table 3 note). Second, dropping the 16 matched-pair published papers as a robustness check (Section 6.3) removes only those 16 observations and leaves  $N = 275$ , because their NBER drafts are already excluded; the corresponding NBER counterparts are not separately deletable. Third, the caliper test (Section 5.1) operates on the  $|z|$ -statistics extracted from the journal-published version of each matched pair when computing the Published-row caliper, and on the NBER draft for the unmatched NBER-only papers; no test statistic is counted twice. The same rule governs every subsequent table in the paper.

In the regression analysis, the 134 published papers are coded  $\text{Published} = 1$  and the 157 NBER-only working papers are coded  $\text{Published} = 0$ . The direction distributions of the two groups are nearly identical—about 80% of papers in each group have a directional finding—suggesting the two samples are broadly comparable in composition before any regression adjustment.

### 3.6 Summary Statistics

Table 3 presents summary statistics for the 291-paper analysis sample. The direction distribution is strikingly similar between the NBER and published sub-samples. Among NBER working papers, 20.4% are coded null or mixed, compared with 20.9% among published papers—a difference of less than 1 percentage point. Positive-finding papers constitute 35.1% and negative-finding papers 44.0% of the published sample, versus 35.7% and 43.9% of the NBER sample. The two samples look nearly the same in their direction distributions.

[INSERT TABLE 3 HERE]

## 4. Empirical Strategy

### 4.1 Research Design and Identification

Our empirical strategy compares papers circulated as NBER working papers against papers published in the *Journal of Finance*, the *Review of Financial Studies*, or the *Journal of Financial Economics*, asking whether the direction distribution of findings differs between the two groups. This is not a within-paper design: only 16 of 134 published papers are matched to an NBER working paper, so the two samples are largely different papers selected through different processes.

A causal editorial-selection interpretation would require both that NBER working papers approximate the direction distribution of papers actually submitted to top finance journals, and that the NBER and published samples be comparable on other factors related jointly to finding direction and publication. Similar working-paper-based approaches are taken by Franco, Malhotra and Simonovits (2014), who use NSF pre-registrations as a pre-publication baseline, and DellaVigna and Linos (2022), who use institutional RCT records. Our design is weaker than either: NBER affiliation is selective on researcher quality and institutional prestige; most finance papers circulate through SSRN rather than NBER; and NBER posting dates do not pin down whether a paper circulated before or after journal submission. The direction of any resulting bias is unclear—NBER selection could either overstate or understate the true editorial gap—so we interpret all results in this paper as cross-sectional associations between corpus membership and paper characteristics, not causal editorial-filter effects. The full identification discussion is in Section 7.4.

### 4.2 Primary Specification

Our primary test estimates a probit model at the paper level:

$$\Pr(\text{Published}_i = 1) = \Phi(\alpha + \beta_1 \text{Positive}_i + \beta_2 \text{Negative}_i + \gamma' \mathbf{X}_i), \quad (1)$$

where  $\Phi(\cdot)$  is the standard normal CDF,  $\text{Published}_i$  equals one if in-scope paper  $i$  appears as an article in the *Journal of Finance*, the *Review of Financial Studies*, or the *Journal of Financial Economics*,  $\text{Positive}_i$  ( $\text{Negative}_i$ ) is an indicator equal to one if the LLM-assigned direction code is +1 (−1), and the omitted category is papers with direction code 0 (null or mixed). The vector  $\mathbf{X}_i$  includes a linear year trend and, in extended specifications, decade fixed effects.

Our primary coefficients of interest are  $\beta_1$  and  $\beta_2$ . A finding of  $\beta_1 > 0$  ( $\beta_2 > 0$ ) indicates that positive-finding (negative-finding) papers are more likely to be in the published sample than null-finding papers, conditional on controls. Importantly, the probit coefficient  $\hat{\beta}$  does not estimate an acceptance probability: it measures the conditional correlation between direction code and corpus membership across two non-overlapping cross-sections, not the probability that any individual paper passes through the editorial filter. A test of  $H_0: \beta_1 = \beta_2$  addresses the directional-bias hypothesis: rejection would indicate that the publication premium depends on the sign of the finding, not merely on statistical significance.

Because the probit model is nonlinear, we report marginal effects. For the bivariate (Column 2) specification, which has no covariates beyond the direction indicators, we report the null-paper baseline difference

$$\Delta_d = \Phi(\hat{\alpha} + \hat{\beta}_d) - \Phi(\hat{\alpha}), \quad d \in \{+, -\}.$$

For controlled specifications (Columns 3–7), the probit index also includes  $\hat{\gamma}'\mathbf{X}_i$ , so the corresponding expression at a covariate baseline  $\mathbf{X}^{\text{base}}$  is

$$\Delta_d(\mathbf{X}^{\text{base}}) = \Phi(\hat{\alpha} + \hat{\beta}_d + \hat{\gamma}'\mathbf{X}^{\text{base}}) - \Phi(\hat{\alpha} + \hat{\gamma}'\mathbf{X}^{\text{base}}).$$

In the controlled specifications we report two summaries. (a) Marginal effects at the centered baseline, defined by  $\text{year}_c = 0$  (the sample median year),  $\log(1 + N_{\text{authors}}) = 0$  (single-author papers), Quasi-Exp. ID = 0, and Top Institution = 0; this is the baseline reported in

the bivariate Column (2) expression above, where the centered controls all enter at zero.

(b) Average partial effects (APE), computed by averaging over the in-sample empirical distribution of  $\mathbf{X}_i$ :

$$\text{APE}_d = \frac{1}{N} \sum_{i=1}^N \left[ \Phi(\hat{\alpha} + \hat{\beta}_d + \hat{\gamma}'\mathbf{X}_i) - \Phi(\hat{\alpha} + \hat{\gamma}'\mathbf{X}_i) \right].$$

The two summaries are close in our sample because the centered controls mostly take values near zero. We also report raw corpus-membership rates directly from the data as a transparent, model-free summary of the unconditional patterns.

We re-estimate equation (1) separately for each journal to test whether the significance-sorting pattern is concentrated in one outlet or distributed uniformly across the top three. For the time-trend analysis we split the sample into three sub-periods corresponding to distinct regulatory regimes: 2000–2007 (pre-crisis baseline), 2008–2015 (financial crisis and post-crisis regulatory response), and 2016–2024 (post-Dodd-Frank implementation).

## 5. Results

### 5.1 Within-Paper Test Statistic Distributions

We test for unusual clustering of test statistics near the conventional 5% significance threshold using the caliper test of Brodeur, Lé, Sangnier and Zylberberg (2016), run separately on two row types: regression rows where both a coefficient and a standard error are reported (verifiable by their ratio), and standalone  $t$ - or  $z$ -statistics appearing mainly in robustness and supplementary tables. Automated text parsing of all 134 published and 157 NBER PDFs yields 7,761 test statistics from 114 published papers (20 contribute none: 13 with no extractable rows, 7 with values failing numeric validation) and 8,351 test statistics from 117 NBER papers (40 contribute none for similar reasons). The pre-publication sample lets us compare the two distributions directly rather than inferring bias from the published distribution alone.

*Standalone-statistic comparability.* The extraction treats every standalone number as a two-sided  $|z|$ -test against  $H_0 = 0$  evaluated against the 1.96 cutoff (the dominant convention in published-finance tables), but is not certified row-by-row. A small share of standalone rows are flagged as  $F$ -stats,  $\chi^2$ -stats, one-sided tests, or tests against non-zero nulls—these would not be comparable to the 1.96 cutoff but we do not separately filter them out. A row-level provenance audit would benefit a future revision; the caliper conclusions should be read with this caveat.

One concern is that published journal PDFs might be harder to parse than NBER draft PDFs, artificially inflating the caliper gap. We address this concern with several checks but acknowledge that none of them rules it out completely; each speaks to a different aspect of differential missingness.

First, for rows where both a coefficient and a standard error are extracted, the reported test statistic matches their ratio to within 0.1 in 98.3% of published rows and 99.7% of NBER rows—no meaningful gap between corpora. By construction this check applies only to Coef+SE rows; the standalone  $z$ -statistics that drive the headline caliper finding (ratio 2.54 published vs. 1.58 NBER) cannot be cross-validated. Within the verifiable Coef+SE rows, neither corpus shows excess bunching ( $R = 0.82$  published, 0.95 NBER); the gap is concentrated precisely in the unverifiable row type. Second, per-paper extraction rates (the fraction of in-scope papers yielding at least one usable  $|z|$ -statistic) differ by both corpus and direction: 91.5%/83.0%/75.0% of published negative-/positive-/null-finding papers contribute statistics, vs. 82.6%/75.0%/56.2% for NBER. Null-finding papers are over-represented in the missing-extraction group in both corpora, consistent with their reporting fewer significance-laden tables. The negative-finding result—the one driving the headline—has the highest extraction rates and the most balanced corpus comparison, and is therefore the least likely of the three to reflect selective extraction. Third, papers with no usable statistics are spread roughly evenly across journals and years within journal (within-cell counts are small). Fourth, restricting to papers with at least 20 extracted statistics

moves the main caliper ratios by less than  $\pm 0.05$ . Two channels remain outside our checks: within-paper differential missingness across row types, and table-format heterogeneity that systematically suppresses one side of the threshold in published papers. We flag these as residual concerns rather than claim the parsing channel is fully closed.

The results differ sharply by statistic type. Primary Coef+SE rows show no threshold bunching in either corpus ( $R = 0.82$  published,  $0.95$  NBER). Bunching appears only in standalone test statistics, which come mainly from supplementary tables, and published papers show a higher bunching ratio ( $R = 2.54$ ) than NBER papers ( $R = 1.58$ ). We read this as *selective reporting*: authors are more likely to highlight a supplementary result when it clears a significance threshold, not because they are manipulating their primary estimates. Primary tables report all estimates regardless of significance and show no bunching. The same pattern—though weaker—appearing in NBER working papers suggests it reflects an authorial reporting choice rather than a publication filter imposed by journals. We cannot rule out that some standalone statistics come from searches over many specifications, but the absence of bunching in primary regression rows is evidence against widespread specification search (see Table 5 below and Appendix E for full counts).

We use two complementary tests. The *caliper test* (Brodeur, Lé, Sangnier and Zylberberg, 2016) compares the number of  $|z|$ -statistics on either side of the conventional significance cutoff  $|z| = 1.96$  (equivalently, the two-sided  $p = 0.05$  threshold). Throughout the paper we report the ratio

$$R = \frac{\#\{|z| \in [1.645, 1.96)\}}{\#\{|z| \in [1.96, 2.275]\}},$$

the count of statistics in the narrow window *just below* the  $|z| = 1.96$  cutoff (where  $0.05 < p \leq 0.10$ ) divided by the count *just above* it (where  $0.023 \leq p \leq 0.05$ ). Under the null of no threshold sorting, the two zones should be equally populated and  $R = 1$ . A value  $R > 1$  indicates excess mass on the just-below-1.96 side (i.e., test statistics that fall just short of conventional significance); a value  $R < 1$  indicates excess mass on the just-above-1.96 side (i.e., results that just clear the bar). We adopt this  $|z|$ -space convention—the inverse of the

$p$ -space convention used by some authors—throughout the paper. The *density discontinuity test* asks whether there is a sharp jump in the frequency of test statistics right at the  $|z| = 1.96$  threshold, which is also consistent with selective reporting.

Table 4 reports the caliper ratios and significance tests for the full sample, broken out by journal and by finding direction. Figures 1 and 3 plot the full-sample and by-journal test-statistic distributions; the direction-specific distributions are summarized in Table 6.

[INSERT TABLE 4 HERE]

[INSERT FIGURE 1 HERE]

Under the naive chi-square, four sub-sample patterns stand out. (i) Across all published papers, statistics in the just-insignificant window outnumber those in the just-significant window by about 20% ( $R = 1.20$ ,  $p = 0.002$  naive). (ii) Among the three journals the asymmetry is most visible in the *Journal of Financial Economics* ( $R = 1.49$ ,  $p < 0.001$  naive); *Journal of Finance* and *Review of Financial Studies* are essentially flat ( $R = 0.88$  and  $0.97$ , both  $p > 0.40$ ). (iii) Among direction categories the asymmetry is largest in *negative-finding papers* ( $R = 1.52$ ,  $p < 0.001$  naive), compared with  $R = 0.94$  positive and  $R = 1.05$  null/mixed. (iv) The discontinuity test conflicts with the caliper for *Journal of Finance* (caliper  $p = 0.423$  but discontinuity  $p < 0.001$ ); given only 26 JF papers, we read this as statistical noise and give priority to the caliper throughout. The corresponding paper-clustered bootstrap  $p$ -values for the negative-finding subgroup ( $p = 0.135$ ) and the *Journal of Financial Economics* are not significant either; we therefore read the per-subgroup framing as describing where the naive-test point estimates are largest, not where statistical significance is robustly established.

This direction-level result is novel: positive-finding papers are not over-represented in the published corpus (next section), yet within published negative-finding papers test statistics show unusually strong near-threshold clustering—a pattern absent from positive-finding papers. One interpretation is that authors of anti-intervention papers face greater scrutiny

and search more aggressively for significant specifications; we treat this as suggestive rather than definitive, since cross-direction differences in research design could also contribute.

A simple decomposition explains how the full-sample ratio of 1.20 arises. Published papers devote 28.3% of their table rows to standalone  $z$ -statistics (those reported without a paired coefficient and standard error), compared with 18.1% in NBER working papers; and standalone statistics cluster far more sharply in published papers (ratio 2.54) than in NBER papers (1.58). The full-sample difference therefore reflects both a within-type effect—more bunching per standalone statistic in published papers—and a composition effect—a higher standalone share. In the verifiable coefficient rows, neither corpus shows excess clustering (ratios of 0.82 and 0.95, respectively). Table 5 presents these numbers together; Appendix E provides the full counts.

**[INSERT TABLE 5 HERE]**

Most prior work on significance bias can only study published papers and must infer what the research pipeline looked like before editorial decisions were made. Our NBER sample lets us compare directly.

Table 6 shows the results side by side.

**[INSERT TABLE 6 HERE]**

The point estimates suggest a gap: published  $R = 1.20$  vs. NBER  $R = 1.01$ . The two groups have nearly identical overall significance shares (57% published, 56% NBER), so the ratio difference is not driven by noise levels. A naive  $2 \times 2$  chi-square treating each statistic as independent gives  $\chi^2 = 3.74$  ( $p = 0.053$ ) for the full sample and  $\chi^2 = 10.66$  ( $p = 0.001$ ) for negative-finding papers; the paper-clustered bootstrap in Table 6 does not clear 0.05 for either ( $p = 0.150$  full,  $p = 0.135$  negative). We read the published–NBER gap as suggestive rather than definitively significant; the point estimates are larger for published papers and especially so for negative-finding papers, but the cross-corpus difference is not

robust to paper-level clustering on our preferred test. Within these caveats, the size and direction of the gap are concentrated in negative-finding papers: published negative-finding papers cluster on the just-below-1.96 side at more than 50% above the just-above-1.96 count; NBER negative-finding papers show essentially no such asymmetry ( $R = 1.00$ ). Positive- and null-finding papers show no gap in either direction (Figure 2).

**[INSERT FIGURE 2 HERE]**

Two interpretations are plausible. Published negative-finding papers may differ from NBER negative-finding papers on unobserved structural dimensions (more robustness checks, different table formats, different vintage) that mechanically produce more statistics near the threshold. This is a comparison between two direction-matched cross-sections, not a draft-to-publication comparison; only 16 of 134 published papers are matched to a specific NBER draft. Alternatively, authors arguing a policy is harmful may face greater reviewer scrutiny and try more specifications until results clear the threshold. We cannot distinguish these with the data available.

Because within-paper statistics are not independent, Table 6 reports three test versions: naive chi-square (independent statistics), permutation (within-paper), and paper-clustered bootstrap (our preferred test). The permutation test is significant for the full sample and negative-finding papers; the bootstrap does not reach conventional significance in either case. NBER papers show no clustering under any test.

**[INSERT FIGURE 3 HERE]**

Because we run 18 tests across subsamples, journals, and directions, some significant results could appear by chance. We designated the two full-sample caliper tests (published and NBER) as primary ex ante. Under naive chi-square, the published full-sample  $p = 0.002$  clears the Bonferroni threshold  $0.05/18 \approx 0.003$ ; under the paper-clustered bootstrap (preferred),  $p = 0.150$  does not clear 0.05 even before correction. NBER is null under both

( $p = 0.849$  naive,  $p = 0.432$  bootstrap). Among exploratory subgroups, the *Journal of Financial Economics* and negative-finding results clear the naive Bonferroni cutoff ( $p < 0.001$  each) but the negative-finding bootstrap  $p = 0.135$  does not. We treat the naive-test results as the less-conservative reading and the bootstrap as our preferred reading.

We include all regression coefficients in the test rather than just the headline estimate, since researchers can adjust any reported result toward a significance threshold. As a check, dropping standard control-variable rows ( $\approx 9\%$  of observations) leaves the main results unchanged (Appendix D).

## 5.2 Supporting Context: Cross-Sectional Direction Sorting

*Terminology convention.* For brevity we use “publication rate” as a compact label for the within-cell *corpus-membership share*—published / (published + NBER-only) within the indicated cell of our 291-paper sample. These shares are mechanically determined by sample composition (134 published vs. 157 NBER-only) and do not estimate journal acceptance probabilities; the latter would require observing the submission pool and editorial decisions, neither of which is in our data. Readers should interpret every occurrence as a corpus-membership share rather than as a transition probability.

We next ask whether finding direction predicts membership in the published corpus relative to the NBER-only comparison group, which provides a cross-sectional test of whether the published-versus-pre-publication mix of findings differs by direction. Table 7 reports the share of papers in each direction category that appear in the published corpus. About 46% of papers in our 291-paper sample fall on the published side regardless of finding direction: 45.6% of papers finding regulation beneficial (47 of 103), 46.1% of papers finding it harmful (59 of 128), and 46.7% of papers with inconclusive findings (28 of 60). The three corpus-membership rates are virtually identical and statistically indistinguishable ( $p = 0.992$ ). This rules out a direction effect on *corpus membership* in our two-corpus comparison; it does not, by itself, rule out direction effects at the journal acceptance stage, which would require

observing the submission pool and editorial decisions (neither is in our data). We interpret the result as evidence that finding direction does not predict published-vs. NBER-only membership, not as a direct test of journal acceptance.

**[INSERT TABLE 7 HERE]**

Table 8 reports probit regression results across seven specifications, each predicting whether a paper ends up published in one of the three journals. Column (2) is our primary specification, matching equation (1) in Section 4.1: it separates positive (+1) and negative (−1) findings against null/mixed papers (0), and the directional-bias hypothesis is the test of  $H_0: \beta_1 = \beta_2$ . Column (1) reports a more parsimonious binary specification—directional vs. null/mixed—as a complementary check on whether *any* directional sign carries a publication premium; we report it alongside Column (2) but treat it as supporting rather than primary, because the binary version answers a different question (significance bias) than the directional-bias test that motivates the paper. Columns (3)–(7) extend the primary Column (2) specification by adding controls for publication year, decade fixed effects, team size, identification method, and author institution. We classify a paper as *quasi-experimental* (“QE”) if its primary identification strategy is regression discontinuity (RDD), instrumental variables (IV), difference-in-differences (DiD), a natural experiment with an exogenous shock, or an event study that the AI judges to carry a causal interpretation (i.e., the event is used to identify a causal effect, not merely to characterize an event-window return pattern). Descriptive event-window analyses without a causal claim are coded as non-QE (`ols_reduced`), matching the rubric used by the classifier (`scripts/09_code_id_strategy.py`). We use this five-method definition consistently throughout the paper.

**[INSERT TABLE 8 HERE]**

Our primary Column (2) shows no relationship between either signed direction indicator and publication: both coefficients are near zero and insignificant (positive:  $p = 0.898$ ;

negative:  $p = 0.942$ ). A joint test of whether both are zero gives  $p = 0.992$ . The directional-bias test of whether the two coefficients are equal to each other gives  $\chi^2 = 0.005$ ,  $p = 0.944$ , providing no evidence that journals favor one direction over the other.

The supporting binary specification in Column (1) reaches the same conclusion in a more parsimonious form: the pooled directional coefficient is  $-0.020$  ( $p = 0.914$ ), indistinguishable from zero, so papers with any directional finding are no more or less likely than null-finding papers to sit on the published side of our corpus. The design rules out large sorting effects—roughly 18 percentage points or more—but cannot rule out smaller ones; we read both columns as evidence against strong directional sorting, not as proof of complete editorial neutrality.

Adding controls does not change the picture. The Column (3) year trend is positive and significant: more recently dated in-scope papers disproportionately sit on the published side. Right-censoring of recent NBER papers would push the year coefficient in the *opposite* direction (inflating the NBER share at higher years), so censoring is not the explanation. Compositional channels (post-2008 regulatory topics where journal publication outpaces NBER circulation; better recent-year indexing on the published side) are more plausible, but our two-corpus design cannot separate them. Restricting to papers circulated before 2021 leaves the direction coefficients unchanged ( $p > 0.50$ ), so the positive year trend does not interact with the direction null. Decade fixed effects (Column (4)) and team size controls (Column (5)) likewise leave the direction results unaffected.

Column (6) adds the QE indicator (RDD, IV, DiD, natural experiment, or event study with causal interpretation). QE papers are significantly more likely to be published ( $p = 0.003$ ), consistent with journals rewarding rigorous identification, though the indicator is AI-coded and may capture correlated quality attributes; Section 5.3 shows the premium survives sequential controls for prestige and team size ( $r = 0.015$  between QE and institutional rank). Direction coefficients remain near zero after adding QE ( $p > 0.44$ ). Column (7) adds a top-institution indicator (UTD Top 100 Business School Research Rankings 2021–2025,

Naveen Jindal School of Management, University of Texas at Dallas 2026, via OpenAlex). Top Institution is strongly positive (0.661,  $p < 0.001$ ), the QE coefficient is essentially unchanged (0.473\*\*\*), and direction coefficients remain insignificant ( $\hat{\beta}_1 = -0.160$ ,  $\hat{\beta}_2 = -0.159$ ).

Table 9 breaks down publication rates by direction and research design. QE papers publish at much higher rates than non-QE within every direction: gaps of 20.9 pp positive (55.6% vs. 34.7%), 13.2 pp negative (52.2% vs. 39.0%), and 38.6 pp null (75.0% vs. 36.4%, but the null×QE cell has only 16 papers). Within non-QE papers the three direction rates are similar (spread 4.3 pp); within QE papers they are more dispersed (spread 22.8 pp), driven mainly by the small null cell. Two checks bear on whether the QE premium varies with direction: a joint probit with both indicators leaves QE strongly positive ( $p = 0.004$ ) and both direction coefficients near zero (positive  $p = 0.548$ , negative  $p = 0.499$ ); adding direction×QE interactions yields  $p = 0.158$  and  $p = 0.093$  (the latter marginal at the 10% level). We read this as ruling out strong heterogeneity but not pure uniformity.

**[INSERT TABLE 9 HERE]**

Dividing each direction-specific corpus-membership share by the null-paper share gives ratios of 0.98 (positive) and 0.99 (negative) —essentially 1.0 in both cases. These ratios are not directly comparable to the 1.5–2.0× figures reported by Brodeur, Lé, Sangnier and Zylberberg (2016), which measure how often researchers choose test statistics that land just above the significance threshold relative to just below it. Our ratios instead measure whether papers with different conclusions are equally represented in the published corpus.

The one factor that does predict published-side corpus membership with both statistical and practical significance is research design. Conditional on observed controls (year, team size, top-institution status, finding direction), QE papers are 18–20 pp more likely to appear in the published corpus than papers with the same observed characteristics that use standard regression. “Otherwise similar” refers to these observed dimensions only; unobserved quality differences cannot be ruled out, and we read this association as cross-sectional rather than a causal estimate of editorial weighting.

Untabulated journal-level probits show no statistically significant direction coefficient at any of the three journals individually, consistent with the pooled null but carrying limited inferential weight given small sub-samples (the *Journal of Finance* contributes only 26 in-scope published articles).

### 5.3 The Identification Quality Premium

The quasi-experimental coefficient—QE papers (RDD, IV, DiD, natural experiment, or event study with causal interpretation) are roughly 18–20 percentage points more likely to appear in the published corpus conditional on observed controls—is a substantive cross-sectional association in its own right, not merely a control variable. We develop it here. We flag at the outset that the QE indicator is LLM-coded from abstracts and has not been manually validated on the full sample; a 30-paper audit reported in Appendix G suggests the headline magnitude is fragile to stricter coding criteria, so the discussion in this subsection should be read as describing the sign and order of magnitude of the QE–corpus-membership association rather than a precisely measured premium.

A natural concern is that the premium proxies for confounds: top institutions may produce more quasi-experimental work, or larger teams may have better access to data and computing resources that enable credible identification. To assess this, Table 10 adds controls sequentially.

**[INSERT TABLE 10 HERE]**

The QE coefficient is essentially unchanged across all six specifications, ranging from 0.441 to 0.508, all significant at  $p < 0.01$ . The stability reflects a striking feature of the data: QE status and institutional prestige are nearly orthogonal ( $r = 0.015$ ), and top-institution researchers use QE methods at virtually the same rate as other-institution researchers (48.9% vs. 47.3%). Untabulated  $2 \times 2$  corpus-membership shares confirm the premium operates within both institution tiers (top: 72.7% QE vs. 52.2% non-QE, gap 20.5 pp; other: 48.4% vs. 30.2%, gap 18.2 pp), with the two effects additive across the 42-pp spread

between top-QE (72.7%) and other-non-QE (30.2%) papers; cell counts 32/44, 24/46, 46/95, 32/106 respectively. We cannot fully rule out residual confounding from unobserved quality dimensions that travel with QE method choice (data access, framing, robustness practices), but the near-zero correlation with institutional rank argues against a pure researcher-quality story; if the premium were entirely about quality it would be expected to correlate at least moderately with prestige.

## 6. Robustness

This section checks whether the main findings hold under a range of alternative assumptions. Table 11 summarizes results across three areas: sample and scope restrictions, alternative model specifications, and concerns about classification quality.

### 6.1 Caliper Window Sensitivity

The caliper test compares how many test statistics fall just below the 5% significance threshold versus just above it, so the result depends on exactly how wide a band is used on each side. Untabulated, when the window starts at  $|z| = 1.645$  or higher—covering three of the four specifications we examine, including the standard Brodeur, Lé, Sangnier and Zylberberg (2016) window—published papers consistently show excess mass just below the threshold (ratios 1.20 at the standard window, 1.33 at the narrow  $[1.70, 2.22]$  window, and 1.33 at Brodeur’s  $[1.65, 2.25]$  window), while NBER ratios stay between 0.97 and 1.10. The exception is the widest window ( $[1.50, 2.42]$ ): the published ratio drops to 0.92 ( $p = 0.088$ ), mechanically driven by rising test-statistic density between  $|z| = 1.50$  and  $|z| = 1.645$ . This rules out the alternative explanation that published papers simply have a general shift toward larger test statistics; the pattern is specifically concentrated in the narrow band just below the 5% significance cutoff.

## 6.2 Sample and Scope Robustness

Our sample includes papers on central bank actions such as quantitative easing, discount window lending, and unconventional monetary policy alongside papers on explicit legislation, securities regulation, and banking supervision. One might worry that monetary policy papers behave differently—perhaps because their findings are harder to classify directionally, or because they attract a different editorial audience. To check, we drop the 21 monetary policy papers and re-estimate using only the 270 remaining papers on financial legislation and regulation. The results are unchanged: both direction coefficients are close to zero and statistically insignificant ( $\hat{\beta}_1 = -0.028$ ,  $p = 0.896$ ;  $\hat{\beta}_2 = -0.068$ ,  $p = 0.741$ ), confirming that the null finding is not driven by the inclusion of monetary policy papers.

As a further check, we ask whether finding direction predicts placement in one journal rather than another. Holding the full 291-paper sample fixed, we redefine the dependent variable in two ways and re-run the probit:  $\text{Pub}_{\text{Journal of Finance}} = 1$  if a paper appears in the *Journal of Finance* (0 otherwise, including *Review of Financial Studies*, *Journal of Financial Economics*, and NBER), and  $\text{Pub}_{\text{Review of Financial Studies}+\text{Journal of Financial Economics}} = 1$  if a paper appears in the *Review of Financial Studies* or the *Journal of Financial Economics* (0 otherwise, including *Journal of Finance* and NBER). Each specification asks whether directional papers are disproportionately likely to land in the indicated outlet relative to everything else in the corpus. In neither case does finding direction predict outlet placement: the  $\text{Pub}_{\text{Journal of Finance}}$  coefficients are 0.400 and 0.350 ( $p > 0.20$ ,  $N = 291$ ), and the  $\text{Pub}_{\text{Review of Financial Studies}+\text{Journal of Financial Economics}}$  coefficients are  $-0.147$  and  $-0.130$  ( $p > 0.48$ ,  $N = 291$ ). This is a target-journal-membership test on the full sample, not a journal-vs-NBER subsample comparison (a JF-vs-NBER subsample would be  $26 + 157 = 183$  observations); it asks whether any single outlet sorts on direction more sharply than the corpus as a whole, and the answer is no.

### 6.3 Alternative Specifications

Our baseline model separates papers into positive, negative, and null findings. One concern is that labeling a finding as positive or negative requires a judgment call about what counts as a beneficial outcome for a given intervention. As a check, we sidestep this issue entirely by collapsing the three categories into a simple binary: did the paper find any directional result at all, or was the finding null? The result is unchanged—the coefficient is  $-0.020$  (s.e.  $0.182$ ,  $p = 0.914$ )—and adding year and author-count controls leaves it essentially the same at  $-0.014$  ( $p = 0.941$ ). The null result does not depend on how finely we classify finding direction.

We also check whether the result depends on the choice of statistical model. Probit models make distributional assumptions that may not hold perfectly in small samples. Re-estimating the same model using ordinary least squares (a linear probability model with robust standard errors) gives coefficients of  $-0.010$  ( $p = 0.900$ ) for positive-finding papers and  $-0.006$  ( $p = 0.942$ ) for negative-finding papers—directly readable as percentage-point differences in publication rates relative to null papers. Both are negligible, confirming that the null result is not an artifact of the probit specification.

Finally, we check whether the results are driven by the small set of 16 published papers that we can confirm previously circulated as NBER working papers. Because our NBER comparison group is defined as the 157 in-scope working papers that did *not* reach the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics*, the matched-pair NBER drafts are by construction already excluded from the estimation sample (see Section 3.5); dropping the 16 matched published papers therefore leaves  $N = 275$  (118 published + 157 NBER). The direction coefficients remain small and statistically insignificant on this subsample ( $\hat{\beta}_1 = -0.122$ ,  $p = 0.561$ ;  $\hat{\beta}_2 = -0.054$ ,  $p = 0.787$ , year-trend included), confirming that the results do not depend on the matched-pair subset.

## 6.4 Classification and Quality Concerns

One concern with our main result is that null-finding papers may simply use weaker research designs, and editors may be rejecting them for quality reasons rather than because of their conclusions. To address this, we restrict the sample to the 139 papers that use credible causal identification methods—natural experiments, instrumental variables, difference-in-differences, or regression discontinuity designs—holding research design quality roughly constant across finding directions.

Within this subsample, neither direction coefficient is significantly positive: the positive-finding coefficient is  $-0.598$  ( $p = 0.121$ ) and the negative-finding coefficient is  $-0.656$  ( $p = 0.081$ ). The second figure is marginally significant at the 10% level, but with a *negative* sign, meaning quasi-experimental negative-finding papers are if anything *less* likely to be published than null papers—the opposite of what a direction-premium story would predict. We do not over-interpret this: it is fully consistent with the pooled null result and with random sampling variation in a small subsample. The six direction-by-corpus cells in the QE subsample are not uniformly small—Table 9 implies roughly 30/24 (positive published/NBER-only), 36/33 (negative), and 12/4 (null/mixed)—but the null/mixed baseline cells are very small (12 and 4 papers), and the negative coefficient is computed relative to that baseline. The marginal  $p = 0.081$  on the negative coefficient is therefore best read as instability driven by the small null/mixed-QE reference group rather than as evidence of an anti-intervention discount. The classification of papers as quasi-experimental is also based on AI coding of abstracts and has not been manually verified. The key takeaway is that neither coefficient supports the existence of a positive direction premium among the most rigorously designed papers in the sample.

A second concern is that AI direction coding may differ across abstract styles. The AI assigns high-confidence codes to 81.3% of published abstracts versus 62.4% of NBER abstracts, a 18.9 pp gap; NBER abstracts are actually longer on average (158 vs. 109 words), so the gap is not driven by text length. To separate writing style from substantive

finding differences, we count two term families separately: *style-only hedges* (“may,” “might,” “could,” “possibly,” “perhaps,” “suggests,” “appears,” “tends to,” “seems,” “unclear”) and *null-result phrases* (“no significant,” “do not find,” “no evidence,” “insignificant,” “null”). Per 100 words, NBER vs. published abstracts contain 0.36 vs. 0.22 style-only hedges ( $1.67\times$ ) and 0.05 vs. 0.04 null-result phrases ( $1.15\times$ ). The style-only gap dominates; the null-result gap is small. We read the confidence gap as primarily reflecting writing style.

The bias direction on our direction premium depends on which miscoding channel dominates: NBER directional-as-null inflates the premium; published directional-as-null deflates it; published null-as-directional also inflates it. The LLM replication exercise (Section 3.4) shows null is the least stable category: 15 of 20 null-coded validation papers re-code as directional in replication, with the flips concentrated on the published side (13 of 15 published nulls vs. 2 of 5 NBER nulls). We therefore do not claim the estimates are conservative.

As a partial robustness check, adding  $\log(1 + \text{abstract length})$  as a control leaves direction coefficients essentially zero ( $\hat{\beta}_1 = 0.008$ ,  $p = 0.968$ ;  $\hat{\beta}_2 = 0.011$ ,  $p = 0.958$ ), ruling out the length channel but not the per-100-word hedging-style gap, which could affect coding conditional on length. A proper test of the style channel would require re-coding from style-normalized abstracts; we do not attempt this here.

[INSERT TABLE 11 HERE]

## 7. Discussion

### 7.1 Implications for Research and Regulation

Our findings argue against the aggregate concern that top finance journals systematically favor studies that find an effect of government intervention over those that find none. Positive-, negative-, and null-finding papers appear in the published corpus at nearly identical rates ( $\approx 46\%$ ), and finding direction is unrelated to corpus membership across all seven probit

specifications. This is an aggregate statement: the sample spans mandatory disclosure rules, corporate governance mandates, short-selling restrictions, and Dodd-Frank provisions, but with 134 published articles we cannot support reliable within-topic comparisons. Our study is also underpowered to detect modest imbalances—the smallest reliably detectable gap corresponds to an 18-percentage-point publication-rate difference—so policymakers should treat our null as evidence against *large-scale* direction-based editorial distortion, not a clean bill of health; bias on specific regulations could exist below our detection threshold.

A separate and distinct caveat applies to the within-paper caliper analysis: the bootstrap confidence interval on the negative-finding caliper ratio is wide ( $[0.81, 2.67]$ ), which limits how confidently we can pin down the magnitude of within-paper threshold pressure in negative-finding published papers. We flag this here because the two caveats are about different estimands—the first concerns *direction-based corpus membership* measured by the probit; the second concerns *within-paper selective reporting near  $|z| = 1.96$*  measured by the caliper test—and the wide caliper interval does not by itself qualify the no-direction-sorting result.

A clearer cross-sectional association does emerge: QE papers (RDD, IV, DiD, natural experiment, or event study with causal interpretation; see Section 5.2) are substantially more likely to appear in the published corpus, conditional on year, team size, top-institution status, and finding direction. For regulators, this association is reassuring rather than distorting if it reflects editorial or pipeline selection on methodological quality, but our cross-sectional design cannot exclude that quasi-experimental papers also differ on unobserved dimensions (data, framing, robustness practices) that travel with method choice. The descriptive statement we can defend is that the published evidence base is disproportionately drawn from quasi-experimental work; whether that pattern reflects an editorial quality filter or unobserved correlated quality is a question our design leaves open.

One further caution applies to how individual results are read. The excess mass we observe on the just-below-1.96 side of the threshold ( $|z| \in [1.645, 1.96)$ ,  $p \in (0.05, 0.10]$ )—

concentrated in negative-finding papers—suggests that *within-paper* specification choices may play a role even when the publication decision itself is direction-neutral. Those relying on this literature should apply extra scrutiny to anti-intervention findings whose test statistics sit close to the conventional significance boundary on either side (roughly  $|z| \in [1.645, 2.275]$ ), while treating publication in a top-three finance journal as a credible signal of overall methodological quality.

## 7.2 How Finance Compares to Economics

Economics and related social sciences typically show a substantial publication premium for statistically significant findings. Brodeur, Lé, Sangnier and Zylberberg (2016) document excess test statistics just above the significance threshold in top general-interest economics journals; Franco, Malhotra and Simonovits (2014) find that NSF-funded null results are published at roughly one-third the rate of significant ones; DellaVigna and Linos (2022) report similar disparities in field experiments, and Brodeur, Cook, Hartley and Heyes (2024) show that pre-registration reduces but does not eliminate this advantage.

Our results stand out against this backdrop. Null-finding papers in the finance-regulation literature reach the published corpus at the same rate as positive- or negative-finding papers, with no three-to-one gap. The published-sample caliper ratio  $R = 1.20$  (just-below/just-above convention; equivalently  $\approx 0.83$  under Brodeur, Lé, Sangnier and Zylberberg 2016’s convention) compared with  $R = 1.01$  for NBER is consistent with editorial or revision-stage selection introducing the asymmetry, but because the two corpora are largely non-overlapping cross-sections we cannot rule out a compositional explanation in which published and unmatched NBER papers differ on unobserved dimensions correlated with near-threshold density. Notably, the published-sample asymmetry is in the opposite direction from the classical Brodeur pattern: excess on the just-insignificant side rather than just clearing the threshold.

Two channels could explain the contrast with economics. First, editorial norms at the

*Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* may be less sensitive to finding direction than norms at general-interest economics journals. Second, since NBER working papers come disproportionately from well-resourced authors, null results in this domain may circulate and reach editors more reliably than null results from less-connected researchers; the absence of a selection gap may partly reflect NBER-baseline composition rather than editorial neutrality per se.

Our caliper finding fits alongside prior finance work. Harvey and Liu (2021) document near-threshold clustering in published finance research and estimate what the pre-publication distribution must have looked like; our NBER data are consistent with their inferred shape, but the cross-sectional design does not let us pin the asymmetry on the revision stage rather than on compositional differences. Chen (2021) note that very large published test statistics ( $|t| > 4$ ) appear too frequently to be explained by marginal-significance reporting alone, implying most large effects in the published literature are real. One interpretation of why near-threshold asymmetry is concentrated in negative-finding papers is that anti-intervention results face elevated reviewer scrutiny, leading authors to report many borderline specifications. A compositional alternative—that negative-finding published papers differ from their NBER counterparts on unobserved dimensions correlated with the number of specifications reported—is equally consistent with the data.

### 7.3 Alternative Interpretations

Equal publication rates across directions are consistent with two explanations the data cannot separate: genuine editorial neutrality (editors do not favor papers that find an effect), and a supply-side equilibrium in which researchers who produce null findings submit them to top finance journals at higher rates than in fields where null results are known to face rejection. Equal observed rates would then reflect balanced submission rather than balanced acceptance. The near-identical share of directional findings in NBER (79.6%) and published (79.1%) samples is consistent with the supply-side story. Similarly, the quasi-experimental

coefficient (0.467<sup>\*\*\*</sup>) is a cross-sectional association conditional on observed controls; research design quality travels with unobserved attributes (resources, framing, data access) that our model does not measure, so the coefficient reflects a bundle of quality-related attributes correlated with identification strategy rather than a causal estimate of editorial weighting.

## 7.4 Limitations

Our design has six limitations. *First*, only 16 of 134 published papers (11.9%) match a specific NBER counterpart, so we compare two largely non-overlapping cross-sections rather than tracking papers through the editorial process. All direction coefficients are cross-sectional associations, not causal editorial-filter effects. *Second*, our outcome variable is top-three placement (*Journal of Finance*, *Review of Financial Studies*, *Journal of Financial Economics*); NBER papers published in other peer-reviewed journals are correctly classified as not-in-top-three under this outcome, so there is no misclassification. Generalizing to broader peer-reviewed publication requires data on second-tier placements that we do not have, so we cannot sign the bias on that broader outcome. *Third*, direction coding from NBER abstracts is noisier than from published ones (high-confidence rate 62.4% vs. 81.3%), and null is the least stable category: 15 of 20 null-coded validation papers re-code as directional in replication, with the flips concentrated in published (13 of 15) rather than NBER (2 of 5). The direction of the resulting bias depends on which side is more frequently miscoded, so we cannot sign it; a joint equality test across all three direction categories ( $\chi^2 = 0.017$ ,  $p = 0.992$ ) confirms the overall null without leaning on any single category as the reference. *Fourth*, the quasi-experimental premium could mask a direction effect if rigorous studies systematically find directional results, but direction coefficients are unchanged with and without the QE control, and Table 9 shows flat publication rates across directions within both QE and non-QE strata. The QE indicator itself is LLM-coded from abstracts and not separately validated on the full sample. *Fifth*, NBER researchers cluster at leading universities; if their papers are inherently more publishable our pre-publication baseline sets a high bar and

would understate any true direction premium. The near-identical directional share in NBER (79.6%) and published (79.1%) samples is the most direct evidence against this concern. *Sixth*, the caliper analysis runs 18 comparisons. Under Bonferroni-corrected naive chi-square ( $p < 0.003$ ), three caliper tests (full published  $p = 0.002$ , *Journal of Financial Economics*  $p < 0.001$ , negative-finding  $p < 0.001$ ) and two density tests (*Journal of Finance*, null/mixed,  $p < 0.001$  each) survive; under our preferred paper-clustered bootstrap, the published full sample and negative-finding subgroup ( $p = 0.150$  and  $0.135$  respectively) do not clear  $0.05$  even before correction. We treat the Bonferroni-surviving set as the less-conservative reading and the bootstrap as the preferred reading; the *Journal of Finance* density anomaly is included for completeness but conflicts with the *Journal of Finance* caliper ratio ( $0.88$ ), and we give priority to the caliper. NBER caliper tests do not approach significance under any test.

## 8. Conclusion

This paper asks whether top finance journals are more likely to publish government-intervention studies that find an effect than studies that find none. Using a sample of 134 published articles from the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* and 157 NBER working papers as a pre-publication benchmark, we find no evidence of direction-based editorial selection at a scale our design can detect. Positive-, negative-, and null-finding papers occupy the published side of our combined 291-paper corpus at nearly identical corpus-membership shares —  $45.6\%$ ,  $46.1\%$ , and  $46.7\%$  respectively (within-direction shares, see the terminology convention at the start of Section 5.2) — and finding direction is statistically unrelated to published-side corpus membership across all seven model specifications. The design rules out gaps larger than 18 percentage points at 80% power; smaller asymmetries remain possible.

The distribution of test statistics tells a more nuanced story. As shown in Table 5, the full-sample caliper ratio of  $1.20$  reflects two reinforcing effects: published papers cluster

standalone  $z$ -statistics more sharply near the significance threshold than NBER papers do (ratio 2.54 vs. 1.58), and they devote a larger share of table rows to such statistics (28.3% vs. 18.1%). Verifiable coefficient rows show no systematic threshold manipulation in either corpus. The bunching is most pronounced in negative-finding papers (ratio 1.52), though this pattern is suggestive rather than statistically definitive.


Together these two results point to different stages of the research pipeline. The absence of detectable large-scale direction-based sorting between published and pre-publication papers suggests that editorial selection in this domain is not strongly tilted toward finding an effect. The within-paper bunching just below the  $|z| = 1.96$  threshold, concentrated in published negative-finding papers, is more consistent with revision-stage pressure than with outright fabrication of results. We flag this as suggestive rather than definitive: the bunching pattern clears a Bonferroni cutoff under a naive chi-square that treats every test statistic as independent, but does not clear conventional significance under our preferred paper-clustered bootstrap ( $p = 0.135$ – $0.150$ , see Section 5.1).

The two largest cross-sectional predictors of corpus membership are research-design quality and institutional prestige, not finding direction. Conditional on observed controls, QE papers are 18–20 pp more likely to appear in the published corpus and top-institution papers 24–26 pp more likely (Table 10); in Column (6) the institutional-prestige marginal effect (26.2 pp) exceeds QE’s (18.8 pp). Both indicators exceed any direction-based predictor and are nearly orthogonal to each other ( $r = 0.015$ ). These associations are consistent with journals selecting on a bundle of quality attributes correlated with identification strategy and pedigree. For policymakers, the takeaway is that the published record of government-intervention research in top finance journals does not appear large-scale tilted toward positive or negative findings by editorial selection. The more pressing concern is within-paper: those relying on this evidence base should treat marginal significance in anti-intervention findings with additional scrutiny, while treating publication in a top-three finance journal as a credible signal of methodological quality.


## References

- Agarwal, S., Lucca, D., Seru, A., Trebbi, F., 2014. Inconsistent regulators: Evidence from banking. *Quarterly Journal of Economics* 129, 889–938.
- Andrews, I., Kasy, M., 2019. Identification of and correction for publication bias. *American Economic Review* 109, 2766–2794.
- Ash, E., Hansen, S., 2023. Text algorithms in economics. *Annual Review of Economics* 15, 659–688.
- Berger, A.N., Bouwman, C.H.S., 2017. Bank liquidity creation, monetary policy, and financial crises. *Journal of Financial Stability* 30, 139–155.
- Brodeur, A., Cook, N., Hartley, J., Heyes, A., 2024. Do pre-registration and pre-analysis plans reduce p-hacking and publication bias? Evidence from 15,992 test statistics and suggestions for improvement. *IZA Discussion Paper* 15476.
- Brodeur, A., Lé, M., Sangnier, M., Zylberberg, Y., 2016. Star wars: The empirics strike back. *American Economic Journal: Applied Economics* 8, 1–32.
- Chen, A.Y., 2021. The limits of p-hacking: Some thought experiments. *Journal of Finance* 76, 2447–2480.
- DellaVigna, S., Linos, E., 2022. Rcts to scale: Comprehensive evidence from two nudge units. *Econometrica* 90, 81–116.
- Franco, A., Malhotra, N., Simonovits, G., 2014. Publication bias in the social sciences: Unlocking the file drawer. *Science* 345, 1502–1505.
- Harvey, C.R., Liu, Y., 2021. Uncovering the iceberg from its tip: A model of publication bias and p-hacking. *SSRN Working Paper* 3865813.


- Harvey, C.R., Liu, Y., Zhu, H., 2016. ... and the cross-section of expected returns. *Review of Financial Studies* 29, 5–68.
- Hou, K., Xue, C., Zhang, L., 2020. Replicating anomalies. *Review of Financial Studies* 33, 2019–2133.
- Jensen, T.I., Kelly, B.T., Pedersen, L.H., 2023. Is there a replication crisis in finance? *Journal of Finance* 78, 2465–2518.
- Korinek, A., 2023. Language models and cognitive automation for economic research. NBER Working Paper 30957.
- La Porta, R., Lopez-de Silanes, F., Shleifer, A., Vishny, R.W., 1998. Law and finance. *Journal of Political Economy* 106, 1113–1155.
- Landis, J.R., Koch, G.G., 1977. The measurement of observer agreement for categorical data. *Biometrics* 33, 159–174.
- Naveen Jindal School of Management, University of Texas at Dallas, 2026. The UTD Top 100 Business School Research Rankings. <https://jsom.utdallas.edu/the-utd-top-100-business-school-research-rankings/>. 2021–2025 worldwide ranking period. Accessed June 2026.



**Figure 1.** Distribution of  $|z|$ -statistics from all 7,761 regression rows across 114 published papers. The dashed vertical lines mark  $|z| = 1.645$  (10% two-sided),  $|z| = 1.96$  (5%), and  $|z| = 2.576$  (1%). The red reference curve is the half-normal density—the implied distribution of  $|Z|$  when  $Z \sim N(0, 1)$ , equivalently the change-of-variables image of the  $\chi^2(1)$  distribution of  $Z^2$  onto the  $|z|$  scale—so it is commensurate with the histogram and represents the null benchmark of no true effects and no selective reporting. Within the caliper window, statistics in  $[1.645, 1.96)$ —the just-insignificant zone, where  $0.05 < p \leq 0.10$ —outnumber those in  $[1.96, 2.275]$ —the just-significant zone, where  $p \leq 0.05$ —by about 20% ( $R = 1.20$ ,  $p = 0.002$ ). The excess on the just-below-1.96 side is consistent with selective reporting of borderline results that fall just short of conventional significance.



**Figure 2.** Each panel compares published journal papers (left) to NBER working papers (right). Bars show how frequently test statistics fall in each range. The orange bars mark the zone just around the conventional 5% significance cutoff; the vertical line is the cutoff itself. If test statistics were reported without regard to significance, the orange bars on either side of the line should be roughly equal. *Top panel:* All 7,761 test statistics from published papers and 8,351 from NBER papers. Published papers show visibly more statistics just below the cutoff than just above it; NBER papers do not. *Bottom panel:* The same comparison restricted to papers that conclude a government regulation was harmful (3,491 published; 4,375 NBER). The asymmetry is more pronounced in published papers and essentially absent in NBER papers; see Section 5.1 for the cross-corpus significance discussion.



**Figure 3.** Test-statistic distributions broken out by journal. All  $p$ -values reported in this caption are caliper-test  $p$ -values ( $R = \#\{|z| \in [1.645, 1.96]\} / \#\{|z| \in [1.96, 2.275]\}$ ); the density discontinuity test (Disc.  $p$  in Table 4) gives a different reading for some panels and we note conflicts explicitly below. *Top:* Journal of Finance (1,308 statistics; no caliper-test clustering, caliper  $p = 0.423$ ). Note that the discontinuity test for JF gives  $p < 0.001$ ; the body text explains why we give priority to the caliper here (Section 5.1). *Middle:* Review of Financial Studies (2,581 statistics; no caliper-test clustering, caliper  $p = 0.752$ ; discontinuity  $p = 0.180$ , no conflict). *Bottom:* Journal of Financial Economics (3,872 statistics; significant caliper-test clustering on the just-below-1.96 side, caliper  $p < 0.001$ ; discontinuity  $p = 0.038$ , same direction). Under the caliper test, near-threshold clustering is present only in the JFE; the discontinuity test flags the JF as well, and we read that JF flag as noise (see Section 5.1).

**Table 1.** Sample Selection

| Step                                           | Papers | Excluded |
|------------------------------------------------|--------|----------|
| Starting corpus (NBER + JF/RFS/JFE, 2000–2024) | 39,136 |          |
| Pass 1: Keyword filter                         | 1,849  | 37,287   |
| Pass 2: LLM classification & manual review     | 291    | 1,558    |
| <i>Final sample</i>                            | 291    |          |
| Published articles (JF / RFS / JFE)            | 134    |          |
| Journal of Finance                             | 26     |          |
| Review of Financial Studies                    | 48     |          |
| Journal of Financial Economics                 | 60     |          |
| NBER working papers                            | 157    |          |

*Notes.* The starting corpus includes all 30,212 NBER working papers (every program, 2000–2024) and all research articles from the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics*, 2000–2024. The keyword filter screens for government intervention terminology (see Appendix A). Pass 2 combines three sub-steps. (a) The LLM in-scope prompt is run on all 1,849 keyword candidates, flagging 217 as in-scope. (b) The author manually reviews each LLM in-scope decision and overrides 22 published articles (*Journal of Finance*: 18, *Review of Financial Studies*: 2, *Journal of Financial Economics*: 2) plus 10 NBER papers as out of scope, retaining 51 of the 61 LLM in-scope NBER papers and 134 of the 156 LLM in-scope published articles. (c) An NBER-side refilter re-codes 803 NBER papers that the LLM had flagged as low-confidence out-of-scope; 126 are returned to the in-scope pool and 106 survive the same manual review criteria. Net of all three sub-steps the working-paper sample is  $51 + 106 = 157$  NBER and the published sample is 134 articles, summing to the 291-paper analysis sample.

**Table 2.** Prompt-Perturbation Robustness: Agreement with Original Direction Codes

| Prompt variant                   | $N$ | Agreed | Agreement % | $\kappa_w$ |
|----------------------------------|-----|--------|-------------|------------|
| Baseline (original wording)      | 58  | 53     | 91.4        | 0.904      |
| Variant A (“primary conclusion”) | 58  | 48     | 82.8        | 0.742      |
| Variant B (“classify direction”) | 58  | 49     | 84.5        | 0.814      |

*Notes.* Each variant re-codes the same 58-paper stratified sample using a smaller model with the same classification rules.  $\kappa_w$  is computed against the original direction codes used throughout the paper.

**Table 3.** Summary Statistics

| Sample                                | <i>N</i> | Positive (%) | Negative (%) | Null (%) |
|---------------------------------------|----------|--------------|--------------|----------|
| All papers                            | 291      | 35.4         | 44.0         | 20.6     |
| NBER working papers                   | 157      | 35.7         | 43.9         | 20.4     |
| Published (JF/RFS/JFE)                | 134      | 35.1         | 44.0         | 20.9     |
| <i>By journal:</i>                    |          |              |              |          |
| <i>Journal of Finance</i>             | 26       | 42.3         | 46.2         | 11.5     |
| <i>Review of Financial Studies</i>    | 48       | 29.2         | 54.2         | 16.7     |
| <i>Journal of Financial Economics</i> | 60       | 36.7         | 35.0         | 28.3     |

*Notes.* The sample consists of 157 NBER working papers identified from the full NBER corpus (all programs, 2000–2024) via topic filter, and 134 published journal articles (*Journal of Finance*, *Review of Financial Studies*, *Journal of Financial Economics*) classified as in-scope government-intervention papers by our LLM pipeline and manual validation. Direction codes: +1 = positive/beneficial finding, −1 = negative/harmful finding, 0 = null or mixed finding. The NBER and Published samples are mutually exclusive under the accounting rule stated in Section 3.5: the 291 papers consist of 157 NBER working papers whose published version did not subsequently appear in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics* during 2000–2024, and 134 distinct published journal articles, with the 16 matched-pair NBER drafts represented only by their journal-published incarnation. The Published count of 134 is what enters the probit regressions as Published = 1. Column percentages may not sum to 100 due to rounding.

**Table 4.** Within-Paper Test Statistic Distribution Tests

| Subset                       | $N$   | % Sig | Caliper ratio | Caliper $\chi^2$ | Caliper $p$ | Disc. $p$ |
|------------------------------|-------|-------|---------------|------------------|-------------|-----------|
| Full sample                  | 7,761 | 57.1% | 1.20***       | 9.48             | 0.002       | 0.040     |
| <i>By journal:</i>           |       |       |               |                  |             |           |
| JF                           | 1,308 | 65.5% | 0.88          | 0.64             | 0.423       | <0.001    |
| RFS                          | 2,581 | 59.2% | 0.97          | 0.10             | 0.752       | 0.180     |
| JFE                          | 3,872 | 52.9% | 1.49***       | 23.56            | <0.001      | 0.038     |
| <i>By finding direction:</i> |       |       |               |                  |             |           |
| Positive                     | 2,633 | 60.7% | 0.94          | 0.33             | 0.565       | 0.456     |
| Negative                     | 3,491 | 54.9% | 1.52***       | 22.54            | <0.001      | 0.183     |
| Null/mixed                   | 1,637 | 56.1% | 1.05          | 0.11             | 0.740       | <0.001    |

*Notes:*  $N$  is the number of test statistics extracted from regression tables, after removing implausible values and non-numeric entries. % Sig is the share of test statistics that exceed the conventional 5% significance threshold. *Caliper ratio* is  $R = \#\{|z| \in [1.645, 1.96)\} / \#\{|z| \in [1.96, 2.275]\}$ , the count of  $|z|$ -statistics in the just-insignificant window divided by the count in the just-significant window.  $R = 1.0$  means equal frequency;  $R > 1.0$  indicates clustering on the just-below-1.96 side (results that fall just short of the conventional significance bar), and  $R < 1.0$  indicates clustering on the just-above-1.96 side. *Caliper  $p$*  is the associated  $p$ -value testing whether the two zones are equally populated. *Disc.  $p$*  tests whether there is a sharp drop in the frequency of test statistics right at the 5% threshold (Brodeur, Lé, Sangnier and Zylberberg, 2016). \*\*\*  $p < 0.01$ , \*\*  $p < 0.05$ , \*  $p < 0.10$  (caliper ratio).

**Table 5.** Caliper Ratio Decomposition by Statistic Type and Corpus

| Corpus    | Row type       | $N$ stats | Ratio   | Share |
|-----------|----------------|-----------|---------|-------|
| Published | All rows       | 7,761     | 1.20*** | 100%  |
|           | Coef.+SE rows  | 5,563     | 0.82*** | 71.7% |
|           | Standalone $z$ | 2,198     | 2.54*** | 28.3% |
| NBER      | All rows       | 8,351     | 1.01    | 100%  |
|           | Coef.+SE rows  | 6,840     | 0.95    | 81.9% |
|           | Standalone $z$ | 1,511     | 1.58**  | 18.1% |

*Notes.* *Coef.+SE rows* are regression rows where both a coefficient and a standard error are reported, so the test statistic can be independently verified. *Standalone  $z$*  rows report only a  $t$ - or  $z$ -statistic with no coefficient or standard error. *Share* is the fraction of each corpus’s statistics that fall into that row type. The full-sample ratio of 1.20 reflects both the higher standalone bunching in published papers and their larger standalone share relative to the NBER corpus. Significance markers are two-sided (\*\*\*  $p < 0.01$ , \*\*  $p < 0.05$ , \*  $p < 0.10$  for the equality test). Direction arrows:  $\uparrow$  flags  $R > 1$  (excess on the just-below-1.96 side, the focal pattern);  $\downarrow$  flags  $R < 1$  (excess on the just-above-1.96 side, the classical Brodeur “just-cleared-the-bar” direction). The published Coef+SE row is significant in the  $\downarrow$  direction ( $p = 0.007$ ): in verifiable rows, published papers show slightly more statistics just above 1.96 than just below. See Section 5.1 for the two-direction interpretation and Appendix E for the underlying counts.

**Table 6.** Within-Paper Caliper Test: Published vs. NBER Working Papers

| Sample    | Subset           | $N$ stats | Papers | Ratio   | $p$ -value |        |       |
|-----------|------------------|-----------|--------|---------|------------|--------|-------|
|           |                  |           |        |         | Naive      | Perm.  | Boot. |
| Published | Full sample      | 7,761     | 108    | 1.20*** | 0.002      | 0.001  | 0.150 |
| NBER      | Full sample      | 8,351     | 106    | 1.01    | 0.849      | 0.424  | 0.432 |
| Published | Positive-finding | 2,633     | 37     | 0.94    | 0.565      | 0.732  | 0.623 |
| NBER      | Positive-finding | 2,658     | 39     | 0.99    | 0.918      | 0.554  | 0.492 |
| Published | Negative-finding | 3,491     | 50     | 1.52*** | <0.001     | <0.001 | 0.135 |
| NBER      | Negative-finding | 4,375     | 50     | 1.00    | 0.964      | 0.532  | 0.474 |
| Published | Null-finding     | 1,637     | 21     | 1.05    | 0.740      | 0.388  | 0.383 |
| NBER      | Null-finding     | 1,318     | 17     | 1.14    | 0.442      | 0.239  | 0.260 |

*Notes:* The *Published* sample contains 7,761 test statistics from 114 papers in the JF, RFS, and JFE; 108 of these papers contribute at least one statistic near the significance threshold and are used in the permutation and bootstrap tests. The *NBER* sample contains 8,351 test statistics from 117 working papers (out of 157 in-scope); 106 contribute near-threshold statistics. The *Ratio* column is  $R = \#\{|z| \in [1.645, 1.96)\} / \#\{|z| \in [1.96, 2.275]\}$ .  $R = 1.0$  means the two zones are equally populated;  $R > 1.0$  indicates clustering on the just-below-1.96 side (just-insignificant results), and  $R < 1.0$  indicates clustering on the just-above-1.96 side (just-significant results). Three significance tests are reported. *Naive*: treats every statistic as an independent observation. *Perm.*: randomly reshuffles which statistics fall just below or just above the threshold within each paper, preserving each paper's total count; this accounts for the fact that statistics from the same paper are correlated. *Boot.*: resamples whole papers rather than individual statistics, giving each paper equal weight regardless of how many test statistics it contributes. \*\*\*  $p < 0.01$ , \*\*  $p < 0.05$ , \*  $p < 0.10$  (caliper ratio). Boot.  $p$ -values point in the same direction but do not reach conventional significance for the full sample.

**Table 7.** Corpus-Membership Rates by Finding Direction

| Finding Direction | Sample composition |           | Total | Corpus-Membership Rate | 95% CI       |
|-------------------|--------------------|-----------|-------|------------------------|--------------|
|                   | NBER-only          | Published |       |                        |              |
| Positive (+1)     | 56                 | 47        | 103   | 45.6%                  | [36.3, 55.2] |
| Negative (−1)     | 69                 | 59        | 128   | 46.1%                  | [37.7, 54.7] |
| Null/Mixed (0)    | 32                 | 28        | 60    | 46.7%                  | [34.6, 59.1] |
| Total             | 157                | 134       | 291   | 46.0%                  |              |

*Notes.* Sample covers 2000–2024. *Finding Direction* is the main conclusion of the paper as coded by the AI classifier: positive (+1) if the paper finds a government intervention was beneficial, negative (−1) if harmful, and null/mixed (0) if the evidence is inconclusive. *NBER-only* counts working papers that were never published in the JF, RFS, or JFE; *Published* counts papers that were; *Total* is both groups combined. *Corpus-Membership Rate* is the share of papers in each direction category that end up published. This reflects the relative sizes of our two samples and should not be interpreted as a journal acceptance rate. *95% CI* is the confidence interval around each publication rate. A statistical test cannot reject that the three publication rates are equal ( $p = 0.992$ ), and the positive vs. negative comparison gives  $p = 0.944$ .

**Table 8.** Direction–Publication Association: Probit Estimates

|                                | (1)<br>Binary     | (2)<br>Bivariate<br>(Primary) | (3)<br>+ Year       | (4)<br>+ Decade FE | (5)<br>+ Authors    | (6)<br>+ ID Strategy | (7)<br>+ Top Inst.  |
|--------------------------------|-------------------|-------------------------------|---------------------|--------------------|---------------------|----------------------|---------------------|
| Directional                    | −0.020<br>(0.182) |                               |                     |                    |                     |                      |                     |
| Positive (+1)                  |                   | −0.026<br>(0.204)             | −0.002<br>(0.205)   | −0.006<br>(0.205)  | −0.005<br>(0.205)   | −0.135<br>(0.212)    | −0.160<br>(0.215)   |
| Negative (−1)                  |                   | −0.014<br>(0.196)             | −0.011<br>(0.198)   | −0.017<br>(0.198)  | −0.020<br>(0.198)   | −0.156<br>(0.206)    | −0.159<br>(0.208)   |
| Year trend                     |                   |                               | 0.033***<br>(0.012) |                    | 0.034***<br>(0.012) | 0.028**<br>(0.012)   | 0.034***<br>(0.013) |
| $\log(1 + N_{\text{authors}})$ |                   |                               |                     |                    | −0.186<br>(0.279)   | −0.255<br>(0.283)    | −0.264<br>(0.291)   |
| Quasi-Exp. ID                  |                   |                               |                     |                    |                     | 0.467***<br>(0.157)  | 0.473***<br>(0.160) |
| Top Institution                |                   |                               |                     |                    |                     |                      | 0.661***<br>(0.167) |
| Decade FEs                     | No                | No                            | No                  | Yes                | No                  | No                   | No                  |
| Intercept                      | −0.084            | −0.084                        | −0.097              | −0.371             | 0.145               | 0.114                | −0.075              |
| $N$                            | 291               | 291                           | 291                 | 291                | 291                 | 291                  | 291                 |
| Pseudo $R^2$                   | 0.000             | 0.000                         | 0.020               | 0.013              | 0.021               | 0.043                | 0.083               |

*Notes.* Each column reports probit estimates of the probability that a paper in our dataset ends up published in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics*. Column (2) is the primary specification: it separates positive (+1) and negative (−1) findings against null/mixed papers (0), matching equation (1) in Section 4.1; direction codes are assigned by the AI pipeline from paper abstracts. Column (1) reports a parsimonious binary alternative—“Directional” versus null/mixed—as a supporting check; Columns (3)–(7) extend Column (2) by adding controls. Column (5) adds the number of authors (log scale) as a proxy for team size. Column (6) adds the quasi-experimental (QE) indicator defined in the preceding text (RDD, IV, DiD, natural experiment, or event study with causal interpretation), as classified by the AI from abstracts. Column (7) adds an indicator for whether the lead author’s institution appears in the UTD Top 100 Business School Research Rankings (2021–2025; Naveen Jindal School of Management, University of Texas at Dallas 2026), retrieved via OpenAlex. Standard errors in parentheses. \*\*\*  $p < 0.01$ , \*\*  $p < 0.05$ , \*  $p < 0.10$ . The baseline publication probability implied by the intercept alone is  $\Phi(-0.084) \approx 46.7\%$ , equal to the raw null-paper publication rate (28 of 60 null-finding papers, 46.7%); positive- and negative-finding papers publish at 45.6% and 46.1%, respectively, for an overall rate of 46.0%. The marginal effects of all direction indicators are negligible across all columns. The near-zero pseudo- $R^2$  in Columns (1)–(2) is exactly what we would expect if finding direction has no predictive power for publication; it reflects a null relationship, not poor model quality. A formal test of whether positive and negative coefficients are equal to each other in Column (2) gives  $\chi^2 = 0.005$ ,  $p = 0.944$ , providing no evidence that journals treat the two directions differently. The sample is large enough to detect coefficient differences of roughly 0.465 units with 80% power (standard error of the difference = 0.166), corresponding to a publication-rate advantage of approximately  $1.4\times$  the null-paper rate (roughly 18 percentage points on the probability scale). The null result therefore does not rule out moderate directional preferences, only large ones.

**Table 9.** Corpus-Membership Rates by Finding Direction and Identification Strategy

|                | Non-QE         | QE                         | Total           |
|----------------|----------------|----------------------------|-----------------|
| Positive (+1)  | 34.7% (17/49)  | 55.6% (30/54)              | 45.6% (47/103)  |
| Negative (−1)  | 39.0% (23/59)  | 52.2% (36/69)              | 46.1% (59/128)  |
| Null/mixed (0) | 36.4% (16/44)  | 75.0% (12/16) <sup>†</sup> | 46.7% (28/60)   |
| Total          | 36.8% (56/152) | 56.1% (78/139)             | 46.0% (134/291) |

*Notes.* Each cell shows the share of papers published in a top-3 finance journal (JF, RFS, or JFE), with paper counts in parentheses. QE: papers whose primary identification strategy is regression discontinuity (RDD), instrumental variables (IV), difference-in-differences (DiD), a natural experiment with an exogenous shock, or an event study with causal interpretation (see definition following Table 8). Non-QE: papers that use standard regression, panel fixed effects without an instrument, descriptive methods, or event-window analyses without a causal claim, as classified by the AI from abstracts. †: Only 16 papers fall in the null/mixed  $\times$  QE cell; the 75% rate is based on a small sample and should be treated with caution. Within the Non-QE column, publication rates are similar across the three finding directions (34.7%, 39.0%, 36.4%); within the QE column, the rates are more dispersed (55.6%, 52.2%, 75.0%), driven mainly by the small null/mixed cell. A formal interaction test of whether the research-design premium varies by finding direction does not reject pure uniformity at conventional levels, but one interaction is marginal ( $p = 0.093$ ); we therefore characterize the evidence as ruling out strong heterogeneity rather than confirming complete uniformity (see text).

**Table 10.** Identification Quality Premium: Sequential Controls

|                                | (1)                 | (2)                 | (3)                 | (4)                 | (5)                 | (6)                 |
|--------------------------------|---------------------|---------------------|---------------------|---------------------|---------------------|---------------------|
| Quasi-Exp. ID                  | 0.490***<br>(0.149) | 0.501***<br>(0.151) | 0.497***<br>(0.150) | 0.508***<br>(0.152) | 0.441***<br>(0.155) | 0.473***<br>(0.160) |
| Top Institution                |                     | 0.607***<br>(0.164) |                     | 0.606***<br>(0.164) | 0.659***<br>(0.167) | 0.661***<br>(0.167) |
| $\log(1 + N_{\text{authors}})$ |                     |                     | -0.176<br>(0.280)   | -0.160<br>(0.281)   | -0.265<br>(0.286)   | -0.264<br>(0.291)   |
| Year trend                     |                     |                     |                     |                     | 0.036***<br>(0.013) | 0.034***<br>(0.013) |
| Direction dummies              | No                  | No                  | No                  | No                  | No                  | Yes                 |
| Marg. effect, Quasi-Exp. (pp)  | 19.4                | 19.9                | 19.7                | 20.1                | 17.5                | 18.8                |
| Marg. effect, Top Inst. (pp)   | —                   | 24.1                | —                   | 24.1                | 26.1                | 26.2                |
| Pseudo $R^2$                   | 0.027               | 0.062               | 0.028               | 0.062               | 0.081               | 0.083               |
| $N$                            | 291                 | 291                 | 291                 | 291                 | 291                 | 291                 |

*Notes.* Each column reports a probit of publication on the indicator variables listed, estimated on the full sample of 291 papers. *Quasi-Exp. ID* equals one if the paper’s primary identification strategy is RDD, IV, DiD, a natural experiment, or an event study with causal interpretation (full definition following Table 8), as coded by the AI from abstracts. *Top Institution* equals one if the first author’s institution appears in the UTD Top 100 Business School Research Rankings. Column (6) includes positive and negative finding-direction dummies (omitted from the table for brevity; both are individually and jointly insignificant,  $p > 0.44$ ). Marginal effects (in percentage points) are evaluated at the sample mean of the linear index. Standard errors in parentheses. \*\*\* $p < 0.01$ , \*\* $p < 0.05$ , \* $p < 0.10$ .

**Table 11.** Robustness Checks

| Specification                                                                         | Variable    | Coefficient | p-value | N   |
|---------------------------------------------------------------------------------------|-------------|-------------|---------|-----|
| <i>Panel A: Sample and Scope</i>                                                      |             |             |         |     |
| Excl. monetary policy                                                                 | Positive    | −0.028      | 0.896   | 270 |
|                                                                                       | Negative    | −0.068      | 0.741   | 270 |
| <i>Journal of Finance</i> membership                                                  | Positive    | 0.400       | 0.217   | 291 |
|                                                                                       | Negative    | 0.350       | 0.271   | 291 |
| <i>Review of Financial Studies</i> + <i>Journal of Financial Economics</i> membership | Positive    | −0.147      | 0.483   | 291 |
|                                                                                       | Negative    | −0.130      | 0.519   | 291 |
| <i>Panel B: Alternative Specifications</i>                                            |             |             |         |     |
| Binary directional                                                                    | Directional | −0.020      | 0.914   | 291 |
| Binary + year + authors                                                               | Directional | −0.014      | 0.941   | 291 |
| LPM (OLS, HC3)                                                                        | Positive    | −0.010      | 0.900   | 291 |
|                                                                                       | Negative    | −0.006      | 0.942   | 291 |
| <i>Panel C: Classification and Quality Concerns</i>                                   |             |             |         |     |
| Quasi-exp. papers only                                                                | Positive    | −0.598      | 0.121   | 139 |
|                                                                                       | Negative    | −0.656*     | 0.081   | 139 |
| + log(abstract length)                                                                | Positive    | 0.008       | 0.968   | 291 |
| [abstract length control]                                                             | Negative    | 0.011       | 0.958   | 291 |

*Notes.* Each row reports the probit latent-index coefficient on finding direction, except the LPM row, which reports an OLS coefficient estimated on the same sample. Probit coefficients are on the latent-index ( $\Phi^{-1}$ ) scale and do not by themselves give a percentage-point change in publication probability; the implied publication-probability change is  $\Phi(\hat{\alpha} + \hat{\beta}_d + \hat{\gamma}'\mathbf{X}^{\text{base}}) - \Phi(\hat{\alpha} + \hat{\gamma}'\mathbf{X}^{\text{base}})$  at a covariate baseline  $\mathbf{X}^{\text{base}}$  (the centered baseline defined in Section 4.1 is used elsewhere in the paper). The LPM (OLS) row, by contrast, has a direct percentage-point interpretation: an LPM coefficient of −0.010 means a 1.0 percentage-point decrease in publication probability for the indicated direction category relative to null/mixed papers. All specifications include a linear year trend unless noted. Standard errors omitted for brevity. *Panel A:* The first row drops monetary policy papers. The “*Journal of Finance* membership” and “*Review of Financial Studies*+*Journal of Financial Economics* membership” rows hold the full 291-paper sample fixed and redefine the dependent variable as an indicator for placement in the indicated outlet (so the reference group includes the other journals and NBER); these are target-journal-membership tests, not journal-vs-NBER subsample regressions. Direction coefficients are small and insignificant in all cases. *Panel B:* “Binary directional” collapses positive and negative findings into a single directional category; adding author-count controls leaves results unchanged. “LPM” re-estimates via OLS with robust standard errors, giving percentage-point differences in publication rates directly; results remain negligible. *Panel C:* “Quasi-exp. only” restricts to the 139 papers using credible causal methods (natural experiments, IV, DiD, RD), holding research design quality constant across direction categories. “+ log(abstract length)” controls for abstract length only; the row tests the length channel of the writing-style concern and does not control for per-word hedging style. See Section 7.4 and the body discussion above for the broader scope. \*\*\* $p < 0.01$ , \*\* $p < 0.05$ , \* $p < 0.10$ .

## **Appendix A. Keyword Taxonomy (Online Appendix OA.1)**

The complete seven-category keyword taxonomy used at the first-pass classification stage (financial regulation, securities law, banking policy, monetary/fiscal policy, market microstructure rules, corporate governance law, and general intervention signals) and the co-occurrence rules for ambiguous standalone terms are provided in Online Appendix OA.1 to conserve space.

## **Appendix B. LLM Prompts (Online Appendix OA.2)**

The verbatim in-scope classification prompt and the two direction-coding prompts (abstract-only and title-only fallback), including system instructions and JSON return-schema specifications, are reproduced in Online Appendix OA.2. The deployed pipeline is implemented in `scripts/03_classify_inscope.py` (in-scope) and `scripts/04_code_direction.py` (direction); these scripts and the prompts they invoke are available in the replication package.

## Appendix C. Replication Confusion Matrix

**Table A1.** LLM Direction Coding: Replication Confusion Matrix

| Original code  | Replication code |          |           | Row total |
|----------------|------------------|----------|-----------|-----------|
|                | −1               | 0        | +1        |           |
| Negative (−1)  | <b>20</b>        | 0        | 0         | 20        |
| Null/Mixed (0) | 9                | <b>5</b> | 6         | 20        |
| Positive (+1)  | 2                | 0        | <b>18</b> | 20        |
| Column total   | 31               | 5        | 24        | 60        |

*Notes.* To assess how consistently the AI classifies paper findings, we drew a random sample of 60 papers (20 from each direction category) and ran the direction-coding prompt a second time using the same model (claude-haiku-4-5-20251001) in an independent second call. Each cell shows how many papers the two runs agreed or disagreed on. Rows are the original codes used in the main analysis; columns are the codes from the second independent run. Bold entries on the main diagonal are papers where both runs assigned the same code—these are the agreements. Off-diagonal entries are disagreements. Overall, the two runs agree on 43 of 60 papers (71.7%). Cohen’s weighted  $\kappa_w = 0.674$  (linear weights), which the Landis and Koch (1977) scale classifies as “substantial agreement.” Agreement is perfect for negative-finding papers (20/20 = 100%) and high for positive-finding papers (18/20 = 90%), but much lower for null/mixed papers (5/20 = 25%). When the two runs disagree on a null-coded paper, the replication almost always re-classifies it as directional, suggesting that abstracts with inconclusive findings are harder for the AI to classify consistently. The corpus breakdown of these null-category flips matters for signing the measurement-error bias. The 20 originally null-coded validation papers comprise 5 NBER working papers and 15 published articles (*Journal of Finance*: 5, *Review of Financial Studies*: 5, *Journal of Financial Economics*: 5). Of the 15 papers re-coded as directional in replication, 2 are NBER and 13 are published. The direction of bias in our estimated direction premium depends on which side dominates: NBER directional papers miscoded as null would inflate the premium (and make our null finding conservative), while published directional papers miscoded as null would deflate the premium (and make the null finding easier to obtain). The replication evidence shows null-coding instability is concentrated on the *published* side, so we do not claim our estimates are conservative; the validation establishes instability of the null category but does not establish a clean sign for the residual measurement error.

## Appendix D. Caliper Robustness to Coefficient Inclusion

A potential concern is that pooling all regression-table rows—including standard controls—could inflate the caliper count without reflecting publication pressure on headline results. Table A2 addresses this by re-running the caliper test after dropping rows whose variable name matches common control patterns: firm size, leverage, profitability (ROA/ROE/EBITDA),

book-to-market, Tobin’s  $q$ , return volatility, momentum, cash flow, market capitalization, intercepts, year dummies, and fixed-effect indicators, as well as macro controls (GDP, inflation, interest rates). The retained “non-control” rows are predominantly treatment indicators, DiD interactions, and main-result variables. The restriction drops fewer than 10% of rows (7,055 of 7,761 published and 7,599 of 8,351 NBER statistics are retained). Caliper ratios are essentially unchanged: the published full-sample ratio remains 1.203 and the negative-finding ratio *increases* to 1.625 (from 1.522), while NBER ratios stay near 1.0. As in the main analysis, the published-side results clear the naive chi-square at very small  $p$ -values but do not reach 0.05 on the paper-clustered bootstrap ( $p = 0.161$  for the non-control full sample,  $p = 0.121$  for the non-control negative-finding subsample); the \*\*\* stars in the table reflect the naive chi-square only. The pattern that the caliper result is not an artifact of pooling control coefficients therefore holds at the level of the point estimates and the naive test, but does not strengthen the result on the preferred test.

## **Appendix E. Caliper Test Stratified by Row Type: Coef+SE vs. Standalone $z$ -Statistics**

Regression tables present statistical results in two ways. Some rows show a coefficient alongside its standard error, which allows a reader to independently verify the test statistic by dividing one by the other. Other rows report only a  $t$ - or  $z$ -statistic directly, with no coefficient or standard error shown, making it impossible to cross-check the number. Both types appear routinely in published finance papers. Table A3 asks whether the threshold-clustering pattern differs between these two types of rows—that is, whether the excess mass just below the 5% significance cutoff is concentrated in the rows readers can verify or in the rows they cannot.

**Table A2.** Caliper Test: All Statistics vs. Non-Control Statistics

| Sample    | Subset           | Filter      | $N$   | Papers | Ratio    | Naive $p$ | Perm. $p$ | Boot. $p$ |
|-----------|------------------|-------------|-------|--------|----------|-----------|-----------|-----------|
| Published | Full sample      | All stats   | 7,761 | 108    | 1.203*** | 0.002     | 0.001     | 0.150     |
| Published | Full sample      | Non-control | 7,055 | 108    | 1.203*** | 0.003     | 0.001     | 0.161     |
| Published | Negative-finding | All stats   | 3,491 | 50     | 1.522*** | <0.001    | <0.001    | 0.135     |
| Published | Negative-finding | Non-control | 3,139 | 50     | 1.625*** | <0.001    | <0.001    | 0.121     |
| Published | Positive-finding | All stats   | 2,633 | 37     | 0.941    | 0.565     | 0.732     | 0.623     |
| Published | Positive-finding | Non-control | 2,409 | 37     | 0.927    | 0.481     | 0.769     | 0.637     |
| NBER      | Full sample      | All stats   | 8,351 | 106    | 1.012    | 0.849     | 0.424     | 0.432     |
| NBER      | Full sample      | Non-control | 7,599 | 104    | 0.981    | 0.767     | 0.625     | 0.561     |
| NBER      | Negative-finding | All stats   | 4,375 | 50     | 0.996    | 0.964     | 0.532     | 0.474     |
| NBER      | Negative-finding | Non-control | 3,981 | 50     | 0.955    | 0.633     | 0.693     | 0.614     |

*Notes.* Each paper’s regression tables contain two kinds of rows: (1) rows for the main variable of interest (e.g., a treatment indicator or policy dummy), and (2) rows for standard control variables (e.g., firm size, leverage, profitability, GDP). The “All stats” rows use every regression coefficient; the “Non-control” rows keep only the main-variable rows, dropping controls such as size, leverage, ROA/ROE, book-to-market, GDP, inflation, firm age, return volatility, fixed-effect indicators, and intercepts. The *Ratio* column is  $R = \#\{z \in [1.645, 1.96)\} / \#\{z \in [1.96, 2.275]\}$ .  $R = 1.0$  means equal frequency;  $R > 1.0$  indicates excess clustering on the just-below-1.96 side (just-insignificant results), and  $R < 1.0$  indicates excess clustering on the just-above-1.96 side (just-significant results). *Naive p*: standard chi-squared test. *Perm.*: permutation test that accounts for the fact that statistics from the same paper are correlated (10,000 iterations). *Boot.*: bootstrap that gives each paper equal weight regardless of how many statistics it contributes (10,000 draws). \*\*\*  $p < 0.01$ , \*\*  $p < 0.05$ , \*  $p < 0.10$  refer to the *naive* chi-square only; under the paper-clustered bootstrap the published-side  $p$ -values are 0.121–0.161 and do not reach conventional significance. Point estimates of the ratio are virtually identical whether or not control rows are included, and the clustering among negative-finding papers strengthens (in point estimate) when controls are dropped.

## Appendix F. Sample Frame Validation (Online Appendix OA.3)

Three precision/recall estimates bound the sample-frame quality. (i) *LLM precision on its in-scope decisions* is  $58/68 = 85.3\%$  (manual recheck file). (ii) *LLM recall on the low-confidence out-of-scope NBER pool* is  $1 - 126/803 = 84.3\%$  (NBER-side refilter). (iii) *Keyword-filter false-negative rate*: a 200-paper audit of the 37,656-paper keyword-negative pool (with  $\geq 20$ -word abstracts; random seed 42) flags 4 as in-scope under the same LLM prompt (2.0%, 95% binomial CI 0.5% to 4.0%), implying roughly 753 in-scope misses (95% CI 188–1,506) before adjusting for the LLM’s  $\approx 85\%$  precision and roughly 640 (95% CI 160–1,280) after. Full audit details, per-paper verdicts, and scripts are in Online Appendix OA.3 and the replication package (`scripts/31.keyword_recall_audit.py`,

**Table A3.** Caliper Test Stratified by Row Type

| Corpus    | Row type       | $N$   | Below | Above | Ratio                 | Naive $p$ |
|-----------|----------------|-------|-------|-------|-----------------------|-----------|
| Published | All rows       | 7,761 | 611   | 508   | 1.203*** $\uparrow$   | 0.002     |
| Published | Coef+SE only   | 5,563 | 322   | 394   | 0.817*** $\downarrow$ | 0.007     |
| Published | $z$ -stat only | 2,198 | 289   | 114   | 2.535*** $\uparrow$   | $<0.001$  |
| NBER      | All rows       | 8,351 | 501   | 495   | 1.012                 | 0.849     |
| NBER      | Coef+SE only   | 6,840 | 422   | 445   | 0.948                 | 0.435     |
| NBER      | $z$ -stat only | 1,511 | 79    | 50    | 1.580*** $\uparrow$   | 0.011     |

*Notes.* Regression tables report statistical results in two ways. Some rows report a coefficient and its standard error side by side, which allows the test statistic to be independently verified by dividing one by the other (*Coef+SE only* rows). Other rows report only a  $t$ - or  $z$ -statistic directly, with no coefficient or standard error shown, so there is no way to cross-check the number ( $z$ -stat only rows). The extraction treats every such number as a two-sided  $|z|$ -test against  $H_0 = 0$ ; a small share of standalone rows are flagged in the underlying data as  $F$ - or  $\chi^2$ -statistics, one-sided tests, or tests against non-zero nulls and would not be comparable to the 1.96 cutoff used in the caliper window, but we do not separately filter these out (Section 5.1). The *Ratio* column is  $R = \#\{|z| \in [1.645, 1.96]\} / \#\{|z| \in [1.96, 2.275]\}$ .  $R = 1.0$  means equal frequency;  $R > 1.0$  indicates excess clustering on the just-below-1.96 side, and  $R < 1.0$  indicates excess clustering on the just-above-1.96 side. *Below* and *Above* are the raw counts of test statistics falling in the windows just below and just above the 5% cutoff. *Naive p* tests whether those two counts are equal. Significance markers are two-sided: \*\*\* $p < 0.01$ , \*\* $p < 0.05$ , \* $p < 0.10$  for the test that the below and above counts are equal. Direction arrows indicate which side has the excess:  $\uparrow$  flags  $R > 1$  (excess on the just-below-1.96 side, the focal pattern), and  $\downarrow$  flags  $R < 1$  (excess on the just-above-1.96 side). The published Coef+SE-only row is significant at  $p = 0.007$  in the  $\downarrow$  direction: among verifiable rows in published papers, the sample shows more statistics in  $[1.96, 2.275]$  than in  $[1.645, 1.96]$ . This is the classical Brodeur-style “just-cleared-the-bar” pattern, in the opposite direction from the headline excess driven by standalone  $z$ -statistics; we describe both directions in Section 5.1. Across the table, threshold-side asymmetry is concentrated in the standalone  $z$ -statistics, while the verifiable rows show a smaller asymmetry in the opposite direction.

output/tables/keyword\_recall\_audit.csv).

## Appendix G. Validation of Core LLM-Coded Variables

Two LLM-coded variables drive the main inference: the direction code that defines the probit dependent groups and the QE indicator that produces the headline 18–20 percentage-point research-design premium. This appendix audits each with a stricter coding prompt on a 30-paper balanced sample (15 published + 15 NBER-only). Scripts and per-paper verdicts: scripts/32\_core\_variable\_validation.py, output/tables/null\_stability\_audit.csv, output/tables/qe\_validation\_audit.csv.

### (A) Null-coded paper stability

The Appendix C confusion matrix shows null/mixed (0) is the least stable category, with only 5/20 (25%) agreement between the two original LLM runs. Because the null category is the *baseline* in the probit (positive and negative are measured *relative* to null), instability in null coding could either attenuate or bias both direction coefficients.

We re-code a stratified sample of 30 null-coded papers (15 published, 15 NBER-only; random seed 42) using a stricter prompt that requires the abstract to explicitly state *no significant effect*, mixed evidence, or refusal to take a directional stance for the paper to remain coded 0; hedged language with a substantive directional claim flips the code to +1 or −1. Under this stricter rule, 10 of 15 published null-coded papers (66.7%) and 10 of 15 NBER-only null-coded papers (66.7%) flip to directional.

The flip rate is large but, critically, *symmetric across corpora*: the within-corpus null-instability is the same on both sides. This has a clean implication for our headline directional null result: null-misclassification injects direction-neutral noise into the baseline category but does not bias the cross-corpus comparison of directional categories. In particular, the equality test  $H_0: \beta_1 = \beta_2$  (Section 4.1) does not depend on the integrity of the null baseline—it compares the positive and negative coefficients to each other—and gives  $\chi^2 = 0.005$ ,  $p = 0.944$  in our primary specification. We therefore read the directional null as robust to null-coding instability.

### (B) QE indicator validation

The QE indicator is AI-coded from abstracts and—unlike the direction code—was not separately validated in the original pipeline (we noted this in Section 5.2 and Section 7.4). A referee raised the specific concern that the 18–20 pp QE premium might capture differences in whether published abstracts *explicitly name* their identification strategy (“we use a regression discontinuity design”) rather than differences in actual research-design quality, because published abstracts tend to be sharper and more declarative than NBER drafts.

We test this directly by re-coding a stratified sample of 30 QE-coded papers (15 published, 15 NBER-only; random seed 42) under a strict rule: a paper qualifies as QE only if its abstract *names* an identification strategy (RDD, IV, DiD, natural experiment, or event study with explicit causal claim) AND uses that strategy as the *primary* empirical method, not as a robustness check.

Results. Of 15 QE-coded *published* papers, only 5 (33.3%) survive the strict rule; the other 10 are re-coded as not-QE. Of 15 QE-coded *NBER-only* papers, 8 (53.3%) survive. The strict survival rate is therefore higher on the NBER side than on the published side—the *opposite* of the directional asymmetry the referee’s hypothesis predicted. Extrapolating to the full QE pool (78 published + 61 NBER), the strict-QE published share would be  $26/(26 + 33) = 44.4\%$ , against an original 56.1% and a sample overall published share of 46.0%. Most of the original 18–20 pp QE premium attenuates toward zero under the strict coding rule.

We interpret this as evidence that a meaningful fraction of the documented QE premium reflects looser LLM coding of abstract-style cues—explicit-method-naming language, declarative framing—rather than differences in true research-design quality. We therefore do not treat the magnitude of the QE premium as a precisely measured quantity, and we read its sign (positive) and direction-independent character as more robust than its specific 18–20 pp point estimate. The audit also suggests that QE coding from abstracts is less reliable than direction coding from abstracts and would benefit from a future manual-coding exercise on the full 291-paper sample. The headline finding—that finding direction does not predict corpus membership—does not depend on the QE indicator’s calibration and is unaffected.