

Do Finance Journals Pick Sides? Evidence from Government Intervention Research

Emre Kuvvet*

June 18, 2026

ABSTRACT

We study whether top finance journals favor certain policy conclusions. Using 134 published articles in the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* alongside 157 NBER working papers on government intervention in financial markets (2000–2024), we find no evidence that journals sort on finding direction at scales our design can detect. Whether a paper concludes that a regulation is beneficial, harmful, or inconsequential, its share of the combined 291-paper corpus that sits on the published side is roughly the same, about 46 percent. This is a within-direction corpus-membership share, not a journal acceptance rate (which would require observing submissions and editorial decisions neither of which is in our data). The two largest cross-sectional predictors of corpus membership are research-design quality and institutional prestige, not finding direction: papers using quasi-experimental methods are about 18–20 percentage points more likely to appear in the published corpus, and papers from UTD Top 100 institutions are about 24–26 percentage points more likely, conditional on the same observed controls (year, team size, finding direction). The two premia are essentially orthogonal ($r = 0.015$) and both survive sequential controls; they are consistent with—but do not prove—an editorial preference for credible causal identification and institutional pedigree, since our cross-sectional design cannot rule out that QE and top-institution papers differ on unobserved quality dimensions. We also find that published papers show modest threshold bunching in supplementary test statistics but not in their main regression tables, suggesting that any specification searching is concentrated at the margins of the paper rather than in its core results.

JEL classification: G18, G28, C80, B40.

*Department of Finance, Nova Southeastern University. E-mail: ekuvvet@nova.edu.

Keywords: Publication bias; Government intervention; Finance journals; Selective reporting;
Regulatory finance; p-hacking

1. Introduction

Academic research on government intervention in financial markets occupies a privileged position in public policy. The Federal Reserve Board, the Securities and Exchange Commission, and the Basel Committee on Banking Supervision routinely incorporate peer-reviewed finance research into regulatory impact assessments, rule-makings, and policy speeches. The Dodd-Frank Wall Street Reform and Consumer Protection Act alone generated an estimated 400 new regulatory mandates, many of which were designed in explicit reference to academic evidence on the effects of capital requirements, disclosure rules, and consumer financial protection. The credibility of this evidence base is therefore a first-order policy question: if the academic record systematically misrepresents which interventions have measurable effects—and which do not—regulators face a distorted picture that can generate welfare-reducing policy choices.

Publication bias offers a mechanism by which such distortion arises without any deliberate misconduct. When researchers are reluctant to write up null results, when journals are more likely to accept papers with statistically significant findings, or when both effects operate simultaneously, the published literature understates the true prevalence of null or inconclusive evidence (Franco, Malhotra and Simonovits, 2014). In a domain such as government intervention research, where policy decisions depend on whether a regulation has a measurable causal effect, suppression of null results generates the false impression that most interventions work—or, more precisely, that most interventions have discernible effects in one direction or the other.

Why might directional publication bias arise in the government-intervention literature specifically? Two theoretical channels operate in opposite directions, making the sign of any bias an open empirical question. A *pro-intervention channel* could arise if journals and audiences favor results that confirm a policy’s stated objective (e.g., that capital requirements reduce bank risk), both because positive results are more actionable for regulators and because they more naturally confirm the prior that motivated the policy. An *anti-intervention*

channel could arise if research that documents unintended costs or failures of regulation is more surprising—and thus more publishable—than research confirming that a policy works as intended. Under either channel, the direction of sorting would differ from significance bias (null vs. directional), which reflects a simpler threshold between finding-something and finding-nothing. This paper is designed to detect both types.

This paper studies sorting patterns in a domain the existing literature has not examined: the finance literature’s treatment of government intervention. Prior work on significance bias in finance has focused on the cross-section of asset pricing anomalies, documenting that factor-zoo proliferation reflects selection for high t-statistics (Harvey, Liu and Zhu, 2016; Hou, Xue and Zhang, 2020; Jensen, Kelly and Pedersen, 2023). A parallel economics literature (Brodeur, Lé, Sangnier and Zylberberg, 2016; Andrews and Kasy, 2019) infers bias from the shape of published p-value distributions, finding excess mass just below conventional significance thresholds. Neither strand addresses the policy-intervention literature, nor do they distinguish two conceptually distinct forms of bias with different welfare implications. *Significance bias*—non-null findings over-published relative to null/inconclusive ones—makes every intervention appear detectable. *Directional bias*—pro-intervention findings over-published relative to anti-intervention findings, conditional on being non-null— additionally skews the regulatory record toward a particular conclusion about the sign of policy effects. These two biases map onto distinct empirical tests in our analysis: the across-corpus probit with separate Positive and Negative indicators against the null baseline (equation (1)) speaks to both at once—the $\beta_1 = \beta_2$ test asks about directional bias, while the joint test $\beta_1 = \beta_2 = 0$ asks about significance bias—and the within-paper caliper test (Section 5.1) speaks to within-paper selective reporting near the significance threshold, a related but distinct phenomenon from across-paper editorial selection. We design our study to address all three.

Our research design combines two complementary samples. The *published paper* sample consists of all in-scope articles in the *Journal of Finance*, the *Review of Financial Studies*,

and the *Journal of Financial Economics* from 2000 to 2024. The *working paper* sample is drawn from all NBER working papers published from 2000 to 2024 (30,212 papers across every program), with in-scope papers identified through the same keyword filter and LLM classification applied to the journal sample. No program restriction is imposed; topic relevance alone determines inclusion. The key advantage of this design is that we observe findings from papers that circulated *before* potential journal publication, allowing a direct two-corpus comparison rather than inferring the pre-publication distribution from the shape of published p-values alone.

Because only 12% of published papers are matched to a confirmed NBER counterpart, the two corpora are largely non-overlapping cross-sections rather than the same papers tracked through the editorial filter. We therefore interpret all estimates as cross-sectional associations between paper characteristics and corpus membership, not as causal editorial-filter effects; the full identification discussion is in Section 7.4.

We define a paper as in-scope if its primary research question concerns the causal effect of a specific government law, regulation, or public policy on financial markets, firms, or investors. Identifying these papers across nearly 40,000 documents required an automated approach: a keyword filter followed by a large language model (LLM) classifier that reads each abstract and assigns an in-scope judgment along with a directional code—positive (+1, the intervention was beneficial), negative (−1, harmful), or null/mixed (0).

Our central result is that finding direction does not predict published-versus-NBER-only corpus membership at the scale our design can detect. Within each direction category—papers concluding that a regulation helps, papers concluding it hurts, and papers finding no significant effect—the share that sits on the published side of our combined 291-paper sample is essentially the same, around 46 percent. We emphasize that this 46 percent is a within-direction corpus-membership share, computed as $\text{published}/(\text{published} + \text{NBER-only})$ within each direction category, not a journal acceptance rate; the latter would require observing the submission pool and editorial decisions, neither of which is in our data. This

uniformity in corpus-membership shares holds unconditionally and survives every regression specification we estimate, including controls for publication year, decade, team size, and research design quality. No detectable direction premium is found for positive findings, for negative findings, or when positive and negative are pooled against null.

What does predict published-side corpus membership is how a paper was designed. Papers using quasi-experimental identification strategies—regression discontinuity, instrumental variables, difference-in-differences, natural experiments, or event studies with a causal interpretation (full definition in Section 5.2)—are 18–20 percentage points more likely to sit on the published side of our corpus than observational papers on the same policy questions. This methodological premium is the only statistically robust predictor we find; finding direction adds nothing beyond it. We interpret this cautiously: the quasi-experimental indicator likely captures a bundle of correlated quality attributes—author prestige, institutional resources, topic fashionability—not identification strategy alone.

Within published papers, we find a secondary pattern. Test statistics bunch modestly just below the conventional 5 percent significance threshold, but only in standalone statistics that appear primarily in supplementary and robustness tables. Primary regression tables—where both a coefficient and a standard error are reported and can be independently verified—show no bunching and are statistically indistinguishable from the corresponding NBER working paper tables. This distinction is informative: it suggests that any selective reporting occurs at the margin of revision, in the choice of which additional checks to highlight, rather than in the core empirical results that define the paper’s contribution.

This paper makes three contributions. First, it provides direct evidence that top finance journals do not sort government-intervention research by the direction of its conclusions at the scale detectable by our design. Second, it documents that identification quality and institutional prestige—not finding direction—are the two largest cross-sectional predictors of corpus membership in our sample: conditional on observed controls (year, team size, direction), papers using credible causal designs are about 18–20 percentage points more

likely to appear in the published corpus, and papers from UTD Top 100 institutions are about 24–26 points more likely; the two premia are essentially orthogonal to each other ($r = 0.015$) and both survive sequential controls. We do not interpret this as a causal estimate of editorial weighting, but the pattern speaks to ongoing debates about editorial standards and the replication crisis in finance. Third, it introduces a validated LLM-assisted pipeline for classifying and direction-coding large paper corpora, with code and prompts publicly available.

The remainder of the paper proceeds as follows. Section 2 reviews the related literature. Section 3 describes data construction and the classification pipeline. Section 4 presents the empirical strategy. Section 5 reports main results and economic magnitudes. Section 6 presents robustness checks. Section 7 discusses mechanisms, policy implications, comparisons to other domains, and limitations. Section 8 concludes.

2. Related Literature

We use four terms with distinct meanings throughout, and each maps to a specific empirical test. *Significance bias* refers to the selective suppression of statistically null results, giving the published record an inflated share of non-null findings regardless of their substantive direction; in our analysis it corresponds to a joint test that the Positive and Negative coefficients in equation (1) are both zero (equivalently, the binary “Directional vs. null” coefficient in Column (1) of Table 8). *Directional bias* refers to editorial selection based on the substantive sign of a paper’s conclusion—whether a policy is found beneficial or harmful—conditional on the finding being non-null; it corresponds to the equality test $\beta_1 = \beta_2$ in our primary Column (2) specification. *Selective reporting* refers to within-paper researcher behavior: the choice of which specifications, robustness checks, or standalone test statistics to highlight in the final manuscript based on proximity to a significance threshold; in our analysis it corresponds to the caliper test in Section 5.1. *Editorial selection* is our umbrella term for any mechanism by which the submission-to-publication pipeline shapes the composition of the

published literature. The probit tests speak to significance and directional bias; the caliper test speaks to within-paper selective reporting; the significance-bias literature provides the methodological benchmark for both.

This paper relates to three strands of research. The first is the literature on significance bias in economics and finance. Brodeur, Lé, Sangnier and Zylberberg (2016) document excess mass just above the $|z| = 1.96$ significance cutoff (equivalently, on the $p < 0.05$ side of the threshold) in the top economics journals—the classical p-hacking pattern. Our caliper analysis uses the same windows but the inverse ratio convention; we describe both directions explicitly in Section 5.1. Franco, Malhotra and Simonovits (2014) show, using NSF-funded pre-registered experiments, that null results are published at roughly one-third the rate of significant ones. In finance, Harvey, Liu and Zhu (2016) show that given the scale of data mining in the factor literature, conventional significance thresholds are far too low and most claimed factors are likely false discoveries under multiple testing; Hou, Xue and Zhang (2020) confirm this empirically, finding that between 65% and 82% of published anomalies (depending on the test and sample) fail to clear conventional significance thresholds under uniform replication standards, pointing to upward-biased published effect sizes. These papers test for *significance* bias—whether null results are suppressed relative to significant ones. Our paper asks a different question: whether published findings are sorted by the *direction* of the policy conclusion, which is only meaningful in a domain where findings have a clear normative interpretation. We also contribute a direct pre-publication baseline by using NBER working paper texts, rather than inferring the pre-publication distribution from the shape of published p-values alone.

The second strand is the literature on government intervention in financial markets. A large body of empirical work studies the effects of securities law, banking regulation, and macroprudential policy (La Porta, Lopez-de Silanes, Shleifer and Vishny, 1998; Berger and Bouwman, 2017; Agarwal, Lucca, Seru and Trebbi, 2014), producing heterogeneous findings across interventions and time periods. Whether the published record faithfully represents

the underlying distribution of evidence on these questions has not previously been examined. Our paper fills this gap.

This literature is particularly susceptible to directional publication bias for three reasons. First, the direction of findings is normatively salient in a way that it is not for most empirical finance research. A paper asking whether momentum predicts returns has no obvious ideological valence; a paper asking whether Dodd-Frank improved financial stability does. When findings carry clear policy implications, editors, referees, and authors arrive with directional priors that can generate asymmetric selection (Andrews and Kasy, 2019). Second, the academic finance community has historically leaned skeptical of government intervention, reflecting the efficient-markets tradition. If referees hold this prior, the Andrews and Kasy (2019) model predicts that papers with positive (+1, pro-intervention) findings face higher scrutiny than papers with negative (−1, anti-intervention) findings, generating editorial selection toward over-publication of negative findings relative to positive ones. We are agnostic about how the prior maps to scrutiny of null/mixed (0) papers in our coding, which contain no clear directional signal and are not the focus of the prior-driven asymmetric-selection prediction. Post-2008, the field has become more receptive to intervention research, but the net direction of any ideological filter remains unclear. Third, papers in the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* are routinely cited in regulatory proceedings and policy debates; papers with clear pro- or anti-intervention conclusions are more actionable for policymakers than null results, potentially creating an editorial demand for directional findings independent of their statistical strength. These three mechanisms—normative salience, ideological sorting, and policy-venue citation pressure—jointly predict directional sorting in this literature while leaving the sign of any bias theoretically ambiguous. Our paper provides the first direct test using a pre-publication benchmark, resolving what theory alone cannot.

The third strand is the growing use of large language models for text classification in economics (Ash and Hansen, 2023; Korinek, 2023). We contribute a validated two-pass

pipeline—keyword filter followed by LLM classification—for identifying and direction-coding government-intervention finance papers at scale, and we release the coding rubric and prompts publicly.

3. Data and Sample Construction

3.1 Published Papers

The published-paper sample consists of all articles appearing in the *Journal of Finance*, the *Review of Financial Studies*, and the *Journal of Financial Economics* from 2000 to 2024. These three journals constitute the traditional top tier of academic finance and account for the majority of highly cited finance research on government intervention. We collect article metadata—title, authors, year, volume, issue, DOI, and abstract—for all articles in each journal over this period.

Journal articles are identified and retrieved as follows. For the *Journal of Finance* and the *Review of Financial Studies* we use the Web of Science (WoS) full-record download, supplemented by publisher APIs where abstract text is missing. For the *Journal of Financial Economics*, abstracts were obtained via the Elsevier ScienceDirect Article Retrieval API. Abstract text is available for 100% of research articles in all three journals after manual abstract enrichment. The *Journal of Finance* denominator is restricted to research articles only; the raw WoS record includes editorial matter, meeting announcements, and AFA administrative items that carry no abstracts, and our pipeline excludes these before computing coverage.

3.2 NBER Working Papers

The working-paper sample is drawn from the NBER Working Paper Series via the public API. We collect *all* working papers published from 2000 to 2024 across every NBER program (30,212 papers in total) and identify in-scope government-intervention papers through the keyword filter and LLM classification described in Section 3.3. No program restriction is

applied at the collection stage; the topic filter alone determines inclusion. Government intervention finance research is concentrated in five programs— Corporate Finance (CF), Asset Pricing (AP), Monetary Economics (ME), Economic Fluctuations and Growth (EFG), and Public Economics (PE)—but papers assigned to neighboring programs (e.g., Economic History or International Finance and Macroeconomics) are automatically captured if they pass the scope filter.

We use the NBER working paper sample as a reference distribution for comparison with the published sample. A causal editorial-selection interpretation would require two distinct conditions: (i) NBER selection into working-paper circulation is itself direction-neutral, so that the NBER distribution is not pre-filtered on finding direction; and (ii) conditional on direction, the NBER sample approximates the relevant top-journal submission pool on other factors jointly related to direction and publication (e.g., research-design quality, topic, author network, data access, vintage). Both conditions are strong. NBER affiliation is selective on researcher quality and institutional prestige, and many finance papers circulate through SSRN rather than NBER; the NBER distribution may therefore differ from the full pre-publication distribution for reasons unrelated to editorial selection. We also note that “pre-publication” here means circulated as an NBER working paper at some point—NBER posting dates establish working-paper circulation but do not establish that this circulation preceded journal submission, since papers may be posted to NBER before, alongside, or after a submission decision. We therefore read all results in this paper as cross-sectional associations between corpus membership and paper characteristics, not as causal estimates of editorial filtering; we return to this identification challenge formally in Section 4.1.

3.3 In-Scope Classification Pipeline

Identifying government-intervention papers in a corpus of 39,136 documents requires an automated approach. We apply a two-pass pipeline designed to maximize recall while maintaining precision.

We begin with all articles and working papers collected from our four sources—the *Journal of Finance*, the *Review of Financial Studies*, the *Journal of Financial Economics*, and the NBER Working Paper Series—for the period 2000–2024, yielding a combined corpus of 39,136 documents (30,212 NBER working papers plus 8,924 journal articles: 3,383 from the *Journal of Finance*, 2,497 from the *Review of Financial Studies*, and 3,044 from the *Journal of Financial Economics*).

In the first pass, a paper passes the keyword filter if the combined text of its title and abstract contains at least one term from a pre-specified taxonomy of government intervention keywords spanning seven categories: financial regulation (e.g., “Dodd-Frank,” “capital requirements,” “Basel”), securities law (e.g., “SEC,” “Sarbanes-Oxley,” “insider trading law”), banking policy (e.g., “FDIC,” “deposit insurance,” “TARP”), monetary and fiscal policy (e.g., “quantitative easing,” “fiscal stimulus”), market microstructure rules (e.g., “uptick rule,” “tick size pilot”), corporate governance law (e.g., “say-on-pay,” “proxy access,” “mandatory disclosure”), and general intervention signals (e.g., “government intervention,” “public policy,” “law and finance”). Ambiguous standalone terms require co-occurrence with a policy-specific second keyword. The full taxonomy is reported in Appendix A.

Searching the combined title-and-abstract text rather than the abstract alone improves recall for three reasons. First, authors frequently name the specific law or regulation in the title but refer to it only as “the reform” or “the rule” throughout the abstract, so the keyword would appear only in the title. Second, common abbreviations (e.g., “Dodd-Frank” in the title) and full statutory names (e.g., “Wall Street Reform and Consumer Protection Act” in the abstract) are not always used interchangeably; searching both fields catches either form. Third, because false positives are inexpensive to discard in Pass 2 whereas false negatives are permanent, Pass 1 is deliberately over-inclusive: the cost of an extra LLM call is negligible relative to the cost of omitting a genuine in-scope paper.

The keyword filter reduces the corpus from 39,136 documents to 1,849 candidates (1,313 NBER, 182 *Journal of Finance*, 153 *Review of Financial Studies*, and 201 *Journal of*

Financial Economics).

In the second pass, each keyword-flagged candidate is passed to Claude with the following instruction (full prompt in Appendix B):

“A finance research paper is ‘in scope’ if and only if its PRIMARY research question is the CAUSAL EFFECT of a specific government law, regulation, or public policy on financial markets, firms, or investors. The following are excluded: (1) purely theoretical papers with no empirical test of a policy; (2) papers that only describe a regulation without estimating its effect; (3) papers whose main topic is private firm or market behavior, with policy as incidental context; (4) papers studying central bank communication without a specific policy instrument; and (5) papers whose primary policy variable is a rule imposed by a stock exchange rather than a government or regulatory mandate.”

For each paper, the model returns an in-scope decision, a confidence level (high, medium, or low), and a one-sentence justification. At this stage the model classifies 343 of the 1,849 keyword-flagged candidates as in scope (NBER: 187, *Journal of Finance*: 44, *Review of Financial Studies*: 50, *Journal of Financial Economics*: 62) and the remaining 1,506 as out of scope.

Following the automated pass, all 343 LLM in-scope classifications are reviewed manually by the author against the paper’s abstract, with the PDF introduction consulted whenever the abstract is ambiguous about whether the policy is the primary research question. The review proceeds in two directions.

(i) *Re-examining LLM in-scope decisions.* The most important judgment call concerns papers that use a policy event only as an identification device to study a downstream outcome that the regulation did not target. We operationalize the boundary at the level of the *outcome variable*:

- A paper is *in scope* if its primary outcome is something the regulation explicitly aimed to affect (its stated regulatory objective)—e.g., for Sarbanes-Oxley, audit quality,

internal-control disclosure, financial-reporting accuracy, restatement frequency, or fraud detection. The paper’s contribution is then a direct test of “did SOX achieve its disclosure objective?”

- A paper is *out of scope* if its primary outcome is downstream of the regulated mechanism but is *not* a stated regulatory objective—e.g., for Sarbanes-Oxley, firm investment, M&A activity, cost of capital, or innovation. The paper’s contribution is then “what is the effect of disclosure (or governance, or liability) on outcome Y , using SOX as an instrument?” These are mechanism papers whose policy variable is incidental.

The criterion turns on whether the regulator stated the outcome as a target, not on whether the outcome plausibly responds to the regulation. Borderline cases—where authors describe both a policy-effect and a mechanism interpretation—are resolved by the position of the outcome in the abstract: if the abstract’s headline result names a stated regulatory objective, the paper is in scope; if the abstract frames the regulation as the source of variation for a non-objective outcome, the paper is out of scope. A second class of overrides drops papers that test a policy whose empirical content turns out to be a private market response (e.g., voluntary cross-listing decisions after a directive). This re-examination produces 52 net exclusions relative to the LLM (NBER: 30, *Journal of Finance*: 18, *Review of Financial Studies*: 2, *Journal of Financial Economics*: 2).

(ii) *Reconsidering LLM out-of-scope decisions in the NBER corpus.* Because false-negative LLM exclusions are difficult to detect without a second pass, the author re-codes 803 NBER papers that the LLM flagged as low-confidence out-of-scope. Of these, 126 are returned to the in-scope pool; 106 survive the same Pass (i) criteria and enter the analysis sample. The published-side LLM decisions are not subjected to a symmetric refilter because the journal denominators (8,924 articles) are small enough that the keyword filter retains all plausible candidates at the first stage.

Borderline cases are resolved by the following rule: if the policy variable is the right-hand-side treatment in the paper’s headline regression, the paper is in scope; if the policy is

used as an instrument or a covariate to study a non-policy outcome, it is out of scope. The full per-paper override log—LLM decision, manual decision, and reason when they differ—is available in the replication package.

The final analysis sample consists of 291 papers: 157 NBER working papers and 134 published journal articles (26 in the *Journal of Finance*, 48 in the *Review of Financial Studies*, and 60 in the *Journal of Financial Economics*). Table 1 reports the number of papers surviving each stage of the selection process.

Table 1. Sample Selection

Step	Papers	Excluded
Starting corpus (NBER + JF/RFS/JFE, 2000–2024)	39,136	
Pass 1: Keyword filter	1,849	37,287
Pass 2: LLM classification & manual review	291	1,558
<i>Final sample</i>	291	
Published articles (JF / RFS / JFE)	134	
Journal of Finance	26	
Review of Financial Studies	48	
Journal of Financial Economics	60	
NBER working papers	157	

Notes. The starting corpus includes all 30,212 NBER working papers (every program, 2000–2024) and all research articles from the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics*, 2000–2024. The keyword filter screens for government intervention terminology (see Appendix A). Pass 2 combines three sub-steps. (a) The LLM in-scope prompt is run on all 1,849 keyword candidates, flagging 217 as in-scope. (b) The author manually reviews each LLM in-scope decision and overrides 22 published articles (*Journal of Finance*: 18, *Review of Financial Studies*: 2, *Journal of Financial Economics*: 2) plus 10 NBER papers as out of scope, retaining 51 of the 61 LLM in-scope NBER papers and 134 of the 156 LLM in-scope published articles. (c) An NBER-side refilter re-codes 803 NBER papers that the LLM had flagged as low-confidence out-of-scope; 126 are returned to the in-scope pool and 106 survive the same manual review criteria. Net of all three sub-steps the working-paper sample is $51 + 106 = 157$ NBER and the published sample is 134 articles, summing to the 291-paper analysis sample.

3.4 Direction Coding

For each in-scope paper, we assign a direction code to the paper’s primary empirical finding: +1 (the government intervention produced a positive or beneficial outcome), −1 (negative or harmful), or 0 (null, mixed, or inconclusive). Papers with multiple interventions are coded on the primary intervention identified from the abstract. Both the in-scope classifier and the direction coder use Claude Haiku (claude-haiku-4-5-20251001; Anthropic).

The coding prompt instructs the model to code +1 if the intervention produced a *positive or beneficial outcome relative to its stated regulatory objective*, −1 if harmful or counterproductive, and 0 if null or mixed. Normative framing by the authors is explicitly excluded. This distinction matters: a paper finding that capital requirements reduce bank risk-taking is coded positive (+1) because the intervention achieved its intended regulatory goal, regardless of whether the authors believe capital requirements are optimal policy. The full coding prompt is in Appendix B.

We validate the direction codes in two ways. First, we re-ran the coding on a stratified 60-paper sample (5 papers per source $\times$ direction cell) using the same model (claude-haiku-4-5-20251001) in an independent second call. The two runs agree on 71.7% of papers ($\kappa_w = 0.674$, “substantial agreement”). Agreement is high for directional papers—100% for negative findings and 90% for positive ones—but only 25% for null/mixed papers; among the 20 originally null-coded papers, 15 were re-coded as directional and only 5 stayed null. These flips are not evenly distributed across corpora: 13 of the 15 originally null-coded *published* articles flipped to directional, against 2 of 5 NBER papers. Because the sign of the resulting bias in our direction premium depends on which corpus the miscoded papers come from—NBER directional-as-null inflates the premium, while published directional-as-null deflates it—the replication exercise establishes that null is the least stable category but does *not* establish a clean sign for the residual measurement error. The full confusion matrix is in Appendix C.

Second, we re-ran coding on 58 papers under three alternative prompt wordings. Agree-

ment with the original codes ranges from 83% to 91% ($\kappa_w = 0.742\text{--}0.904$), all above the model-choice κ_w , confirming the results are not sensitive to how the question is phrased (Table 2).

Table 2. Prompt-Perturbation Robustness: Agreement with Original Direction Codes

Prompt variant	N	Agreed	Agreement %	κ_w
Baseline (original wording)	58	53	91.4	0.904
Variant A (“primary conclusion”)	58	48	82.8	0.742
Variant B (“classify direction”)	58	49	84.5	0.814

Notes. Each variant re-codes the same 58-paper stratified sample using a smaller model with the same classification rules. κ_w is computed against the original direction codes used throughout the paper.

Human-coded validation of a 40-paper subsample is included in the replication package and remains a key additional check; the model-choice $\kappa_w = 0.674$ is the primary reliability evidence reported here.

3.5 Linking NBER Papers to Published Papers

We matched each of the 134 published journal articles to the NBER working paper archive using title similarity, author overlap, and external databases including OpenAlex, CrossRef, and Semantic Scholar. The pipeline identified 16 matched pairs—about 12% of the published sample. The remaining 118 published papers have no NBER counterpart, reflecting that most corporate finance papers circulate through SSRN rather than NBER.

Accounting rule for the 16 matched pairs. We adopt a single auditable rule for the 16 matched-pair NBER drafts and apply it consistently across every table and regression in the paper:

- (i) The 134 *Journal of Finance/Review of Financial Studies/Journal of Financial Economics* articles enter the analysis sample with Published = 1 (the journal-published version is the one counted).

- (ii) The NBER comparison group of 157 papers contains *only* NBER working papers whose published version did not subsequently appear in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics* during 2000–2024. The 16 NBER drafts of papers that did reach a top-three journal are excluded from the 157 by construction.
- (iii) No paper-author-content unit is therefore double-counted. The 291-paper sample partitions cleanly into $134 + 157$ disjoint units, with the 16 matched-pair NBER drafts appearing in neither cell—they are represented only by their journal-published incarnation in the 134-paper Published group.

This rule has three immediate implications. First, the dependent variable in the probit ($\text{Published} \in \{0, 1\}$) is well-defined and mutually exclusive across the 291 observations (matching the Table 3 note). Second, dropping the 16 matched-pair published papers as a robustness check (Section 6.3) removes only those 16 observations and leaves $N = 275$, because their NBER drafts are already excluded; the corresponding NBER counterparts are not separately deletable. Third, the caliper test (Section 5.1) operates on the $|z|$ -statistics extracted from the journal-published version of each matched pair when computing the Published-row caliper, and on the NBER draft for the unmatched NBER-only papers; no test statistic is counted twice. The same rule governs every subsequent table in the paper.

In the regression analysis, the 134 published papers are coded $\text{Published} = 1$ and the 157 NBER-only working papers are coded $\text{Published} = 0$. The direction distributions of the two groups are nearly identical—about 80% of papers in each group have a directional finding—suggesting the two samples are broadly comparable in composition before any regression adjustment.

3.6 Summary Statistics

Table 3 presents summary statistics for the 291-paper analysis sample. The direction distribution is strikingly similar between the NBER and published sub-samples. Among

NBER working papers, 20.4% are coded null or mixed, compared with 20.9% among published papers—a difference of less than 1 percentage point. Positive-finding papers constitute 35.1% and negative-finding papers 44.0% of the published sample, versus 35.7% and 43.9% of the NBER sample. The two samples look nearly the same in their direction distributions.

Table 3. Summary Statistics

Sample	<i>N</i>	Positive (%)	Negative (%)	Null (%)
All papers	291	35.4	44.0	20.6
NBER working papers	157	35.7	43.9	20.4
Published (JF/RFS/JFE)	134	35.1	44.0	20.9
<i>By journal:</i>				
<i>Journal of Finance</i>	26	42.3	46.2	11.5
<i>Review of Financial Studies</i>	48	29.2	54.2	16.7
<i>Journal of Financial Economics</i>	60	36.7	35.0	28.3

Notes. The sample consists of 157 NBER working papers identified from the full NBER corpus (all programs, 2000–2024) via topic filter, and 134 published journal articles (*Journal of Finance*, *Review of Financial Studies*, *Journal of Financial Economics*) classified as in-scope government-intervention papers by our LLM pipeline and manual validation. Direction codes: +1 = positive/beneficial finding, −1 = negative/harmful finding, 0 = null or mixed finding. The NBER and Published samples are mutually exclusive under the accounting rule stated in Section 3.5: the 291 papers consist of 157 NBER working papers whose published version did not subsequently appear in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics* during 2000–2024, and 134 distinct published journal articles, with the 16 matched-pair NBER drafts represented only by their journal-published incarnation. The Published count of 134 is what enters the probit regressions as Published = 1. Column percentages may not sum to 100 due to rounding.

4. Empirical Strategy

4.1 Research Design and Identification

Our empirical strategy compares papers circulated as NBER working papers against papers published in the *Journal of Finance*, the *Review of Financial Studies*, or the *Journal of Financial Economics*, asking whether the direction distribution of findings differs between the two groups. This is not a within-paper design: only 16 of 134 published papers are matched to an NBER working paper, so the two samples are largely different papers selected

through different processes.

A causal editorial-selection interpretation would require both that NBER working papers approximate the direction distribution of papers actually submitted to top finance journals, and that the NBER and published samples be comparable on other factors related jointly to finding direction and publication. Similar working-paper-based approaches are taken by Franco, Malhotra and Simonovits (2014), who use NSF pre-registrations as a pre-publication baseline, and DellaVigna and Linos (2022), who use institutional RCT records. Our design is weaker than either: NBER affiliation is selective on researcher quality and institutional prestige; most finance papers circulate through SSRN rather than NBER; and NBER posting dates do not pin down whether a paper circulated before or after journal submission. The direction of any resulting bias is unclear—NBER selection could either overstate or understate the true editorial gap—so we interpret all results in this paper as cross-sectional associations between corpus membership and paper characteristics, not causal editorial-filter effects. The full identification discussion is in Section 7.4.

4.2 Primary Specification

Our primary test estimates a probit model at the paper level:

$$\Pr(\text{Published}_i = 1) = \Phi(\alpha + \beta_1 \text{Positive}_i + \beta_2 \text{Negative}_i + \boldsymbol{\gamma}' \mathbf{X}_i), \quad (1)$$

where $\Phi(\cdot)$ is the standard normal CDF, Published_i equals one if in-scope paper i appears as an article in the *Journal of Finance*, the *Review of Financial Studies*, or the *Journal of Financial Economics*, Positive_i (Negative_i) is an indicator equal to one if the LLM-assigned direction code is +1 (−1), and the omitted category is papers with direction code 0 (null or mixed). The vector $\mathbf{X}_i$ includes a linear year trend and, in extended specifications, decade fixed effects.

Our primary coefficients of interest are β_1 and β_2 . A finding of $\beta_1 > 0$ ($\beta_2 > 0$) indicates that positive-finding (negative-finding) papers are more likely to be in the published sample

than null-finding papers, conditional on controls. Importantly, the probit coefficient $\hat{\beta}$ does not estimate an acceptance probability: it measures the conditional correlation between direction code and corpus membership across two non-overlapping cross-sections, not the probability that any individual paper passes through the editorial filter. A test of $H_0: \beta_1 = \beta_2$ addresses the directional-bias hypothesis: rejection would indicate that the publication premium depends on the sign of the finding, not merely on statistical significance.

Because the probit model is nonlinear, we report marginal effects. For the bivariate (Column 2) specification, which has no covariates beyond the direction indicators, we report the null-paper baseline difference

$$\Delta_d = \Phi(\hat{\alpha} + \hat{\beta}_d) - \Phi(\hat{\alpha}), \quad d \in \{+, -\}.$$

For controlled specifications (Columns 3–7), the probit index also includes $\hat{\gamma}'\mathbf{X}_i$, so the corresponding expression at a covariate baseline $\mathbf{X}^{\text{base}}$ is

$$\Delta_d(\mathbf{X}^{\text{base}}) = \Phi(\hat{\alpha} + \hat{\beta}_d + \hat{\gamma}'\mathbf{X}^{\text{base}}) - \Phi(\hat{\alpha} + \hat{\gamma}'\mathbf{X}^{\text{base}}).$$

In the controlled specifications we report two summaries. (a) Marginal effects at the centered baseline, defined by $\text{year}_c = 0$ (the sample median year), $\log(1 + N_{\text{authors}}) = 0$ (single-author papers), Quasi-Exp. ID = 0, and Top Institution = 0; this is the baseline reported in the bivariate Column (2) expression above, where the centered controls all enter at zero. (b) Average partial effects (APE), computed by averaging over the in-sample empirical distribution of $\mathbf{X}_i$:

$$\text{APE}_d = \frac{1}{N} \sum_{i=1}^N \left[\Phi(\hat{\alpha} + \hat{\beta}_d + \hat{\gamma}'\mathbf{X}_i) - \Phi(\hat{\alpha} + \hat{\gamma}'\mathbf{X}_i) \right].$$

The two summaries are close in our sample because the centered controls mostly take values near zero. We also report raw corpus-membership rates directly from the data as a

transparent, model-free summary of the unconditional patterns.

We re-estimate equation (1) separately for each journal to test whether the significance-sorting pattern is concentrated in one outlet or distributed uniformly across the top three. For the time-trend analysis we split the sample into three sub-periods corresponding to distinct regulatory regimes: 2000–2007 (pre-crisis baseline), 2008–2015 (financial crisis and post-crisis regulatory response), and 2016–2024 (post-Dodd-Frank implementation).

5. Results

5.1 Within-Paper Test Statistic Distributions

We test whether published papers show unusual clustering of test statistics just below the conventional 5% significance threshold, using the caliper test of Brodeur, Lé, Sangnier and Zylberberg (2016). We run this test separately on two types of reported statistics. The first type covers regression rows where both a coefficient and a standard error are reported, allowing independent verification of the test statistic. The second type covers standalone test statistics reported without a coefficient or standard error—numbers that appear mainly in robustness and supplementary tables. We extract both types from the PDFs of all published and NBER papers in the sample.

We extract regression-table test statistics from the PDFs of all 134 published papers and 157 NBER working papers using automated text parsing. The published sample yields 7,761 test statistics from 114 papers; the remaining 20 published papers contribute no usable statistics (13 have no extractable rows; 7 have rows whose values fail numeric validation). The NBER sample yields 8,351 test statistics from 117 papers; the remaining 40 contribute no usable statistics for similar reasons. Having a pre-publication sample of test statistics allows us to compare the two distributions directly, rather than inferring bias solely from the shape of the published distribution.

One concern is that published journal PDFs might be harder to parse than NBER draft

PDFs, artificially inflating the caliper gap. We address this concern with several checks but acknowledge that none of them rules it out completely; each speaks to a different aspect of differential missingness.

First, for rows where both a coefficient and a standard error are extracted, we verify that the reported test statistic equals their ratio to within 0.1. This consistency check passes for 98.3% of published rows and 99.7% of NBER rows, with no meaningful difference between corpora. By construction this check applies only to Coef+SE rows; the standalone z -statistics that drive the headline caliper finding (ratio 2.54 published vs. 1.58 NBER) report no coefficient or standard error and therefore cannot be cross-validated against another extracted field. What we can say is that within the Coef+SE rows, where extraction is verifiable, neither corpus shows excess bunching ($R = 0.82$ published, 0.95 NBER); the gap between corpora is concentrated in the unverifiable row type, which is precisely where parsing concerns are hardest to rule out.

Second, the per-paper extraction rate—the fraction of in-scope papers that yield at least one usable $|z|$ -statistic—differs by both corpus and direction. Among published papers, 91.5% of negative-finding, 83.0% of positive-finding, and 75.0% of null-finding papers contribute at least one extracted statistic; the corresponding NBER rates are 82.6%, 75.0%, and 56.2%. Null-finding papers are over-represented in the missing-extraction group in both corpora, which is consistent with the fact that null-finding papers report fewer significance-laden regression tables to begin with. The headline caliper result is concentrated in negative-finding papers, where extraction rates are highest (above 80% in both corpora) and most similar across corpora; the negative-finding result is therefore the least likely of the three direction-specific results to be driven by selective extraction.

Third, papers yielding no usable statistics are spread roughly evenly across journals and across years within journal (within-cell counts are small, so we report this descriptively rather than as a formal test).

Fourth, restricting the analysis to papers with at least 20 extracted statistics leaves the

main caliper ratios within ± 0.05 of their full-sample values, indicating that the result is not driven by a small set of high-volume papers.

Two parsing-bias channels remain genuinely outside our checks: differential missingness within a paper’s tables (a paper might contribute its Coef+SE rows but not its standalone z -stats, or vice versa), and table-format heterogeneity that systematically suppresses just-above-1.96 statistics relative to just-below-1.96 ones in published papers. We are not able to rule these out with the data we have; we flag them as residual concerns rather than claim the parsing channel is fully closed.

The results differ sharply by statistic type. For primary regression tables—where authors report all coefficient estimates for a given specification—neither published papers nor NBER papers show threshold bunching (caliper ratios of 0.82 and 0.95, respectively). Threshold bunching appears only in standalone test statistics, which come mainly from supplementary and robustness tables where authors choose which results to display. Published papers show a higher bunching ratio in these rows (2.54) than NBER papers (1.58).

We interpret this pattern as *selective reporting*: authors are more likely to highlight a supplementary result when it clears a significance threshold, not because they are manipulating their primary estimates. Standard regression tables report all estimates regardless of significance, and those rows show no bunching. The fact that the same pattern—though weaker—appears in NBER working papers as well suggests it reflects an authorial reporting choice rather than a publication filter imposed by journals. We cannot rule out that some standalone statistics come from searches over many specifications, but the absence of bunching in primary regression rows is evidence against widespread specification search (see Table 5 below and Appendix E for full counts).

We use two complementary tests. The *caliper test* (Brodeur, Lé, Sangnier and Zylberberg, 2016) compares the number of $|z|$ -statistics on either side of the conventional significance cutoff $|z| = 1.96$ (equivalently, the two-sided $p = 0.05$ threshold). Throughout the paper we

report the ratio

$$R = \frac{\#\{|z| \in [1.645, 1.96)\}}{\#\{|z| \in [1.96, 2.275]\}},$$

the count of statistics in the narrow window *just below* the $|z| = 1.96$ cutoff (where $0.05 < p \leq 0.10$) divided by the count *just above* it (where $0.023 \leq p \leq 0.05$). Under the null of no threshold sorting, the two zones should be equally populated and $R = 1$. A value $R > 1$ indicates excess mass on the just-below-1.96 side (i.e., test statistics that fall just short of conventional significance); a value $R < 1$ indicates excess mass on the just-above-1.96 side (i.e., results that just clear the bar). We adopt this $|z|$ -space convention—the inverse of the p -space convention used by some authors—throughout the paper. The *density discontinuity test* asks whether there is a sharp jump in the frequency of test statistics right at the $|z| = 1.96$ threshold, which is also consistent with selective reporting.

Table 4 reports the caliper ratios and significance tests for the full sample, broken out by journal and by finding direction. Figures 1–4 plot the corresponding test-statistic distributions.

Table 4. Within-Paper Test Statistic Distribution Tests

Subset	N	% Sig	Caliper ratio	Caliper χ^2	Caliper p	Disc. p
Full sample	7,761	57.1%	1.20***	9.48	0.002	0.040
<i>By journal:</i>						
JF	1,308	65.5%	0.88	0.64	0.423	<0.001
RFS	2,581	59.2%	0.97	0.10	0.752	0.180
JFE	3,872	52.9%	1.49***	23.56	<0.001	0.038
<i>By finding direction:</i>						
Positive	2,633	60.7%	0.94	0.33	0.565	0.456
Negative	3,491	54.9%	1.52***	22.54	<0.001	0.183
Null/mixed	1,637	56.1%	1.05	0.11	0.740	<0.001

Notes: N is the number of test statistics extracted from regression tables, after removing implausible values and non-numeric entries. % Sig is the share of test statistics that exceed the conventional 5% significance threshold. Caliper ratio is $R = \#\{|z| \in [1.645, 1.96)\} / \#\{|z| \in [1.96, 2.275]\}$, the count of $|z|$ -statistics in the just-insignificant window divided by the count in the just-significant window. $R = 1.0$ means equal frequency; $R > 1.0$ indicates clustering on the just-below-1.96 side (results that fall just short of the conventional significance bar), and $R < 1.0$ indicates clustering on the just-above-1.96 side. Caliper p is the associated p -value testing whether the two zones are equally populated. Disc. p tests whether there is a sharp drop in the frequency of test statistics right at the 5% threshold (Brodeur, Lé, Sangnier and Zylberberg, 2016). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$ (caliper ratio).




Figure 1. Distribution of $|z|$ -statistics from all 7,761 regression rows across 114 published papers. The dashed vertical lines mark $|z| = 1.645$ (10% two-sided), $|z| = 1.96$ (5%), and $|z| = 2.576$ (1%). The red reference curve is the half-normal density—the implied distribution of $|Z|$ when $Z \sim N(0, 1)$, equivalently the change-of-variables image of the $\chi^2(1)$ distribution of Z^2 onto the $|z|$ scale—so it is commensurate with the histogram and represents the null benchmark of no true effects and no selective reporting. Within the caliper window, statistics in $[1.645, 1.96)$ —the just-insignificant zone, where $0.05 < p \leq 0.10$ —outnumber those in $[1.96, 2.275]$ —the just-significant zone, where $p \leq 0.05$ —by about 20% ($R = 1.20$, $p = 0.002$). The excess on the just-below-1.96 side is consistent with selective reporting of borderline results that fall just short of conventional significance.

Four findings stand out. First, across all published papers, test statistics in the just-insignificant window $|z| \in [1.645, 1.96)$ outnumber those in the just-significant window $|z| \in [1.96, 2.275]$ by about 20% ($R = 1.20$, $p = 0.002$). Second, this bunching is concentrated in the *Journal of Financial Economics*: its caliper ratio is 1.49 ($p < 0.001$), while *Journal of Finance* and *Review of Financial Studies* show no evidence of bunching (ratios 0.88 and 0.97, both $p > 0.40$). Third, breaking the sample by a paper’s main conclusion reveals that bunching is almost entirely driven by *negative-finding papers*—those that conclude a regulation was harmful or ineffective. Their caliper ratio is 1.52 ($p < 0.001$), compared to 0.94 in positive-finding papers and 1.05 in null-finding papers. Fourth, the discontinuity test shows a conflicting result for *Journal of Finance*: the caliper ratio (0.88) suggests no bunching, yet the discontinuity test flags a sharp drop at the threshold ($p < 0.001$). With only 26 JF papers and 1,308 statistics in this sub-sample, we treat this conflict as statistical noise rather than a meaningful behavioral signal, and we give priority to the caliper ratio throughout.

The direction-level result is novel. As shown in the next section, positive-finding papers are not over-represented in the published corpus relative to negative-finding ones. Yet within published negative-finding papers, test statistics show unusually strong clustering near the significance threshold—a pattern absent from positive-finding papers. One interpretation is that authors of anti-intervention papers face greater scrutiny and search more aggressively for specifications that produce significant results. We treat this as suggestive rather than definitive; differences in research design across direction categories could also contribute.

A simple decomposition explains how the full-sample ratio of 1.20 arises. Published papers devote 28.3% of their table rows to standalone z -statistics (those reported without a paired coefficient and standard error), compared with 18.1% in NBER working papers; and standalone statistics cluster far more sharply in published papers (ratio 2.54) than in NBER papers (1.58). The full-sample difference therefore reflects both a within-type effect—more bunching per standalone statistic in published papers—and a composition effect—a higher

standalone share. In the verifiable coefficient rows, neither corpus shows excess clustering (ratios of 0.82 and 0.95, respectively). Table 5 presents these numbers together; Appendix E provides the full counts.

Table 5. Caliper Ratio Decomposition by Statistic Type and Corpus

Corpus	Row type	N stats	Ratio	Share
Published	All rows	7,761	1.20*** [↑]	100%
	Coef.+SE rows	5,563	0.82*** [↓]	71.7%
	Standalone z	2,198	2.54*** [↓]	28.3%
NBER	All rows	8,351	1.01	100%
	Coef.+SE rows	6,840	0.95	81.9%
	Standalone z	1,511	1.58*** [↑]	18.1%

Notes. *Coef.+SE rows* are regression rows where both a coefficient and a standard error are reported, so the test statistic can be independently verified. *Standalone z* rows report only a t - or z -statistic with no coefficient or standard error. *Share* is the fraction of each corpus’s statistics that fall into that row type. The full-sample ratio of 1.20 reflects both the higher standalone bunching in published papers and their larger standalone share relative to the NBER corpus. Significance markers are two-sided ($***p < 0.01$, $**p < 0.05$, $*p < 0.10$ for the equality test). Direction arrows: [↑] flags $R > 1$ (excess on the just-below-1.96 side, the focal pattern); [↓] flags $R < 1$ (excess on the just-above-1.96 side, the classical Brodeur “just-cleared-the-bar” direction). The published Coef+SE row is significant in the [↓] direction ($p = 0.007$): in verifiable rows, published papers show slightly more statistics just above 1.96 than just below. See Section 5.1 for the two-direction interpretation and Appendix E for the underlying counts.

Most prior work on significance bias can only study published papers and must infer what the research pipeline looked like before editorial decisions were made. Our NBER sample lets us compare directly.

Table 6 shows the results side by side.

Table 6. Within-Paper Caliper Test: Published vs. NBER Working Papers

Sample	Subset	N stats	Papers	Ratio	p -value		
					Naive	Perm.	Boot.
Published	Full sample	7,761	108	1.20***	0.002	0.001	0.150
NBER	Full sample	8,351	106	1.01	0.849	0.424	0.432
Published	Positive-finding	2,633	37	0.94	0.565	0.732	0.623
NBER	Positive-finding	2,658	39	0.99	0.918	0.554	0.492
Published	Negative-finding	3,491	50	1.52***	<0.001	<0.001	0.135
NBER	Negative-finding	4,375	50	1.00	0.964	0.532	0.474
Published	Null-finding	1,637	21	1.05	0.740	0.388	0.383
NBER	Null-finding	1,318	17	1.14	0.442	0.239	0.260

Notes: The *Published* sample contains 7,761 test statistics from 114 papers in the JF, RFS, and JFE; 108 of these papers contribute at least one statistic near the significance threshold and are used in the permutation and bootstrap tests. The *NBER* sample contains 8,351 test statistics from 117 working papers (out of 157 in-scope); 106 contribute near-threshold statistics. The *Ratio* column is $R = \#\{|z| \in [1.645, 1.96)\} / \#\{|z| \in [1.96, 2.275]\}$. $R = 1.0$ means the two zones are equally populated; $R > 1.0$ indicates clustering on the just-below-1.96 side (just-insignificant results), and $R < 1.0$ indicates clustering on the just-above-1.96 side (just-significant results). Three significance tests are reported. *Naive*: treats every statistic as an independent observation. *Perm.*: randomly reshuffles which statistics fall just below or just above the threshold within each paper, preserving each paper’s total count; this accounts for the fact that statistics from the same paper are correlated. *Boot.*: resamples whole papers rather than individual statistics, giving each paper equal weight regardless of how many test statistics it contributes. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$ (caliper ratio). Boot. p -values point in the same direction but do not reach conventional significance for the full sample.

The point estimates suggest a gap: the published caliper ratio (1.20) is larger than the NBER ratio (1.01), and the overall fraction of statistically significant results is nearly identical in both groups (57% published, 56% NBER), so the difference in ratios is not driven by working papers being noisier or less rigorous overall. How confident we can be in the gap depends on which test we use. A naive 2×2 chi-square comparing the cross-corpus difference in just-below/just-above counts—treating each statistic as independent—gives $\chi^2 = 3.74$, $p = 0.053$ for the full sample; the same test on negative-finding papers gives $\chi^2 = 10.66$, $p = 0.001$ (published: 318 below / 209 above; NBER: 239/240). Both figures, however, treat within-paper statistics as independent and so overstate precision. The within-sample paper-clustered bootstrap reported in Table 6 does not reach conventional significance for either the

published full sample ($p = 0.150$) or the published negative-finding subsample ($p = 0.135$). We therefore read the published–NBER gap as suggestive rather than definitively significant: the point estimates are larger for published papers, and especially so for negative-finding papers, but the cross-corpus difference is not robust to paper-level clustering on our preferred test.

Within these caveats, the size and direction of the gap are concentrated in negative-finding papers. Among published negative-finding papers, test statistics in the caliper window cluster on the just-below-1.96 side at more than 50% above the just-above-1.96 count; the corresponding NBER negative-finding papers show essentially no such asymmetry ($R = 1.00$). Papers with positive or inconclusive findings show no gap in either direction. Figure 2 displays the distributions side by side.

Two interpretations are plausible. First, published negative-finding papers may differ from negative-finding papers in the NBER comparison group on unobserved structural dimensions—more robustness checks, different table formats, different vintage—that happen to produce more statistics near the threshold. This is a comparison between two direction-matched cross-sections (published negative-finding vs. NBER-only negative-finding papers), not a draft-to-publication comparison of the same papers; only 16 of 134 published papers in our sample are matched to a specific NBER draft, so most published papers have no observable working-paper counterpart against which to read off within-paper changes. Second, authors arguing that a policy is harmful may face greater scrutiny from reviewers and try more specifications until their results are strong enough to publish. We cannot distinguish between these explanations with the data available, and our cross-sectional design cannot speak to within-paper changes between draft and publication for the unmatched majority of the sample.

Because test statistics from the same paper are not independent, we report three versions of the significance test for each entry in Table 6: a standard test that treats all statistics as independent, a permutation test that accounts for within-paper correlation, and a bootstrap




Figure 2. Each panel compares published journal papers (left) to NBER working papers (right). Bars show how frequently test statistics fall in each range. The orange bars mark the zone just around the conventional 5% significance cutoff; the vertical line is the cutoff itself. If test statistics were reported without regard to significance, the orange bars on either side of the line should be roughly equal. *Top panel:* All 7,761 test statistics from published papers and 8,351 from NBER papers. Published papers show visibly more statistics just below the cutoff than just above it; NBER papers do not. *Bottom panel:* The same comparison restricted to papers that conclude a government regulation was harmful (3,491 published; 4,375 NBER). The asymmetry is more pronounced in published papers and essentially absent in NBER papers; see Section 5.1 for the cross-corpus significance discussion.

test that treats each paper as a single observation. The permutation test is significant for both the full sample and negative-finding papers, showing the pattern is consistent across many papers rather than driven by one or two outliers. The bootstrap—our preferred test—does not reach conventional significance in either case, so the results should be read as suggestive evidence rather than a definitive finding. The NBER working papers show no clustering under any test.

Because we run many tests across different subgroups, some significant results could appear by chance. We address this in two ways, but the conclusion depends on which test one prioritizes. First, we designated two tests as primary before analyzing the data: the full-sample caliper test for published papers and for NBER papers. Under the naive chi-square—which treats every statistic as independent—the published full-sample $p = 0.002$ clears a Bonferroni threshold of $0.05/18 \approx 0.003$, and the NBER full-sample $p = 0.849$ is (trivially) consistent with no clustering. Under the paper-clustered bootstrap, however, which we prefer because it does not assume within-paper independence, the published full-sample p is 0.150 and does not reach conventional significance even before correction; the NBER full-sample $p = 0.432$ is again null. We therefore do not claim the primary published-side result “survives” multiple-testing correction under our preferred test; what we can say is that it survives correction under the naive test that ignores within-paper clustering, and that the corpus where clustering is absent (NBER) remains absent under every test.

Among the exploratory subgroups, the negative-finding and *Journal of Financial Economics* results likewise reach the Bonferroni cutoff under the naive test ($p < 0.001$ each) but do not under the bootstrap ($p = 0.135$ for negative-finding; the *Journal of Financial Economics* bootstrap is not reported separately). We read these as suggestive evidence whose significance depends materially on whether within-paper correlation is accounted for, not as results that have been shown to be robust to multiple-testing correction on the preferred inferential standard.

We include all regression coefficients in the test rather than just the headline estimate,




Figure 3. Test-statistic distributions broken out by journal. All p -values reported in this caption are caliper-test p -values ($R = \#\{|z| \in [1.645, 1.96]\} / \#\{|z| \in [1.96, 2.275]\}$); the density discontinuity test (Disc. p in Table 4) gives a different reading for some panels and we note conflicts explicitly below. *Top:* Journal of Finance (1,308 statistics; no caliper-test clustering, caliper $p = 0.423$). Note that the discontinuity test for JF gives $p < 0.001$; the body text explains why we give priority to the caliper here (Section 5.1). *Middle:* Review of Financial Studies (2,581 statistics; no caliper-test clustering, caliper $p = 0.752$; discontinuity $p = 0.180$, no conflict). *Bottom:* Journal of Financial Economics (3,872 statistics; significant caliper-test clustering on the just-below-1.96 side, caliper $p < 0.001$; discontinuity $p = 0.038$, same direction). Under the caliper test, near-threshold clustering is present only in the JFE; the discontinuity test flags the JF as well, and we read that JF flag as noise (see Section 5.1).




Figure 4. Test-statistic distributions broken out by the paper’s main conclusion. All p -values in this caption are caliper-test p -values; the density discontinuity test gives a different reading for the null/mixed panel and we note it explicitly. *Top:* Papers finding that a government intervention was beneficial (2,633 statistics; no caliper-test clustering, caliper $p = 0.565$; discontinuity $p = 0.456$, no conflict). *Middle:* Papers finding that an intervention was harmful or ineffective (3,491 statistics; significant caliper-test clustering on the just-below-1.96 side, caliper $p < 0.001$; discontinuity $p = 0.183$). *Bottom:* Papers with inconclusive or mixed findings (1,637 statistics; no caliper-test clustering, caliper $p = 0.740$). Note that the discontinuity test for null/mixed papers gives $p < 0.001$; as with the JF density anomaly, we prioritize the caliper ratio for the substantive interpretation (Section 5.1). Under the caliper test, near-threshold clustering is concentrated in negative-finding papers; the discontinuity test also flags the null/mixed panel.

because researchers can adjust any reported result—not only the main coefficient—when revising toward a significance threshold. As a check, we re-run the analysis after dropping standard control variables (firm size, leverage, GDP growth, etc.). Removing roughly 9% of observations leaves the main results unchanged, confirming they are not driven by the inclusion of many control-variable rows (see Appendix D).

5.2 Supporting Context: Cross-Sectional Direction Sorting

Terminology convention. Throughout the rest of this section and in Sections 6 and 7, we use “publication rate” and related phrases (e.g., “a paper reaches a top journal at $X\%$ ”) as compact labels for the within-cell *corpus-membership share*—published / (published + NBER-only) within the indicated cell of our 291-paper sample. These shares are mechanically determined by the chosen sample composition (134 published vs. 157 NBER-only) and do not estimate journal acceptance, transition, or editorial-filter probabilities, which would require observing the submission pool and editorial decisions, neither of which is in our data. We retain the shorter “publication rate” phrasing in tables and prose for readability; readers should interpret every occurrence as a corpus-membership share rather than as a transition probability.

We next ask whether finding direction predicts membership in the published corpus relative to the NBER-only comparison group, which provides a cross-sectional test of whether the published-versus-pre-publication mix of findings differs by direction. Table 7 reports the share of papers in each direction category that appear in the published corpus. About 46% of papers in our 291-paper sample fall on the published side regardless of finding direction: 45.6% of papers finding regulation beneficial (47 of 103), 46.1% of papers finding it harmful (59 of 128), and 46.7% of papers with inconclusive findings (28 of 60). The three corpus-membership rates are virtually identical and statistically indistinguishable ($p = 0.992$). This rules out a direction effect on *corpus membership* in our two-corpus comparison; it does not, by itself, rule out direction effects at the journal acceptance stage, which would require

observing the submission pool and editorial decisions (neither is in our data). We interpret the result as evidence that finding direction does not predict published-vs. NBER-only membership, not as a direct test of journal acceptance.

Table 7. Corpus-Membership Rates by Finding Direction

Finding Direction	Sample composition			Corpus-Membership Rate	95% CI
	NBER-only	Published	Total		
Positive (+1)	56	47	103	45.6%	[36.3, 55.2]
Negative (−1)	69	59	128	46.1%	[37.7, 54.7]
Null/Mixed (0)	32	28	60	46.7%	[34.6, 59.1]
Total	157	134	291	46.0%	

Notes. Sample covers 2000–2024. *Finding Direction* is the main conclusion of the paper as coded by the AI classifier: positive (+1) if the paper finds a government intervention was beneficial, negative (−1) if harmful, and null/mixed (0) if the evidence is inconclusive. *NBER-only* counts working papers that were never published in the JF, RFS, or JFE; *Published* counts papers that were; *Total* is both groups combined. *Corpus-Membership Rate* is the share of papers in each direction category that end up published. This reflects the relative sizes of our two samples and should not be interpreted as a journal acceptance rate. *95% CI* is the confidence interval around each publication rate. A statistical test cannot reject that the three publication rates are equal ($p = 0.992$), and the positive vs. negative comparison gives $p = 0.944$.

Table 8 reports probit regression results across seven specifications, each predicting whether a paper ends up published in one of the three journals. Column (2) is our primary specification, matching equation (1) in Section 4.1: it separates positive (+1) and negative (−1) findings against null/mixed papers (0), and the directional-bias hypothesis is the test of $H_0: \beta_1 = \beta_2$. Column (1) reports a more parsimonious binary specification—directional vs. null/mixed—as a complementary check on whether *any* directional sign carries a publication premium; we report it alongside Column (2) but treat it as supporting rather than primary, because the binary version answers a different question (significance bias) than the directional-bias test that motivates the paper. Columns (3)–(7) extend the primary Column (2) specification by adding controls for publication year, decade fixed effects, team size, identification method, and author institution. We classify a paper as *quasi-experimental*

(“QE”) if its primary identification strategy is regression discontinuity (RDD), instrumental variables (IV), difference-in-differences (DiD), a natural experiment with an exogenous shock, or an event study that the AI judges to carry a causal interpretation (i.e., the event is used to identify a causal effect, not merely to characterize an event-window return pattern). Descriptive event-window analyses without a causal claim are coded as non-QE (`ols_reduced`), matching the rubric used by the classifier (`scripts/09_code_id_strategy.py`). We use this five-method definition consistently throughout the paper.

Table 8. Direction–Publication Association: Probit Estimates

	(1) Binary	(2) Bivariate (Primary)	(3) + Year	(4) + Decade FE	(5) + Authors	(6) + ID Strategy	(7) + Top Inst.
Directional	−0.020 (0.182)						
Positive (+1)		−0.026 (0.204)	−0.002 (0.205)	−0.006 (0.205)	−0.005 (0.205)	−0.135 (0.212)	−0.160 (0.215)
Negative (−1)		−0.014 (0.196)	−0.011 (0.198)	−0.017 (0.198)	−0.020 (0.198)	−0.156 (0.206)	−0.159 (0.208)
Year trend			0.033*** (0.012)		0.034*** (0.012)	0.028** (0.012)	0.034*** (0.013)
$\log(1 + N_{\text{authors}})$					−0.186 (0.279)	−0.255 (0.283)	−0.264 (0.291)
Quasi-Exp. ID						0.467*** (0.157)	0.473*** (0.160)
Top Institution							0.661*** (0.167)
Decade FEs	No	No	No	Yes	No	No	No
Intercept	−0.084	−0.084	−0.097	−0.371	0.145	0.114	−0.075
N	291	291	291	291	291	291	291
Pseudo R^2	0.000	0.000	0.020	0.013	0.021	0.043	0.083

Notes. Each column reports probit estimates of the probability that a paper in our dataset ends up published in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics*. Column (2) is the primary specification: it separates positive (+1) and negative (−1) findings against null/mixed papers (0), matching equation (1) in Section 4.1; direction codes are assigned by the AI pipeline from paper abstracts. Column (1) reports a parsimonious binary alternative—“Directional” versus null/mixed—as a supporting check; Columns (3)–(7) extend Column (2) by adding controls. Column (5) adds the number of authors (log scale) as a proxy for team size. Column (6) adds the quasi-experimental (QE) indicator defined in the preceding text (RDD, IV, DiD, natural experiment, or event study with causal interpretation), as classified by the AI from abstracts. Column (7) adds an indicator for whether the lead author’s institution appears in the UTD Top 100 Business School Research Rankings (2021–2025; Naveen Jindal School of Management, University of Texas at Dallas 2026), retrieved via OpenAlex. Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. The baseline publication probability implied by the intercept alone is $\Phi(-0.084) \approx 46.7\%$, equal to the raw null-paper publication rate (28 of 60 null-finding papers, 46.7%); positive- and negative-finding papers publish at 45.6% and 46.1%, respectively, for an overall rate of 46.0%. The marginal effects of all direction indicators are negligible across all columns. The near-zero pseudo- R^2 in Columns (1)–(2) is exactly what we would expect if finding direction has no predictive power for publication; it reflects a null relationship, not poor model quality. A formal test of whether positive and negative coefficients are equal to each other in Column (2) gives $\chi^2 = 0.005$, $p = 0.944$, providing no evidence that journals treat the two directions differently. The sample is large enough to detect coefficient differences of roughly 0.465 units with 80% power (standard error of the difference = 0.166), corresponding to a publication-rate advantage of approximately $1.4\times$ the null-paper rate (roughly 18 percentage points on the probability scale). The null result therefore does not rule out moderate directional preferences, only large ones.

Our primary Column (2) shows no relationship between either signed direction indicator and publication: both coefficients are near zero and insignificant (positive: $p = 0.898$; negative: $p = 0.942$). A joint test of whether both are zero gives $p = 0.992$. The directional-bias test of whether the two coefficients are equal to each other gives $\chi^2 = 0.005$, $p = 0.944$, providing no evidence that journals favor one direction over the other.

The supporting binary specification in Column (1) reaches the same conclusion in a more parsimonious form: the pooled directional coefficient is -0.020 ($p = 0.914$), indistinguishable from zero, so papers with any directional finding are no more or less likely than null-finding papers to sit on the published side of our corpus. The design rules out large sorting effects—roughly 18 percentage points or more—but cannot rule out smaller ones; we read both columns as evidence against strong directional sorting, not as proof of complete editorial neutrality.

Adding controls does not change the picture. The year trend (Column (3)) is positive and significant, indicating that, conditional on direction, more recently dated papers in the in-scope government-intervention literature are disproportionately on the published side of the sample. We note that any right-censoring of recent NBER papers (working papers that have not yet had time to be published) would push in the *opposite* direction—attenuating or reversing a positive year trend, since recent NBER papers stuck in review would inflate the NBER share at higher years—so censoring is not the explanation for the positive coefficient. More plausible explanations are compositional: the in-scope set may shift over the 2000–2024 window toward topics (e.g., post-2008 capital and resolution regulation) where journal publication outpaces NBER circulation in our sample, or the published-side data may simply be more completely indexed for recent years than the NBER side. Our two-corpus design does not let us separate these compositional channels. As a robustness check, restricting the sample to papers circulated before 2021 leaves the direction coefficients unchanged ($p > 0.50$ in both cases), so the positive year trend does not interact with the direction null. Decade fixed effects (Column (4)) and team size controls (Column (5)) likewise leave the direction

results unaffected.

Column (6) adds the QE indicator defined above (RDD, IV, DiD, natural experiment, or event study with causal interpretation). Papers with such designs are significantly more likely to be published ($p = 0.003$), consistent with journals rewarding rigorous causal identification. This variable is classified by the AI from paper abstracts and has not been manually validated, so it may also capture other quality attributes—such as institutional affiliation or topic fashionability—that happen to correlate with research design. Section 5.3 examines this concern directly, showing that the premium survives sequential controls for institutional prestige and team size, and that quasi-exp status and institutional rank are nearly orthogonal ($r = 0.015$). Direction coefficients remain near zero and statistically insignificant after adding the QE control ($p > 0.44$), which reduces but does not eliminate the omitted-QE-variable concern: a hidden direction effect masked by QE would have been expected to surface once QE is partialled out, and we do not see one. This does not rule out an underlying direction effect—the design has the same 18-percentage-point power floor noted in Section 7.4, so moderate direction-based imbalances would remain below the detection threshold—but it does close off the specific channel in which the absence of a direction premium in earlier columns might have reflected omitted research-design quality.

Column (7) adds a top-institution indicator equal to one if the first author’s institutional affiliation appears in the UTD Top 100 Business School Research Rankings (2021–2025 worldwide ranking; Naveen Jindal School of Management, University of Texas at Dallas 2026), as retrieved from OpenAlex. Papers from UTD Top 100 institutions are significantly more likely to be published: the top-institution coefficient is 0.661 (s.e. 0.167, $p < 0.001$), indicating a strong institutional prestige premium on top of the quasi-experimental design premium. Direction coefficients remain unchanged and insignificant ($\hat{\beta}_1 = -0.160$, $\hat{\beta}_2 = -0.159$; both $p > 0.44$), and the quasi-experimental premium (0.473***, s.e. 0.160) is virtually unchanged from Column (6) of Table 8 (0.467***, s.e. 0.157). The null result for finding direction is robust to controlling for institutional prestige.

Table 9 breaks down publication rates by both finding direction and research design. The clearest pattern is that papers using quasi-experimental methods are published at much higher rates than papers that rely on standard regression approaches, regardless of what the paper finds: the gap is about 20 percentage points for positive-finding papers (55.6% vs. 34.7%), 13 points for negative-finding papers (52.2% vs. 39.0%), and 39 points for null-finding papers (75.0% vs. 36.4%), though the null $\times$ QE cell contains only 16 papers so the 75% figure should be interpreted cautiously.

Within each research-design group, publication rates are similar across finding directions for non-QE papers but more dispersed for QE papers. Among papers without quasi-experimental designs, positive, negative, and null papers are published at 34.7%, 39.0%, and 36.4% respectively—a spread of only 4.3 percentage points. Among quasi-experimental papers the rates are 55.6%, 52.2%, and 75.0%—a spread of 22.8 percentage points, driven mainly by the high publication rate in the null cell, which contains only 16 papers and should be interpreted cautiously. We therefore do not describe the QE rates as “nearly identical” across direction categories.

Two additional checks bear on whether the research-design premium itself varies with finding direction. First, a probit that includes both direction and research-design indicators jointly finds a strong quasi-experimental coefficient ($p = 0.004$) while both direction coefficients remain indistinguishable from zero (positive: $p = 0.548$; negative: $p = 0.499$); this leaves the research-design premium intact after conditioning on direction. Second, adding interaction terms between direction and research design—which would detect whether the research-design premium differs by finding direction—yields one interaction at $p = 0.158$ and one at $p = 0.093$ (marginally significant at the 10% level the paper uses elsewhere). We read these checks as ruling out strong heterogeneity in the QE premium across direction categories, but not as conclusive evidence of pure uniformity; the $p = 0.093$ interaction leaves room for a modest direction-specific component that our sample size cannot reliably distinguish from noise.

Table 9. Corpus-Membership Rates by Finding Direction and Identification Strategy

	Non-QE	QE	Total
Positive (+1)	34.7% (17/49)	55.6% (30/54)	45.6% (47/103)
Negative (−1)	39.0% (23/59)	52.2% (36/69)	46.1% (59/128)
Null/mixed (0)	36.4% (16/44)	75.0% (12/16) [†]	46.7% (28/60)
Total	36.8% (56/152)	56.1% (78/139)	46.0% (134/291)

Notes. Each cell shows the share of papers published in a top-3 finance journal (JF, RFS, or JFE), with paper counts in parentheses. QE: papers whose primary identification strategy is regression discontinuity (RDD), instrumental variables (IV), difference-in-differences (DiD), a natural experiment with an exogenous shock, or an event study with causal interpretation (see definition following Table 8). Non-QE: papers that use standard regression, panel fixed effects without an instrument, descriptive methods, or event-window analyses without a causal claim, as classified by the AI from abstracts. †: Only 16 papers fall in the null/mixed $\times$ QE cell; the 75% rate is based on a small sample and should be treated with caution. Within the Non-QE column, publication rates are similar across the three finding directions (34.7%, 39.0%, 36.4%); within the QE column, the rates are more dispersed (55.6%, 52.2%, 75.0%), driven mainly by the small null/mixed cell. A formal interaction test of whether the research-design premium varies by finding direction does not reject pure uniformity at conventional levels, but one interaction is marginal ($p = 0.093$); we therefore characterize the evidence as ruling out strong heterogeneity rather than confirming complete uniformity (see text).

How large is the direction effect in practical terms? The raw publication rates tell the story directly: positive-finding papers appear in top journals 45.6% of the time, negative-finding papers 46.1%, and null-finding papers 46.7%. Dividing each by the null rate gives ratios of 0.98 and 0.99 — essentially 1.0 in both cases. A paper’s directional conclusion has no meaningful bearing on whether it gets published.

These ratios are not the same as the 1.5–2.0 $\times$ figures reported by Brodeur, Lé, Sangnier and Zylberberg (2016), and the two should not be directly compared. Brodeur’s estimates measure a different phenomenon: how often researchers choose test statistics that land just above the 5% significance threshold relative to just below it, within the set of already-published papers. Our ratios instead measure whether papers with different conclusions are equally likely to reach print in the first place. We compare our results to Brodeur on the within-paper metric in Section 5.1, where our published caliper ratio of 1.20 is somewhat lower than Brodeur’s range. The absence of detectable large-scale direction-based sorting in the policy-finance literature is a distinct and noteworthy finding.

The one factor that does predict corpus membership with both statistical and practical

significance is research design. Conditional on the controls included in our model (publication year, team size, top-institution status, and finding direction), QE papers (the five-method definition given above: RDD, IV, DiD, natural experiment, or event study with causal interpretation) are 18–20 percentage points more likely to appear in the published corpus than papers with the same observed characteristics that use standard regression or descriptive approaches. We use the phrase “otherwise similar” to refer specifically to these observed dimensions; unobserved differences in quality, data access, or narrative framing between quasi-experimental and other papers cannot be ruled out, and we read this association as cross-sectional rather than as a causal estimate of how editors weight identification strategy.

Untabulated journal-level probits show no statistically significant direction coefficient at any of the three journals individually, consistent with the pooled null result. These estimates carry limited inferential weight given small sub-samples—the *Journal of Finance* contributes only 26 in-scope published articles over the full sample period—and should be treated as descriptive.

5.3 The Identification Quality Premium

The quasi-experimental coefficient—QE papers (RDD, IV, DiD, natural experiment, or event study with causal interpretation) are roughly 18–20 percentage points more likely to appear in the published corpus conditional on observed controls—is a substantive cross-sectional association in its own right, not merely a control variable. We develop it here.

A natural concern is that the premium proxies for confounds: top institutions may produce more quasi-experimental work, or larger teams may have better access to data and computing resources that enable credible identification. To assess this, Table 10 adds controls sequentially.

The quasi-experimental premium is essentially unchanged across all six specifications. Adding top-institution status (Column 2) does not shrink the quasi-exp coefficient—it rises marginally from 0.490 to 0.501. Adding team size (Column 3) leaves it equally stable.

Table 10. Identification Quality Premium: Sequential Controls

	(1)	(2)	(3)	(4)	(5)	(6)
Quasi-Exp. ID	0.490*** (0.149)	0.501*** (0.151)	0.497*** (0.150)	0.508*** (0.152)	0.441*** (0.155)	0.473*** (0.160)
Top Institution		0.607*** (0.164)		0.606*** (0.164)	0.659*** (0.167)	0.661*** (0.167)
$\log(1 + N_{\text{authors}})$			-0.176 (0.280)	-0.160 (0.281)	-0.265 (0.286)	-0.264 (0.291)
Year trend					0.036*** (0.013)	0.034*** (0.013)
Direction dummies	No	No	No	No	No	Yes
Marg. effect, Quasi-Exp. (pp)	19.4	19.9	19.7	20.1	17.5	18.8
Marg. effect, Top Inst. (pp)	—	24.1	—	24.1	26.1	26.2
Pseudo R^2	0.027	0.062	0.028	0.062	0.081	0.083
N	291	291	291	291	291	291

Notes. Each column reports a probit of publication on the indicator variables listed, estimated on the full sample of 291 papers. *Quasi-Exp. ID* equals one if the paper’s primary identification strategy is RDD, IV, DiD, a natural experiment, or an event study with causal interpretation (full definition following Table 8), as coded by the AI from abstracts. *Top Institution* equals one if the first author’s institution appears in the UTD Top 100 Business School Research Rankings. Column (6) includes positive and negative finding-direction dummies (omitted from the table for brevity; both are individually and jointly insignificant, $p > 0.44$). Marginal effects (in percentage points) are evaluated at the sample mean of the linear index. Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Adding both prestige proxies together (Column 4), then a year trend (Column 5), then direction dummies (Column 6) produces a range of 0.441–0.508, all significant at $p < 0.01$.

The stability is explained by a striking feature of the data: quasi-exp and institutional prestige are nearly orthogonal. The correlation between the two indicators is $r = 0.015$. Researchers at top institutions use quasi-experimental methods at virtually the same rate as those at other institutions (48.9% vs. 47.3%). This means controlling for prestige provides almost no new information about method choice, and vice versa.

Table 11 presents raw publication rates by the four combinations of method and institution. The quasi-exp premium operates within both institution tiers: among top-institution authors, quasi-exp papers publish at 72.7% vs. 52.2% for non-QE papers (a gap of 20.5

pp); among other institutions, the gap is 48.4% vs. 30.2% (18.2 pp). The institutional prestige premium likewise holds within both method groups. The two effects are additive: top-institution quasi-exp papers sit on the published side of the corpus 72.7% of the time, non-top non-QE papers 30.2% of the time—a 42-percentage-point corpus-membership spread reflecting roughly equal contributions from method and pedigree.

Table 11. Publication Rates by Identification Strategy and Institutional Tier

	Non-QE	Quasi-Experimental
Top Institution	52.2% (24/46)	72.7% (32/44)
Other Institution	30.2% (32/106)	48.4% (46/95)
QE premium (pp)	—	+20.5 (top) / +18.2 (other)

Notes. Each cell shows the share of papers published in a top-3 finance journal, with counts in parentheses. *Quasi-Experimental* (QE) denotes papers using the five-method definition in Section 5.2: RDD, IV, DiD, natural experiment, or event study with causal interpretation. *Top Institution* is the UTD Top 100 indicator. The correlation between the two indicators is $r = 0.015$.

What can we conclude? Institutional prestige and team size—the most natural observable confounds—do not explain the quasi-experimental premium; the premium is not a prestige premium in disguise on the observed dimensions. What remains is consistent with journals in this domain having internalized credible causal identification as a quality criterion, but our cross-sectional design does not let us identify that channel directly. Unobserved quality dimensions that travel with quasi-experimental method choice—data access, framing, robustness practices—could equally well explain the cross-sectional association that we measure.

We cannot fully rule out residual confounding: quasi-experimental papers may be higher quality on unobservable dimensions — better natural experiments, richer data, more careful robustness checking — that correlate with method choice but not with institutional affiliation. But the separability from observable prestige proxies is itself informative. If the premium were purely about researcher quality, it would be expected to correlate at least moderately with institutional rank; the near-zero correlation argues against a pure researcher-quality

story.

6. Robustness

This section checks whether the main findings hold under a range of alternative assumptions. Table 13 summarizes results across three areas: sample and scope restrictions, alternative model specifications, and concerns about classification quality. Caliper window sensitivity is addressed separately in Table 12.

6.1 Caliper Window Sensitivity

The caliper test compares how many test statistics fall just below the 5% significance threshold versus just above it, so the result depends on exactly how wide a band is used on each side. Table 12 repeats the test under four different window choices.

When the window starts at $|z| = 1.645$ or higher—covering three of the four specifications, including the standard Brodeur, Lé, Sangnier and Zylberberg (2016) window—published papers consistently show excess mass just below the threshold: the ratio of below-threshold to above-threshold test statistics ranges from 1.20 to 1.33, while the same ratio for NBER working papers stays near 1.0.

The exception is the widest window, which extends the lower bound down to $|z| = 1.50$. Here the ratio for published papers drops to 0.92 ($p = 0.088$), suggesting fewer statistics below the threshold than above. The reason is mechanical: at $|z|$ values between 1.50 and 1.645, the overall density of test statistics is still rising—there is no bunching there—so pulling that region into the window dilutes and reverses the bunching signal found closer to the threshold. This rules out the alternative explanation that published papers simply have a general shift toward larger test statistics; the pattern is specifically concentrated in the narrow band just below the 5% significance cutoff.

Table 12. Caliper Window Robustness (Summary)

Window	Sample	Ratio	Naive p	Perm. p
Standard (1.645–2.275)	Published	1.20***	0.002	0.001
	NBER	1.01	0.849	0.424
Narrow (1.70–2.22)	Published	1.33***	<0.001	<0.001
	NBER	1.01	0.917	0.468
Wide (1.50–2.42)	Published	0.92	0.088	0.960
	NBER	0.97	0.549	0.737
Brodeur2 (1.65–2.25)	Published	1.33***	<0.001	<0.001
	NBER	1.10	0.153	0.081

Notes. The caliper ratio is the count of $|z|$ -statistics with $|z|$ strictly below 1.96 (the just-insignificant side) divided by the count with $|z|$ at or above 1.96 (the just-significant side), within the specified window. $R > 1.0$ means more statistics fall on the just-below-1.96 side than the just-above-1.96 side—an excess of borderline-but-not-quite-significant results. The wide window (1.50–2.42) reverses sign for published papers because test-statistic density is still rising between $|z| = 1.50$ and $|z| = 1.645$, pulling the ratio below 1.0; this is a mechanical artifact of the wider lower bound, not evidence against bunching. The three windows whose lower bounds start at or above $|z| = 1.645$ all show consistent excess mass on the just-below-1.96 side of the published distribution. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$ (caliper ratio).

6.2 Sample and Scope Robustness

Our sample includes papers on central bank actions such as quantitative easing, discount window lending, and unconventional monetary policy alongside papers on explicit legislation, securities regulation, and banking supervision. One might worry that monetary policy papers behave differently—perhaps because their findings are harder to classify directionally, or because they attract a different editorial audience. To check, we drop the 21 monetary policy papers and re-estimate using only the 270 remaining papers on financial legislation and regulation. The results are unchanged: both direction coefficients are close to zero and statistically insignificant ($\hat{\beta}_1 = -0.028$, $p = 0.896$; $\hat{\beta}_2 = -0.068$, $p = 0.741$), confirming that the null finding is not driven by the inclusion of monetary policy papers.

As a further check, we ask whether finding direction predicts placement in one journal rather than another. Holding the full 291-paper sample fixed, we redefine the dependent variable in two ways and re-run the probit: $\text{Pub}_{\text{Journal of Finance}} = 1$ if a paper appears in the *Journal of Finance* (0 otherwise, including *Review of Financial Studies*, *Journal of Financial*

Economics, and NBER), and $\text{Pub}_{\text{Review of Financial Studies}+\text{Journal of Financial Economics}} = 1$ if a paper appears in the *Review of Financial Studies* or the *Journal of Financial Economics* (0 otherwise, including *Journal of Finance* and NBER). Each specification asks whether directional papers are disproportionately likely to land in the indicated outlet relative to everything else in the corpus. In neither case does finding direction predict outlet placement: the $\text{Pub}_{\text{Journal of Finance}}$ coefficients are 0.400 and 0.350 ($p > 0.20$, $N = 291$), and the $\text{Pub}_{\text{Review of Financial Studies}+\text{Journal of Financial Economics}}$ coefficients are -0.147 and -0.130 ($p > 0.48$, $N = 291$). This is a target-journal-membership test on the full sample, not a journal-vs-NBER subsample comparison (a JF-vs-NBER subsample would be $26 + 157 = 183$ observations); it asks whether any single outlet sorts on direction more sharply than the corpus as a whole, and the answer is no.

6.3 Alternative Specifications

Our baseline model separates papers into positive, negative, and null findings. One concern is that labeling a finding as positive or negative requires a judgment call about what counts as a beneficial outcome for a given intervention. As a check, we sidestep this issue entirely by collapsing the three categories into a simple binary: did the paper find any directional result at all, or was the finding null? The result is unchanged—the coefficient is -0.020 (s.e. 0.182, $p = 0.914$)—and adding year and author-count controls leaves it essentially the same at -0.014 ($p = 0.941$). The null result does not depend on how finely we classify finding direction.

We also check whether the result depends on the choice of statistical model. Probit models make distributional assumptions that may not hold perfectly in small samples. Re-estimating the same model using ordinary least squares (a linear probability model with robust standard errors) gives coefficients of -0.010 ($p = 0.900$) for positive-finding papers and -0.006 ($p = 0.942$) for negative-finding papers—directly readable as percentage-point differences in publication rates relative to null papers. Both are negligible, confirming that

the null result is not an artifact of the probit specification.

Finally, we check whether the results are driven by the small set of 16 published papers that we can confirm previously circulated as NBER working papers. Because our NBER comparison group is defined as the 157 in-scope working papers that did *not* reach the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics*, the matched-pair NBER drafts are by construction already excluded from the estimation sample (see Section 3.5); dropping the 16 matched published papers therefore leaves $N = 275$ (118 published + 157 NBER). The direction coefficients remain small and statistically insignificant on this subsample ($\hat{\beta}_1 = -0.122$, $p = 0.561$; $\hat{\beta}_2 = -0.054$, $p = 0.787$, year-trend included), confirming that the results do not depend on the matched-pair subset.

6.4 Classification and Quality Concerns

One concern with our main result is that null-finding papers may simply use weaker research designs, and editors may be rejecting them for quality reasons rather than because of their conclusions. To address this, we restrict the sample to the 139 papers that use credible causal identification methods—natural experiments, instrumental variables, difference-in-differences, or regression discontinuity designs—holding research design quality roughly constant across finding directions.

Within this subsample, neither direction coefficient is significantly positive: the positive-finding coefficient is -0.598 ($p = 0.121$) and the negative-finding coefficient is -0.656 ($p = 0.081$). The second figure is marginally significant at the 10% level, but with a *negative* sign, meaning quasi-experimental negative-finding papers are if anything *less* likely to be published than null papers—the opposite of what a direction-premium story would predict. We do not over-interpret this: it is fully consistent with the pooled null result and with random sampling variation in a small subsample. The six direction-by-corpus cells in the QE subsample are not uniformly small—Table 9 implies roughly 30/24 (positive published/NBER-only), 36/33 (negative), and 12/4 (null/mixed)—but the null/mixed baseline cells are very

small (12 and 4 papers), and the negative coefficient is computed relative to that baseline. The marginal $p = 0.081$ on the negative coefficient is therefore best read as instability driven by the small null/mixed-QE reference group rather than as evidence of an anti-intervention discount. The classification of papers as quasi-experimental is also based on AI coding of abstracts and has not been manually verified. The key takeaway is that neither coefficient supports the existence of a positive direction premium among the most rigorously designed papers in the sample.

A second concern is that the AI may classify the same paper differently depending on how the abstract is written. Published abstracts tend to be sharper and more declarative; working paper abstracts often present multiple findings, hedge conclusions, and include sensitivity discussions. This writing-style difference could cause the AI to assign high-confidence direction codes more often to published papers than to NBER papers—and that asymmetry could inflate our estimated direction premium even if none exists.

The data support this concern. The AI assigns a high-confidence code to 81.3% of published abstracts but only 62.4% of NBER abstracts, a gap of 18.9 percentage points. NBER abstracts are actually longer on average (158 words versus 109 words for published abstracts), so the lower confidence rate is not simply about having less text to work with.

To pin down whether the confidence gap reflects writing style or substantive differences in reported findings, we count occurrences of two distinct families of terms separately, since a single mixed dictionary would conflate the two. *Style-only hedges* are phrases that signal author tentativeness without naming a finding—“may,” “might,” “could,” “possibly,” “perhaps,” “suggests,” “appears,” “tends to,” “seems,” “unclear.” *Null-result phrases* explicitly describe the substantive finding—“no significant,” “do not find,” “no evidence,” “insignificant,” “statistically insignificant,” “null,” “null result.” Per 100 words of abstract text, NBER abstracts contain 0.36 style-only hedges versus 0.22 in published abstracts ($1.67\times$), and 0.05 null-result phrases versus 0.04 in published abstracts ($1.15\times$). The style-only gap is the larger of the two and persists after length adjustment; the null-result gap is small

and consistent with the already-known direction-distribution difference between the two corpora. We therefore read the confidence gap as primarily reflecting writing style rather than substantive finding differences, while acknowledging that a residual finding-driven component cannot be ruled out from these counts alone.

This writing-style gap creates a measurement risk in two directions, and the sign of the resulting bias in our direction premium depends on which direction dominates. (a) Hedged NBER abstracts may cause the AI to classify genuinely directional working papers as null, which would inflate the NBER null share, deflate the NBER directional share, and push the estimated direction premium *upward*. (b) Polished published abstracts may cause the AI to classify some genuinely null published articles as directional, which would inflate the published directional share and also push the premium *upward*; conversely, miscoding genuinely directional published articles as null (a third channel) would push the premium *downward*.

The LLM replication exercise (Section 3.4) shows the null category to be the least stable: of the 20 papers the AI coded as null in the validation sample, 15 were re-coded as directional in an independent replication. Crucially, these 15 flips are not symmetric across corpora: 13 occurred among the 15 originally null-coded published articles (87% flip rate), against 2 of 5 NBER papers (40%). Channel (a), which would make our null result conservative, is therefore not the dominant empirical pattern in the replication: the corpus where null coding swings most is the *published* side, not NBER. We accordingly do not claim the estimates are conservative; the replication exercise establishes instability of the null category but not a clean sign for the residual measurement bias.

As a partial robustness check, we add $\log(1 + \text{abstract length})$ as a control variable. The direction coefficients remain essentially zero ($\hat{\beta}_1 = 0.008$, $p = 0.968$; $\hat{\beta}_2 = 0.011$, $p = 0.958$), indicating that the length component of the writing-style gap is not what is driving the null direction result. This check addresses only the length channel; it does not control for the residual per-100-word hedging-style gap documented above (NBER 0.36 style hedges

per 100 words vs. 0.22 in published abstracts), which could affect LLM direction coding even conditional on length. A proper test of the style channel would require either re-coding direction from length-and-style-normalized abstracts or controlling for per-word style intensity, neither of which we attempt here; we therefore read this row as evidence against a length-driven story specifically, not as a closure of the broader writing-style concern.

Table 13. Robustness Checks

Specification	Variable	Coefficient	p-value	N
<i>Panel A: Sample and Scope</i>				
Excl. monetary policy	Positive	−0.028	0.896	270
	Negative	−0.068	0.741	270
<i>Journal of Finance</i> membership	Positive	0.400	0.217	291
	Negative	0.350	0.271	291
<i>Review of Financial Studies</i> + <i>Journal of Financial Economics</i> membership	Positive	−0.147	0.483	291
	Negative	−0.130	0.519	291
<i>Panel B: Alternative Specifications</i>				
Binary directional	Directional	−0.020	0.914	291
Binary + year + authors	Directional	−0.014	0.941	291
LPM (OLS, HC3)	Positive	−0.010	0.900	291
	Negative	−0.006	0.942	291
<i>Panel C: Classification and Quality Concerns</i>				
Quasi-exp. papers only	Positive	−0.598	0.121	139
	Negative	−0.656*	0.081	139
+ log(abstract length)	Positive	0.008	0.968	291
[abstract length control]	Negative	0.011	0.958	291

Notes. Each row reports the probit latent-index coefficient on finding direction, except the LPM row, which reports an OLS coefficient estimated on the same sample. Probit coefficients are on the latent-index (Φ^{-1}) scale and do not by themselves give a percentage-point change in publication probability; the implied publication-probability change is $\Phi(\hat{\alpha} + \hat{\beta}_d + \hat{\gamma}'\mathbf{X}^{\text{base}}) - \Phi(\hat{\alpha} + \hat{\gamma}'\mathbf{X}^{\text{base}})$ at a covariate baseline $\mathbf{X}^{\text{base}}$ (the centered baseline defined in Section 4.1 is used elsewhere in the paper). The LPM (OLS) row, by contrast, has a direct percentage-point interpretation: an LPM coefficient of −0.010 means a 1.0 percentage-point decrease in publication probability for the indicated direction category relative to null/mixed papers. All specifications include a linear year trend unless noted. Standard errors omitted for brevity. *Panel A:* The first row drops monetary policy papers. The “*Journal of Finance* membership” and “*Review of Financial Studies*+*Journal of Financial Economics* membership” rows hold the full 291-paper sample fixed and redefine the dependent variable as an indicator for placement in the indicated outlet (so the reference group includes the other journals and NBER); these are target-journal-membership tests, not journal-vs-NBER subsample regressions. Direction coefficients are small and insignificant in all cases. *Panel B:* “Binary directional” collapses positive and negative findings into a single directional category; adding author-count controls leaves results unchanged. “LPM” re-estimates via OLS with robust standard errors, giving percentage-point differences in publication rates directly; results remain negligible. *Panel C:* “Quasi-exp. only” restricts to the 139 papers using credible causal methods (natural experiments, IV, DiD, RD), holding research design quality constant across direction categories. “+ log(abstract length)” controls for abstract length only; the row tests the length channel of the writing-style concern and does not control for per-word hedging style. See Section 7.4 and the body discussion above for the broader scope. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

7. Discussion

7.1 Implications for Research and Regulation

Our findings argue against the concern that top finance journals systematically favor studies that find an effect of government intervention over those that find none, at the aggregate level of the in-scope government-intervention literature we sample. Positive-, negative-, and null-finding papers appear in the published corpus at nearly identical rates (all around 46%), and finding direction is unrelated to corpus membership across our seven probit specifications. We emphasize that this is an aggregate statement: the sample spans mandatory disclosure rules (Sarbanes-Oxley, Reg FD), corporate governance mandates (say-on-pay, proxy access, board independence), short-selling restrictions, and Dodd-Frank capital and resolution provisions, but with 134 published articles the sample does not support reliable within-topic publication-rate comparisons. Whether the aggregate null also holds at the level of any specific regulation remains an open question that our design cannot resolve.

One qualification limits how strongly this direction-publication conclusion can be stated: our study is underpowered to detect modest imbalances. The smallest directional gap we could reliably identify corresponds to an 18-percentage-point difference in publication rates, so smaller direction-based asymmetries in corpus membership remain possible. Policymakers who rely on this literature should treat our null as evidence against *large-scale* direction-based editorial distortion at the aggregate level, not as a clean bill of health; residual direction-based bias on specific regulations could exist below our detection threshold.

A separate and distinct caveat applies to the within-paper caliper analysis: the bootstrap confidence interval on the negative-finding caliper ratio is wide ($[0.81, 2.67]$), which limits how confidently we can pin down the magnitude of within-paper threshold pressure in negative-finding published papers. We flag this here because the two caveats are about different estimands—the first concerns *direction-based corpus membership* measured by the probit; the second concerns *within-paper selective reporting near $|z| = 1.96$* measured by the

caliper test—and the wide caliper interval does not by itself qualify the no-direction-sorting result.

A clearer cross-sectional association does emerge: QE papers (RDD, IV, DiD, natural experiment, or event study with causal interpretation; see Section 5.2) are substantially more likely to appear in the published corpus, conditional on year, team size, top-institution status, and finding direction. For regulators, this association is reassuring rather than distorting if it reflects editorial or pipeline selection on methodological quality, but our cross-sectional design cannot exclude that quasi-experimental papers also differ on unobserved dimensions (data, framing, robustness practices) that travel with method choice. The descriptive statement we can defend is that the published evidence base is disproportionately drawn from quasi-experimental work; whether that pattern reflects an editorial quality filter or unobserved correlated quality is a question our design leaves open.

One further caution applies to how individual results are read. The excess mass we observe on the just-below-1.96 side of the threshold ($|z| \in [1.645, 1.96)$, $p \in (0.05, 0.10]$)—concentrated in negative-finding papers—suggests that *within-paper* specification choices may play a role even when the publication decision itself is direction-neutral. Those relying on this literature should apply extra scrutiny to anti-intervention findings whose test statistics sit close to the conventional significance boundary on either side (roughly $|z| \in [1.645, 2.275]$), while treating publication in a top-three finance journal as a credible signal of overall methodological quality.

7.2 How Finance Compares to Economics

Studies in economics and related social sciences find that papers reporting statistically significant results are published at substantially higher rates than papers reporting null results. Brodeur, Lé, Sangnier and Zylberberg (2016) show that the top general-interest economics journals (*AER*, *JPE*, *QJE*) contain substantially more test statistics that land just above the conventional significance threshold than just below it—a sign that researchers and editors

favor results that clear the significance bar. Franco, Malhotra and Simonovits (2014) tracked NSF-funded social science experiments from registration through to publication and found that only about 21% of null results were eventually published, compared with roughly 62% of significant results—a roughly three-to-one publication advantage for significant findings. DellaVigna and Linos (2022) report similar disparities in a sample of field experiments, and Brodeur, Cook, Hartley and Heyes (2024) show that pre-registering a study in advance reduces but does not eliminate this advantage.

Our results stand out against this backdrop. Null-finding papers in the finance-regulation literature are published at essentially the same rate as papers that find a positive or negative effect—no three-to-one gap, and no evidence that editors systematically reward significant findings. Our published sample also shows clustering near the $|z| = 1.96$ threshold (caliper ratio $R = 1.20$ in our just-below/just-above convention; equivalently, ≈ 0.83 in the just-above/just-below convention more familiar from Brodeur, Lé, Sangnier and Zylberberg 2016), while the same measure applied to NBER working papers is $R = 1.01$ (essentially symmetric). The published–NBER gap is consistent with editorial or revision-stage selection introducing the asymmetry, but because the two corpora are largely non-overlapping cross-sections, we cannot rule out a compositional explanation—published and unmatched NBER papers may differ on unobserved dimensions (research design, table format, vintage) that themselves correlate with near-threshold density. We discuss this identification limit in Section 7.4. Notably, the published-sample asymmetry is in the opposite direction from the classical Brodeur pattern: excess mass falls on the just-insignificant side rather than just clearing the threshold (see Section 5.1 for our convention and Appendix E for the row-type decomposition).

Two explanations could account for why finance-regulation research looks different from other social science domains. First, editorial norms at the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* may genuinely be less sensitive to whether a paper finds an effect than norms at general-interest economics journals. Second,

because our pre-publication baseline consists of NBER working papers—which tend to come from faculty at well-resourced research universities—null results in this domain may circulate and reach editors more reliably than null results produced by less-connected researchers. If so, the absence of a selection gap in our data could partly reflect the composition of the NBER group rather than editorial neutrality per se. Whether the same pattern holds for lower-ranked finance journals or for finance sub-fields without explicit policy content (such as corporate governance or asset pricing) remains an open question.

Our caliper finding also fits alongside prior work within finance. Harvey and Liu (2021) document near-threshold clustering of test statistics in published finance research and estimate what the pre-publication distribution must have looked like to produce it; our NBER data provide a direct look at that pre-publication distribution and are consistent with their inferred shape, but—given the largely non-overlapping cross-section design—do not by themselves pin the asymmetry on the revision stage. Chen (2021) make a complementary observation: the very largest published test statistics ($|t| > 4$, corresponding to highly significant results) appear too frequently to be explained by selective reporting of marginal findings—implying that most large, convincing effects in the published literature are real. These observations fit together if genuine large effects account for the strong results in the published literature, while revision-stage selection is one candidate channel for the near- $|z| = 1.96$ asymmetry we document—in our data, an excess of statistics in $[1.645, 1.96)$ (just short of conventional significance) relative to $[1.96, 2.275]$ (just clearing it). We note again that the cross-sectional design does not let us distinguish revision-stage selection from compositional differences between the published and NBER samples. One interpretation of why this asymmetry is concentrated in negative-finding papers is that anti-intervention results face elevated scrutiny during peer review, leading authors to report many borderline specifications—some that clear the threshold and many that come close but do not—rather than reducing the reported set to only the just-significant ones. A compositional alternative is that negative-finding published papers differ from negative-finding NBER-only papers

along unobserved dimensions—research design complexity, empirical detail, or vintage—that correlate with the number of specifications reported. This is a cross-sectional comparison of two direction-matched corpora, not a draft-to-publication comparison of the same papers; our design cannot distinguish between these two explanations.

7.3 Alternative Interpretations

Finding equal publication rates across finding directions is consistent with two different explanations, and our data cannot tell them apart.

The first is genuine editorial neutrality: editors evaluate papers on their merits and do not systematically favor those that find an effect. The second is a supply-side equilibrium: if top finance journals have a reputation for publishing null results, researchers who produce null findings may be more willing to submit them in this field than in fields where null results are known to face rejection. Equal publication rates would then emerge not because editors are neutral but because the papers arriving on editors’ desks are already balanced across directions. The near-identical share of directional findings in NBER working papers (79.6%) and published articles (79.1%) is consistent with this story.

The second finding—that papers using credible causal methods are more likely to appear in the published corpus—also warrants caution in interpretation. The quasi-experimental coefficient (0.467^{***}) is a cross-sectional association conditional on year, team size, top-institution status, and finding direction. Research design quality also tends to travel with other attributes (better-resourced research groups, more carefully crafted papers, richer data) that our model does not observe, so the coefficient should be read as evidence that the published corpus disproportionately contains a bundle of quality-related attributes correlated with identification strategy, not as a causal estimate of how editors weight identification strategy in the acceptance decision.

7.4 Limitations

Our design has six limitations worth flagging. First, only 16 of 134 published papers (11.9%) can be matched to a specific NBER counterpart, so we are comparing two largely non-overlapping groups of papers rather than tracking the same paper through the editorial process. We cannot directly observe whether any paper was rejected because of its findings; we can only compare the distribution of findings across the two groups. All direction coefficients should be read as cross-sectional associations, not causal effects of editorial filtering.

Second, we define publication as appearing in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics*, so the outcome variable measures top-three placement rather than broader peer-reviewed publication. Under our stated outcome, NBER papers that ended up in other peer-reviewed journals are *correctly* coded as not in the top three—there is no misclassification. The only sense in which this is a limitation is for a reader who wants to generalize to broader peer-reviewed publication; in that case our outcome is a strict subset of the broader one, and the direction of any resulting bias in the relative direction estimates depends on whether positive-, negative-, or null-finding NBER papers are differentially likely to publish in second-tier outlets. We do not have data on second-tier journal placements, so we cannot sign this bias and we do not claim our top-three null is conservative for the broader outcome.

Third, our AI pipeline codes direction from abstracts, and NBER abstracts are written more tentatively than published ones (high-confidence rate: 62.4% vs. 81.3%). Our validation exercise shows that null is the least stable category: 15 of 20 null-coded validation papers were re-coded as directional in an independent replication, with the flips concentrated on the published side (13 of 15 published nulls flipped vs. 2 of 5 NBER nulls). Because the sign of the resulting bias on the direction premium depends on whether the miscoding is concentrated in NBER (which would inflate the premium) or in published (which would deflate it), the validation evidence does not let us sign the residual measurement error. The

null-finding publication rate ($\approx 46.7\%$) is the least reliable of the three rates, but a joint equality test across all three categories ($\chi^2 = 0.017$, $p = 0.992$) confirms the overall null without relying on any single category as the reference.

Fourth, the 18-percentage-point quasi-experimental premium could mask a direction effect if rigorous studies also tend to find directional results. This does not appear to be the case: direction coefficients are unchanged whether or not research design quality is controlled, and Table 9 shows flat publication rates across directions within both rigorous and non-rigorous studies. The AI classification of research design is less reliable than the direction coding — many abstracts do not name their methods explicitly — and has not been separately validated.

Fifth, NBER researchers come from a select group at leading universities. If their papers are inherently more publishable, our pre-publication baseline sets a higher bar than the average unpublished paper, and we would understate any true direction premium. The near-identical share of directional findings in the NBER (79.6%) and published (79.1%) samples is the most reassuring evidence against this concern.

Sixth, the caliper analysis runs 18 comparisons across subsamples, journals, and directions. Applying the Bonferroni correction ($p < 0.003$ for 18 tests) to the *naive* chi-square p -values, three caliper tests survive: the full published sample ($p = 0.002$), the *Journal of Financial Economics* ($p < 0.001$), and negative-finding papers ($p < 0.001$); two density tests—*Journal of Finance* and null/mixed ($p < 0.001$ each)—also survive. However, the naive chi-square treats every test statistic as independent and we have flagged it as overstating precision relative to our preferred paper-clustered bootstrap. Under the bootstrap, the published full sample ($p = 0.150$) and negative-finding sample ($p = 0.135$) do not clear 0.05 even before correction, so we cannot claim these results “survive” multiple-testing correction on our preferred test. The Bonferroni-surviving set above should therefore be read as the set that survives under the less conservative inferential standard, not as a definitive multiple-testing-robust finding. The *Journal of Finance* density result is included for completeness, but as

noted in Section 5.1, it conflicts with the *Journal of Finance* caliper ratio (0.88, $p = 0.423$); given the small *Journal of Finance* sample (26 papers, 1,308 statistics), we give priority to the caliper and do not interpret the density finding as evidence of bunching. The NBER caliper tests do not approach significance under any of the three test versions ($p > 0.42$ under both naive chi-square and bootstrap), and the main direction regression is pre-specified in a single test ($p > 0.90$), where the concern is power rather than false discovery.

8. Conclusion

This paper asks whether top finance journals are more likely to publish government-intervention studies that find an effect than studies that find none. Using a sample of 134 published articles from the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* and 157 NBER working papers as a pre-publication benchmark, we find no evidence of direction-based editorial selection at a scale our design can detect. Positive-, negative-, and null-finding papers occupy the published side of our combined 291-paper corpus at nearly identical corpus-membership shares — 45.6%, 46.1%, and 46.7% respectively (within-direction shares, see the terminology convention at the start of Section 5.2) — and finding direction is statistically unrelated to published-side corpus membership across all seven model specifications. The design rules out gaps larger than 18 percentage points at 80% power; smaller asymmetries remain possible.

The distribution of test statistics tells a more nuanced story. As shown in Table 5, the full-sample caliper ratio of 1.20 reflects two reinforcing effects: published papers cluster standalone z -statistics more sharply near the significance threshold than NBER papers do (ratio 2.54 vs. 1.58), and they devote a larger share of table rows to such statistics (28.3% vs. 18.1%). Verifiable coefficient rows show no systematic threshold manipulation in either corpus. The bunching is most pronounced in negative-finding papers (ratio 1.52), though this pattern is suggestive rather than statistically definitive.

Together these two results point to different stages of the research pipeline. The absence

of detectable large-scale direction-based sorting between published and pre-publication papers suggests that editorial selection in this domain is not strongly tilted toward finding an effect. The within-paper bunching just below the $|z| = 1.96$ threshold, concentrated in published negative-finding papers, is more consistent with revision-stage pressure than with outright fabrication of results. We flag this as suggestive rather than definitive: the bunching pattern clears a Bonferroni cutoff under a naive chi-square that treats every test statistic as independent, but does not clear conventional significance under our preferred paper-clustered bootstrap ($p = 0.135\text{--}0.150$, see Section 5.1).

Among observed paper characteristics, the two largest cross-sectional predictors of corpus membership are research-design quality and institutional prestige, not finding direction. Conditional on year, team size, and direction, QE papers (RDD, IV, DiD, natural experiment, or event study with causal interpretation) are 18–20 percentage points more likely to appear in the published corpus, and papers from UTD Top 100 institutions are 24–26 percentage points more likely (Table 10); in the fully specified Column (6) the institutional-prestige marginal effect (26.2 pp) actually exceeds the QE marginal effect (18.8 pp), so we do not claim research design is uniformly the single strongest predictor. What we can say is that both QE and top-institution status are substantively large cross-sectional predictors that exceed any direction-based predictor in our sample, and the two are essentially orthogonal to each other ($r = 0.015$). These associations are consistent with journals selecting on a bundle of quality attributes— rigorous identification, careful data construction, well-crafted writing—of which identification strategy and pedigree are two components.

For policymakers and researchers who use this literature, the main takeaway is that the published record of government-intervention research in top finance journals does not appear to be large-scale tilted toward positive or negative findings by editorial selection. The more pressing concern is within-paper: test statistics clustered just below the significance threshold in published negative-finding papers suggest that authors may be adjusting their specifications during revision to clear the significance bar. Those relying on this evidence

base should treat marginal significance in anti-intervention findings with additional scrutiny, while treating publication in a top-three finance journal as a credible signal of methodological quality.

References

- Agarwal, S., Lucca, D., Seru, A., Trebbi, F., 2014. Inconsistent regulators: Evidence from banking. *Quarterly Journal of Economics* 129, 889–938.
- Andrews, I., Kasy, M., 2019. Identification of and correction for publication bias. *American Economic Review* 109, 2766–2794.
- Ash, E., Hansen, S., 2023. Text algorithms in economics. *Annual Review of Economics* 15, 659–688.
- Berger, A.N., Bouwman, C.H.S., 2017. Bank liquidity creation, monetary policy, and financial crises. *Journal of Financial Stability* 30, 139–155.
- Brodeur, A., Cook, N., Hartley, J., Heyes, A., 2024. Do pre-registration and pre-analysis plans reduce p-hacking and publication bias? Evidence from 15,992 test statistics and suggestions for improvement. *IZA Discussion Paper* 15476.
- Brodeur, A., Lé, M., Sangnier, M., Zylberberg, Y., 2016. Star wars: The empirics strike back. *American Economic Journal: Applied Economics* 8, 1–32.
- Chen, A.Y., 2021. The limits of p-hacking: Some thought experiments. *Journal of Finance* 76, 2447–2480.
- DellaVigna, S., Linos, E., 2022. Rcts to scale: Comprehensive evidence from two nudge units. *Econometrica* 90, 81–116.
- Franco, A., Malhotra, N., Simonovits, G., 2014. Publication bias in the social sciences: Unlocking the file drawer. *Science* 345, 1502–1505.
- Harvey, C.R., Liu, Y., 2021. Uncovering the iceberg from its tip: A model of publication bias and p-hacking. *SSRN Working Paper* 3865813.

- Harvey, C.R., Liu, Y., Zhu, H., 2016. ... and the cross-section of expected returns. *Review of Financial Studies* 29, 5–68.
- Hou, K., Xue, C., Zhang, L., 2020. Replicating anomalies. *Review of Financial Studies* 33, 2019–2133.
- Jensen, T.I., Kelly, B.T., Pedersen, L.H., 2023. Is there a replication crisis in finance? *Journal of Finance* 78, 2465–2518.
- Korinek, A., 2023. Language models and cognitive automation for economic research. NBER Working Paper 30957.
- La Porta, R., Lopez-de Silanes, F., Shleifer, A., Vishny, R.W., 1998. Law and finance. *Journal of Political Economy* 106, 1113–1155.
- Landis, J.R., Koch, G.G., 1977. The measurement of observer agreement for categorical data. *Biometrics* 33, 159–174.
- Naveen Jindal School of Management, University of Texas at Dallas, 2026. The UTD Top 100 Business School Research Rankings. <https://jsom.utdallas.edu/the-utd-top-100-business-school-research-rankings/>. 2021–2025 worldwide ranking period. Accessed June 2026.

Appendix A. Keyword Taxonomy

Table A1 lists all terms used in Pass 1 of the classification pipeline.

Table A1. In-Scope Classification: Keyword Taxonomy

Category	Keywords
Financial regulation	regulation, regulatory, Dodd-Frank, Basel, Glass-Steagall, capital requirements, leverage ratio, bank capital
Securities law	SEC, securities and exchange, disclosure requirement, Reg FD, Sarbanes-Oxley, SOX, insider trading law, short-selling ban, short sale ban, short-sale restriction
Banking policy	FDIC, deposit insurance, lender of last resort, bailout, TARP, stress test, bank resolution, government guarantee
Monetary/fiscal policy	quantitative easing, QE, Fed intervention, government subsidy, public guarantee, state guarantee, fiscal stimulus
Market microstructure rules	circuit breaker, tick size, trading halt, uptick rule, limit-down, limit-up, trading curb
Corporate governance law	board independence requirement, say-on-pay, proxy access, mandatory audit, governance mandate, clawback provision, mandatory disclosure
General intervention signals	government intervention, government policy, public policy, law and finance, legal protection, shareholder rights, creditor rights, investor protection, financial regulation reform

Ambiguous standalone terms (*regulation*, *regulatory*, *SEC*) require co-occurrence with at least one additional in-scope keyword before the paper is flagged as a candidate.

Appendix B. LLM Prompts

2.1 In-Scope Classification Prompt

The prompt receives the paper’s title and abstract. System instruction: *“You are a research assistant helping to classify finance research papers. You must return valid JSON only — no other text.”* User message:

A finance research paper is ‘in scope’ for a study of directional bias in government intervention research if and only if its PRIMARY research question is the CAUSAL EFFECT of a specific government law, regulation, or public policy on financial markets, firms, or investors.

Exclusion criteria (mark `in_scope=false`):

- Purely theoretical papers (no empirical test of a policy)
- Papers that only DESCRIBE a regulation without estimating its effect
- Papers whose main topic is private firm/market behaviour, with policy as incidental context
- Papers studying central bank communication WITHOUT a specific policy instrument
- Papers whose primary policy variable is a rule imposed by a stock exchange (NYSE, NASDAQ, etc.) without a specific government or regulatory mandate (e.g., exchange-designed trading halts, voluntary tick size conventions, exchange fee structures)

Return JSON with exactly these keys: `in_scope` (true or false), `confidence` (“high” | “medium” | “low”), `reason` (one sentence).

2.2 Direction Coding Prompt

The deployed pipeline has two paths, selected per paper by the classifier in `scripts/04_code_direction.py`. When the abstract is present and contains at

least 20 words, the model receives the *abstract-only* prompt below; the title is *not* concatenated with the abstract or otherwise passed. When the abstract is missing or shorter than 20 words, the model instead receives the *title-only* fallback prompt also shown below. All 291 papers in the final analysis sample have abstracts of at least 20 words, so the title-only fallback was not invoked for any analysis-sample paper; the validation samples and AFA pipeline use the same two-path logic.

System instruction (both prompts): *“You are a research assistant classifying finance paper findings. Return valid JSON only — no other text.”*

Abstract-only prompt (used for all 291 analysis-sample papers).

The following is the abstract of a finance paper that studies the effect of a government law, regulation, or public policy.

Classify the MAIN empirical finding:

+1 The government intervention produced a POSITIVE / BENEFICIAL outcome (e.g. improved market quality, reduced systemic risk, better investor protection, increased firm value, lower cost of capital)

-1 The government intervention produced a NEGATIVE / HARMFUL outcome (e.g. reduced liquidity, higher costs, unintended consequences, welfare losses)

0 NULL, MIXED, or INCONCLUSIVE finding (no statistically significant effect, evidence on both sides, or the paper explicitly refuses to take a directional stance)

Coding rules:

— Code the PRIMARY intervention and PRIMARY outcome (named in title or first sentence)

— If paper is purely theoretical with no empirical test, set direction=0, confidence=low

— Do NOT infer direction from normative framing; code the sign of the estimated

coefficient

Abstract:

```
""{abstract}""
```

Return JSON with exactly: direction (1 | -1 | 0), confidence (“high” | “medium” | “low”), rationale (one sentence).

The instruction “named in title or first sentence” is preserved from the deployed prompt verbatim; because only the abstract is passed in this path, the model interprets “title” as title-like leading material that some abstracts include in their first sentence (e.g., naming the specific statute being studied). Body Section 3.4 paraphrases the spirit of the rule—coding the sign of the estimated effect relative to the stated regulatory goal—which the prompt conveys through the example outcomes for +1 (regulatory objectives achieved: improved market quality, reduced systemic risk, etc.) and −1 (regulatory side-effects: reduced liquidity, higher costs, etc.).

Title-only fallback (not used for any analysis-sample paper).

The following is the title of a finance paper that studies the effect of a government law, regulation, or public policy. No abstract is available.

Based on the title alone, infer the likely direction of the main empirical finding:

+1 The intervention likely produced a POSITIVE / BENEFICIAL outcome

-1 The intervention likely produced a NEGATIVE / HARMFUL outcome

0 NULL, MIXED, or cannot be determined from the title alone

Coding rules:

- Default to 0 (null) if the title is neutral or provides no directional signal
- Only assign +1 or −1 if the title clearly implies a directional finding
- Set confidence to “low” unless the directional signal is unambiguous

Title:

""{title}""

Return JSON with exactly: direction (1 | -1 | 0), confidence (“high” | “medium”
| “low”), rationale (one sentence).

Appendix C. Replication Confusion Matrix

Table A2. LLM Direction Coding: Replication Confusion Matrix

Original code	Replication code			Row total
	−1	0	+1	
Negative (−1)	20	0	0	20
Null/Mixed (0)	9	5	6	20
Positive (+1)	2	0	18	20
Column total	31	5	24	60

Notes. To assess how consistently the AI classifies paper findings, we drew a random sample of 60 papers (20 from each direction category) and ran the direction-coding prompt a second time using the same model (claude-haiku-4-5-20251001) in an independent second call. Each cell shows how many papers the two runs agreed or disagreed on. Rows are the original codes used in the main analysis; columns are the codes from the second independent run. Bold entries on the main diagonal are papers where both runs assigned the same code—these are the agreements. Off-diagonal entries are disagreements. Overall, the two runs agree on 43 of 60 papers (71.7%). Cohen’s weighted $\kappa_w = 0.674$ (linear weights), which the Landis and Koch (1977) scale classifies as “substantial agreement.” Agreement is perfect for negative-finding papers (20/20 = 100%) and high for positive-finding papers (18/20 = 90%), but much lower for null/mixed papers (5/20 = 25%). When the two runs disagree on a null-coded paper, the replication almost always re-classifies it as directional, suggesting that abstracts with inconclusive findings are harder for the AI to classify consistently. The corpus breakdown of these null-category flips matters for signing the measurement-error bias. The 20 originally null-coded validation papers comprise 5 NBER working papers and 15 published articles (*Journal of Finance*: 5, *Review of Financial Studies*: 5, *Journal of Financial Economics*: 5). Of the 15 papers re-coded as directional in replication, 2 are NBER and 13 are published. The direction of bias in our estimated direction premium depends on which side dominates: NBER directional papers miscoded as null would inflate the premium (and make our null finding conservative), while published directional papers miscoded as null would deflate the premium (and make the null finding easier to obtain). The replication evidence shows null-coding instability is concentrated on the *published* side, so we do not claim our estimates are conservative; the validation establishes instability of the null category but does not establish a clean sign for the residual measurement error.

Appendix D. Caliper Robustness to Coefficient Inclusion

A potential concern is that pooling all regression-table rows—including standard controls—could inflate the caliper count without reflecting publication pressure on headline results. Table A3 addresses this by re-running the caliper test after dropping rows whose variable name matches common control patterns: firm size, leverage, profitability (ROA/ROE/EBITDA), book-to-market, Tobin’s q , return volatility, momentum, cash flow, market capitalization, intercepts, year dummies, and fixed-effect indicators, as well as macro controls (GDP, inflation, interest rates). The retained “non-control” rows are predominantly treatment indicators, DiD interactions, and main-result variables. The restriction drops fewer than 10% of rows (7,055 of 7,761 published and 7,599 of 8,351 NBER statistics are retained). Caliper ratios are essentially unchanged: the published full-sample ratio remains 1.203 and the negative-finding ratio *increases* to 1.625 (from 1.522), while NBER ratios stay near 1.0. As in the main analysis, the published-side results clear the naive chi-square at very small p -values but do not reach 0.05 on the paper-clustered bootstrap ($p = 0.161$ for the non-control full sample, $p = 0.121$ for the non-control negative-finding subsample); the *** stars in the table reflect the naive chi-square only. The pattern that the caliper result is not an artifact of pooling control coefficients therefore holds at the level of the point estimates and the naive test, but does not strengthen the result on the preferred test.

Appendix E. Caliper Test Stratified by Row Type: Coef+SE vs. Standalone z -Statistics

Regression tables present statistical results in two ways. Some rows show a coefficient alongside its standard error, which allows a reader to independently verify the test statistic by dividing one by the other. Other rows report only a t - or z -statistic directly, with no coefficient or standard error shown, making it impossible to cross-check the number. Both types appear routinely in published finance papers. Table A4 asks whether the threshold-

Table A3. Caliper Test: All Statistics vs. Non-Control Statistics

Sample	Subset	Filter	N	Papers	Ratio	Naive p	Perm. p	Boot. p
Published	Full sample	All stats	7,761	108	1.203***	0.002	0.001	0.150
Published	Full sample	Non-control	7,055	108	1.203***	0.003	0.001	0.161
Published	Negative-finding	All stats	3,491	50	1.522***	<0.001	<0.001	0.135
Published	Negative-finding	Non-control	3,139	50	1.625***	<0.001	<0.001	0.121
Published	Positive-finding	All stats	2,633	37	0.941	0.565	0.732	0.623
Published	Positive-finding	Non-control	2,409	37	0.927	0.481	0.769	0.637
NBER	Full sample	All stats	8,351	106	1.012	0.849	0.424	0.432
NBER	Full sample	Non-control	7,599	104	0.981	0.767	0.625	0.561
NBER	Negative-finding	All stats	4,375	50	0.996	0.964	0.532	0.474
NBER	Negative-finding	Non-control	3,981	50	0.955	0.633	0.693	0.614

Notes. Each paper’s regression tables contain two kinds of rows: (1) rows for the main variable of interest (e.g., a treatment indicator or policy dummy), and (2) rows for standard control variables (e.g., firm size, leverage, profitability, GDP). The “All stats” rows use every regression coefficient; the “Non-control” rows keep only the main-variable rows, dropping controls such as size, leverage, ROA/ROE, book-to-market, GDP, inflation, firm age, return volatility, fixed-effect indicators, and intercepts. The *Ratio* column is $R = \#\{|z| \in [1.645, 1.96)\} / \#\{|z| \in [1.96, 2.275]\}$. $R = 1.0$ means equal frequency; $R > 1.0$ indicates excess clustering on the just-below-1.96 side (just-insignificant results), and $R < 1.0$ indicates excess clustering on the just-above-1.96 side (just-significant results). *Naive p*: standard chi-squared test. *Perm.*: permutation test that accounts for the fact that statistics from the same paper are correlated (10,000 iterations). *Boot.*: bootstrap that gives each paper equal weight regardless of how many statistics it contributes (10,000 draws). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$ refer to the *naive* chi-square only; under the paper-clustered bootstrap the published-side p -values are 0.121–0.161 and do not reach conventional significance. Point estimates of the ratio are virtually identical whether or not control rows are included, and the clustering among negative-finding papers strengthens (in point estimate) when controls are dropped.

clustering pattern differs between these two types of rows—that is, whether the excess mass just below the 5% significance cutoff is concentrated in the rows readers can verify or in the rows they cannot.

Table A4. Caliper Test Stratified by Row Type

Corpus	Row type	N	Below	Above	Ratio	Naive p
Published	All rows	7,761	611	508	1.203*** $\uparrow$	0.002
Published	Coef+SE only	5,563	322	394	0.817*** $\downarrow$	0.007
Published	z -stat only	2,198	289	114	2.535*** $\uparrow$	<0.001
NBER	All rows	8,351	501	495	1.012	0.849
NBER	Coef+SE only	6,840	422	445	0.948	0.435
NBER	z -stat only	1,511	79	50	1.580*** $\uparrow$	0.011

Notes. Regression tables report statistical results in two ways. Some rows report a coefficient and its standard error side by side, which allows the test statistic to be independently verified by dividing one by the other (*Coef+SE only* rows). Other rows report only a t - or z -statistic directly, with no coefficient or standard error shown, so there is no way to cross-check the number (z -stat only rows). The *Ratio* column is $R = \#\{|z| \in [1.645, 1.96)\} / \#\{|z| \in [1.96, 2.275]\}$. $R = 1.0$ means equal frequency; $R > 1.0$ indicates excess clustering on the just-below-1.96 side, and $R < 1.0$ indicates excess clustering on the just-above-1.96 side. *Below* and *Above* are the raw counts of test statistics falling in the windows just below and just above the 5% cutoff. *Naive p* tests whether those two counts are equal. Significance markers are two-sided: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$ for the test that the below and above counts are equal. Direction arrows indicate which side has the excess: $\uparrow$ flags $R > 1$ (excess on the just-below-1.96 side, the focal pattern), and $\downarrow$ flags $R < 1$ (excess on the just-above-1.96 side). The published Coef+SE-only row is significant at $p = 0.007$ in the $\downarrow$ direction: among verifiable rows in published papers, the sample shows more statistics in $[1.96, 2.275]$ than in $[1.645, 1.96)$. This is the classical Brodeur-style “just-cleared-the-bar” pattern, in the opposite direction from the headline excess driven by standalone z -statistics; we describe both directions in Section 5.1. Across the table, threshold-side asymmetry is concentrated in the standalone z -statistics, while the verifiable rows show a smaller asymmetry in the opposite direction.