

Online Appendix to “Do Finance Journals Pick Sides? Evidence from Government Intervention Research”

Overview

This Online Appendix collects three supporting materials referenced in the main paper: the full keyword taxonomy used at the first-pass classification stage (Online Appendix OA.1), the verbatim LLM prompts used to assign in-scope and direction codes (Online Appendix OA.2), and the sample-frame precision and recall audit (Online Appendix OA.3). Section, table, and figure numbers in this Online Appendix begin with “OA”. References to sections of the main paper retain their main-paper numbering.

OA.1. Keyword Taxonomy

Table OA.1 lists all terms used in Pass 1 of the classification pipeline.

Table OA.1. In-Scope Classification: Keyword Taxonomy

Category	Keywords
Financial regulation	regulation, regulatory, Dodd-Frank, Basel, Glass-Steagall, capital requirements, leverage ratio, bank capital
Securities law	SEC, securities and exchange, disclosure requirement, Reg FD, Sarbanes-Oxley, SOX, insider trading law, short-selling ban, short sale ban, short-sale restriction
Banking policy	FDIC, deposit insurance, lender of last resort, bailout, TARP, stress test, bank resolution, government guarantee
Monetary/fiscal policy	quantitative easing, QE, Fed intervention, government subsidy, public guarantee, state guarantee, fiscal stimulus
Market microstructure rules	circuit breaker, tick size, trading halt, uptick rule, limit-down, limit-up, trading curb
Corporate governance law	board independence requirement, say-on-pay, proxy access, mandatory audit, governance mandate, clawback provision, mandatory disclosure
General intervention signals	government intervention, government policy, public policy, law and finance, legal protection, shareholder rights, creditor rights, investor protection, financial regulation reform

Ambiguous standalone terms (*regulation*, *regulatory*, *SEC*) require co-occurrence with at least one additional in-scope keyword before the paper is flagged as a candidate.

OA.2. LLM Prompts

OA.2.1 In-Scope Classification Prompt

The prompt receives the paper’s title and abstract. System instruction: “You are a research assistant helping to classify finance research papers. You must return valid JSON only, no other text.” User message:

A finance research paper is ‘in scope’ for a study of directional bias in government intervention research if and only if its PRIMARY research question is the CAUSAL EFFECT of a specific government law, regulation, or public policy on financial markets, firms, or investors.

Exclusion criteria (mark in_scope=false):

- , Purely theoretical papers (no empirical test of a policy)
- , Papers that only DESCRIBE a regulation without estimating its effect
- , Papers whose main topic is private firm/market behaviour, with policy as incidental context
- , Papers studying central bank communication WITHOUT a specific policy instrument
- , Papers whose primary policy variable is a rule imposed by a stock exchange (NYSE, NASDAQ, etc.) without a specific government or regulatory mandate (e.g., exchange-designed trading halts, voluntary tick size conventions, exchange fee structures)

Return JSON with exactly these keys: in_scope (true or false), confidence (“high” | “medium” | “low”), reason (one sentence).

OA.2.2 Direction Coding Prompt

The deployed pipeline has two paths, selected per paper by the direction-coding classifier. When the abstract is present and contains at least 20 words, the model receives the *abstract-only* prompt below; the title is not concatenated with the abstract

or otherwise passed. When the abstract is missing or shorter than 20 words, the model instead receives the *title-only* fallback prompt also shown below. All 291 papers in the final analysis sample have abstracts of at least 20 words, so the title-only fallback was not invoked for any analysis-sample paper; the validation samples also use the same two-path logic.

System instruction (both prompts): “You are a research assistant classifying finance paper findings. Return valid JSON only, no other text.”

Abstract-only prompt (used for all 291 analysis-sample papers).

The following is the abstract of a finance paper that studies the effect of a government law, regulation, or public policy.

Classify the MAIN empirical finding:

+1 The government intervention produced a POSITIVE / BENEFICIAL outcome (e.g. improved market quality, reduced systemic risk, better investor protection, increased firm value, lower cost of capital)

-1 The government intervention produced a NEGATIVE / HARMFUL outcome (e.g. reduced liquidity, higher costs, unintended consequences, welfare losses)

0 NULL, MIXED, or INCONCLUSIVE finding (no statistically significant effect, evidence on both sides, or the paper explicitly refuses to take a directional stance)

Coding rules:

, Code the PRIMARY intervention and PRIMARY outcome (named in title or first sentence)

, If paper is purely theoretical with no empirical test,

set direction=0, confidence=low

, Do NOT infer direction from normative framing; code the sign of the estimated coefficient

Abstract:

[ABSTRACT TEXT HERE]

Return JSON with exactly: direction (1 | -1 | 0), confidence (“high” | “medium” | “low”), rationale (one sentence).

The instruction “named in title or first sentence” is preserved from the deployed prompt verbatim; because only the abstract is passed in this path, the model interprets “title” as title-like leading material that some abstracts include in their first sentence (e.g., naming the specific statute being studied). Section 3.4 of the main paper paraphrases the spirit of the rule, coding the sign of the estimated effect relative to the stated regulatory goal, which the prompt conveys through the example outcomes for +1 (regulatory objectives achieved: improved market quality, reduced systemic risk, etc.) and −1 (regulatory side-effects: reduced liquidity, higher costs, etc.).

Title-only fallback prompt (documented for pipeline completeness; not invoked for any paper in the 291-paper analysis sample).

The following is the title of a finance paper that studies the effect of a government law, regulation, or public policy. No abstract is available.

Based on the title alone, infer the likely direction of the main empirical finding:

+1 The intervention likely produced a POSITIVE / BENEFICIAL outcome

-1 The intervention likely produced a NEGATIVE / HARMFUL outcome

0 NULL, MIXED, or cannot be determined from the title alone

Coding rules:

, Default to 0 (null) if the title is neutral or provides no directional signal

, Only assign +1 or −1 if the title clearly implies a directional finding

, Set confidence to “low” unless the directional signal is unambiguous

Title:

[TITLE TEXT HERE]

Return JSON with exactly: direction (1 | -1 | 0), confidence (“high” | “medium”
| “low”), rationale (one sentence).

OA.3. Sample Frame Validation: Precision and Recall

The selection pipeline narrows 39,136 documents to 1,849 keyword candidates and then to 291 final papers (Table 1 of the main paper). Differential exclusion at any stage, between NBER and journal papers, or between null and directional papers, could bias the within-direction corpus-membership shares that our main inference relies on. This appendix reports three precision/recall estimates that bound this concern.

LLM precision on its in-scope decisions

We re-evaluated 68 papers that the LLM flagged as in-scope to test whether the in-scope label was correct under the boundary rule of Section 3.3 of the main paper. The author overrode 10 of the 68, giving an implied LLM precision (against the manual gold standard) of $58/68 = 85.3\%$. Across the full keyword-positive pool, the LLM flagged 217 papers as in-scope at the first pass; manual review trimmed this to 185 (32 of the initial 217 were overridden: 10 NBER and 22 published, with the published-side split JF: 18, RFS: 2, JFE: 2). The NBER refilter then promoted 126 additional papers to in-scope, of which 20 were overridden on second-pass review and 106 survived, yielding the final 291-paper sample ($217 - 32 + 106 = 291$, equivalently $217 + 126 - 52 = 291$ with the combined override count). Every per-paper override is recorded in the replication package; the boundary criterion (policy as primary research question vs. as identification device for a non-policy outcome) is operationalized in

Section 3.3 of the main paper.

LLM recall on the keyword-positive pool

The NBER refilter re-coded 803 NBER papers that the LLM had flagged as low-confidence out-of-scope; 126 were promoted to in-scope and 106 survived the same manual review used elsewhere. Treating manual + refilter as the gold standard restricted to the NBER subset, the LLM’s first-pass recall on this low-confidence pool is $1 - 126/803 = 84.3\%$, comparable to its precision. The published-side LLM decisions were not subjected to a symmetric refilter because the journal denominators (8,924 articles) are small enough that the keyword filter retains all plausible candidates at the first stage, but the symmetry between the precision and recall point estimates suggests no strong directional asymmetry in LLM errors within the keyword-positive pool.

Keyword-filter recall: audit of keyword-negative papers

Stage 1 (the keyword filter) is by design over-inclusive, but its recall on truly in-scope papers cannot be inferred from Table 1 of the main paper alone. We audit the keyword-negative pool by sampling 200 papers (random seed 42) from the 37,287 documents that fail the keyword filter and have abstracts of at least 20 words, and re-running the in-scope LLM prompt on each. Of the 200 sampled papers, the LLM flagged 4 as in-scope: false-negative rate 2.0% (95% binomial CI: 0.5% to 4.0%). Extrapolated to the full 37,287-paper keyword-negative pool with usable abstracts, this implies roughly 746 in-scope papers (95% CI: 186 to 1,491) that the keyword filter misses; adjusting for the $\approx 85\%$ LLM precision documented above brings the implied count of manually-validated in-scope misses to roughly 634 (95% CI: 158 to 1,267).

The 4 audited keyword-misses concern papers studying specific government policies

(a tax rule change in commercial mortgage markets, voluntary restraint arrangements with Japan, bank nationalization, foreign royalty tax rates on R&D) whose abstracts happen not to contain any term in our seven-category taxonomy. Per-source false-negative rates in the small audit sample are not significantly differentiated (JF: 1/14, RFS: 1/7, JFE: 0/15, NBER: 2/164), but the journal side carries a slightly higher point estimate. If a true differential exists, the published-side denominator is modestly understated relative to the NBER denominator; the implied bias on the within-direction corpus-membership shares depends additionally on how the missed papers distribute across direction categories, which we cannot estimate from this small sample.

Interpretation

Our main inference rests on the equality of direction shares across the NBER and published samples within the selected corpus. A direction-neutral keyword miss leaves these shares unchanged; a direction-correlated keyword miss would bias them. The 200-paper audit is too small to detect such differential exclusion if it exists.