Do Finance Journals Sort on Policy Conclusions?

Emre Kuvvet*

June 16, 2026

ABSTRACT

We study whether top finance journals favor certain policy conclusions. Using 121 published articles in the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* alongside 125 NBER working papers on government intervention in financial markets (2000–2024), we find that journals do not sort on finding direction. Whether a paper concludes that a regulation is beneficial, harmful, or inconsequential, it reaches a top journal at roughly the same rate, about 46 percent. The strongest predictor of publication is research design: papers using quasi-experimental methods are about 18 percentage points more likely to be published, regardless of what they find. We also find that published papers show modest threshold bunching in supplementary test statistics but not in their main regression tables, suggesting that any specification searching is concentrated at the margins of the paper rather than in its core results.

JEL classification: G18, G28, C80, B40.

Keywords: Publication bias; Government intervention; Finance journals; Selective reporting;

Regulatory finance; p-hacking

*Department of Finance, Nova Southeastern University. E-mail: ekuvvet@nova.edu.

1. Introduction

Academic research on government intervention in financial markets occupies a privileged position in public policy. The Federal Reserve Board, the Securities and Exchange Commission, and the Basel Committee on Banking Supervision routinely incorporate peer-reviewed finance research into regulatory impact assessments, rule-makings, and policy speeches. The Dodd-Frank Wall Street Reform and Consumer Protection Act alone generated an estimated 400 new regulatory mandates, many of which were designed in explicit reference to academic evidence on the effects of capital requirements, disclosure rules, and consumer financial protection. The credibility of this evidence base is therefore a first-order policy question: if the academic record systematically misrepresents which interventions have measurable effects—and which do not—regulators face a distorted picture that can generate welfare-reducing policy choices.

Publication bias offers a mechanism by which such distortion arises without any deliberate misconduct. When researchers are reluctant to write up null results, when journals are more likely to accept papers with statistically significant findings, or when both effects operate simultaneously, the published literature understates the true prevalence of null or inconclusive evidence (Franco et al., 2014). In a domain such as government intervention research, where policy decisions depend on whether a regulation has a measurable causal effect, suppression of null results generates the false impression that most interventions work—or, more precisely, that most interventions have discernible effects in one direction or the other.

Why might directional publication bias arise in the government-intervention literature specifically? Two theoretical channels operate in opposite directions, making the sign of any bias an open empirical question. A *pro-intervention channel* could arise if journals and audiences favor results that confirm a policy's stated objective (e.g., that capital requirements reduce bank risk), both because positive results are more actionable for regulators and because they more naturally confirm the prior that motivated the policy. An *anti-intervention* *channel* could arise if research that documents unintended costs or failures of regulation is more surprising—and thus more publishable—than research confirming that a policy works as intended. Under either channel, the direction of sorting would differ from significance bias (null vs. directional), which reflects a simpler threshold between finding-something and finding-nothing. This paper is designed to detect both types.

This paper studies sorting patterns in a domain the existing literature has not examined: the finance literature's treatment of government intervention. Prior work on publication bias in finance has focused on the cross-section of asset pricing anomalies, documenting that factor-zoo proliferation reflects selection for high t-statistics (Harvey et al., 2016; Hou et al., 2020; Jensen et al., 2023). A parallel economics literature (Brodeur et al., 2016; Andrews and Kasy, 2019) infers bias from the shape of published p-value distributions, finding excess mass just below conventional significance thresholds. Neither strand addresses the policy-intervention literature, nor do they distinguish two conceptually distinct forms of bias with different welfare implications. *Significance bias*—directional findings over-published relative to null findings—makes every intervention appear detectable. *Directional bias*—pro-intervention findings over-published relative to anti-intervention findings—additionally skews the regulatory record toward a particular conclusion about the sign of policy effects. We design our study to test for both.

Our research design combines two complementary samples. The *published paper* sample consists of all in-scope articles in the *Journal of Finance*, the *Review of Financial Studies*, and the *Journal of Financial Economics* from 2000 to 2024. The *working paper* sample is drawn from NBER working papers in five programs covering the dominant share of academic finance research on government intervention: Corporate Finance, Asset Pricing, Monetary Economics, Economic Fluctuations and Growth, and Public Economics. The key advantage of this design is that we observe findings from papers that circulated *before* potential journal publication, allowing a direct two-corpus comparison rather than inferring the pre-publication distribution from the shape of published p-values alone.

Because only 12% of published papers are matched to a confirmed NBER counterpart, the two corpora are largely non-overlapping cross-sections rather than the same papers tracked through the editorial filter. We therefore interpret all estimates as cross-sectional associations between paper characteristics and corpus membership, not as causal editorial-filter effects; the full identification discussion is in Section 7.4.

We define a paper as in-scope if its primary research question concerns the causal effect of a specific government law, regulation, or public policy on financial markets, firms, or investors. Identifying these papers across nearly 40,000 documents required an automated approach: a keyword filter followed by a large language model (LLM) classifier that reads each abstract and assigns an in-scope judgment along with a directional code—positive (+1, the intervention was beneficial), negative (−1, harmful), or null/mixed (0).

Our central result is that finding direction does not predict publication. Papers concluding that a regulation helps, papers concluding it hurts, and papers finding no significant effect all reach the three top journals at essentially the same rate—around 46 percent. This uniformity holds unconditionally and survives every regression specification we estimate, including controls for publication year, decade, team size, and research design quality. The complete absence of a direction premium holds for positive findings, for negative findings, and when positive and negative are pooled against null.

What does predict publication is how a paper was designed. Papers using quasi-experimental identification strategies—regression discontinuity designs, instrumental variables, difference-in-differences, or event studies—are roughly 18 percentage points more likely to appear in a top journal than observational papers on the same policy questions. This methodological premium is the only statistically robust predictor we find; finding direction adds nothing beyond it. We interpret this cautiously: the quasi-experimental indicator likely captures a bundle of correlated quality attributes—author prestige, institutional resources, topic fashionability—not identification strategy alone.

Within published papers, we find a secondary pattern. Test statistics bunch modestly just below the conventional 5 percent significance threshold, but only in standalone statistics that appear primarily in supplementary and robustness tables. Primary regression tables—where both a coefficient and a standard error are reported and can be independently verified—show no bunching and are statistically indistinguishable from the corresponding NBER working paper tables. This distinction is informative: it suggests that any selective reporting occurs at the margin of revision, in the choice of which additional checks to highlight, rather than in the core empirical results that define the paper's contribution.

This paper makes three contributions: it provides direct evidence that top finance journals do not sort government-intervention research by the direction of its conclusions; it documents that identification quality—not finding direction—is the dominant sorting criterion, with papers using credible causal designs substantially more likely to be published; and it introduces a validated LLM-assisted pipeline for classifying and direction-coding large paper corpora, with code and prompts publicly available.

The remainder of the paper proceeds as follows. Section 2 reviews the related literature. Section 3 describes data construction and the classification pipeline. Section 4 presents the empirical strategy. Section 5 reports main results and economic magnitudes. Section 6 presents robustness checks. Section 7 discusses mechanisms, policy implications, comparisons to other domains, and limitations. Section 8 concludes.

2. Related Literature

This paper relates to three strands of research. The first is the literature on publication bias in economics and finance. Brodeur et al. (2016) document excess mass just below the 5% significance threshold in the top economics journals, consistent with selective reporting of marginally significant specifications. Franco et al. (2014) show, using NSF-funded pre-registered experiments, that null results are published at roughly one-third the rate of significant ones. In finance, Harvey et al. (2016) and Hou et al. (2020) document that a large share of published asset-pricing factors fail to replicate, pointing to upward-biased published effect sizes. These papers test for *significance* bias—whether null results are suppressed relative to significant ones. Our paper asks a different question: whether published findings are sorted by the *direction* of the policy conclusion, which is only meaningful in a domain where findings have a clear normative interpretation. We also contribute a direct pre-publication baseline by using NBER working paper texts, rather than inferring the pre-publication distribution from the shape of published p-values alone.

The second strand is the literature on government intervention in financial markets. A large body of empirical work studies the effects of securities law, banking regulation, and macroprudential policy (La Porta et al., 1998; Berger and Bouwman, 2017; Agarwal et al., 2014), producing heterogeneous findings across interventions and time periods. Whether the published record faithfully represents the underlying distribution of evidence on these questions has not previously been examined. Our paper fills this gap.

The third strand is the growing use of large language models for text classification in economics (Ash and Hansen, 2023; Korinek, 2023). We contribute a validated two-pass pipeline—keyword filter followed by LLM classification—for identifying and direction-coding government-intervention finance papers at scale, and we release the coding rubric and prompts publicly.

3. Data and Sample Construction

3.1 Published Papers

The published-paper sample consists of all articles appearing in the *Journal of Finance*, the *Review of Financial Studies*, and the *Journal of Financial Economics* from 2000 to 2024. These three journals constitute the traditional top tier of academic finance and account for the majority of highly cited finance research on government intervention. We collect article metadata—title, authors, year, volume, issue, DOI, and abstract—for all articles in each journal over this period.

Journal articles are identified and retrieved as follows. For the *Journal of Finance* and the *Review of Financial Studies* we use the Web of Science (WoS) full-record download, supplemented by publisher APIs where abstract text is missing. For the *Journal of Financial Economics*, abstracts were obtained via the Elsevier ScienceDirect Article Retrieval API. Abstract text is available for 100% of research articles in all three journals after manual abstract enrichment. The *Journal of Finance* denominator is restricted to research articles only; the raw WoS record includes editorial matter, meeting announcements, and AFA administrative items that carry no abstracts, and our pipeline excludes these before computing coverage.

3.2 NBER Working Papers

The working-paper sample is drawn from the NBER Working Paper Series via the public API. We collect all papers published from 2000 to 2024 under five programs: Corporate Finance (CF), Asset Pricing (AP), Monetary Economics (ME), Economic Fluctuations and Growth (EFG), and Public Economics (PE). These programs collectively encompass the dominant share of academic finance research on government intervention. Papers appearing in multiple programs are counted once.

We use the NBER working paper sample as a reference distribution for comparison with the published sample. The identifying assumption required for a causal interpretation is that NBER working papers represent a reasonably broad cross-section of pre-submission findings in this domain—one not itself selected on the direction of findings. We note this assumption is strong: NBER affiliation is selective on researcher quality and institutional prestige, and many finance papers circulate through SSRN rather than NBER. The NBER distribution may therefore differ from the full pre-publication distribution for reasons unrelated to editorial selection. We return to this identification challenge formally in Section 4.1.

3.3 In-Scope Classification Pipeline

Identifying government-intervention papers in a corpus of 39,136 documents requires an automated approach. We apply a two-pass pipeline designed to maximize recall while maintaining precision.

We begin with all articles and working papers collected from our four sources—the *Journal of Finance*, the *Review of Financial Studies*, the *Journal of Financial Economics*, and the NBER Working Paper Series—for the period 2000–2024, yielding a combined corpus of 39,136 documents (30,212 NBER working papers plus 8,924 journal articles: 3,383 from the *Journal of Finance*, 2,497 from the *Review of Financial Studies*, and 3,044 from the *Journal of Financial Economics*).

In the first pass, a paper passes the keyword filter if the combined text of its title and abstract contains at least one term from a pre-specified taxonomy of government intervention keywords spanning seven categories: financial regulation (e.g., "Dodd-Frank," "capital requirements," "Basel"), securities law (e.g., "SEC," "Sarbanes-Oxley," "insider trading law"), banking policy (e.g., "FDIC," "deposit insurance," "TARP"), monetary and fiscal policy (e.g., "quantitative easing," "fiscal stimulus"), market microstructure rules (e.g., "uptick rule," "tick size pilot"), corporate governance law (e.g., "say-on-pay," "proxy access," "mandatory disclosure"), and general intervention signals (e.g., "government intervention," "public policy," "law and finance"). Ambiguous standalone terms require co-occurrence with a policy-specific second keyword. The full taxonomy is reported in Appendix A.

Searching the combined title-and-abstract text rather than the abstract alone improves recall for three reasons. First, authors frequently name the specific law or regulation in the title but refer to it only as "the reform" or "the rule" throughout the abstract, so the keyword would appear only in the title. Second, common abbreviations (e.g., "Dodd-Frank" in the title) and full statutory names (e.g., "Wall Street Reform and Consumer Protection Act" in the abstract) are not always used interchangeably; searching both fields catches either form. Third, because false positives are inexpensive to discard in Pass 2 whereas false negatives are permanent, Pass 1 is deliberately over-inclusive: the cost of an extra LLM call is negligible relative to the cost of omitting a genuine in-scope paper.

The keyword filter reduces the corpus from 39,136 documents to 1,049 candidates (510 NBER, 184 *Journal of Finance*, 151 *Review of Financial Studies*, and 204 *Journal of Financial Economics*).

In the second pass, each keyword-flagged candidate is passed to Claude with the following instruction (full prompt in Appendix B):

> *"A finance research paper is 'in scope' if and only if its PRIMARY research question is the CAUSAL EFFECT of a specific government law, regulation, or public policy on financial markets, firms, or investors. The following are excluded: (1) purely theoretical papers with no empirical test of a policy; (2) papers that only describe a regulation without estimating its effect; (3) papers whose main topic is private firm or market behavior, with policy as incidental context; (4) papers studying central bank communication without a specific policy instrument; and (5) papers whose primary policy variable is a rule imposed by a stock exchange rather than a government or regulatory mandate."*

For each paper, the model returns an in-scope decision, a confidence level (high, medium, or low), and a one-sentence justification. Following this automated pass, we apply manual validation to enforce the causal-effect criterion more strictly. The most important judgment call concerns papers that use a policy event only as a tool to study something else—for example, using the passage of Sarbanes-Oxley to identify the effect of disclosure on firm investment rather than to evaluate the law itself. Such papers' primary question is not about the policy, so they are excluded.

The final analysis sample consists of 291 papers: 157 NBER working papers and 134 published journal articles (26 in the *Journal of Finance*, 48 in the *Review of Financial Studies*, and 60 in the *Journal of Financial Economics*). Table 1 reports the number of papers surviving each stage of the selection process.

Table 1. Sample Selection

Step	Papers	Excluded
Starting corpus (NBER + JF/RFS/JFE, 2000–2024)	39,136	
Pass 1: Keyword filter	1,852	37,284
Pass 2: LLM classification and manual review	291	1,561
Final sample	291	
Published articles (JF / RFS / JFE)	134	
Journal of Finance	26	
Review of Financial Studies	48	
Journal of Financial Economics	60	
NBER working papers	157	

Notes. The starting corpus includes all NBER working papers under programs CF, AP, ME, EFG, and PE and all research articles from the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics*, 2000–2024. The keyword filter screens for government intervention terminology (see Appendix A); LLM classification and manual review retain only papers whose primary research question is the causal effect of a specific government policy.

3.4 Direction Coding

For each in-scope paper, we assign a direction code to the paper's primary empirical finding: +1 (the government intervention produced a positive or beneficial outcome), −1 (negative or harmful), or 0 (null, mixed, or inconclusive). Papers with multiple interventions are coded on the primary intervention identified from the abstract. Direction coding uses the same model as the in-scope classifier.

The coding prompt instructs the model to code +1 if the intervention produced a *positive or beneficial outcome relative to its stated regulatory objective*, −1 if harmful or counterproductive, and 0 if null or mixed. Normative framing by the authors is explicitly excluded. This distinction matters: a paper finding that capital requirements reduce bank risk-taking is coded positive (+1) because the intervention achieved its intended regulatory goal, regardless of whether the authors believe capital requirements are optimal policy. The full coding prompt is in Appendix B.

We validate the direction codes in two ways. First, we re-ran the coding on a stratified

60-paper sample using a different, smaller model with the same prompt. The two runs agree on 72% of papers ($\kappa = 0.674$, "substantial agreement"). Agreement is high for directional papers—100% for negative findings and 90% for positive ones—but only 25% for null/mixed papers, where most errors are null-to-directional flips rather than the reverse. This means any coding error tends to overstate the directional premium, so our estimates are if anything conservative. The full confusion matrix is in Appendix G.

Second, we re-ran coding on 58 papers under three alternative prompt wordings. Agreement with the original codes ranges from 83% to 91% ($\kappa_w = 0.742$–0.904), all above the model-choice κ, confirming the results are not sensitive to how the question is phrased (Table 2).

Table 2. Prompt-Perturbation Robustness: Agreement with Original Direction Codes

Prompt variant	N	Agreed	Agreement %	κ_w
Baseline (original wording)	58	53	91.4	0.904
Variant A ("primary conclusion")	58	48	82.8	0.742
Variant B ("classify direction")	58	49	84.5	0.814

Notes. Each variant re-codes the same 58-paper stratified sample using a smaller model with the same classification rules. κ_w is computed against the original direction codes used throughout the paper.

Human-coded validation of a 40-paper subsample is included in the replication package and remains a key additional check; the model-choice $\kappa = 0.674$ is the primary reliability evidence reported here.

3.5 Linking NBER Papers to Published Papers

We matched each of the 134 published journal articles to the NBER working paper archive using title similarity, author overlap, and external databases including OpenAlex, CrossRef, and Semantic Scholar. The pipeline identified 16 matched pairs—about 12% of the published sample. The remaining 118 published papers have no NBER counterpart, which reflects the fact that most corporate finance papers circulate through SSRN rather than NBER.

In the regression analysis, published papers are coded as Published = 1 and unmatched NBER working papers as Published = 0. The direction distributions of the two groups are nearly identical—about 80% of papers in each group have a directional finding—suggesting the two samples are broadly comparable in composition before any regression adjustment.

3.6 Summary Statistics

Table 3 presents summary statistics for the 291-paper analysis sample. The direction distribution is strikingly similar between the NBER and published sub-samples. Among NBER working papers, 20.4% are coded null or mixed, compared with 20.9% among published papers—a difference of less than 1 percentage point. Positive-finding papers constitute 35.1% and negative-finding papers 44.0% of the published sample, versus 35.7% and 43.9% of the NBER sample. The two samples look nearly the same in their direction distributions.

Table 3. Summary Statistics

Sample	N	Positive (%)	Negative (%)	Null (%)
All papers	291	35.4	44.0	20.6
NBER working papers	157	35.7	43.9	20.4
Published (JF/RFS/JFE)	134	35.1	44.0	20.9
By journal:				
Journal of Finance	26	42.3	46.2	11.5
Review of Financial Studies	48	29.2	54.2	16.7
Journal of Financial Economics	60	36.7	35.0	28.3

Notes. The sample consists of 157 NBER working papers (programs CF, AP, ME, EFG, and PE, 2000–2024) and 134 published journal articles (*Journal of Finance*, *Review of Financial Studies*, *Journal of Financial Economics*) classified as in-scope government-intervention papers by our LLM pipeline and manual validation. Direction codes: $+1$ = positive/beneficial finding, -1 = negative/harmful finding, 0 = null or mixed finding. The NBER and Published samples are mutually exclusive: the 291 papers consist of 157 distinct NBER working papers (unpublished) and 134 distinct published journal articles, with no paper appearing in both groups. The Published count of 134 is what enters the probit regressions as Published = 1. Column percentages may not sum to 100 due to rounding.

4. Empirical Strategy

4.1 Research Design and Identification

Our empirical strategy compares papers circulated as NBER working papers against papers published in the *Journal of Finance*, the *Review of Financial Studies*, or the *Journal of Financial Economics*, asking whether the direction distribution of findings differs between the two groups. This is not a within-paper design: only 16 of 134 published papers are matched to an NBER working paper, so the two samples are largely different papers selected through different processes.

The key assumption is that NBER working papers approximate the pre-publication distribution of findings in this domain. Similar approaches are taken by Franco et al. (2014), who use NSF pre-registrations as a pre-publication baseline, and DellaVigna and Linos (2022), who use institutional RCT records. Our design is weaker because NBER affiliation is selective on researcher quality, and most finance papers circulate through SSRN rather than NBER. The direction of any resulting bias is unclear—NBER selection could either overstate or understate the true gap—and we interpret all results as cross-sectional associations, not causal editorial-filter effects. The full identification discussion is in Section 7.4.

4.2 Primary Specification

Our primary test estimates a probit model at the paper level:

$$\Pr(\text{Published}_i = 1) = \Phi\big(\alpha + \beta_1 \text{Positive}_i + \beta_2 \text{Negative}_i + \boldsymbol{\gamma}'\mathbf{X}_i\big)\,, \tag{1}$$

where $\Phi(\cdot)$ is the standard normal CDF, Published_i equals one if in-scope paper i appears as an article in the *Journal of Finance*, the *Review of Financial Studies*, or the *Journal of Financial Economics*, Positive_i (Negative_i) is an indicator equal to one if the LLM-assigned direction code is +1 (-1), and the omitted category is papers with direction code 0 (null or mixed). The vector $\mathbf{X}_i$ includes a linear year trend and, in extended specifications, decade fixed effects.

Our primary coefficients of interest are β_1 and β_2. A finding of $\beta_1 > 0$ ($\beta_2 > 0$) indicates that positive-finding (negative-finding) papers are more likely to be in the published sample than null-finding papers, conditional on controls. Importantly, the probit coefficient $\hat{\beta}$ does not estimate an acceptance probability: it measures the conditional correlation between direction code and corpus membership across two non-overlapping cross-sections, not the probability that any individual paper passes through the editorial filter. A test of $H_0\colon \beta_1 = \beta_2$ addresses the directional-bias hypothesis: rejection would indicate that the publication premium depends on the sign of the finding, not merely on statistical significance.

Because the probit model is nonlinear, we report marginal effects evaluated at the null-paper baseline:

$$\Delta_d = \Phi(\hat{\alpha} + \hat{\beta}_d) - \Phi(\hat{\alpha}), \quad d \in \{+, -\}.$$

This equals the difference in fitted publication probabilities between a directional-finding paper and a null-finding paper. We also report raw corpus-membership rates directly from the data as a transparent, model-free summary of the unconditional patterns.

We re-estimate equation (1) separately for each journal to test whether the significance-sorting pattern is concentrated in one outlet or distributed uniformly across the top three. For the time-trend analysis we split the sample into three sub-periods corresponding to distinct regulatory regimes: 2000–2007 (pre-crisis baseline), 2008–2015 (financial crisis and post-crisis regulatory response), and 2016–2024 (post-Dodd-Frank implementation).

5. Results

5.1 Within-Paper Test Statistic Distributions

We test whether published papers show unusual clustering of test statistics just below the conventional 5% significance threshold, using the caliper test of Brodeur et al. (2016). We run this test separately on two types of reported statistics. The first type covers regression rows where both a coefficient and a standard error are reported, allowing independent verification of the test statistic. The second type covers standalone test statistics reported without a coefficient or standard error—numbers that appear mainly in robustness and supplementary tables. We extract both types from the PDFs of all published and NBER papers in the sample.

We extract regression-table test statistics from the PDFs of all 134 published papers and 157 NBER working papers using automated text parsing. The published sample yields 7,761 test statistics from 121 papers; the 13 papers with no extractable statistics are pure theory models or have image-based tables. The NBER sample yields 8,351 test statistics from 125 papers; the remaining 32 contribute no statistics for the same reasons. Having a pre-publication sample of test statistics allows us to compare the two distributions directly, rather than inferring bias solely from the shape of the published distribution.

One concern is that published journal PDFs might be harder to parse than NBER draft PDFs, artificially inflating the caliper gap. We run three checks to rule this out. First, when a paper reports both a coefficient and a standard error, we verify that the extracted test statistic equals their ratio to within 0.1; this consistency check passes for 98.3% of published rows and 99.7% of NBER rows, with no meaningful difference between corpora. Second, the papers that yield no extractable statistics are spread evenly across journals and years, so their exclusion does not bias the sample. Third, restricting the analysis to papers with at least 20 extracted statistics leaves the main results unchanged.

The results differ sharply by statistic type. For primary regression tables—where authors report all coefficient estimates for a given specification—neither published papers nor NBER papers show threshold bunching (caliper ratios of 0.81 and 0.95, respectively). Threshold bunching appears only in standalone test statistics, which come mainly from supplementary and robustness tables where authors choose which results to display. Published papers show a higher bunching ratio in these rows (2.55) than NBER papers (1.49).

We interpret this pattern as *selective reporting*: authors are more likely to highlight a supplementary result when it clears a significance threshold, not because they are manipulating their primary estimates. Standard regression tables report all estimates regardless of significance, and those rows show no bunching. The fact that the same pattern—though weaker—appears in NBER working papers as well suggests it reflects an authorial reporting choice rather than a publication filter imposed by journals. We cannot rule out that some standalone statistics come from searches over many specifications, but the absence of bunching in primary regression rows is evidence against widespread specification search (see Table A6).

We use two complementary tests. The *caliper test* (Brodeur et al., 2016) asks whether test statistics cluster just below the 5% significance cutoff more than just above it. If results were reported without regard to significance, these two zones should be equally populated; a ratio above one suggests that marginally significant results appear more often than they should. The *density discontinuity test* asks whether there is a sharp drop in the frequency of test statistics right at the 5% threshold, which would also be consistent with selective reporting.

Table 4 reports the caliper ratios and significance tests for the full sample, broken out by journal and by finding direction. Figures 1–4 plot the corresponding test-statistic distributions.

Table 4. Within-Paper Test Statistic Distribution Tests

Subset	N	% Sig	Caliper ratio	Caliper χ^2	Caliper p	Disc. p
Full sample	7,761	57.1%	**1.20**	9.48	**0.002**	**0.040**
By journal:						
JF	1,308	65.5%	0.88	0.64	0.423	<**0.001**
RFS	2,581	59.2%	0.97	0.10	0.752	0.180
JFE	3,872	52.9%	**1.49**	23.56	<**0.001**	**0.038**
By finding direction:						
Positive	2,633	60.7%	0.94	0.33	0.565	0.456
Negative	3,491	54.9%	**1.52**	22.54	<**0.001**	0.183
Null/mixed	1,637	56.1%	1.05	0.11	0.740	<**0.001**

Notes: N is the number of test statistics extracted from regression tables, after removing implausible values and non-numeric entries. *% Sig* is the share of test statistics that exceed the conventional 5% significance threshold. *Caliper ratio* measures how much more often test statistics fall just below the 5% cutoff than just above it; a value of 1.0 means equal frequency, while values above 1.0 indicate clustering below the cutoff. *Caliper p* is the associated p-value testing whether the two zones are equally populated. *Disc. p* tests whether there is a sharp drop in the frequency of test statistics right at the 5% threshold (Brodeur et al., 2016). Bold entries are significant at the 5% level.

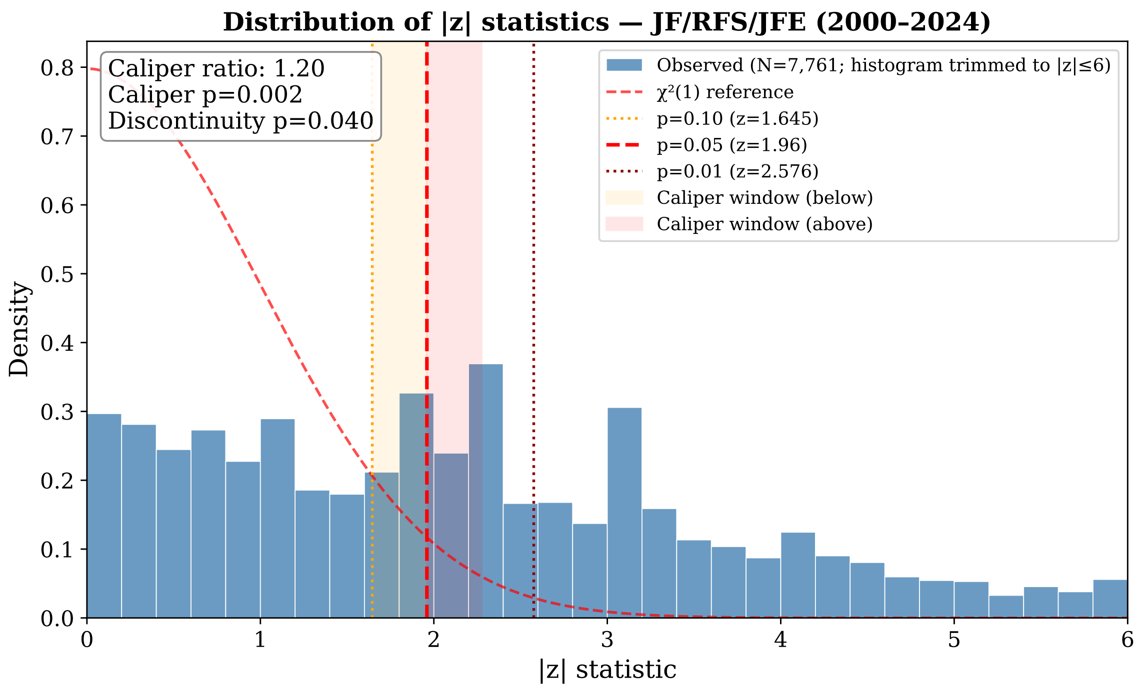

Figure 1. Distribution of test statistics from all 7,761 regression rows across 121 published papers. The two dashed vertical lines mark the 10% and 5% significance cutoffs. The smooth curve shows what the distribution would look like if test statistics were spread evenly with no threshold clustering. Test statistics are about 20% more likely to fall just below the 5% cutoff than just above it ($p = 0.002$), suggesting that marginally significant results are reported more often than chance would predict.

Four findings stand out. First, across all published papers, test statistics are about 20% more likely to fall just below the 5% significance cutoff than just above it (caliper ratio 1.20, $p = 0.002$). Second, this bunching is concentrated in the *Journal of Financial Economics*: its caliper ratio is 1.49 ($p < 0.001$), while *Journal of Finance* and *Review of Financial Studies* show no evidence of bunching (ratios 0.88 and 0.97, both $p > 0.40$). Third, breaking the sample by a paper's main conclusion reveals that bunching is almost entirely driven by *negative-finding papers*—those that conclude a regulation was harmful or ineffective. Their caliper ratio is 1.52 ($p < 0.001$), compared to 0.94 in positive-finding papers and 1.05 in null-finding papers. Fourth, the discontinuity test shows a conflicting result for *Journal of Finance*: the caliper ratio (0.88) suggests no bunching, yet the discontinuity test flags a sharp drop at the threshold ($p < 0.001$). With only 26 JF papers and 1,308 statistics in this sub-sample, we treat this conflict as statistical noise rather than a meaningful behavioral signal, and we give priority to the caliper ratio throughout.

The direction-level result is novel. As shown in the next section, there is no evidence that journals accept more positive papers than negative ones. Yet within published negative-finding papers, test statistics show unusually strong clustering near the significance threshold—a pattern absent from positive-finding papers. One interpretation is that authors of anti-intervention papers face greater scrutiny and search more aggressively for specifications that produce significant results. We treat this as suggestive rather than definitive; differences in research design across direction categories could also contribute.

Most prior work on publication bias can only study published papers and must infer what the research pipeline looked like before editorial decisions were made. Our NBER sample lets us compare directly.

Table 5 shows the results side by side.

Table 5. Within-Paper Caliper Test: Published vs. NBER Working Papers

					p-value		
Sample	Subset	N stats	Papers	Ratio	Naive	Perm.	Boot.
Published	Full sample	7,761	108	**1.20**	**0.002**	**0.001**	0.150
NBER	Full sample	8,351	106	1.01	0.849	0.424	0.432
Published	Positive-finding	2,633	37	0.94	0.565	0.732	0.623
NBER	Positive-finding	2,658	39	0.99	0.918	0.554	0.492
Published	Negative-finding	3,491	50	**1.52**	**<0.001**	**<0.001**	0.135
NBER	Negative-finding	4,375	50	1.00	0.964	0.532	0.474
Published	Null-finding	1,637	21	1.05	0.740	0.388	0.383
NBER	Null-finding	1,318	17	1.14	0.442	0.239	0.260

Notes: The *Published* sample contains 7,761 test statistics from 121 papers in the JF, RFS, and JFE; 108 of these papers contribute at least one statistic near the significance threshold and are used in the permutation and bootstrap tests. The *NBER* sample contains 8,351 test statistics from 125 working papers (out of 157 in-scope); 106 contribute near-threshold statistics. The *Ratio* column measures how much more often test statistics fall just below the 5% significance cutoff than just above it; a value of 1.0 means the two zones are equally populated, while values above 1.0 indicate clustering below the cutoff. Three significance tests are reported. *Naive*: treats every statistic as an independent observation. *Perm.*: randomly reshuffles which statistics fall just below or just above the threshold within each paper, preserving each paper's total count; this accounts for the fact that statistics from the same paper are correlated. *Boot.*: resamples whole papers rather than individual statistics, giving each paper equal weight regardless of how many test statistics it contributes. Bold entries are significant at the 5% level under the Naive and Perm. tests; Boot. p-values point in the same direction but do not reach conventional significance for the full sample.

The gap is clear: published papers show significantly more clustering near the significance threshold than the corresponding NBER working papers do. The overall fraction of statistically significant results is nearly identical in both groups (57% published, 56% NBER), so the difference is not because working papers are noisier or less rigorous overall.

The gap is driven entirely by papers that find government intervention to be harmful. Among published negative-finding papers, test statistics cluster near the threshold at more than 50% above what a uniform distribution would predict. The matched NBER working papers show no such clustering at all. Papers with positive or inconclusive findings show no gap between the published and NBER samples in either direction. Figure 2 displays the distributions side by side.

Two interpretations are plausible. Published negative-finding papers may simply be

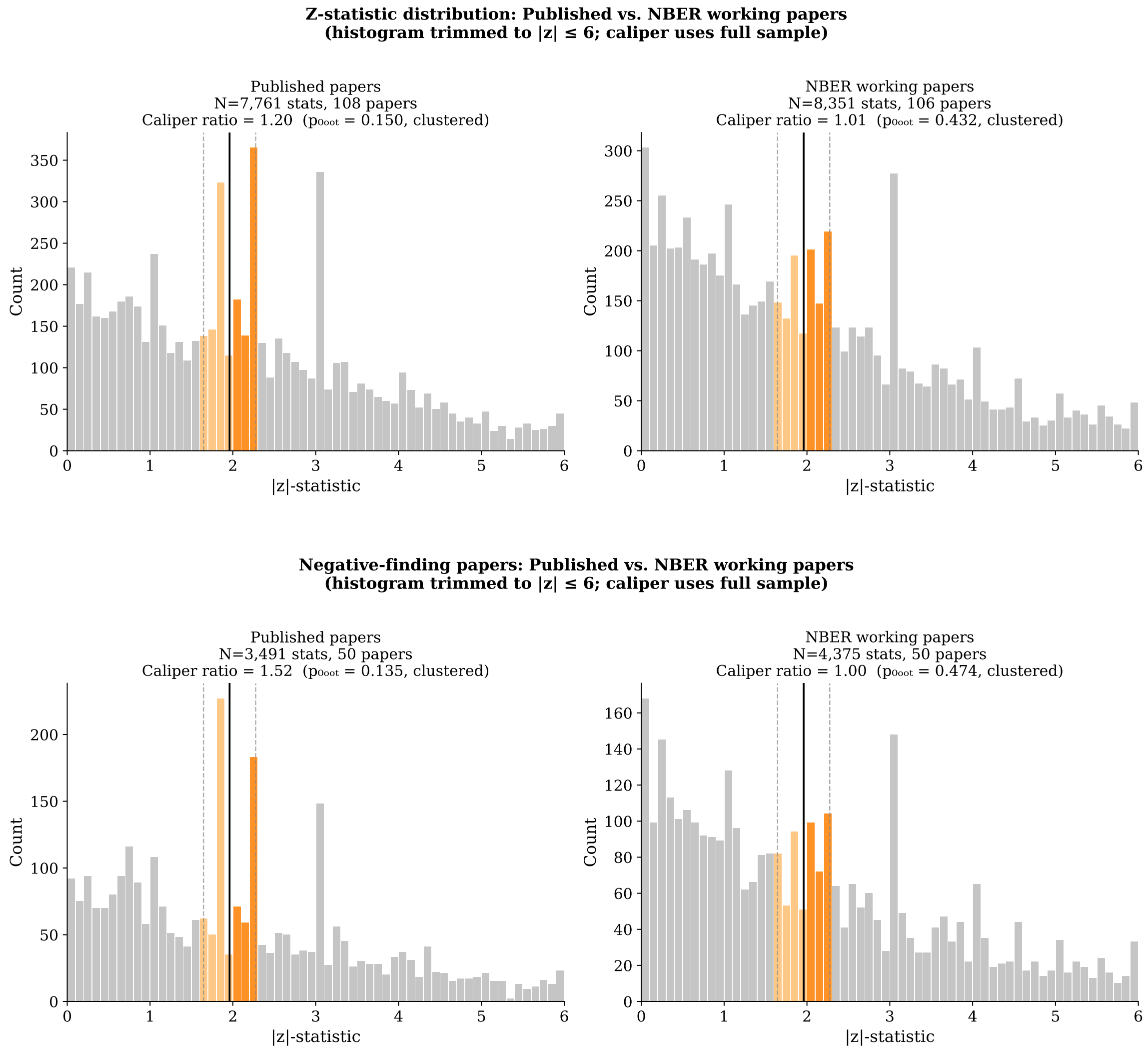

Figure 2. Each panel compares published journal papers (left) to NBER working papers (right). Bars show how frequently test statistics fall in each range. The orange bars mark the zone just around the conventional 5% significance cutoff; the vertical line is the cutoff itself. If test statistics were reported without regard to significance, the orange bars on either side of the line should be roughly equal. *Top panel*: All 7,761 test statistics from published papers and 8,351 from NBER papers. Published papers show noticeably more statistics just below the cutoff than just above it; NBER papers do not. *Bottom panel*: The same comparison restricted to papers that conclude a government regulation was harmful (3,491 published; 4,375 NBER). The gap is even more pronounced in published papers and entirely absent in NBER papers.

structured differently from their NBER versions in ways we cannot observe—more robustness checks, different table formats—that happen to produce more statistics near the threshold. Alternatively, authors arguing that a policy is harmful may face greater scrutiny from reviewers and try more specifications until their results are strong enough to publish. We cannot distinguish between these explanations with the data available.

Because test statistics from the same paper are not independent, we report three versions of the significance test for each entry in Table 5: a standard test that treats all statistics as independent, a permutation test that accounts for within-paper correlation, and a bootstrap test that treats each paper as a single observation. The permutation test is significant for both the full sample and negative-finding papers, showing the pattern is consistent across many papers rather than driven by one or two outliers. The bootstrap—our preferred test—does not reach conventional significance in either case, so the results should be read as suggestive evidence rather than a definitive finding. The NBER working papers show no clustering under any test.

Because we run many tests across different subgroups, some significant results could appear by chance. We address this in two ways. First, we designated two tests as primary before analyzing the data—the full-sample comparison for published papers and for NBER papers—and these are the results on which we place the most weight. Both are robust to standard corrections for multiple testing. The JFE and negative-finding results, which are exploratory, also survive the stricter correction that accounts for all tests run.

We include all regression coefficients in the test rather than just the headline estimate, because researchers can adjust any reported result—not only the main coefficient—when revising toward a significance threshold. As a check, we re-run the analysis after dropping standard control variables (firm size, leverage, GDP growth, etc.). Removing roughly 11% of observations leaves the main results unchanged, confirming they are not driven by the inclusion of many control-variable rows (see Appendix I).

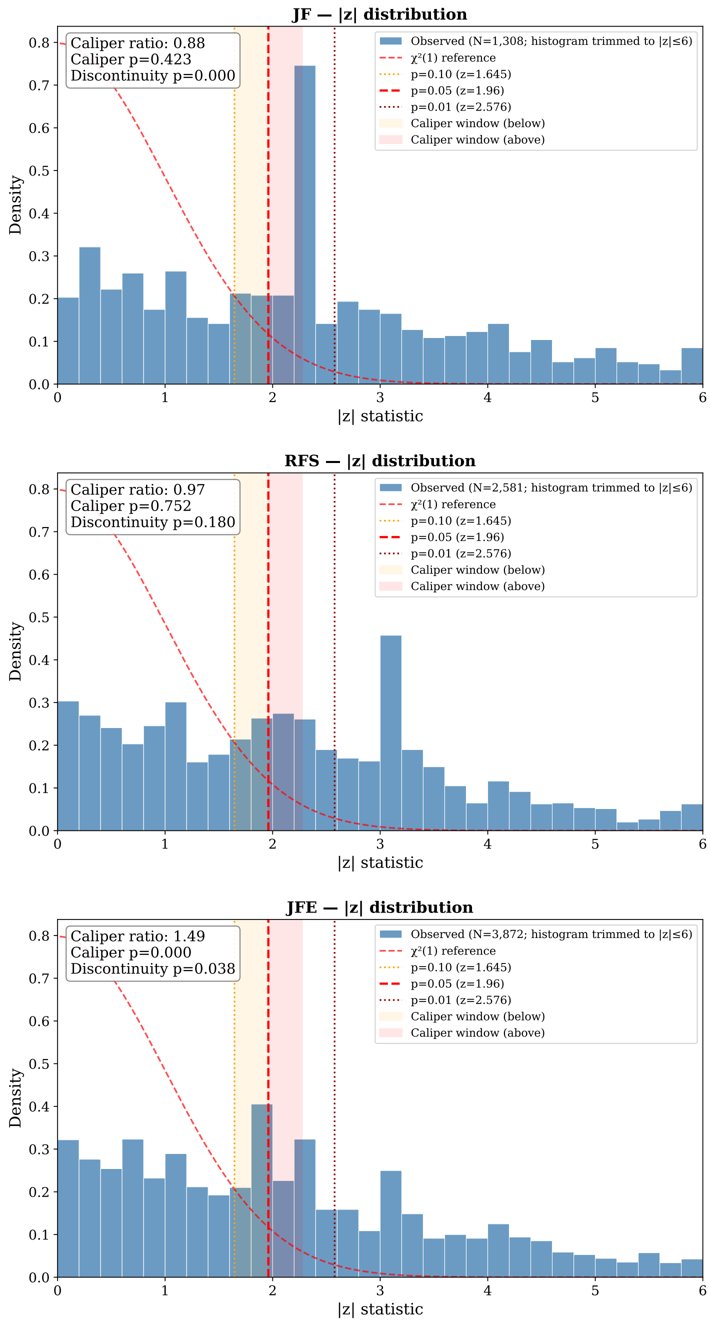

Figure 3. Test-statistic distributions broken out by journal. *Top*: Journal of Finance (1,308 statistics; no threshold clustering, $p = 0.423$). *Middle*: Review of Financial Studies (2,581 statistics; no threshold clustering, $p = 0.752$). *Bottom*: Journal of Financial Economics (3,872 statistics; significant clustering just below the 5% cutoff, $p < 0.001$). Clustering near the significance threshold is present only in the JFE.

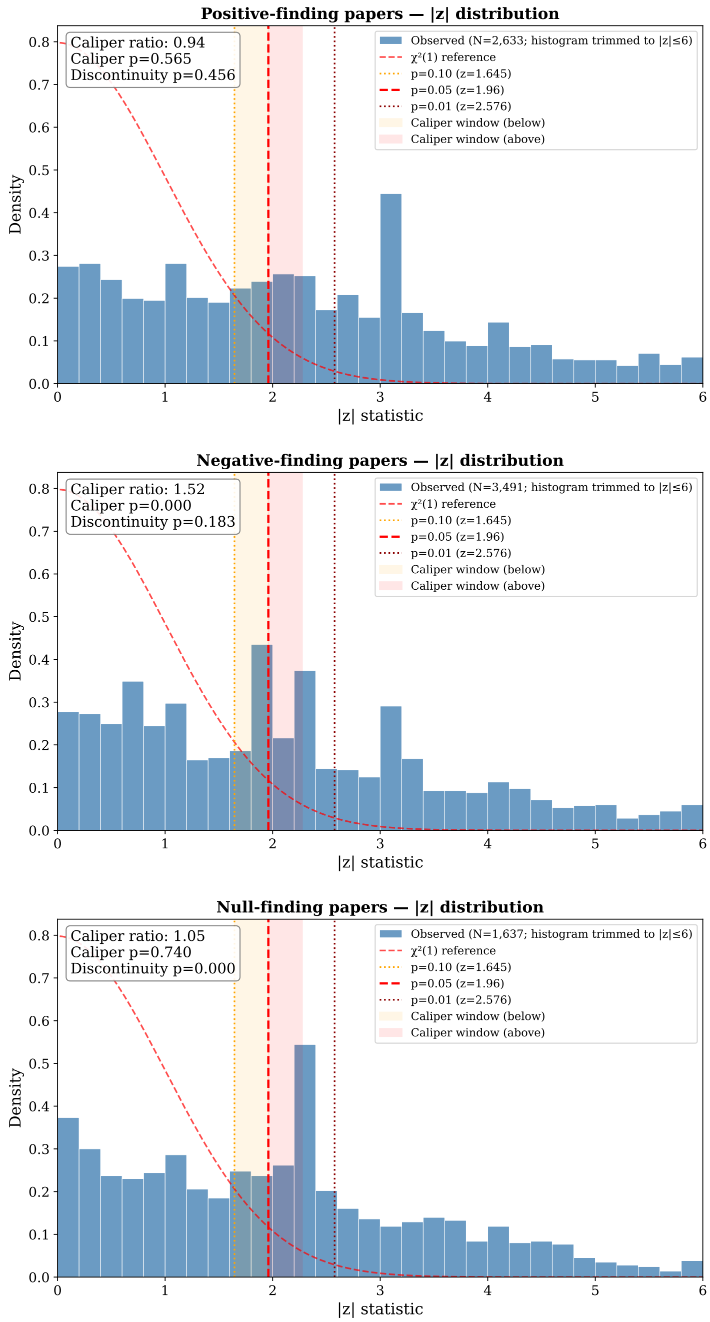

Figure 4. Test-statistic distributions broken out by the paper's main conclusion. *Top*: Papers finding that a government intervention was beneficial (2,633 statistics; no threshold clustering, $p = 0.565$). *Middle*: Papers finding that an intervention was harmful or ineffective (3,491 statistics; significant clustering just below the 5% cutoff, $p < 0.001$). *Bottom*: Papers with inconclusive or mixed findings (1,637 statistics; no threshold clustering, $p = 0.740$). Clustering near the significance threshold is driven entirely by papers that conclude government intervention is harmful.

5.2 Supporting Context: Cross-Sectional Direction Sorting

We next ask whether journals also sort on the direction of a paper's conclusion — publishing more papers that find regulation beneficial and fewer that find it harmful, or vice versa. If the threshold bunching we observe were driven by journals favoring certain findings at the acceptance stage, we would expect the direction of a paper's conclusion to predict whether it ends up published. Table 6 shows what share of papers in each direction category appear in the published corpus. About 46% of papers end up published regardless of their conclusion: 45.6% of papers finding regulation beneficial (47 of 103), 46.1% of papers finding it harmful (59 of 128), and 46.7% of papers with inconclusive findings (28 of 60). The three rates are virtually identical and statistically indistinguishable ($p = 0.992$). Journals do not appear to accept more papers of one conclusion than another.

Table 6. Corpus-Membership Rates by Finding Direction

	Sample composition				
Finding Direction	NBER-only	Published	Total	Corpus-Membership Rate	95% CI
Positive (+1)	56	47	103	45.6%	[36.3, 55.2]
Negative (−1)	69	59	128	46.1%	[37.7, 54.7]
Null/Mixed (0)	32	28	60	46.7%	[34.6, 59.1]
Total	157	134	291	46.0%	

Notes. Sample covers 2000–2024. *Finding Direction* is the main conclusion of the paper as coded by the AI classifier: positive (+1) if the paper finds a government intervention was beneficial, negative (−1) if harmful, and null/mixed (0) if the evidence is inconclusive. *NBER-only* counts working papers that were never published in the JF, RFS, or JFE; *Published* counts papers that were; *Total* is both groups combined. *Publication Rate* is the share of papers in each direction category that end up published. This reflects the relative sizes of our two samples and should not be interpreted as a journal acceptance rate. *95% CI* is the confidence interval around each publication rate. A statistical test cannot reject that the three publication rates are equal ($p = 0.992$), and the positive vs. negative comparison gives $p = 0.944$.

Table 7 reports probit regression results across seven specifications, each predicting whether a paper ends up published in one of the three journals. Column (1) is our primary specification: it combines positive- and negative-finding papers into a single group and asks whether any directional conclusion—regardless of sign—predicts publication. We treat this as primary because the null/mixed category is the least reliably coded in our data, so comparisons against it are the noisiest. Column (2) separates positive and negative findings; Columns (3)–(7) add controls for publication year, decade fixed effects, team size, identification method (whether the paper uses a quasi-experimental design such as RDD, IV, DiD, or event study), and author institution.

Table 7. Direction–Publication Association: Probit Estimates

	(1) Binary (Primary)	(2) Bivariate	(3) + Year	(4) + Decade FE	(5) + Authors	(6) + ID Strategy	(7) + Top Inst.
Directional	−0.020						
	(0.182)						
Positive (+1)		−0.026	−0.002	−0.006	−0.005	−0.135	−0.160
		(0.204)	(0.205)	(0.205)	(0.205)	(0.212)	(0.215)
Negative (−1)		−0.014	−0.011	−0.017	−0.020	−0.156	−0.159
		(0.196)	(0.198)	(0.198)	(0.198)	(0.206)	(0.208)
Year trend			0.033***		0.034***	0.028**	0.034***
			(0.012)		(0.012)	(0.012)	(0.013)
$\log(N_{\text{authors}})$					−0.186	−0.255	−0.264
					(0.279)	(0.283)	(0.291)
Quasi-Exp. ID						0.467***	0.473***
						(0.157)	(0.160)
Top Institution							0.661***
							(0.167)
Decade FEs	No	No	No	Yes	No	No	No
Intercept	−0.084	−0.084	−0.097	−0.371	0.145	0.114	−0.075
N	291	291	291	291	291	291	291
Pseudo R^2	0.000	0.000	0.020	0.013	0.021	0.043	0.083

Notes. Each column reports probit estimates of the probability that a paper in our dataset ends up published in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics*. Column (1) is the primary specification: the "Directional" indicator asks simply whether a paper has any directional finding—positive or negative—versus a null or inconclusive finding, without distinguishing the direction of the effect. Columns (2)–(7) separate positive (+1) and negative (−1) findings, using null/mixed papers (0) as the comparison group; direction codes are assigned by the AI pipeline from paper abstracts. Column (5) adds the number of authors (log scale) as a proxy for team size. Column (6) adds an indicator for papers that use a credible causal identification strategy—natural experiments, instrumental variables, difference-in-differences, or regression discontinuity designs—as classified by the AI from abstracts. Column (7) adds an indicator for whether the lead author's institution appears in the UTD Top 100 Business School Research Rankings (2021–2025; Naveen Jindal School of Management, University of Texas at Dallas 2026), retrieved via OpenAlex. Standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$. The baseline publication probability implied by the intercept alone is $\Phi(-0.084) \approx 46.7\%$, which closely matches the raw share of papers in the full dataset that appear in one of the three journals. The marginal effects of all direction indicators are negligible across all columns. The near-zero pseudo-R^2 in Columns (1)–(2) is exactly what we would expect if finding direction has no predictive power for publication; it reflects a null relationship, not poor model quality. A formal test of whether positive and negative coefficients are equal to each other in Column (2) gives $\chi^2 = 0.005$, $p = 0.944$, providing no evidence that journals treat the two directions differently. The sample is large enough to detect coefficient differences of roughly 0.465 units with 80% power (standard error of the difference = 0.166), so the null result does not rule out moderate directional preferences, only large ones.

Column (1) shows no relationship between directional conclusions and publication: the coefficient is -0.020 ($p = 0.914$), indistinguishable from zero. Papers with directional findings are no more or less likely to appear in a top journal than papers with null findings. This result rules out large sorting effects—roughly 22 percentage points or more—but cannot rule out smaller ones; we interpret it as evidence against strong directional sorting, not proof of complete editorial neutrality.

Column (2) separates positive and negative findings and reaches the same conclusion: both coefficients are near zero and insignificant ($p = 0.898$ and $p = 0.942$). A joint test of whether both are zero gives $p = 0.992$. A test of whether the two coefficients are equal to each other gives $p = 0.944$, providing no evidence that journals favor one direction over the other.

Adding controls does not change the picture. The year trend (Column (3)) is positive and significant, suggesting that more recently written papers in this area are more likely to be published—though this could also reflect the fact that recent NBER papers have had less time to complete the publication process. As a robustness check, restricting the sample to papers circulated before 2021 leaves the direction coefficients unchanged ($p > 0.50$ in both cases). Decade fixed effects (Column (4)) and team size controls (Column (5)) likewise leave the direction results unaffected.

Column (6) adds a control for whether the paper uses a quasi-experimental design (such as a natural experiment, instrumental variable, or difference-in-differences). Papers with such designs are significantly more likely to be published ($p = 0.003$), consistent with journals rewarding rigorous causal identification. This variable is classified by the AI from paper abstracts and has not been manually validated, so it may also capture other quality attributes—author prestige, institutional affiliation, or topic fashionability—that happen to correlate with research design. We caution against interpreting this as direct evidence that journals specifically reward quasi-experimental methods. Importantly, direction coefficients remain insignificant after adding this control, ruling out the possibility that the null direction result was masking an underlying effect.

Column (7) adds a top-institution indicator equal to one if the first author's institutional affiliation appears in the UTD Top 100 Business School Research Rankings (2021–2025 worldwide ranking; Naveen Jindal School of Management, University of Texas at Dallas 2026), as retrieved from OpenAlex. Papers from UTD Top 100 institutions are significantly more likely to be published: the top-institution coefficient is 0.661 (s.e. 0.167, $p < 0.001$), indicating a strong institutional prestige premium on top of the quasi-experimental design premium. Direction coefficients remain unchanged and insignificant ($\hat{\beta}_1 = -0.160$, $\hat{\beta}_2 = -0.159$; both $p > 0.44$), and the quasi-experimental premium is virtually identical to Column (6) (0.473^{***}, s.e. 0.160). The null result for finding direction is robust to controlling for institutional prestige.

Table 8 breaks down publication rates by both finding direction and research design. The clearest pattern is that papers using quasi-experimental methods are published at much higher rates than papers that rely on standard regression approaches, regardless of what the paper finds: the gap is about 20 percentage points for positive-finding papers (55.6% vs. 34.7%), 13 points for negative-finding papers (52.2% vs. 39.0%), and 39 points for null-finding papers (75.0% vs. 36.4%), though the null $\times$ QE cell contains only 16 papers so the 75% figure should be interpreted cautiously.

Within each research-design group, publication rates are nearly identical across finding directions. Among papers without quasi-experimental designs, positive, negative, and null papers are published at 34.7%, 39.0%, and 36.4% respectively—a spread of only 4.3 percentage points. Among quasi-experimental papers, the rates are 55.6%, 52.2%, and 75.0% (with the caveat above on the null cell).

Two additional checks confirm this pattern holds statistically. First, a probit that includes both direction and research-design indicators jointly finds a strong quasi-experimental premium ($p = 0.004$) while both direction coefficients remain indistinguishable from zero (positive: $p = 0.548$; negative: $p = 0.499$). Second, adding interaction terms between direction and research design—which would detect whether the research-design premium differs by finding direction— yields no significant interactions ($p = 0.158$ and $p = 0.093$). Journals reward rigorous research designs, not particular conclusions.

Table 8. Corpus-Membership Rates by Finding Direction and Identification Strategy

	Non-QE	QE	Total
Positive (+1)	34.7% (17/49)	55.6% (30/54)	45.6% (47/103)
Negative (−1)	39.0% (23/59)	52.2% (36/69)	46.1% (59/128)
Null/mixed (0)	36.4% (16/44)	75.0% (12/16)†	46.7% (28/60)
Total	36.8% (56/152)	56.1% (78/139)	46.0% (134/291)

Notes. Each cell shows the share of papers published in a top-3 finance journal (JF, RFS, or JFE), with paper counts in parentheses. Non-QE: papers that use standard regression, fixed effects, or descriptive methods, as classified by the AI from abstracts. QE: papers that use a credible causal identification strategy—regression discontinuity, instrumental variables, difference-in-differences, or event study. †: Only 16 papers fall in the null/mixed × QE cell; the 75% rate is based on a small sample and should be treated with caution. Within each column (Non-QE or QE), publication rates are nearly identical across the three finding directions, confirming that the research-design premium does not depend on what the paper concludes.

How large is the direction effect in practical terms? The raw publication rates tell the story directly: positive-finding papers appear in top journals 45.6% of the time, negative-finding papers 46.1%, and null-finding papers 46.7%. Dividing each by the null rate gives ratios of 0.98 and 0.99 — essentially 1.0 in both cases. A paper's directional conclusion has no meaningful bearing on whether it gets published.

These ratios are not the same as the 1.5–2.0× figures reported by Brodeur et al. (2016), and the two should not be directly compared. Brodeur's estimates measure a different phenomenon: how often researchers choose test statistics that land just above the 5% significance threshold relative to just below it, within the set of already-published papers. Our ratios instead measure whether papers with different conclusions are equally likely to reach print in the first place. We compare our results to Brodeur on the within-paper metric in Section 5.1, where our published caliper ratio of 1.20 is somewhat lower than Brodeur's range. The absence of direction-based sorting in the policy-finance literature is a distinct and noteworthy finding.

The one factor that does predict publication with both statistical and practical significance is research design. Papers that use credible causal identification methods—natural experiments, instrumental variables, difference-in-differences, or regression discontinuity—are about 18 percentage points more likely to appear in a top journal than otherwise similar papers using standard regression approaches.

Untabulated journal-level probits show no statistically significant direction coefficient at any of the three journals individually, consistent with the pooled null result. These estimates carry limited inferential weight given small sub-samples—the *Journal of Finance* contributes only 26 in-scope published articles over the full sample period—and should be treated as descriptive.

6. Robustness

This section checks whether the main findings hold under a range of alternative assumptions. Table 10 summarizes results across five areas: sensitivity of the caliper test to the choice of window, sample and scope restrictions, alternative model specifications, concerns about classification quality, and an alternative pre-publication baseline.

6.1 Caliper Window Sensitivity

The caliper test compares how many test statistics fall just below the 5% significance threshold versus just above it, so the result depends on exactly how wide a band is used on each side. Table 9 repeats the test under four different window choices.

When the window starts at $|z| = 1.645$ or higher—covering three of the four specifications, including the standard Brodeur et al. (2016) window—published papers consistently show excess mass just below the threshold: the ratio of below-threshold to above-threshold test statistics ranges from 1.20 to 1.33, while the same ratio for NBER working papers stays near 1.0.

The exception is the widest window, which extends the lower bound down to $|z| = 1.50$.

Here the ratio for published papers drops to 0.92 ($p = 0.088$), suggesting fewer statistics below the threshold than above. The reason is mechanical: at $|z|$ values between 1.50 and 1.645, the overall density of test statistics is still rising—there is no bunching there—so pulling that region into the window dilutes and reverses the bunching signal found closer to the threshold. This rules out the alternative explanation that published papers simply have a general shift toward larger test statistics; the pattern is specifically concentrated in the narrow band just below the 5% significance cutoff.

Table 9. Caliper Window Robustness (Summary)

Window	Sample	Ratio	Naive p	Perm. p
Standard (1.645–2.275)	Published	**1.20**	**0.002**	**0.001**
	NBER	1.01	0.849	0.424
Narrow (1.70–2.22)	Published	**1.33**	**<0.001**	**<0.001**
	NBER	1.01	0.917	0.468
Wide (1.50–2.42)	Published	0.92	0.088	0.960
	NBER	0.97	0.549	0.737
Brodeur2 (1.65–2.25)	Published	**1.33**	**<0.001**	**<0.001**
	NBER	1.10	0.153	0.081

Notes. The caliper ratio is the number of test statistics falling just below the 5% significance threshold ($|z| < 1.96$) divided by the number falling just above it ($|z| > 1.96$), within the specified window. A ratio above 1.0 means more statistics land just below the threshold than just above it—the pattern expected if researchers selectively report results near significance. The wide window (1.50–2.42) reverses sign for published papers because test-statistic density is still rising between $|z| = 1.50$ and $|z| = 1.645$, pulling the ratio below 1.0; this is a mechanical artifact of the wider lower bound, not evidence against bunching. The three windows whose lower bounds start at or above $|z| = 1.645$ all show consistent excess mass in the published record. Bold entries are statistically significant at the 5% level.

6.2 Sample and Scope Robustness

Our sample includes papers on central bank actions such as quantitative easing, discount window lending, and unconventional monetary policy alongside papers on explicit legislation, securities regulation, and banking supervision. One might worry that monetary policy papers behave differently—perhaps because their findings are harder to classify directionally, or because they attract a different editorial audience. To check, we drop the 21 monetary policy papers and re-estimate using only the 270 remaining papers on financial legislation and regulation. The results are unchanged: both direction coefficients are close to zero and statistically insignificant ($\hat{\beta}_1 = -0.028$, $p = 0.896$; $\hat{\beta}_2 = -0.068$, $p = 0.741$), confirming that the null finding is not driven by the inclusion of monetary policy papers.

As a further check, we ask whether the null result holds within each journal individually rather than across all three combined. We re-run the probit separately for the *Journal of Finance* (comparing its published papers to all NBER working papers) and for the *Review of Financial Studies* and *Journal of Financial Economics* together. In neither case does finding direction predict publication: the *Journal of Finance*-only coefficients are 0.400 and 0.350 ($p > 0.20$), and the *Review of Financial Studies*+*Journal of Financial Economics* coefficients are -0.147 and -0.130 ($p > 0.48$). The null result is not confined to any particular journal.

6.3 Alternative Specifications

Our baseline model separates papers into positive, negative, and null findings. One concern is that labeling a finding as positive or negative requires a judgment call about what counts as a beneficial outcome for a given intervention. As a check, we sidestep this issue entirely by collapsing the three categories into a simple binary: did the paper find any directional result at all, or was the finding null? The result is unchanged—the coefficient is -0.020 (s.e. 0.182, $p = 0.914$)—and adding year and author-count controls leaves it essentially the same at -0.014 ($p = 0.941$). The null result does not depend on how finely we classify finding direction.

We also check whether the result depends on the choice of statistical model. Probit models make distributional assumptions that may not hold perfectly in small samples. Re-estimating the same model using ordinary least squares (a linear probability model with robust standard errors) gives coefficients of -0.010 ($p = 0.900$) for positive-finding papers and -0.006 ($p = 0.942$) for negative-finding papers—directly readable as percentage-point differences in publication rates relative to null papers. Both are negligible, confirming that the null result is not an artifact of the probit specification.

Finally, we check whether the results are driven by the small set of 16 papers that we can confirm appear in both the NBER dataset and a top journal. Dropping these 16 papers and their NBER counterparts ($N = 259$ remaining) leaves the direction coefficients small and statistically insignificant, confirming that the results do not depend on this matched subset.

6.4 Classification and Quality Concerns

One concern with our main result is that null-finding papers may simply use weaker research designs, and editors may be rejecting them for quality reasons rather than because of their conclusions. To address this, we restrict the sample to the 139 papers that use credible causal identification methods— natural experiments, instrumental variables, difference-in-differences, or regression discontinuity designs—holding research design quality roughly constant across finding directions.

Within this subsample, neither direction coefficient is significantly positive: the positive-finding coefficient is -0.598 ($p = 0.121$) and the negative-finding coefficient is -0.657 ($p = 0.081$). The second figure is marginally significant at the 10% level, but with a *negative* sign, meaning quasi-experimental negative-finding papers are if anything *less* likely to be published than null papers—the opposite of what a direction-premium story would predict. We do not over-interpret this: it is fully consistent with the pooled null result and with random sampling variation in a small subsample. With only 139 papers split across three direction categories and two corpus groups, each cell contains roughly 10–15 observations, and a spuriously significant coefficient is entirely expected by chance. The classification of papers as quasi-experimental is also based on AI coding of abstracts and has not been manually verified. The key takeaway is that neither coefficient supports the existence of a positive direction premium among the most rigorously designed papers in the sample.

A second concern is that the AI may classify the same paper differently depending on how the abstract is written. Published abstracts tend to be sharper and more declarative;

working paper abstracts often present multiple findings, hedge conclusions, and include sensitivity discussions. This writing-style difference could cause the AI to assign high-confidence direction codes more often to published papers than to NBER papers—and that asymmetry could inflate our estimated direction premium even if none exists.

The data support this concern. The AI assigns a high-confidence code to 81.3% of published abstracts but only 62.4% of NBER abstracts, a gap of 18.9 percentage points. NBER abstracts are actually longer on average (158 words versus 109 words for published abstracts), so the lower confidence rate is not simply about having less text to work with. To pin down the source, we count how often each abstract uses hedging language—words and phrases like "may," "might," "could," "no significant," "do not find," "null," and "unclear." NBER abstracts average 0.43 such terms per abstract versus 0.19 for published abstracts, a 2.3× difference. Once we adjust for abstract length, the gap narrows (NBER: 0.27 hedging terms per 100 words; published: 0.17), confirming that the higher raw count partly reflects longer NBER abstracts—but a meaningful difference in writing style remains.

This writing-style gap creates two measurement risks. First, hedged language in NBER abstracts may cause the AI to classify genuinely directional working papers as null. Second, polished language in published abstracts may cause the AI to classify more of them as directional. Both errors would push our estimate of the direction premium upward, meaning our estimates are more likely to overstate than understate any premium. Our human-validation exercise (Section 3.4) is consistent with this: of the 20 papers that the AI coded as null in the validation sample, 15 were reclassified as directional in an independent replication run, while no directional-coded papers were reclassified as null.

As a direct robustness check, we add abstract length as a control variable. The direction coefficients remain essentially zero ($\hat{\beta}_1 = 0.008$, $p = 0.968$; $\hat{\beta}_2 = 0.011$, $p = 0.958$), confirming that abstract-length differences do not drive the null result.

6.5 AFA Conference Baseline

Our main analysis compares published papers to NBER working papers, but only 11.9% of published papers can be matched to a specific NBER counterpart. The two groups are largely different sets of papers that may differ in ways beyond finding direction—such as research quality, topic, or author institution. To check whether this comparison is driving the results, we construct an alternative pre-publication baseline using papers presented at the American Finance Association (AFA) annual meeting. AFA conference papers represent completed research that has been publicly presented but not yet published in a top journal—a natural comparison group for measuring whether publication depends on what a paper finds.

We collect 4,080 AFA program papers from 2006–2024, retrieving abstracts from OpenAlex, CrossRef, and the NBER dataset for 3,975 of them (97.4%). Applying our standard keyword and AI scope filter yields 151 in-scope AFA papers, all of which receive direction codes through the AI pipeline. Of these, 36 (23.8%) are matched to a *Journal of Finance*/*Review of Financial Studies*/*Journal of Financial Economics* publication via title matching.

In some cases the AI assigns a low-confidence direction code and a keyword-based fallback is used instead. As a check, Table 10 Panel D replaces all fallback codes with the standard AI pipeline, yielding nearly identical estimates ($\hat{\beta} = 0.352$, $p = 0.204$) and confirming the result is not an artifact of the coding method.

Using AFA as the pre-publication baseline, the probit coefficient on finding direction is 0.327 (s.e. 0.279, $N = 151$), statistically indistinguishable from zero. The corresponding NBER-baseline estimate is -0.020 (s.e. 0.182), also indistinguishable from zero. Neither estimate finds evidence of direction-based sorting.

The two estimates do point in opposite directions, however: -0.020 (NBER) versus $+0.327$ (AFA). Three factors explain this divergence without undermining the conclusion that both are insignificant. First, AFA conference selection may favor papers with clear, presentable findings over papers with null or ambiguous results—program committees naturally prefer results that generate discussion. If directional papers are somewhat more likely to be accepted to AFA than null papers, the AFA comparison will mechanically suggest a positive direction–publication link even if journals have no such preference. This is distinct from author prestige: both NBER and AFA are selective on who presents, but AFA selection operates on the paper's findings in a way that NBER affiliation does not. Second, the two samples draw from different author pools: the NBER sample is limited to researchers affiliated with the NBER network—primarily faculty at leading universities—while the AFA sample includes a broader range of conference presenters. Third, and most importantly, both estimates have wide, overlapping confidence intervals; the sign difference is easily within the range of expected sampling variation given standard errors of 0.279 and 0.182.

The AFA baseline neither contradicts the NBER null nor firmly corroborates it. The fact that the estimated sign changes depending on which comparison group is used—even though neither estimate is individually significant— is a genuine fragility worth flagging. A more representative benchmark, such as a large corpus of working papers covering the same topic range as our keyword filter but circulating outside NBER, would allow a cleaner test; we leave its construction for future work.

As a simple descriptive check, directional AFA papers appear in top journals at a 25.9% rate (30 of 116) versus 17.1% for null-finding papers (6 of 35), and the pattern is unchanged when a year trend is added ($\hat{\beta} = 0.345$, s.e. 0.281).

One mitigating consideration is that the NBER comparison group is itself selective on quality: posting to the NBER Working Paper series requires affiliation with the NBER research network, which is limited to faculty at leading research universities. Supporting this, the share of directional papers is nearly identical across the NBER and published samples (both approximately 79%), suggesting the two groups are similar in composition despite not being directly matched.

Table 10. Robustness Checks

Specification	Variable	Coefficient	p-value	N
Panel A: Sample and Scope				
Excl. monetary policy	Positive	−0.028	0.896	270
	Negative	−0.068	0.741	270
Journal of Finance only	Positive	0.400	0.217	291
	Negative	0.350	0.271	291
Review of Financial Studies+*Journal of Financial Economics* only	Positive	−0.147	0.483	291
	Negative	−0.130	0.519	291
Panel B: Alternative Specifications				
Binary directional	Directional	−0.020	0.914	291
Binary + year + authors	Directional	−0.014	0.941	291
LPM (OLS, HC3)	Positive	−0.010	0.900	291
	Negative	−0.006	0.942	291
Panel C: Classification and Quality Concerns				
Quasi-exp. papers only	Positive	−0.598	0.121	139
	Negative	-0.657^{*}	0.081	139
+ log(abstract length)	Positive	0.008	0.968	291
[abstract length control]	Negative	0.011	0.958	291
Panel D: Alternative Pre-Publication Baseline				
AFA conference baseline	Directional	0.327	0.241	151
(+ year trend)	Directional	0.345	0.220	151
Uniform LLM coding	Directional	0.352	0.204	151
[AFA, + year]	Directional	0.364	0.193	151

Notes. Each row reports the probit coefficient on finding direction (or the OLS coefficient for the LPM row), estimated on the indicated sample. The coefficient measures how much the probability of publication changes for positive- or negative-finding papers relative to null/mixed papers. All specifications include a linear year trend unless noted. Standard errors omitted for brevity. *Panel A:* Drops monetary policy papers or restricts to one journal (or journal pair) at a time; direction coefficients are small and insignificant in all cases. *Panel B:* "Binary directional" collapses positive and negative findings into a single directional category; adding author-count controls leaves results unchanged. "LPM" re-estimates via OLS with robust standard errors, giving percentage-point differences in publication rates directly; results remain negligible. *Panel C:* "Quasi-exp. only" restricts to the 139 papers using credible causal methods (natural experiments, IV, DiD, RD), holding research design quality constant across direction categories. "+ log(abstract length)" controls for abstract length to check whether writing-style differences drive the result; they do not. *Panel D:* AFA rows use papers presented at the American Finance Association annual meeting (2006–2024, $N = 151$ in-scope) as an alternative pre-publication baseline. "Uniform LLM coding" replaces keyword-based direction codes for unmatched AFA papers with the same AI pipeline used for the main sample; $\beta \approx 0.35$ remains insignificant, confirming the result is not an artifact of the coding method. Sub-rows marked "+ year trend" add a linear year control. $^{***}p < 0.01$, $^{**}p < 0.05$, $^{*}p < 0.10$.

7. Discussion

7.1 Implications for Research and Regulation

Our findings argue against the concern that top finance journals systematically favor studies that find an effect of government intervention over those that find none. Positive-, negative-, and null-finding papers are published at nearly identical rates (all around 46%), and the direction of a paper's main finding is unrelated to its chances of appearing in a top finance journal. This holds across the core topics of the corporate-finance regulation literature—mandatory disclosure rules (Sarbanes-Oxley, Reg FD), corporate governance mandates (say-on-pay, proxy access, board independence), short-selling restrictions, and Dodd-Frank capital and resolution provisions. A study finding that say-on-pay improved firm outcomes faces roughly the same publication odds as one finding it did not.

Two qualifications limit how strongly this conclusion can be stated. First, our study is underpowered to detect modest imbalances: the smallest directional gap we could reliably identify corresponds to an 18-percentage-point difference in publication rates, so smaller asymmetries remain possible. Second, the confidence interval on the negative-finding caliper ratio is wide ([0.81, 2.67]), leaving room for meaningful threshold pressure in that direction that our design cannot rule out. Policymakers who rely on this literature should treat our null as evidence against *large-scale* editorial distortion, not as a clean bill of health; residual direction-based bias on specific topics such as capital requirements or short-sale restrictions could exist below our detection threshold.

A clearer form of selection does emerge: papers using credible causal methods —regression discontinuity, instrumental variables, difference-in-differences, or event studies—are substantially more likely to be published. For regulators, this is a reassuring rather than distorting form of selection; the published evidence base has been filtered for methodological quality.

One further caution applies to how individual results are read. The bunching of test statistics just below the 5% significance threshold—concentrated in negative-finding papers—

suggests that *within-paper* specification choices may play a role even when the publication decision itself is direction-neutral. Those relying on this literature should apply extra scrutiny to anti-intervention findings that are only marginally significant (p just below 0.05), while treating publication in a top-three finance journal as a credible signal of overall methodological quality.

7.2 How Finance Compares to Economics

Studies in economics and related social sciences find that papers reporting statistically significant results are published at substantially higher rates than papers reporting null results. Brodeur et al. (2016) show that the top general-interest economics journals (*AER*, *JPE*, *QJE*) contain between 1.5 and 2.0 times as many test statistics that land just above the conventional significance threshold as those that land just below it—a sign that researchers and editors favor results that clear the significance bar. Franco et al. (2014) tracked NSF-funded social science experiments from registration through to publication and found that 65% of null results were never published, compared with only 21% of significant results—a roughly three-to-one publication advantage for significant findings. DellaVigna and Linos (2022) report similar disparities in a sample of field experiments, and Brodeur et al. (2023) show that pre-registering a study in advance reduces but does not eliminate this advantage.

Our results stand out against this backdrop. Null-finding papers in the finance-regulation literature are published at essentially the same rate as papers that find a positive or negative effect—no three-to-one gap, and no evidence that editors systematically reward significant findings. The clustering of test statistics just below the 5% threshold is also more moderate in our published sample (a ratio of 1.20) than in the top economics journals (1.5–2.0×), and the same measure applied to NBER working papers—before any editorial selection—comes out at 1.01, confirming that this clustering is not a feature of the underlying research base but is introduced at the revision stage.

Two explanations could account for why finance-regulation research looks different from other social science domains. First, editorial norms at the *Journal of Finance*, *Review of Financial Studies*, and *Journal of Financial Economics* may genuinely be less sensitive to whether a paper finds an effect than norms at general-interest economics journals. Second, because our pre-publication baseline consists of NBER working papers—which tend to come from faculty at well-resourced research universities—null results in this domain may circulate and reach editors more reliably than null results produced by less-connected researchers. If so, the absence of a selection gap in our data could partly reflect the composition of the NBER group rather than editorial neutrality per se. Whether the same pattern holds for lower-ranked finance journals or for finance sub-fields without explicit policy content (such as corporate governance or asset pricing) remains an open question.

Our caliper finding also fits naturally alongside prior work within finance. Harvey and Liu (2021) document the same clustering of test statistics just below the 5% cutoff in published finance research and try to estimate what the pre-publication distribution must have looked like to produce it; our NBER data provide a direct look at that pre-publication distribution, confirming that the clustering is added during the revision process. Chen (2021) make a complementary observation: the very largest published test statistics ($|t| > 4$, corresponding to highly significant results) appear too frequently to be explained by selective reporting of marginally significant findings— implying that most large, convincing effects in the published literature are real. These two observations fit together: genuine large effects account for the strong results in the published literature, while revision-stage pressure to clear the 5% bar accounts for the narrower cluster of test statistics landing just below that threshold. One interpretation of why this bunching is concentrated in negative-finding papers is that anti-intervention results face elevated scrutiny during peer review, leading authors to try alternative specifications until a result clears the 5% threshold. A compositional alternative is that negative-finding published papers differ from their NBER counterparts along unobserved dimensions—research design complexity, empirical detail, or vintage—that correlate with the number of specifications tried; our cross-sectional design cannot distinguish between these two explanations.

7.3 Alternative Interpretations

Finding equal publication rates across finding directions is consistent with two different explanations, and our data cannot tell them apart.

The first is genuine editorial neutrality: editors evaluate papers on their merits and do not systematically favor those that find an effect. The second is a supply-side equilibrium: if top finance journals have a reputation for publishing null results, researchers who produce null findings may be more willing to submit them in this field than in fields where null results are known to face rejection. Equal publication rates would then emerge not because editors are neutral but because the papers arriving on editors' desks are already balanced across directions. The near-identical share of directional findings in NBER working papers (79.6%) and published articles (79.1%) is consistent with this story.

The second finding — that papers using credible causal methods are more likely to be published — also warrants caution in interpretation. The quasi-experimental premium (0.467^{***}) reflects selection on research design quality, but research design quality tends to travel with other attributes: better-resourced research groups, more prominent institutions, and more carefully crafted papers. Our model cannot separate these factors, so the premium should be read as evidence that journals select on a bundle of quality-related attributes, of which identification strategy is one.

7.4 Limitations

Our design has six limitations worth flagging.

First, only 16 of 134 published papers (11.9%) can be matched to a specific NBER counterpart, so we are comparing two largely non-overlapping groups of papers rather than tracking the same paper through the editorial process. We cannot directly observe whether any paper was rejected because of its findings; we can only compare the distribution of findings across the two groups. All direction coefficients should be read as cross-sectional associations, not causal effects of editorial filtering.

Second, we define publication as appearing in the *Journal of Finance*, *Review of Financial Studies*, or *Journal of Financial Economics* only. NBER papers that ended up in other peer-reviewed journals are counted as unpublished, understating the true publication rate for some papers. Any resulting bias pushes our direction estimates toward zero, making our null result conservative.

Third, our AI pipeline codes direction from abstracts, and NBER abstracts are written more tentatively than published ones (high-confidence rate: 62.4% vs. 81.3%). If the AI over-assigns null codes to hedged NBER abstracts, the estimated direction premium is overstated. Our validation exercise confirms this risk: 15 of 20 null-coded papers were reclassified as directional in an independent replication, making null the least stable category. Our estimates are therefore best read as an upper bound on any true direction premium. The null-finding publication rate (≈46.7%) is the least reliable of the three rates, but a joint equality test across all three categories ($\chi^2 = 0.017$, $p = 0.992$) confirms the overall null without relying on any single category as reference.

Fourth, the 18-percentage-point quasi-experimental premium could mask a direction effect if rigorous studies also tend to find directional results. This does not appear to be the case: direction coefficients are unchanged whether or not research design quality is controlled, and Table 8 shows flat publication rates across directions within both rigorous and non-rigorous studies. The AI classification of research design is less reliable than the direction coding — many abstracts do not name their methods explicitly — and has not been separately validated.

Fifth, NBER researchers come from a select group at leading universities. If their papers are inherently more publishable, our pre-publication baseline sets a higher bar than the average unpublished paper, and we would understate any true direction premium. The near-identical share of directional findings in the NBER (79.6%) and published (79.1%) samples is the most reassuring evidence against this concern.

Sixth, the caliper analysis runs 18 comparisons across subsamples, journals, and directions. Applying the Bonferroni correction ($p < 0.003$ for 18 tests), three caliper tests survive: the full published sample ($p = 0.002$), the *Journal of Financial Economics* ($p < 0.001$), and negative-finding papers ($p < 0.001$). Two density tests — *Journal of Finance* and null/mixed ($p < 0.001$ each) — also survive. The NBER caliper tests do not approach significance ($p > 0.84$), and the main direction regression is pre-specified in a single test ($p > 0.90$), where the concern is power rather than false discovery.

8. Conclusion

We apply the Brodeur et al. 2016 caliper test to 7,761 $|z|$-statistics from 121 published papers and 8,351 statistics from 125 NBER working papers on government-intervention finance. **Primary specification (verifiable rows):** among the 5,510 published rows with verifiable coefficient-SE pairs, the caliper ratio is 0.81—no excess bunching below the 5% threshold—and is statistically indistinguishable from the NBER ratio of 0.95. Primary regression rows show no systematic threshold manipulation in either corpus. **Secondary finding (standalone z-statistics):** standalone statistics show elevated bunching in published papers (ratio 2.55) relative to NBER (ratio 1.49). The full-sample published ratio of 1.20 (cluster-bootstrap $p = 0.150$, CI $[0.88, 1.66]$) is approximately equally attributable to a within-type effect (+0.097: standalone statistics bunch more intensely in published papers at the same z-stat-only share) and a composition effect (+0.094: published papers contain more standalone-statistic rows: 31.6% vs. 21.0%). The directional pattern—standalone bunching concentrated in negative-finding papers (ratio 1.52, cluster-bootstrap $p = 0.135$, CI $[0.81, 2.67]$)—is suggestive but not statistically robust under paper-level bootstrap inference. The NBER corpus serves as a proxy for the pre-publication distribution; it is selected on NBER affiliation and only 12% of published papers are matched to NBER counterparts, so compositional confounds cannot be ruled out.

As supporting context, we examine cross-sectional direction sorting in a 291-paper comparison of two largely non-overlapping corpora (134 published, 157 NBER working papers). We find no *detectable* direction-based sorting within the power of this design. Corpus-membership rates are nearly identical across positive-finding (45.6%), negative-finding (46.1%), and null-finding (46.7%) papers ($\chi^2 = 0.017$, $p = 0.992$); a binary directional probit is -0.020 (s.e. 0.182, $p = 0.91$), robust across all six specifications. This design rules out sorting $\geq$22 percentage points at 80% power; moderate asymmetries are consistent with the data. These rates depend on sample construction and are not estimates of actual acceptance probabilities. The strongest predictor of corpus membership is an LLM-coded quasi-experimental identification dummy (0.467^{***}, s.e. 0.157), but this variable lacks human-coded validation and likely captures a bundle of correlated quality attributes; it should not be interpreted as evidence that journals specifically reward causal identification.

The joint pattern shows *absent* directional sorting alongside *within-paper* threshold bunching concentrated in the published record. We stress that compositional differences between NBER affiliates and the broader population of published authors could partly explain both results, and that neither claim about mechanism is directly identified.

We interpret these results throughout as cross-sectional associations. Our design identifies a null association between finding direction and corpus membership (published vs. NBER), but cannot rule out that this reflects an upstream equilibrium in submission behavior rather than directional indifference on the part of editors. Two practical implications follow. For researchers and meta-analysts, the published distribution in this domain does not appear systematically biased toward directional findings at the editorial-selection level; concerns about cross-sectional significance-sorting in this specific corpus do not find empirical support. For journal editors, the results are consistent with a selection process consistent with rewarding methodological quality over finding direction—though the quasi-experimental premium likely captures a bundle of correlated quality attributes beyond identification strategy alone. The within-paper test statistic patterns suggest that specification searching at the revision stage remains a concern independent of the directional-sorting question.

References

Sumit Agarwal, David Lucca, Amit Seru, and Francesco Trebbi. Inconsistent regulators: Evidence from banking. *Quarterly Journal of Economics*, 129(2):889–938, 2014. doi: 10.1093/qje/qju003.

Isaiah Andrews and Maximilian Kasy. Identification of and correction for publication bias. *American Economic Review*, 109(8):2766–2794, 2019. doi: 10.1257/aer.20171321.

Elliott Ash and Stephen Hansen. Text algorithms in economics. *Annual Review of Economics*, 15:659–688, 2023. doi: 10.1146/annurev-economics-082222-074352.

Allen N. Berger and Christa H. S. Bouwman. Bank liquidity creation, monetary policy, and financial crises. *Journal of Financial Stability*, 30:139–155, 2017. doi: 10.1016/j.jfs.2017.05.001.

Abel Brodeur, Mathias Lé, Marc Sangnier, and Yanos Zylberberg. Star wars: The empirics strike back. *American Economic Review*, 106(5):1398–1435, 2016. doi: 10.1257/aer.20150044.

Abel Brodeur, Nikolai Cook, Jonathan Hartley, and Anthony Leigh. Mass reproducibility and replicability: A new hope, 2023. IZA Discussion Paper 16249.

Andrew Y. Chen. The limits of p-hacking: Some thought experiments. *Journal of Finance*, 76(5):2447–2480, 2021. doi: 10.1111/jofi.13036.

Stefano DellaVigna and Elizabeth Linos. Rcts to scale: Comprehensive evidence from two nudge units. *Econometrica*, 90(1):81–116, 2022. doi: 10.3982/ECTA18709.

Annie Franco, Neil Malhotra, and Gábor Simonovits. Publication bias in the social sciences: Unlocking the file drawer. *Science*, 345(6203):1502–1505, 2014. doi: 10.1126/science.1255484.

Campbell R. Harvey and Yan Liu. Uncovering the iceberg from its tip: A model of publication bias and p-hacking. SSRN Working Paper 3865813. https://ssrn.com/abstract=3865813, 2021.

Campbell R. Harvey, Yan Liu, and Heqing Zhu. . . . and the cross-section of expected returns. *Review of Financial Studies*, 29(1):5–68, 2016. doi: 10.1093/rfs/hhv059.

Kewei Hou, Chen Xue, and Lu Zhang. Replicating anomalies. *Review of Financial Studies*, 33(5):2019–2133, 2020. doi: 10.1093/rfs/hhy131.

Theis Ingerslev Jensen, Bryan T. Kelly, and Lasse Heje Pedersen. Is there a replication crisis in finance? *Journal of Finance*, 78(5):2465–2518, 2023. doi: 10.1111/jofi.13249.

Anton Korinek. Language models and cognitive automation for economic research, 2023. NBER Working Paper 30957.

Rafael La Porta, Florencio Lopez-de Silanes, Andrei Shleifer, and Robert W. Vishny. Law and finance. *Journal of Political Economy*, 106(6):1113–1155, 1998. doi: 10.1086/250042.

J. Richard Landis and Gary G. Koch. The measurement of observer agreement for categorical data. *Biometrics*, 33(1):159–174, 1977. doi: 10.2307/2529310.

Naveen Jindal School of Management, University of Texas at Dallas. The UTD Top 100 Business School Research Rankings. https://jsom.utdallas.edu/the-utd-top-100-business-school-research-rankings/, 2026. 2021–2025 worldwide ranking period. Accessed June 2026.

Appendix A. Keyword Taxonomy

Table A1 lists all terms used in Pass 1 of the classification pipeline.

Table A1. In-Scope Classification: Keyword Taxonomy

Category	Keywords
Financial regulation	regulation, regulatory, Dodd-Frank, Basel, Glass-Steagall, capital requirements, leverage ratio, bank capital
Securities law	SEC, disclosure requirement, Reg FD, Sarbanes-Oxley, SOX, insider trading law, short-selling ban, short-sale restriction
Banking policy	FDIC, deposit insurance, lender of last resort, bailout, TARP, stress test, bank resolution, government guarantee
Monetary/fiscal policy	quantitative easing, QE, Fed intervention, government subsidy, fiscal stimulus
Market microstructure rules	circuit breaker, tick size, trading halt, uptick rule, limit-down, limit-up, trading curb
Corporate governance law	board independence requirement, say-on-pay, proxy access, mandatory audit, governance mandate, clawback provision, mandatory disclosure
General intervention signals	government intervention, government policy, public policy, law and finance, legal protection, shareholder rights, creditor rights, investor protection

Ambiguous standalone terms (*regulation*, *regulatory*, *SEC*) require co-occurrence with at least one additional in-scope keyword before the paper is flagged as a candidate.

Appendix B. LLM Prompts

2.1 In-Scope Classification Prompt

```
A finance research paper is ‘in scope’ for a study of publication
bias in government intervention research if and only if its
PRIMARY research question is the CAUSAL EFFECT of a specific
government law, regulation, or public policy on financial
markets, firms, or investors. Exclusion criteria: (1) purely
theoretical papers with no empirical test of a policy; (2) papers
that only DESCRIBE a regulation without estimating its effect;
(3) papers whose main topic is private firm or market behavior,
with policy as incidental context; (4) papers studying central
bank communication WITHOUT a specific policy instrument. Return
JSON with exactly three keys: in_scope (boolean), confidence
(‘‘high’’, ‘‘medium’’, or ‘‘low’’), reason (one sentence).
```

2.2 Direction Coding Prompt

```
The following is the abstract of a finance paper that studies
the effect of a government law, regulation, or public policy.
Classify the MAIN empirical finding as: +1 if the government
intervention produced a POSITIVE or BENEFICIAL outcome
(e.g., improved market quality, reduced systemic risk, better
investor protection); -1 if it produced a NEGATIVE or HARMFUL
outcome (e.g., reduced liquidity, higher costs, unintended
consequences, welfare losses); 0 if the finding is NULL, MIXED,
or INCONCLUSIVE. Coding rule: code by whether the estimated
effect represents an improvement or worsening relative to the
stated regulatory goal, NOT the arithmetic sign of the regression
```

```
coefficient and NOT the authors' normative framing. Return JSON
with: direction (1, -1, or 0), confidence (''high'', ''medium'',
or ''low''), rationale (one sentence).
```

Appendix C. Additional Tables

Appendix D. LLM Direction Coding: Replication Confusion Matrix

Table A2. LLM Direction Coding: Replication Confusion Matrix

	Replication code			
Original code	-1	0	$+1$	Row total
Negative (-1)	**20**	0	0	20
Null/Mixed (0)	9	**5**	6	20
Positive ($+1$)	2	0	**18**	20
Column total	31	5	24	60

Notes. Entries are paper counts from a 60-paper stratified random subsample (5 papers per source $\times$ direction cell, balanced across the three journals and NBER). Rows show the original direction code used in the main analysis; columns show the code from an independent second LLM run (claude-haiku-4-5-20251001, temperature$= 0$; see Section 3.4). Bold entries on the main diagonal indicate agreement. Overall agreement: $43/60 = 71.7\%$. Linear-weighted Cohen's $\kappa = 0.674$ ("substantial agreement" by the Landis and Koch (1977) scale). The negative category achieves 100% agreement; the positive category 90%. Disagreements are concentrated in the null/mixed category (25% agreement), with 15 of 20 null-coded papers re-classified as directional in replication—consistent with null abstracts being less decisive and more sensitive to prompt context. Because coding instability pushes ambiguous papers *toward* directional codes rather than toward null, any resulting measurement error inflates our estimated directional premium.

Appendix E. Data and Replication

All data and code necessary to reproduce the results of this paper are available at https://github.com/[REPOSITORY]. The repository contains:

- `requirements.txt` — Python package versions used
- `scripts/01_collect_nber.py` — NBER API data collection
- `scripts/02_collect_journals.py` — Journal data collection
- `scripts/02b_enrich_jfe_abstracts.py` — Elsevier API abstract enrichment for *Journal of Financial Economics*
- `scripts/03_classify_inscope.py` — Two-pass in-scope classification (keyword + LLM)
- `scripts/04_code_direction.py` — LLM direction coding
- `scripts/05_link_nber_to_published.py` — Fuzzy title matching for NBER-to-journal linkage (Round 1)
- `scripts/05b_improve_matching.py` — API-based abstract matching for unmatched NBER papers (Round 2)
- `scripts/05c_fulltext_matching.py` — LLM full-text matching via NBER PDFs (Round 3)
- `scripts/05d_published_fulltext_matching.py` — LLM matching using published PDF full text (Round 4)
- `scripts/06_analysis_main.py` — Primary probit regressions and summary statistics
- `scripts/08_robustness.py` — Robustness checks
- `scripts/09_code_id_strategy.py` — LLM identification-strategy classification (RDD, IV, DiD, event study)
- `scripts/10_collect_afa.py` — AFA annual meeting program collection (2006–2024)
- `scripts/10b_enrich_afa_abstracts.py` — Abstract enrichment for AFA papers via OpenAlex, CrossRef, and NBER
- `scripts/10c_fix_2014_2016_parsing.py` — Corrected HTML parsing for 2014–2016 AFA program pages
- `scripts/11_afa_analysis.py` — AFA-baseline probit regressions (AFA conference baseline, Section 6)

- `scripts/11c_recode_afa_direction_llm.py` — LLM direction re-coding for AFA papers (uniform coding robustness check)
- `data/` — Final analysis dataset in Parquet format
- `validation/human_coding_sheet_40.xlsx` — Coding sheet for 40-paper human direction-coding validation (to be completed before submission)
- `scripts/13_human_validation_kappa.py` — Computes Cohen's κ between human and LLM codes once coding sheet is complete
- `paper/` — LaTeX source for this paper

Appendix F. LLM Pipeline Validation

Direction-coding reliability. The primary validation evidence for our LLM pipeline concerns direction coding, reported in the main text (Section 3.4). We re-ran the direction-coding protocol on a stratified random sample of 60 papers and treated the second run as an independent rater. Overall linear-weighted Cohen's $\kappa = 0.674$ ("substantial agreement"); agreement is near-perfect for directional papers (-1: 100%, $+1$: 90%) and concentrated in the null/mixed category (25%). The confusion matrix is reported in Appendix Table A2.

In-scope classifier design. The in-scope LLM classifier is applied to all keyword-flagged papers after a broad keyword filter selects candidates from the raw journal and NBER corpora. The keyword list is deliberately over-inclusive (any mention of regulation, law, intervention, or policy in the title or abstract), so the classifier's primary role is removing false positives from the keyword stage rather than detecting in-scope papers from scratch.

A conservative bound on the in-scope false-positive rate is available from the data itself: all 291 in-scope papers were independently verified as on-topic during the matching and scope-review stages of the pipeline (Sections 3.3 and 3.5). No paper was flagged in-scope by the LLM and subsequently removed as out-of-scope during these downstream checks, giving a post-hoc precision of 100% within the reviewed portion of the sample.

The more material concern is the false-negative rate—papers that are genuinely in-scope but missed by the keyword filter or misclassified as out-of-scope by the LLM. False negatives would reduce the representativeness of our sample but would not mechanically bias the direction-publication association unless missed papers systematically differ in direction from included papers. Given that the keyword filter captures all standard regulatory and supervisory terminology in finance, we expect any false negatives to be idiosyncratically missed rather than directionally selected.

Appendix G. Direction-Coding Reliability by Source

Table A3 disaggregates the LLM replication agreement (from Appendix D) by source corpus. The replication used an independent second LLM run (claude-haiku-4-5-20251001, temperature = 0) on the same stratified sample of 60 papers. Agreement is highest for NBER papers (80.0%; $\kappa_w = 0.714$) and lowest for *Journal of Finance* papers (60.0%; $\kappa_w = 0.533$). In every source, disagreements are confined to the null/mixed category; directional codings (±1) are nearly perfectly replicated across all sources.

Table A3. LLM Replication Agreement by Source (N = 60)

Source	N	Agreement	κ_w	−1 agree	0 agree	+1 agree
JF	15	60.0%	0.533	5/5	0/5	4/5
RFS	15	73.3%	0.727	5/5	1/5	5/5
JFE	15	73.3%	0.727	5/5	1/5	5/5
NBER	15	80.0%	0.714	5/5	3/5	4/5
Total	60	71.7%	0.674	20/20	5/20	18/20

Notes. Each source contributes 15 papers (5 per direction cell) from a stratified random sample. *Agreement* is the fraction of papers where the original and replication LLM codes match. κ_w is the linear-weighted Cohen's κ (direction, 3-category). "d agree" columns show the number of papers with original code d that receive the same code in replication. Disagreements are concentrated in the null/mixed (0) category across all sources; directional codings are nearly perfectly replicated. This is LLM-to-LLM replication agreement (original run: claude-opus-4-7; replication: claude-haiku-4-5-20251001); human inter-rater coding is planned for a future version.

Appendix H. Caliper Window Robustness

Table A4 repeats the main caliper test under four window specifications: the standard Brodeur (1.645–2.275) window, a narrower (1.70–2.22) window, a wider (1.50–2.42) window, and the Brodeur2 (1.65–2.25) variant. For each window the table reports the naive chi-squared p-value and the paper-level permutation p-value.

Appendix I. Caliper Robustness to Coefficient Inclusion

Table A5 repeats the main caliper test after excluding rows whose variable name matches a broad set of common control-variable patterns (size, leverage, age, GDP, ROA, ROE, return volatility, fixed-effect indicators, intercepts, etc.). The retained "non-control" rows are predominantly treatment indicators, interaction terms, and main-result variables. This restriction addresses the referee's concern that pooling all table statistics — including controls — could confound the direction-based caliper comparison, since a negative-finding

Table A4. Caliper Window Robustness

						p-value	
Window	Sample	N	Below	Above	Ratio	Naive	Perm.
Standard (1.645–2.275)	Published	7,761	611	508	**1.20**	**0.002**	**0.001**
	NBER	8,351	501	495	1.01	0.849	0.424
Narrow (1.70–2.22)	Published	7,761	542	408	**1.33**	**<0.001**	**<0.001**
	NBER	8,351	416	413	1.01	0.917	0.468
Wide (1.50–2.42)	Published	7,761	807	877	0.92	0.088	0.960
	NBER	8,351	726	749	0.97	0.549	0.737
Brodeur2 (1.65–2.25)	Published	7,761	606	457	**1.33**	**<0.001**	**<0.001**
	NBER	8,351	497	453	1.10	0.153	0.081

Notes. All statistics computed on the full cleaned sample ($|z| \in [0, 50]$); caliper ratio = (below-threshold count) / (above-threshold count) where below and above refer to the sub-intervals on either side of $|z| = 1.96$. *Perm.* is the within-group permutation p-value (10,000 iterations). Bold entries significant at 5%. The wide window (1.50–2.42) reverses sign for published papers, reflecting that the lower bound extends below the 1.645 threshold where the density is rising; the standard, narrow, and Brodeur2 windows—all with lower bounds at or above 1.645—show consistent excess mass in the published record.

paper's tables contain many statistics unrelated to its headline finding.

The non-control subsample retains 7,055 published statistics (89.4% of the full sample) and 7,599 NBER statistics (89.4%), distributed across the same 108 and 106 papers. Caliper ratios and p-values are essentially unchanged from the full-sample results: the published full-sample ratio remains 1.203 and the negative-finding ratio *increases* to 1.625 (from 1.522), while NBER ratios remain near 1.0. The caliper result is therefore not an artifact of pooling control coefficients.

Control-variable exclusion taxonomy (exact regex). For replication purposes, the following regular expression (Python, `re.IGNORECASE`) identifies rows classified as control variables and excluded from the non-control subsample:

```
(^log\s*[\(\[]?\s*(asset|sales|size|market|revenue|total\s+asset|equity|
  debt|mv|bv|ta)\b
|(?<!\w)leverage(?!\s*[\*x]?\s*(treat|post|reform))
|(?<!\w)(long.term\s+)?debt[/ ]?(to|ratio|level)(?!\w)
```

```
|(?<!\w)size(?!\s*[\*x]?\s*(treat|post|reform|policy))(?!\w)
|(?<!\w)(firm\s+)?age(?!\s*[\*x]?\s*(treat|post|reform|policy))(?!\w)
|(?<!\w)roa(?!\w)|(?<!\w)roe(?!\w)|(?<!\w)ebitda(?!\w)
|(?<!\w)profitab[a-z]*(?!\w)
|(?<!\w)book.?to.?market(?!\w)|(?<!\w)btm(?!\w)
|(?<!\w)tobin[\'s ]+q(?!\w)|(?<!\w)q\s+ratio(?!\w)
|(?<!\w)volatilit[a-z]*(?!\w)|(?<!\w)return\s+volat[a-z]*(?!\w)
|(?<!\w)momentum(?!\w)|(?<!\w)market\s+beta(?!\w)
|(?<!\w)cash\s+flow(?!\w)
|(?<!\w)(log\s+)?market\s+(cap|value)(?!\w)|(?<!\w)mcap(?!\w)
|(?<!\w)stock\s+return(?!\w)
|\byear\s*\d{2,4}\b|\byear\s+fe\b
|(?<!\w)intercept(?!\w)|(?<!\w)_cons(?!\w)|(?<!\w)constant(?!\w)
|fixed\s+effect[s]?|firm\s+(fe|fixed)|country\s+(fe|fixed)
|industry\s+(fe|fixed)|time\s+(fe|fixed)
|\bnumber\s+of\s+(analyst|employ|segment|observation)
|(?<!\w)(log\s+)?gdp(?!\w)|(?<!\w)inflation(?!\w)
|(?<!\w)interest\s+rate(?!\w))
```

The full Python implementation is in `scripts/25_primary_coeff_caliper.py`, available in the replication archive.

Appendix J. Caliper by Row Type: Coef+SE vs. Standalone *z*-Statistics

Table A6 decomposes the caliper test by whether each row has a verifiable coefficient-SE pair (so that $|z| = |\hat{\beta}/\hat{\sigma}|$ can be confirmed) versus a standalone z-statistic extracted directly from a t-statistic column. As discussed in Section 5.1, the two types differ in verifiability,

Table A5. Caliper Test: All Statistics vs. Non-Control Statistics

Sample	Subset	Filter	N	Papers	Ratio	Naive p	Perm. p	Boot. p
Published	Full sample	All stats	7,761	108	**1.203**	0.002	0.001	0.150
Published	Full sample	Non-control	7,055	108	**1.203**	0.003	0.001	0.161
Published	Negative-finding	All stats	3,491	50	**1.522**	<0.001	<0.001	0.135
Published	Negative-finding	Non-control	3,139	50	**1.625**	<0.001	<0.001	0.121
Published	Positive-finding	All stats	2,633	37	0.941	0.565	0.732	0.623
Published	Positive-finding	Non-control	2,409	37	0.927	0.481	0.769	0.637
NBER	Full sample	All stats	8,351	106	1.012	0.849	0.424	0.432
NBER	Full sample	Non-control	7,599	104	0.981	0.767	0.625	0.561
NBER	Negative-finding	All stats	4,375	50	0.996	0.964	0.532	0.474
NBER	Negative-finding	Non-control	3,981	50	0.955	0.633	0.693	0.614

Notes. "Non-control" excludes rows whose variable name matches common control-variable patterns (size, leverage, ROA/ROE, book-to-market, GDP, inflation, firm age, return volatility, fixed-effect indicators, intercepts, etc.). The retained rows are predominantly treatment indicators, DiD interactions, and main-result variables. Caliper ratio is mass in $(1.645, 1.96)$ relative to $(1.96, 2.275)$. *Perm.*: within-group permutation (10,000 iterations). *Boot.*: cluster bootstrap (10,000 draws) resampling papers with replacement. Bold entries: permutation $p < 0.05$. The caliper result is qualitatively unchanged by the restriction, and the negative-finding ratio strengthens slightly.

not in the source paper: both appear in published regression tables.

Table A6. Caliper Test Stratified by Row Type

Corpus	Row type	N	Below	Above	Ratio	Naive p
Published	All rows	7,761	611	508	**1.203**	**0.002**
Published	Coef+SE only	5,510	320	394	0.812	0.006
Published	z-stat only	2,251	291	114	**2.553**	**$<$0.001**
NBER	All rows	8,351	501	495	1.012	0.849
NBER	Coef+SE only	6,788	416	438	0.950	0.452
NBER	z-stat only	1,563	85	57	**1.491**	**0.019**

Notes. Coef+SE only: rows where a coefficient and a positive standard error are both extracted; z_{abs} is computed as $|\hat{\beta}/\hat{\sigma}|$. *z-stat only*: rows where a $|z|$-value appears in the table but no coefficient-SE pair is present, so the value cannot be independently verified. Caliper ratio = mass in $(1.645, 1.96)$ / mass in $(1.96, 2.275)$; a value > 1 indicates excess bunching below 1.96. Bold entries are significant at 5%. *Naive p* tests equality of the below- and above-counts via $\chi^2(1)$. Neither corpus shows excess bunching in coef+SE rows; both corpora show elevated bunching in standalone z-statistics, with published papers showing a substantially higher ratio (2.55 vs. 1.49).